

SQL Server Data Automation Through Frameworks

Building Metadata-Driven Frameworks with
T-SQL, SSIS, and Azure Data Factory

Andy Leonard
Kent Bradshaw

Apress®

SQL Server Data Automation Through Frameworks

**Building Metadata-Driven
Frameworks with T-SQL, SSIS,
and Azure Data Factory**

**Andy Leonard
Kent Bradshaw**

Apress®

SQL Server Data Automation Through Frameworks: Building Metadata-Driven Frameworks with T-SQL, SSIS, and Azure Data Factory

Andy Leonard
Farmville, VA, USA

ISBN-13 (pbk): 978-1-4842-6212-2
<https://doi.org/10.1007/978-1-4842-6213-9>

Kent Bradshaw
Providence Forge, VA, USA

ISBN-13 (electronic): 978-1-4842-6213-9

Copyright © 2020 by Andy Leonard, Kent Bradshaw

This work is subject to copyright. All rights are reserved by the Publisher, whether the whole or part of the material is concerned, specifically the rights of translation, reprinting, reuse of illustrations, recitation, broadcasting, reproduction on microfilms or in any other physical way, and transmission or information storage and retrieval, electronic adaptation, computer software, or by similar or dissimilar methodology now known or hereafter developed.

Trademarked names, logos, and images may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, logo, or image we use the names, logos, and images only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

The use in this publication of trade names, trademarks, service marks, and similar terms, even if they are not identified as such, is not to be taken as an expression of opinion as to whether or not they are subject to proprietary rights.

While the advice and information in this book are believed to be true and accurate at the date of publication, neither the authors nor the editors nor the publisher can accept any legal responsibility for any errors or omissions that may be made. The publisher makes no warranty, express or implied, with respect to the material contained herein.

Managing Director, Apress Media LLC: Welmoed Spaehr
Acquisitions Editor: Jonathan Gennick
Development Editor: Laura Berendson
Coordinating Editor: Jill Balzano

Cover image designed by Freepik (www.freepik.com)

Distributed to the book trade worldwide by Springer Science+Business Media New York, 233 Spring Street, 6th Floor, New York, NY 10013. Phone 1-800-SPRINGER, fax (201) 348-4505, e-mail orders-ny@springer-sbm.com, or visit www.springeronline.com. Apress Media, LLC is a California LLC and the sole member (owner) is Springer Science + Business Media Finance Inc (SSBM Finance Inc). SSBM Finance Inc is a **Delaware** corporation.

For information on translations, please e-mail booktranslations@springernature.com; for reprint, paperback, or audio rights, please e-mail bookpermissions@springernature.com.

Apress titles may be purchased in bulk for academic, corporate, or promotional use. eBook versions and licenses are also available for most titles. For more information, reference our Print and eBook Bulk Sales web page at <http://www.apress.com/bulk-sales>.

Any source code or other supplementary material referenced by the author in this book is available to readers on GitHub via the book's product page, located at www.apress.com/9781484262122. For more detailed information, please visit <http://www.apress.com/source-code>.

Printed on acid-free paper

For Christy
—Andy

For Ann
—Kent

Table of Contents

About the Authors.....	xi
About the Technical Reviewer	xiii
Acknowledgments	xv
Introduction	xvii
Part I: Stored Procedure-Based Database Frameworks	1
Chapter 1: Stored Procedures 101.....	3
The Need for a Framework	3
Demonstration of a Framework	4
An Example Schema.....	5
The Daily Process	6
Executing the Daily Process	12
Including a Monthly Process	12
Summary.....	18
Chapter 2: Automation with Stored Procedures	19
A Daily Process Controller.....	19
Monthly Process Controller.....	23
Summary.....	26
Chapter 3: Stored Procedure Orchestrators	27
The Orchestrator	27
A Quick Process Overview	31
A Look at Troubleshooting Issues.....	31
Summary.....	32

TABLE OF CONTENTS

Chapter 4: A Stored Procedure–Based Metadata-Driven Framework	33
Building the Metadata.....	33
Metadata-Ready Controllers	40
Metadata-Ready Orchestrator.....	47
Troubleshooting with Metadata	50
Summary.....	56
Part II: SSIS Frameworks.....	57
Chapter 5: A Simple, Custom, File-Based SSIS Framework.....	59
An SSIS Framework, Defined and Designed	60
Functionality	60
Empathy.....	60
Simplicity.....	60
Building a File-Based SSIS Framework	61
Obtain the Code	61
Metadata-Driven Execution Management.....	62
The SSISConfig Database	62
Config Schema	63
The Config.Applications Table.....	65
The Config.Packages Table.....	67
A Sample SSIS Solution	75
Configure Visual Studio 2019 for SSIS Development.....	76
Create the Sample SSIS Solution	78
SSIS Framework Metadata Management	85
Add an SSIS Application	85
Add SSIS Packages	88
Assign SSIS Application Packages	91
Conclusion	96

TABLE OF CONTENTS

Chapter 6: Framework Execution Engine	97
Create a Parent SSIS Package	97
Log Execution Values	98
Retrieve SSIS Application Packages from SSISConfig	104
Iterate Application Packages.....	120
Log Application Package Values.....	130
Execute Application Packages	135
Log Execution Results.....	150
Conclusion	158
Chapter 7: Framework Logging	159
Create a Log Schema.....	159
Add Application Instance Logging to Parent.dtsx.....	164
Add Application Package Instance Logging to Parent.dtsx	176
Viewing Execution Report Data.....	191
Configure Fault Tolerance in SSIS Framework Application Metadata	197
Conclusion	206
Chapter 8: Azure-SSIS Integration Runtime	207
Getting Started with Azure	207
Provisioning an Azure Data Factory	208
Provisioning Azure Storage	217
Provisioning an Azure-SSIS Integration Runtime for SSIS Package Files	223
Stopping the Azure-SSIS Integration Runtime	236
Conclusion	240
Chapter 9: Deploy a Simple, Custom, File-Based Azure-SSIS Framework.....	241
Provisioning the SSISConfig Database.....	241
Deploy the Simple, Custom, File-Based Azure-SSIS Framework	249
Provision an Azure File Share	267
Upload SSIS Packages.....	270
Update PackageLocation Values.....	272

TABLE OF CONTENTS

Build the SSIS Framework ADF Execution Engine	273
Retrieve a List of SSIS Packages.....	273
Execute the Retrieved SSIS Packages.....	287
View Test Execution Logs	296
Conclusion	299
Chapter 10: Framework Logging in ADF.....	301
Add the ApplicationName Parameter	301
A Quick Review of SSIS Framework Applications and Packages	302
Add Application Instance Logging.....	303
Modifying the Log.ApplicationInstance Table	303
Adding the Log.InsertApplicationInstance Stored Procedure	304
Logging Application Instance	306
Adding the Log.UpdateApplicationInstanceState Stored Procedure	311
Updating Application Instance.....	312
Let's Test It!	316
Updating Application Instance on Success.....	317
Adding the ApplicationPackageId Field	320
Adding the ApplicationInstanceld Pipeline Variable.....	321
Add the Child Pipeline	325
Let's Test It!	333
Add Application Package Instance Logging	334
Modifying the Log.ApplicationPackageInstance Table.....	335
Adding the Log.InsertApplicationPackageInstance Stored Procedure.....	336
Logging Application Package Instance	338
Adding the Log.UpdateApplicationPackageInstanceState Stored Procedure	344
Updating Application Package Instance	345
Let's Test It!	348
Updating Application Package Instance on Success	350
Conclusion	353

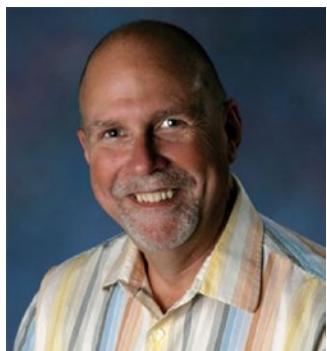
TABLE OF CONTENTS

Chapter 11: Fault Tolerance in the ADF Framework	355
A Brief Introduction to Fault Tolerance.....	355
Add ADF Managed Identity to Contributor Role	356
Add Application Package Fault Tolerance	359
About Fail Application on Package Failure	359
Implementing Application Package Fault Tolerance	359
Checking the Logic.....	362
Implementing the Fault Tolerance Logic.....	363
Let's Test It!	378
Conclusion	383
Index.....	385

About the Authors



Andy Leonard is Chief Data Engineer at Enterprise Data & Analytics, creator and Data Philosopher at DILM (Data Integration Lifecycle Management) Suite, an Azure Data Factory and SQL Server Integration Services trainer and consultant, and a BimlHero. He is a SQL Server database and data warehouse developer, community mentor, engineer, and farmer. Andy is coauthor of *SQL Server Integration Services Design Patterns*, *Data Integration Life Cycle Management with SSIS*, and *The Biml Book*.



Kent Bradshaw is the founder of Tudor Data Solutions, LLC. With over 40 years of IT experience, he is a SQL Server database/ETL developer and database architect with a background in Medicaid claims, public schools, government, retail, and insurance systems. In 2011, Kent founded Tudor Data Solutions, LLC, to pursue new development opportunities which led to his association with Andy Leonard and Enterprise Data & Analytics. In 2017, Kent received the MPP certification for Data Science.

About the Technical Reviewer



André van Meulebrouck has a keen interest in functional programming, especially Haskell and F#.

He also likes data technologies from markup languages to databases and F# type providers.

He lives in Southern California with his wife “Tweety”, and is active in athletics: hiking, mountain biking, and gravity/balance sports like freestyle skating (in-line and ice), skateboarding, surfing, and sandboarding.

To keep his mind sharp, he does compositional origami, plays classical guitar, and enjoys music notation software.

Acknowledgments

This book would not have been possible without the help of Kent Bradshaw, my coauthor, friend, and brother. This book represents Kent's first foray into authoring, and he delivered outstanding work.

Thanks André van Meulebrouck for proofreading, working through the samples, and providing feedback on the manuscript.

I owe my coworkers at Enterprise Data & Analytics (EntDNA.com) a debt of gratitude for their encouragement and for covering for me when I stayed up too late trying to figure out how to make the code work. We have an awesome team, especially my brothers from other mothers – Nick Harris, Shannon Lowder, and Reeves Smith.

Donald Farmer inspires me every time we interact. As Principal Program Manager at Microsoft, Donald worked extensively with SSIS and helped shape the product. Donald continues to shape data software by providing vendors unique strategic guidance at TreeHive Strategy (treehivestrategy.com).

I am certain there are many excellent editors in this business. Jonathan Gennick is the best with whom I have had the privilege to work. Without Jill Balzano's input and project management, I would have been lost!

Finally, I thank my family for their understanding. My children Stevie Ray, Emma, and Riley who live at home at the time of this writing, and Manda and Penny who have children of their own. And Christy, to whom I dedicate this book, my wife, my love. Thank you.

—Andy

Writing this book never really entered my mind until my great friend and EDNA colleague, Andy Leonard, convinced me to do it. I truly appreciate his confidence in me, and it was a very interesting experience. I'm really glad that I could work together with him on this project.

Many thanks to Jonathan Gennick, Jill Balzano, and André van Meulebrouck for all of their help and suggestions. They made things so much easier.

Most important of all, I want to thank my wife, Ann. Much of where I am today is due to her love, support, and encouragement. I couldn't begin to do it without her. Thank you.

—Kent

Introduction

Frameworks have existed for as long as people have been building things. When people began developing solutions using computers, frameworks soon followed. Like woodworking jigs, frameworks exist to simplify work and speed up the process of developing a solution, whether the problem is to complete a woodworking project or deliver an enterprise data integration solution.

Who Is This Book For?

Anyone who is interested in data-related automation or executing multiple chunks of code with a single command will find value in these pages. Three audiences will particularly benefit from this book:

1. Database developers
2. Data engineers and data integration developers
3. Business Intelligence Markup Language (Biml) developers

Database developers will learn about driving database object execution from stored procedures and will see one example of a database framework.

Data engineers and data integration developers will learn how to use SQL Server Integration Services (SSIS) to implement a metadata-driven data integration framework.

Biml developers will see one example of using BimlExpress and metadata to rapidly produce several SSIS packages in selected design patterns *and* an example of interrogating flat files to generate flat file formats and flat file connection managers via Biml.

How This Book Is Structured

In the following sections, we will explore the chapter layout of this book. There are two parts that guide you through the process of building a database framework using stored procedures and an SSIS framework. Each of those parts is composed of several chapters.

Part I: Stored Procedure–Based Database Frameworks

In the first part of this book, we will focus on building a database framework using stored procedures. These chapters are targeted primarily at those who develop database solutions. We do not endeavor to create the most comprehensive primer for every aspect of database development; we cover the basics well.

Chapter 1: Stored Procedures 101

Before we dive into database development, let's cover some basics of stored procedures.

Chapter 2: Automation with Stored Procedures

Once you have a solid grasp on stored procedures, we will use what we've learned about stored procedures to automate execution using the controller pattern, which executes one or more stored procedures.

Chapter 3: Stored Procedure Orchestrators

In this chapter, we examine high-level execution management using the orchestrator pattern, which executes one or more controller stored procedures.

Chapter 4: A Stored Procedure–Based Metadata-Driven Framework

Metadata-driven frameworks are a recurring theme in this book. In this chapter, we store information about stored procedures, controllers, and orchestrators to execute framework applications.

Part II: SSIS Frameworks

If your enterprise is executing a few SSIS packages, there is no need for an SSIS framework. However, if your enterprise is executing several hundred (or several thousand) SSIS packages, you *definitely* need an SSIS framework.

Chapter 5: A Simple, Custom, File-Based SSIS Framework

In this chapter, we define a data integration/engineering framework metadata database that manages configuration and execution. We introduce a version of an SSIS framework that manages these subject areas for SSIS executed on-premises.

Chapter 6: Framework Execution Engine

In this chapter, we add instrumentation in the form of Information and Error events. Built-in SSIS logging will surface these messages, which are useful when troubleshooting SSIS framework application issues.

Chapter 7: Framework Logging

In this chapter, we persist execution metadata to “Instance” tables. The ApplicationInstance table stores one record for each instance of an SSIS framework application execution. The ApplicationPackageInstance table stores one record for each instance of an SSIS framework application package execution. Important execution instance attributes – such as start and end times and execution status – are maintained in the “Instance” tables.

Chapter 8: Azure-SSIS Integration Runtime

In this chapter, we introduce Azure Data Factory (ADF) and the Azure-SSIS integration runtime by walking through the provisioning processes.

Chapter 9: Deploy a Simple, Custom, File-Based Azure-SSIS Framework

In Chapters 5–7, we designed and constructed an SSIS framework aimed at on-premises SSIS execution, configuration, and logging. In this chapter, we begin migrating the SSIS framework designed in Chapters 5–7 by provisioning a new Azure SQL database and then deploying the metadata database to our new Azure SQL database. We provision an Azure File Share and migrate test SSIS packages to the Azure File Share. Finally, we begin building the Azure Data Factory parent pipeline that serves as the execution engine for the ADF version of the SSIS framework.

Chapter 10: Framework Logging in ADF

In this chapter, we add logging functionality to the parent ADF pipeline, much like the functionality we added to the on-premises version of the SSIS framework in Chapter 7.

Chapter 11: Fault Tolerance in the ADF Framework

In this chapter, we complete ADF execution engine functionality by implementing fault tolerance to programmatically stop (or not stop) pipeline execution based on SSISConfig metadata configurations.

Conclusion

This book is for people who want to dive deeper into SQL Server and SSIS automation. We will be discussing and demonstrating database and Integration Services frameworks and covering such topics as SQL Server, SSIS, Azure, and Azure Data Factory.

As you can see, we have a long journey filled with learning and new insight ahead of us. Let's get started!

PART I

Stored Procedure-Based Database Frameworks

CHAPTER 1

Stored Procedures 101

One of the most common issues continually facing IT organizations is finding the proper balance between the effort to develop and deploy processes into production against the efficiency and effectiveness of production control operators. The effort to develop and deploy and the effectiveness of production control seem to be diametrically opposed. Making it easier to develop and deploy processes usually means more work and manual intervention for production control. The real question that needs to be considered is where is it better to “feel the pain”? Pushing the effort toward the development side of the equation can slow down the throughput but minimizes the liability of issues at the production process level. This and the next few chapters are going to concentrate on processes that are executed using stored procedures. In this chapter, you’ll get a basic introduction to stored procedures that is the foundation for the chapters that follow. You’ll see how to create a child stored procedure, and we’ll provide you a template that you can use to create similar procedures in your own work.

The Need for a Framework

When just in development mode, it is awfully easy (and impressive) to construct large, monolithic procedures that do everything from beginning to end. They are the proverbial “black box” where something goes in, many gyrations take place, and then the desired result occurs. Those are great until a problem occurs, or just a change in business requirements means that modifications need to be made. When the procedure does so much, what does it take to test it once the modifications are done? Even though the change only impacts 10 percent of the procedure, the entire procedure has to be tested. What does it take to accomplish that? What if there are several intricate modifications that need to be made? How difficult does it become for more than one developer to work on the changes and coordinate their efforts?

Now, consider that monolith broken up into multiple procedures, each of which performs a unit of work. When a change is made and needs to be tested, that effort can be isolated to just what is necessary to perform that unit of work, and the validation is concentrated on the result of the procedure. And, with it now being multiple procedures, modifications can usually be done simultaneously by multiple developers, and their efforts can be mutually exclusive. Over time, that approach can prove to be much more cost effective and efficient.

That is where a framework helps to organize and manage processes to provide the most flexibility in development and can minimize the maintenance effort (which, sometimes, is not considered until it becomes an obvious issue). A framework provides a consistent methodology for assembling and executing processes. It also promotes writing code in small units of work that can potentially be mixed, matched, and reused. It adds complexity to the development and deployment processes but can reduce the effort for production scheduling. The framework can also provide greater flexibility for managing the execution of the process.

Demonstration of a Framework

To begin the analysis of the framework concept, we need a process. Our example to follow shows a framework built to run a daily process against an example schema. The details of that process don't matter to the example. Just consider that any production system might have something that needs to be done each day, and what follows is a framework by which to make those daily processes happen.

Also, part of the example is a monthly process. Just as a system might need certain tasks to be done each day, it's also common to have certain things that need to be done once monthly. In designing such a system, one must take into account the order in which daily and monthly processes execute when their schedules intersect on – in our example – the first of each month.

For the purpose of this book, a simple process has been developed (NOTE: all of the code described can be downloaded at entdna.com. You can also find a link to the code from the book's catalog page on Apress.com). Downloading the example code enables you to follow along with the upcoming examples on your own machine.

An Example Schema

Listing 1-1 shows code to create a schema called FWDemo that will contain everything needed for the demonstration. Also, there is code to create a table called FWDemo.ProcessLog. Including a pattern for writing to this table throughout all of the procedures certainly adds some complexity and overhead to the procedures, but what it provides in monitoring and troubleshooting more than makes up for the upfront effort.

Listing 1-1. Schema and log table creation

```

print 'FWDemo Schema'
If Not Exists(Select name
              From sys.schemas
              Where name= 'FWDemo')
begin
    print ' - Creating FWDemo schema'
    declare @sql varchar(255) = 'Create Schema FWDemo'
    exec(@sql)
    print ' - FWDemo schema created'
end
Else
print ' - FWDemo schema already exists.'
print ''
GO

IF EXISTS (SELECT * FROM sys.objects
            WHERE object_id = OBJECT_ID(N'FWDemo.ProcessLog')
            AND type in (N'U'))
    DROP TABLE FWDemo.ProcessLog
GO

SET ANSI_NULLS ON
GO

SET QUOTED_IDENTIFIER ON
GO

SET ANSI_PADDING ON
GO

```

```
CREATE TABLE [FWDemo].[ProcessLog](
    [ProcessLogID]      [int] IDENTITY(1,1) NOT NULL,
    [ProcessLogMessage] [nvarchar](255)      NOT NULL,
    [CreateDate]        [smalldatetime]     NOT NULL
)
GO

SET ANSI_PADDING OFF
GO
```

Do you have a SQL Server instance that you can use for learning purposes? Connect to that instance as an administrator, for example, as the sa user. Then, in SQL Server Management Studio (SSMS), open a “New Query” window, copy the code from Listing 1-1, and execute it to create the example schema used in this and subsequent chapters.

The Daily Process

Listing 1-2 shows the code to create two stored procedures that will make up our demonstration Daily Process. We are providing two procedures in our example because it’s common to have more than one, and having two allows us to show how to make the execution of subsequent procedures depend upon the success of earlier ones – because the need to execute a series of procedures and halt or take other actions when an error occurs is the real-life scenario that most of us face.

These procedures (as well as all others that we will be using) can be compiled and executed for your own testing. You will notice that there is some code commented out (lines that are preceded with ‘--’) in each procedure that can be invoked (remove the ‘--’, then recompile) to create an error condition. This ability to create an error condition allows testing for successful and unsuccessful completions that will become more important as we progress through demonstration iterations in later chapters.

For the sake of this exercise, we will declare a business rule for the Daily Process stating that FWDemo.DailyProcess1 must complete successfully before FWDemo.DailyProcess2 can be executed. FWDemo.DailyProcess2 must then complete successfully before the Daily Process can be deemed successfully executed.

Listing 1-2. Daily Process stored procedures

```
If Exists(Select s.name + '.' + p.name
          From sys.procedures p
          Join sys.schemas s
            On s.schema_id = p.schema_id
          Where s.name = 'FWDemo'
            And p.name = 'DailyProcess1')
begin
    print ' - Dropping FWDemo.DailyProcess1 stored procedure'
    Drop Procedure FWDemo.DailyProcess1
    print ' - FWDemo.DailyProcess1 stored procedure dropped'
end
GO

CREATE PROCEDURE FWDemo.DailyProcess1
AS
-----
-- Purpose: This procedure is part of the Stored Procedure Framework Demo.
--
-- NOTE: An Error situation can be created for testing/demo purposes by
--       un-commenting the Error code in the body of the procedure. To return
--       to a procedure with a successful execution, re-comment the code or
--       recompile the original.
--
SET NOCOUNT ON

/***** */
/* Log the START of the procedure to the process log */
/***** */
```

CHAPTER 1 STORED PROCEDURES 101

```
INSERT INTO FWDemo.ProcessLog (
    ProcessLogMessage,
    CreateDate
)
Values ('Procedure FWDemo.DailyProcess1 - STARTING',
        GETDATE()
)

DECLARE @RetStat int

SET @RetStat = 0

/*****************/
/* Force an ERROR CONDITION for this procedure */
/*****************/

--INSERT INTO FWDemo.ProcessLog (
--    ProcessLogMessage,
--    CreateDate
--)
--VALUES ('Procedure FWDemo.DailyProcess1 - Problem Encountered',
--        GETDATE()
--)
--SET @RetStat = 1

/*****************/
/* Log the COMPLETION of the procedure to the process log */
/*****************/

IF @RetStat = 0
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.DailyProcess1 - COMPLETED',
            GETDATE()
    )

```

```
END
ELSE
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.DailyProcess1 - ERROR',
        GETDATE()
    )
END

RETURN @RetStat
GO

If Exists(Select s.name + '.' + p.name
    From sys.procedures p
    Join sys.schemas s
    On s.schema_id = p.schema_id
    Where s.name = 'FWDemo'
    And p.name = 'DailyProcess2')
begin
    print ' - Dropping FWDemo.DailyProcess2 stored procedure'
    Drop Procedure FWDemo.DailyProcess2
    print ' - FWDemo.DailyProcess2 stored procedure dropped'
end
GO

CREATE PROCEDURE FWDemo.DailyProcess2
AS
-----
-----
-- Purpose: This procedure is part of the Stored Procedure Framework Demo.
--
```

CHAPTER 1 STORED PROCEDURES 101

```
-- NOTE: An Error situation can be created for testing/demo purposes by
--       un-commenting the Error code in the body of the procedure. To return
--       to a procedure with a successful execution, re-comment the code or
--       recompile the original.

--



-----



SET NOCOUNT ON

/***** */
/* Log the START of the procedure to the process log */
/***** */

INSERT INTO FWDemo.ProcessLog (
    ProcessLogMessage,
    CreateDate
)
Values ('Procedure FWDemo.DailyProcess2 - STARTING',
        GETDATE()
)

DECLARE @RetStat int

SET @RetStat = 0

/***** */
/* Force an ERROR CONDITION for this procedure */
/***** */

--INSERT INTO FWDemo.ProcessLog (
--    ProcessLogMessage,
--    CreateDate
--)
--VALUES ('Procedure FWDemo.DailyProcess2 - Problem Encountered',
--        GETDATE()
--)
```

```
--SET @RetStat = 1
/*****
/* Log the COMPLETION of the procedure to the process log      */
/****

IF @RetStat = 0
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.DailyProcess2 - COMPLETED',
        GETDATE()
    )
END
ELSE
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.DailyProcess2 - ERROR',
        GETDATE()
    )
END

RETURN @RetStat
GO
```

In a “New Query” window in SSMS, execute the code from Listing 1-2 while connected to the FWDemo schema. The code creates two stored procedures that together make up a daily process. With those procedures in place, you can turn your attention to the next problem, which is to schedule those procedures to actually run each day.

Executing the Daily Process

Now that an environment has been built and a process created, let's turn to the execution. Operations staff will have to set up or schedule the procedures to run and either monitor for any error conditions that are raised or set up precedence rules if such a function exists in any scheduling tool used. In a basic sense, we now have a Daily Process that is ready for production. Listing 1-3 shows the statements that can be used to execute the Daily Process procedures and also a SELECT statement that can be run to view the output written to FWDemo.ProcessLog. You will notice that we are ordering the output in a descending order. This will show the most recent messages at the top and eliminate the need to scroll down to get to the messages for the current execution and much easier once the log starts to become heavily populated.

Listing 1-3. Daily Process execute statements and process log SELECT statement

```
EXECUTE FWDemo.DailyProcess1
```

```
EXECUTE FWDemo.DailyProcess2
```

```
SELECT ProcessLogID  
      ,ProcessLogMessage  
      ,CreateDate  
  FROM FWDemo.ProcessLog  
 ORDER BY ProcessLogID desc
```

Including a Monthly Process

Now it's time to add another layer to our production process. In Listing 1-4, there is code to create two more stored procedures that will make up a Monthly Process. The procedures operate the same as our daily process procedures, and there are some business rules associated with them. First, the monthly process will run on the first day of the month. Second, it will run after the successful execution of the Daily Process, and third, FWDemo.MonthlyProcess1 must complete successfully before FWDemo.MonthlyProcess2 can be executed.

Listing 1-4. Monthly Process stored procedures

```
If Exists(Select s.name + '.' + p.name
      From sys.procedures p
      Join sys.schemas s
        On s.schema_id = p.schema_id
      Where s.name = 'FWDemo'
        And p.name = 'MonthlyProcess1')

begin
  print ' - Dropping FWDemo.MonthlyProcess1 stored procedure'
  Drop Procedure FWDemo.MonthlyProcess1
  print ' - FWDemo.MonthlyProcess1 stored procedure dropped'
end
GO

CREATE PROCEDURE FWDemo.MonthlyProcess1
AS
-----
-- Purpose: This procedure is part of the Stored Procedure Framework Demo.
--
-- NOTE: An Error situation can be created for testing/demo purposes by
--       un-commenting the Error code in the body of the procedure. To return
--       to a procedure with a successful execution, re-comment the code or
--       recompile the original.
--

SET NOCOUNT ON

/***** */
/* Log the START of the procedure to the process log */
/***** */
```

CHAPTER 1 STORED PROCEDURES 101

```
INSERT INTO FWDemo.ProcessLog (
    ProcessLogMessage,
    CreateDate
)
Values ('Procedure FWDemo.MonthlyProcess1 - STARTING',
        GETDATE()
)

DECLARE @RetStat int
SET @RetStat = 0

/*****************/
/* Force an ERROR CONDITION for this procedure */
/*****************/

--INSERT INTO FWDemo.ProcessLog (
--    ProcessLogMessage,
--    CreateDate
--)
--VALUES ('Procedure FWDemo.MonthlyProcess1 - Problem Encountered',
--        GETDATE()
--)
--SET @RetStat = 1

/*****************/
/* Log the COMPLETION of the procedure to the process log */
/*****************/

IF @RetStat = 0
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.MonthlyProcess1 - COMPLETED',
            GETDATE()
    )
END
```

```

ELSE
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.MonthlyProcess1 - ERROR',
        GETDATE()
    )
END

RETURN @RetStat
GO

If Exists(Select s.name + '.' + p.name
    From sys.procedures p
    Join sys.schemas s
    On s.schema_id = p.schema_id
    Where s.name = 'FWDemo'
    And p.name = 'MonthlyProcess2')
begin
    print ' - Dropping FWDemo.MonthlyProcess2 stored procedure'
    Drop Procedure FWDemo.MonthlyProcess2
    print ' - FWDemo.MonthlyProcess2 stored procedure dropped'
end
GO

CREATE PROCEDURE FWDemo.MonthlyProcess2
AS
-----
-----
-- Purpose: This procedure is part of the Stored Procedure Framework Demo.
--
-- NOTE: An Error situation can be created for testing/demo purposes by
--       un-commenting the Error code in the body of the procedure. To return

```

CHAPTER 1 STORED PROCEDURES 101

```
--      to a procedure with a successful execution, re-comment the code or
--      recompile the original.
--
-----
```

SET NOCOUNT ON

```
*****/*
/* Log the START of the procedure to the process log */
*****
```

INSERT INTO FWDemo.ProcessLog (

```
    ProcessLogMessage,
    CreateDate
)
Values ('Procedure FWDemo.MonthlyProcess2 - STARTING',
        GETDATE()
)
```

DECLARE @RetStat int

SET @RetStat = 0

```
*****/*
/* Force an ERROR CONDITION for this procedure */
*****
```

--INSERT INTO FWDemo.ProcessLog (

```
    ProcessLogMessage,
    CreateDate
--)
VALUES ('Procedure FWDemo.MonthlyProcess2 - Problem Encountered',
        GETDATE()
--)
--SET @RetStat = 1
```

```
*****/
```

```
/* Log the COMPLETION of the procedure to the process log      */
/*********************                                         *****/
IF @RetStat = 0
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.MonthlyProcess2 - COMPLETED',
        GETDATE()
    )
END
ELSE
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.MonthlyProcess2 - ERROR',
        GETDATE()
    )
END
RETURN @RetStat
GO
```

The execute statements to be run or scheduled by the operations staff are shown in Listing 1-5. Use the SELECT statement introduced earlier to monitor the progress of the process execution.

Listing 1-5. Monthly Process execute statements

```
EXECUTE FWDemo.MonthlyProcess1
```

```
EXECUTE FWDemo.MonthlyProcess2
```

From an operations perspective, we have now introduced more complexity. Process set up/scheduling now has additional layers of precedence and also a timing factor. Scheduling and monitoring have become more critical, and there are more places that require attention in order for everything to operate smoothly.

Summary

A system has been built that consists of two daily and two monthly processes. Dependencies exist that require the successful execution of all prior processes before initiating the execution of the next process. These dependencies must be applied in the scheduling of the processes in the course of regular production and also in troubleshooting and problem remediation. All of these “moving parts” not only add complexity but also points of vulnerability to the system due to the need for outside attention.

CHAPTER 2

Automation with Stored Procedures

Now that we see all the pieces that need to be coordinated day in and day out for executing the processes defined in Chapter 1, let's introduce the idea of controllers. A controller is the most basic element in our demonstration of the framework concept. The examples that follow demonstrate how controllers can make your job easier while also reducing the overhead that your operations staff must deal with.

A Daily Process Controller

For our purposes, a controller is a stored procedure that itself executes all the elements of a process and can address business rules within the process. Listing 2-1 shows the code to create a stored procedure called FWDemo.DailyProcessController. This stored procedure executes both of our Daily Process stored procedures. It logs its own progress and checks the return status of each of the stored procedures to ensure successful completion of each.

Listing 2-1. Daily Process Controller stored procedure

```
If Exists(Select s.name + '.' + p.name
          From sys.procedures p
          Join sys.schemas s
            On s.schema_id = p.schema_id
          Where s.name = 'FWDemo'
            And p.name = 'DailyProcessController')
begin
```

CHAPTER 2 AUTOMATION WITH STORED PROCEDURES

```
print ' - Dropping FWDemo.DailyProcessController stored procedure'
Drop Procedure FWDemo.DailyProcessController
print ' - FWDemo.DailyProcessController stored procedure dropped'
end
GO

CREATE PROCEDURE FWDemo.DailyProcessController
AS
-----
-- Purpose: This procedure is part of the Stored Procedure Framework Demo.
--           It is the Controller for the Daily Process Stored Procedures.
--
SET NOCOUNT ON
/*****
/* Log the START of the procedure to the process log */
****/

INSERT INTO FWDemo.ProcessLog (
    ProcessLogMessage,
    CreateDate
)
Values ('Procedure FWDemo.DailyProcessController - STARTING',
        GETDATE()
)
DECLARE @RetStat int
SET @RetStat = 0
/*****
/* Execute the DailyProcess1 Stored Procedures */
****/
```

```
EXEC @RetStat = FWDemo.DailyProcess1

IF @RetStat <> 0
    GOTO EndController

/*****************/
/* Execute the DailyProcess2 Stored Procedures */
/*****************/

EXEC @RetStat = FWDemo.DailyProcess2

IF @RetStat <> 0
    GOTO EndController

/*****************/
/* Log the COMPLETION of the procedure to the process log */
/*****************/

EndController:

IF @RetStat = 0
    BEGIN
        INSERT INTO FWDemo.ProcessLog (
            ProcessLogMessage,
            CreateDate
        )
        VALUES ('Procedure FWDemo.DailyProcessController - COMPLETED',
            GETDATE()
        )
    END
ELSE
    BEGIN
        INSERT INTO FWDemo.ProcessLog (
            ProcessLogMessage,
            CreateDate
        )
    
```

```
VALUES ('Procedure FWDemo.DailyProcessController - ERROR',
        GETDATE()
    )
END

RETURN @RetStat
GO
```

With the controller in Listing 2-1, execution by the operations staff has just been made simpler. Instead of scheduling both the stored procedures and dealing with precedence rules from the schedule, only the controller execution needs to be set up, and monitoring for successful completion is now easier. Listing 2-2 shows the code used to execute the Daily Process controller. As always, FWDemo.ProcessLog shows the progress of the execution.

Listing 2-2. Daily Process Controller execute statement

```
EXECUTE FWDemo.DailyProcessController
```

Another advantage of using the controller pattern is the situation where a new stored procedure and/or new business rules need to be added to a process. Without the controller, the new stored procedure is compiled, and it is the responsibility of the operations staff to execute it in its proper position in the process and also account for the new business rule requirements in the setup process. That is a lot of steps that need to happen after development/testing has ended and, thus, an opportunity for error. With a controller in place, development and testing would include not only the new stored procedure(s) but also the necessary modifications to the controller stored procedure to execute the new stored procedures. The business rules then become part of the development/testing effort, and implementation to production becomes as simple as compiling the new stored procedure(s) and the new version of the controller stored procedure. Since the controller execution is already scheduled, no other changes are necessary for the operations staff. We have eliminated a lot of the manual intervention when deploying to production and, therefore, reduced opportunities for errors.

Monthly Process Controller

Next, we will add a controller for our Monthly Process. Listing 2-3 has the code for the stored procedure called FWDemo.MonthlyProcessController. This stored procedure executes both of the monthly process stored procedures and handles the precedence rule within the process while logging its progress in FWDemo.ProgressLog. The operations staff can simply execute this controller to start the process.

Listing 2-3. Monthly Process Controller stored procedure

```
If Exists(Select s.name + '.' + p.name
    From sys.procedures p
    Join sys.schemas s
        On s.schema_id = p.schema_id
    Where s.name = 'FWDemo'
        And p.name = 'MonthlyProcessController')

begin
    print ' - Dropping FWDemo.MonthlyProcessController stored procedure'
    Drop Procedure FWDemo.MonthlyProcessController
    print ' - FWDemo.MonthlyProcessController stored procedure dropped'
end
GO

CREATE PROCEDURE FWDemo.MonthlyProcessController
AS
-----
-- Purpose: This procedure is part of the Stored Procedure Framework Demo.
--          It is the Controller for the Monthly Process Stored Procedures.
--
```

CHAPTER 2 AUTOMATION WITH STORED PROCEDURES

```
SET NOCOUNT ON

/*****  
/* Log the START of the procedure to the process log */  
*****/  
  
INSERT INTO FWDemo.ProcessLog (  
    ProcessLogMessage,  
    CreateDate  
)  
Values ('Procedure FWDemo.MonthlyProcessController - STARTING',  
       GETDATE()  
)  
  
DECLARE @RetStat int  
  
SET @RetStat = 0

/*****  
/* Execute the MonthlyProcess1 Stored Procedures */  
*****/  
  
EXEC @RetStat = FWDemo.MonthlyProcess1  
  
IF @RetStat <> 0  
    GOTO EndController  
  
/*****  
/* Execute the MonthlyProcess2 Stored Procedures */  
*****/  
  
EXEC @RetStat = FWDemo.MonthlyProcess2  
  
IF @RetStat <> 0  
    GOTO EndController  
  
/*****  
/* Log the COMPLETION of the procedure to the process log */  
*****/
```

EndController:

```

IF @RetStat = 0
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.MonthlyProcessController - COMPLETED',
        GETDATE()
    )
END
ELSE
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.MonthlyProcessController - ERROR',
        GETDATE()
    )
END

RETURN @RetStat
GO

```

Listing 2-4 has the statement used to execute the new Monthly Process Controller.

Listing 2-4. Monthly Process Controller execute statement

```
EXECUTE FWDemo.MonthlyProcessController
```

Not addressed here, however, are the business precedence rules between controllers that are 1) the Monthly Process executes on the first day of the month and 2) after the successful completion of the Daily Process. But, these details for the operations staff to address have become much simpler with the introduction of the controller pattern.

Summary

Controllers incorporate dependencies, defined by the business rules, and coordinate the execution of processes. There are two primary advantages that come from this. First, it makes scheduling much easier because the dependencies are built in. Second, it provides ready access and reference to the process flow when in the midst of troubleshooting process issues.

CHAPTER 3

Stored Procedure Orchestrators

Building on the controller pattern in Chapter 2, the next logical progression for the framework is the concept of the orchestrator. It is basically a controller for controllers and can address the issue of precedence rules between controllers.

The Orchestrator

Controllers are themselves a type of process just as the daily and monthly processes that have been built. We may find that dependencies exist between the processing of different controllers. That is where an orchestrator comes in. Orchestrators can address those dependencies in exactly the same way as controllers. Listing 3-1 shows the code to create a stored procedure called FWDemo.ProcessOrchestrator. This stored procedure executes the Daily Controller stored procedure and, when appropriate based on the date, the Monthly Controller stored procedure.

Listing 3-1. Process Orchestrator stored procedure

```
IF EXISTS (SELECT s.name + '.' + p.name
            FROM sys.procedures p
            JOIN sys.schemas s
              ON s.schema_id = p.schema_id
           WHERE s.name = 'FWDemo'
             AND p.name = 'ProcessOrchestrator')
```

CHAPTER 3 STORED PROCEDURE ORCHESTRATORS

```
BEGIN
    PRINT ' - Dropping FWDemo.ProcessOrchestrator stored procedure'
    DROP PROCEDURE FWDemo.ProcessOrchestrator
    PRINT ' - FWDemo.ProcessOrchestrator stored procedure dropped'
END
GO

CREATE PROCEDURE FWDemo.ProcessOrchestrator
    @RunDate smalldatetime = null
AS
-----
-----
-----

-- Purpose: This procedure is part of the Stored Procedure Framework
Demo. -- It is the Orchestrator Procedure that
executes the Controller -- Procedures.

-----
-----
-----


SET NOCOUNT ON

/*****************/
/* Log the START of the procedure to the process log */
/****************/


INSERT INTO FWDemo.ProcessLog (
    ProcessLogMessage,
    CreateDate
)
VALUES ('Procedure FWDemo.ProcessOrchestrator - STARTING',
    GETDATE()
)
IF @RunDate is null
    SET @RunDate = GETDATE()

DECLARE @RetStat int
SET @RetStat = 0
```

```
*****  
/* Execute the Daily Process Controller */  
*****  
  
EXEC @RetStat = FWDemo.DailyProcessController  
  
IF @RetStat <> 0  
    GOTO EndOrchestrator  
  
*****  
/* Execute the MonthlyProcessController IF @RunDate is the first day */  
/* of a month */  
*****  
  
IF DATEPART(DAY, @RunDate) = 1  
    BEGIN  
        EXEC @RetStat = FWDemo.MonthlyProcessController  
  
        IF @RetStat <> 0  
            GOTO EndOrchestrator  
    END  
  
*****  
/* Log the COMPLETION of the procedure to the process log */  
*****  
  
EndOrchestrator:  
  
IF @RetStat = 0  
    BEGIN  
        INSERT INTO FWDemo.ProcessLog (   
            ProcessLogMessage,  
            CreateDate  
        )  
        VALUES ('Procedure FWDemo.ProcessOrchestrator - COMPLETED',  
                GETDATE()  
        )  
    END
```

CHAPTER 3 STORED PROCEDURE ORCHESTRATORS

```
ELSE
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.ProcessOrchestrator - ERROR',
        GETDATE()
    )
END

RETURN @RetStat
GO
```

As with the various stored procedures we have already seen, the orchestrator shows execution progress by logging to FWDemo.ProcessLog. The return status is used to check for successful completion of the controllers. Also, there is code included to determine if it is the first day of the month which is one of the precedence rules for the execution of the Monthly Process. There is an input parameter called “@RunDate” which has a default value of “null.” The null value caused the stored procedure to use the current date for processing. However, to make it easy to experiment with the demonstration system, a date can be supplied at execution time to establish the value of “@RunDate”. Listing 3-2 shows the execute statements for the orchestrator for using both the current date and a date override. For all practical purposes, all that the operations staff needs to do is set up/schedule the execution of FWDemo.ProcessOrchestrator, and everything else is handled within the patterns that have been built.

Listing 3-2. Process Orchestrator execute statements

```
--IF rundate is current date:
EXECUTE FWDemo.ProcessOrchestrator

--IF rundate is not current date:
EXECUTE FWDemo.ProcessOrchestrator
    @RunDate = '<date>'

--NOTE: <date> is in the format of CCYYMMDD
```

A Quick Process Overview

At this point, let's review what has been built up until now. There is a Daily Process comprised of two stored procedures and a Monthly Process also comprised of two stored procedures. The daily process controller, FWDemo.DailyProcessController, executes the daily process and ensures that the stored procedure FWDemo.DailyProcess1 completes successfully before executing the FWDemo.DailyProcess2 stored procedure. The monthly process controller, FWDemo.MonthlyProcessController, executes the monthly process and ensures that the stored procedure FWDemo.MonthlyProcess1 completes successfully before executing the FWDemo.MontlyProcess2 stored procedure. Wrapped around it all is FWDemo.ProcessOrchestrator that executes the daily process controller and, upon successful completion, checks for a run date that is the first day of the month and, if true, executes the monthly process controller. A single execute statement drives the entire process.

A Look at Troubleshooting Issues

Up until now, our attention has been focused on following the successful completion of the processes that have been built. So, it's time for a reality check. How do we handle error situations? What do we do if one of the processes in the sequence has an error and doesn't complete successfully? That can sometimes be accompanied by a phone call in the middle of the night and some action needs to be taken. Diving in and troubleshooting the problem begins. Looking at FWDemo.ProcessLog is a great place to start in order to identify the location and cause of the problem in order for correction to begin. Once the solution is determined and the problem corrected, how do we complete the remaining steps in the sequence? With our current patterns in place, about the only option is to deconstruct the process where the error occurred and finish manually. After that, we find an execution at the next highest level of our patterns and manually execute from there.

For example, the production process has a problem with the stored procedure FWDemo.DailyProcess2. We research and correct the issue. At this point, FWDemo.DailyProcess2 needs to be manually executed. Upon successful completion, we know that everything in FWDemo.DailyProcessController has finished. The next highest level is the FWDemo.ProcessOrchestrator stored procedure. After the execution of the daily process controller, the orchestrator checks to see if the value in the "@RunDate"

parameter equals the first of the month. If that is the case, then we would need to manually execute FWDemo.MonthlyProcessController. If it is not the first of the month, nothing else is necessary. In our system, it is pretty simple. But, in reality, the resolution could have a lot more moving parts and a lot of manual intervention. The more complex the process becomes, the more complex the problem resolution becomes.

Summary

Orchestrators are simply a higher level of controller. It manages the execution of multiple controllers while incorporating any dependencies that may exist between the controllers. It can further simplify scheduling while providing a more complete picture of process requirements that can aid in successful problem remediation.

CHAPTER 4

A Stored Procedure–Based Metadata-Driven Framework

In order to add flexibility and control over the execution of all of the processes within the framework, metadata can be added to the environment. It can identify the stored procedures, controllers, and orchestrators as well as the relationships between each. The metadata will be used as the communication that drives the execution of the processes.

Building the Metadata

If we're going to create metadata, then we need a place in which to store it. We can create a table or a set of tables for that purpose. The metadata for this chapter's example is stored in a set of tables used to control executions. The first table is FWDemo.ProcessProcedure, seen in Listing 4-1.

Listing 4-1. ProcessProcedure table creation

```
IF EXISTS (SELECT * FROM sys.objects  
    WHERE object_id = OBJECT_ID(N'FWDemo.ProcessProcedure')  
        AND type in (N'U'))  
    DROP TABLE FWDemo.ProcessProcedure  
GO  
  
SET ANSI_NULLS ON  
GO
```

```
SET QUOTED_IDENTIFIER ON
GO

SET ANSI_PADDING ON
GO

CREATE TABLE [FWDemo].[ProcessProcedure](
    [ProcessProcedureID]      [int] IDENTITY(1,1) NOT NULL,
    [ProcessProcedureName]    [nvarchar](255)      NOT NULL
)
GO

SET ANSI_PADDING OFF
GO
```

This table identifies all of the stored procedures available for execution in the framework. The stored procedures become the building blocks that are collected together in what we will call framework applications. To establish the application, we use the FWDemo.Application table seen in Listing 4-2.

Listing 4-2. Application table creation

```
IF EXISTS (SELECT * FROM sys.objects
            WHERE object_id = OBJECT_ID(N'FWDemo.Application')
            AND type in (N'U'))
    DROP TABLE FWDemo.Application
GO

SET ANSI_NULLS ON
GO

SET QUOTED_IDENTIFIER ON
GO

SET ANSI_PADDING ON
GO
```

```

CREATE TABLE [FWDemo].[Application](
    [ApplicationID] [int] IDENTITY(1,1) NOT NULL,
    [ApplicationName] [nvarchar](255) NOT NULL
)
GO

SET ANSI_PADDING OFF
GO

```

Now it is necessary to identify which of the stored procedures are going to be included in each of the defined applications. For that, we use the table, seen in Listing 4-3, called FWDemo.ApplicationProcessProcedure. This table contains the metadata that will define the execution of the process. As we discussed in Chapter 1, building small unit of work stored procedures allows for code reuse. Using the metadata tables allows relating a single stored procedure to multiple applications. A data archive stored procedure would be a good example of where this might be helpful.

Listing 4-3. ApplicationProcessProcedure table creation

```

IF EXISTS (SELECT * FROM sys.objects
    WHERE object_id = OBJECT_ID(N'FWDemo.ApplicationProcessProcedure')
        AND type in (N'U'))
    DROP TABLE FWDemo.ApplicationProcessProcedure
GO

SET ANSI_NULLS ON
GO

SET QUOTED_IDENTIFIER ON
GO

SET ANSI_PADDING ON
GO

CREATE TABLE [FWDemo].[ApplicationProcessProcedure](
    [ApplicationProcessProcedureID] [int] IDENTITY(1,1) NOT NULL,
    [ApplicationID] [int] NOT NULL,
    [ProcessProcedureID] [int] NOT NULL,

```

```

    [ExecutionOrder]      [int] NOT NULL,
    [Active]              [bit] NOT NULL
)
GO

SET ANSI_PADDING OFF
GO

```

There is a column named ExecutionOrder which will determine where in the process sequence its associated stored procedure will be executed. Also, the Active column indicates that a stored procedure is enabled for execution. It is used to disable procedures that, for some reason, are not to be executed in the process. This is great if you want to temporarily bypass a procedure and will be discussed later in this chapter.

With the metadata tables created, let's load the data for our processes. Listing 4-4 shows the code that loads all of our stored procedures into the FWDemo.ProcessProcedure table. Only the name is inserted because the value for ProcessProcedureID, being an IDENTITY column, will be automatically generated.

Listing 4-4. Load data to the ProcessProcedure table

```

--Load First Daily Procedure
INSERT INTO FWDemo.ProcessProcedure (ProcessProcedureName)
VALUES ('FWDemo.DailyProcess1')
GO

--Load Second Daily Procedure
INSERT INTO FWDemo.ProcessProcedure (ProcessProcedureName)
VALUES ('FWDemo.DailyProcess2')
GO

--Load First Monthly Procedure
INSERT INTO FWDemo.ProcessProcedure (ProcessProcedureName)
VALUES ('FWDemo.MonthlyProcess1')
GO

--Load Second Monthly Procedure
INSERT INTO FWDemo.ProcessProcedure (ProcessProcedureName)
VALUES ('FWDemo.MonthlyProcess2')
GO

```

Listing 4-5 has the code to define the applications for our demonstration into the table FWDemo.Application. For our basic framework, we are defining our daily and monthly processes. Again, the value of the IDENTITY column ApplicationID will be automatically generated.

Listing 4-5. Load data to the Application table

```
--Load Application for Daily Process
INSERT INTO FWDemo.Application (ApplicationName)
VALUES ('DailyProcessControllerMD')
GO

--Load Application for Monthly Process
INSERT INTO FWDemo.Application (ApplicationName)
VALUES ('MonthlyProcessControllerMD')
GO
```

To put all of the pieces together, Listing 4-6 contains the code to relate the stored procedures to their intended application. Notice the values used for the execution order. Using values in a scale, in this case a scale of ten, leaves gaps between the initially defined execution order. This way, if a new procedure is added later on that needs to execute between existing stored procedures, there is room to add a value, and the existing execution order values do not have to be renumbered. Also, notice that a value of “1” is used for the Active column. This enables the procedure for execution within the application. A value of “0” will disable the procedure in the context of the application.

Listing 4-6. Load metadata for framework

```
DECLARE @AppID      int,
        @ProcID      int

/*****************/
/* Daily Process Application */
/*****************/

SET @AppID = (SELECT ApplicationID
              FROM FWDemo.Application
              WHERE ApplicationName = 'DailyProcessControllerMD')
```

CHAPTER 4 A STORED PROCEDURE–BASED METADATA-DRIVEN FRAMEWORK

```
SET @ProcID = (SELECT ProcessProcedureID
                FROM FWDemo.ProcessProcedure
                WHERE ProcessProcedureName = 'FWDemo.DailyProcess1')

INSERT INTO FWDemo.ApplicationProcessProcedure (
    ApplicationID,
    ProcessProcedureID,
    ExecutionOrder,
    Active
)
VALUES (@AppID,
        @ProcID,
        10,
        1
)

SET @ProcID = (SELECT ProcessProcedureID
                FROM FWDemo.ProcessProcedure
                WHERE ProcessProcedureName = 'FWDemo.DailyProcess2')

INSERT INTO FWDemo.ApplicationProcessProcedure (
    ApplicationID,
    ProcessProcedureID,
    ExecutionOrder,
    Active
)
VALUES (@AppID,
        @ProcID,
        20,
        1
)

/*****************/
/* Monthly Process Application */
/*****************/

SET @AppID = (SELECT ApplicationID
                FROM FWDemo.Application
                WHERE ApplicationName = 'MonthlyProcessControllerMD')
```

```
SET @ProcID = (SELECT ProcessProcedureID
                FROM FWDemo.ProcessProcedure
                WHERE ProcessProcedureName = 'FwDemo.MonthlyProcess1')

INSERT INTO FWDemo.ApplicationProcessProcedure (
    ApplicationID,
    ProcessProcedureID,
    ExecutionOrder,
    Active
)
VALUES (@AppID,
        @ProcID,
        10,
        1
)

SET @ProcID = (SELECT ProcessProcedureID
                FROM FWDemo.ProcessProcedure
                WHERE ProcessProcedureName = 'FwDemo.MonthlyProcess2')

INSERT INTO FWDemo.ApplicationProcessProcedure (
    ApplicationID,
    ProcessProcedureID,
    ExecutionOrder,
    Active
)
VALUES (@AppID,
        @ProcID,
        20,
        1
)
```

Metadata-Ready Controllers

All of the metadata is now in place. But, in order to execute using the metadata, new controllers are needed to use the metadata. Listing 4-7 has the code for the new daily and monthly controllers. You will notice that the names of these new procedures are suffixed with “MD”.

Listing 4-7. New metadata controllers

```

IF EXISTS (SELECT s.name + '.' + p.name
            FROM sys.procedures p
            JOIN sys.schemas s
              ON s.schema_id = p.schema_id
           WHERE s.name = 'FWDemo'
             AND p.name = 'DailyProcessControllerMD')

BEGIN
    PRINT ' - Dropping FWDemo.DailyProcessControllerMD stored procedure'
    DROP PROCEDURE FWDemo.DailyProcessControllerMD
    PRINT ' - FWDemo.DailyProcessControllerMD stored procedure dropped'
END
GO

CREATE PROCEDURE FWDemo.DailyProcessControllerMD
AS
-----
-- Purpose: This procedure is part of the Stored Procedure Framework
Demo.          -- It is the Metadata Controller version for
the Daily Process Stored      -- Procedures.
--
SET NOCOUNT ON

```

```
/****************************************/
/* Log the START of the procedure to the process log */
/****************************************/

INSERT INTO FWDemo.ProcessLog (
    ProcessLogMessage,
    CreateDate
)
VALUES ('Procedure FWDemo.DailyProcessControllerMD - STARTING',
        GETDATE()
)

DECLARE @RetStat int

SET @RetStat = 0

/****************************************/
/* Get and Execute the Active DailyProcess Stored Procedures */
/****************************************/

DECLARE @DailyProcName nvarchar(255)
    ,@AppID      int
    ,@AppPPID    int

DECLARE @DailyCursor as CURSOR

SET @DailyCursor = CURSOR FOR
SELECT pp.ProcessProcedureName, a.ApplicationID,
       app.ApplicationProcessProcedureID
FROM FWDemo.Application a
JOIN FWDemo.ApplicationProcessProcedure app
    ON a.ApplicationID = app.ApplicationID
JOIN FWDemo.ProcessProcedure pp
    ON app.ProcessProcedureID = pp.ProcessProcedureID
WHERE a.ApplicationName = 'DailyProcessControllerMD'
    AND app.Active = 1
ORDER BY app.ExecutionOrder

OPEN @DailyCursor
FETCH NEXT FROM @DailyCursor INTO @DailyProcName, @AppID, @AppPPID
```

CHAPTER 4 A STORED PROCEDURE–BASED METADATA-DRIVEN FRAMEWORK

```
WHILE @@FETCH_STATUS = 0
BEGIN
    INSERT INTO FwDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Executing ' + @DailyProcName + ' ApplicationID = ' +
        convert(varchar(10), @AppID) +
        ' ApplicationProcessProcedureID = ' + convert(varchar(10),
        @AppPPID),
        GETDATE()
    )

    EXEC @RetStat = @DailyProcName

    IF @RetStat <> 0
        BEGIN
            CLOSE @DailyCursor
            DEALLOCATE @DailyCursor
            GOTO EndController
        END

    FETCH NEXT FROM @DailyCursor INTO @DailyProcName, @AppID, @AppPPID
END

CLOSE @DailyCursor
DEALLOCATE @DailyCursor

/*****************/
/* Log the COMPLETION of the procedure to the process log */
/*****************/

EndController:

IF @RetStat = 0
BEGIN
    INSERT INTO FwDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
```

```
VALUES ('Procedure FWDemo.DailyProcessControllerMD - COMPLETED',
        GETDATE()
    )
END
ELSE
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.DailyProcessControllerMD - ERROR',
            GETDATE()
        )
END

RETURN @RetStat
GO

IF EXISTS (SELECT s.name + '.' + p.name
           FROM sys.procedures p
           JOIN sys.schemas s
             ON s.schema_id = p.schema_id
           WHERE s.name = 'FWDemo'
             AND p.name = 'MonthlyProcessControllerMD')
BEGIN
    PRINT ' - Dropping FWDemo.MonthlyProcessControllerMD stored procedure'
    DROP Procedure FWDemo.MonthlyProcessControllerMD
    PRINT ' - FWDemo.MonthlyProcessControllerMD stored procedure dropped'
END
GO

CREATE PROCEDURE FWDemo.MonthlyProcessControllerMD
AS
-----  
-----
```

CHAPTER 4 A STORED PROCEDURE–BASED METADATA-DRIVEN FRAMEWORK

```
--  
-- Purpose: This procedure is part of the Stored Procedure Framework Demo.  
-- It is the Metadata Controller version for the Monthly Process Stored  
-- Procedures.  
--
```

```
SET NOCOUNT ON
```

```
/*****  
/* Log the START of the procedure to the process log */  
*****
```

```
INSERT INTO FWDemo.ProcessLog (  
    ProcessLogMessage,  
    CreateDate  
)  
VALUES ('Procedure FWDemo.MonthlyProcessControllerMD - STARTING',  
       GETDATE()  
)
```

```
DECLARE @RetStat int
```

```
SET @RetStat = 0
```

```
/*****  
/* Get and Execute the Active MonthlyProcess Stored Procedures */  
*****
```

```
DECLARE @MonthlyProcName nvarchar(255)  
    ,@AppID      int  
    ,@AppPPID    int
```

```
DECLARE @MonthlyCursor as CURSOR
```

```
SET @MonthlyCursor = CURSOR FOR  
SELECT pp.ProcessProcedureName, a.ApplicationID,  
       app.ApplicationProcessProcedureID  
FROM FWDemo.Application a
```

```
JOIN FWDemo.ApplicationProcessProcedure app
    ON a.ApplicationID = app.ApplicationID
JOIN FWDemo.ProcessProcedure pp
    ON app.ProcessProcedureID = pp.ProcessProcedureID
WHERE a.ApplicationName = 'MonthlyProcessControllerMD'
    AND app.Active = 1
ORDER BY app.ExecutionOrder

OPEN @MonthlyCursor
FETCH NEXT FROM @MonthlyCursor INTO @MonthlyProcName, @AppID, @AppPPID

WHILE @@FETCH_STATUS = 0
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES (
        'Executing ' + @DailyProcName + ' ApplicationID = ' +
        convert(varchar(10), @AppID) +
        ' ApplicationProcessProcedureID = ' + convert(varchar(10),
        @AppPPID),
        GETDATE()
    )
    EXEC @RetStat = @MonthlyProcName

    IF @RetStat <> 0
        BEGIN
            CLOSE @MonthlyCursor
            DEALLOCATE @MonthlyCursor
            GOTO EndController
        END
    FETCH NEXT FROM @MonthlyCursor INTO @MonthlyProcName, @AppID,
        @AppPPID
END
```

CHAPTER 4 A STORED PROCEDURE–BASED METADATA-DRIVEN FRAMEWORK

```
CLOSE @MonthlyCursor
DEALLOCATE @MonthlyCursor

/*****
/* Log the COMPLETION of the procedure to the process log */
****/

EndController:
IF @RetStat = 0
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.MonthlyProcessControllerMD - COMPLETED',
            GETDATE()
        )
END
ELSE
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.MonthlyProcessControllerMD - ERROR',
            GETDATE()
        )
END

RETURN @RetStat
GO
```

As you can see, these new versions of the controllers use a cursor to retrieve all of the active stored procedures related to the respective applications, sorted by the execution order. The procedure names, along with the values of all of the associated keys, are loaded into variables. The stored procedure name is used to direct the execution

while the key values are used in the messages written to FWDemo.ProcessLog to track execution and aid in monitoring or troubleshooting. The cursor loops through each stored procedure until all have been executed or an error occurs.

Metadata-Ready Orchestrator

With the new versions of the controllers in place, a new version of the orchestrator, also with the name suffixed with “MD”, is needed to execute them. Listing 4-8 has the code for the new orchestrator. Note, separate versions of the controllers and orchestrator were created to keep the original versions in place and operational so that all aspects of the framework demonstration remain intact and functional.

Listing 4-8. New metadata orchestrator

```
IF EXISTS (SELECT s.name + '.' + p.name
           FROM sys.procedures p
           JOIN sys.schemas s
             ON s.schema_id = p.schema_id
            WHERE s.name = 'FWDemo'
              AND p.name = 'ProcessOrchestratorMD')

BEGIN
    PRINT ' - Dropping FWDemo.ProcessOrchestratorMD stored procedure'
    DROP PROCEDURE FWDemo.ProcessOrchestratorMD
    PRINT ' - FWDemo.ProcessOrchestratorMD stored procedure dropped'

END
GO

CREATE PROCEDURE FWDemo.ProcessOrchestratorMD
    @RunDate smalldatetime = null
AS
-----
-----
```

CHAPTER 4 A STORED PROCEDURE–BASED METADATA-DRIVEN FRAMEWORK

```
-- Purpose: This procedure is part of the Stored Procedure Framework Demo.  
-- It is the Orchestrator Procedure that executes the Metadate version  
-- of the Controller Procedures.
```

```
--
```

```
SET NOCOUNT ON
```

```
/******  
/* Log the START of the procedure to the process log */  
******/
```

```
INSERT INTO FWDemo.ProcessLog (  
    ProcessLogMessage,  
    CreateDate  
)  
Values ('Procedure FWDemo.ProcessOrchestratorMD - STARTING',  
    GETDATE()  
)
```

```
IF @RunDate is null  
    SET @RunDate = GETDATE()
```

```
DECLARE @RetStat int
```

```
SET @RetStat = 0
```

```
/******  
/* Execute the DailyProcessControllerMD */  
*****/
```

```
EXEC @RetStat = FWDemo.DailyProcessControllerMD
```

```
IF @RetStat <> 0  
    GOTO EndOrchestratorMD
```

```
/******  
/* Execute the MonthlyProcessControllerMD IF @RunDate is the first day */  
/* of a month */  
*****/
```

```
IF DATEPART(DAY, @RunDate) = 1
BEGIN
    EXEC @RetStat = FWDemo.MonthlyProcessControllerMD

    IF @RetStat <> 0
        GOTO EndOrchestratorMD

END
/************************************************************/
/* Log the COMPLETION of the procedure to the process log */
/************************************************************/
EndOrchestratorMD:

IF @RetStat = 0
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.ProcessOrchestratorMD - COMPLETED',
            GETDATE()
        )
END
ELSE
BEGIN
    INSERT INTO FWDemo.ProcessLog (
        ProcessLogMessage,
        CreateDate
    )
    VALUES ('Procedure FWDemo.ProcessOrchestratorMD - ERROR',
            GETDATE()
        )
END

RETURN @RetStat
GO
```

Just like the original version, the only thing that needs to be set up/scheduled for execution is the orchestrator with the optional @RunDate parameter override. Listing 4-9 has the execution commands.

Listing 4-9. Metadata version of Process Orchestrator execute statements

```
--IF rundate is current date:  
EXECUTE FWDemo.ProcessOrchestratorMD  
  
--IF rundate is not current date:  
EXECUTE FWDemo.ProcessOrchestratorMD  
    @RunDate = '<date>'  
  
--NOTE: <date> is in the format of CCYYMMDD
```

Troubleshooting with Metadata

This is the most basic metadata setup for the system we have built. So how can executions be completed in an error resolution situation like we ended up with in Chapter 3? Let's say that FWDemo.DailyProcess2 was unsuccessful. Listing 4-10 shows the messages from FWDemo.ProcessLog that shows that the error occurred.

Listing 4-10. Output from the Process Log

ProcessLogID	ProcessLogMessage
CreateDate	
39	Procedure FWDemo.ProcessOrchestratorMD - ERROR
2018-09-17 20:47:00	
38	Procedure FWDemo.DailyProcessControllerMD - ERROR
2018-09-17 20:47:00	
37	Procedure FWDemo.DailyProcess2 - Problem Encountered
2018-09-17 20:47:00	
36	Procedure FWDemo.DailyProcess2 - STARTING
2018-09-17 20:46:00	
35	Executing FWDemo.DailyProcess2 ApplicatonID = 1
2018-09-17 20:46:00	
	ApplicatonProcessProcedureID = 2

```

34          Procedure FWDemo.DailyProcess1 - COMPLETED
2018-09-17 20:46:00
33          Procedure FWDemo.DailyProcess1 - STARTING
2018-09-17 20:45:00
32          Executing FWDemo.DailyProcess1 ApplicationID = 1
2018-09-17 20:45:00
                  ApplicationProcessProcedureID = 1
31          Procedure FWDemo.DailyProcessControllerMD - STARTING
2018-09-17 20:45:00
30          Procedure FWDemo.ProcessOrchestratorMD -
STARTING
                  2018-09-17 20:45:00

```

Once the resolution is determined, we can see from the messages that FWDemo.DailyProcess1 has already run successfully and therefore does not have to be rerun. We see in the message that the ApplicationProcessProcedureID value is 1. We can use the code in Listing 4-11 to disable the FWDemo.DailyProcess1 so that it will not run if the application is executed.

Listing 4-11. Code to disable execution

```
--DISABLE stored procedure from execution for RESTART

UPDATE FWDemo.ApplicationProcessProcedure
    SET Active = 0
WHERE ApplicationProcessProcedureID = 1      -- or value from Process Log
```

At this point, the scheduled job can be restarted as before, and the execution will start with FWDemo.DailyProcess2 as shown in Listing 4-12. Once the job is completed, simply reset the Active column for FWDemo.DailyProcess1 with code like that found in Listing 4-13.

Listing 4-12. Process Log after restart

ProcessLogID	ProcessLogMessage
CreateDate	
54	Procedure FWDemo.ProcessOrchestratorMD - COMPLETED
2018-09-17 20:58:00	

CHAPTER 4 A STORED PROCEDURE–BASED METADATA-DRIVEN FRAMEWORK

```
53          Procedure FWDemo.MonthlyProcessControllerMD - COMPLETED
2018-09-17 20:58:00
52          Procedure FWDemo.MonthlyProcess2 - COMPLETED
2018-09-17 20:58:00
51          Procedure FWDemo.MonthlyProcess2 - STARTING
2018-09-17 20:57:00
50          Executing FWDemo.MonthlyProcess2 ApplicatonID = 2
2018-09-17 20:57:00
                  ApplicatonProcessProcedureID = 4
49          Procedure FWDemo.MonthlyProcess1 - COMPLETED
2018-09-17 20:57:00
48          Procedure FWDemo.MonthlyProcess1 - STARTING
2018-09-17 20:56:00
47          Executing FWDemo.MonthlyProcess1 ApplicatonID = 2
2018-09-17 20:56:00
                  ApplicatonProcessProcedureID = 3
46          Procedure FWDemo.MonthlyProcessControllerMD - STARTING
2018-09-17 20:56:00
45          Procedure FWDemo.DailyProcessControllerMD - COMPLETED
2018-09-17 20:56:00
44          Procedure FWDemo.DailyProcess2 - COMPLETED
2018-09-17 20:56:00
43          Procedure FWDemo.DailyProcess2 - STARTING
2018-09-17 20:56:00
42          Executing FWDemo.DailyProcess2 ApplicatonID = 1
2018-09-17 20:55:00
                  ApplicatonProcessProcedureID = 2
41          Procedure FWDemo.DailyProcessControllerMD - STARTING
2018-09-17 20:55:00
40          Procedure FWDemo.ProcessOrchestratorMD - STARTING
2018-09-17 20:55:00
```

Listing 4-13. Code to enable execution of disabled stored procedure

```
--ENABLE stored procedure for execution for NEXT EXECUTION
```

```
UPDATE FWDemo.ApplicationProcessProcedure
      SET Active = 1
 WHERE ApplicationProcessProcedureID = 1      -- or value from Process Log
```

In another example, suppose the error occurred in the FWDemo.MonthlyProcess1 stored procedure. The process log, shown in Listing 4-14, shows that both of the daily process stored procedures completed successfully in the daily process controller.

Listing 4-14. Output from Process Log

ProcessLogID	ProcessLogMessage
CreateDate	
79	Procedure FWDemo.ProcessOrchestratorMD - ERROR 2018-09-18 19:35:00
78	Procedure FWDemo.MonthlyProcessControllerMD - ERROR 2018-09-18 19:35:00
77	Procedure FWDemo.MonthlyProcess1 - Problem Encountered 2018-09-18 19:35:00
76	Procedure FWDemo.MonthlyProcess1 - STARTING 2018-09-18 19:34:00
75	Executing FWDemo.MonthlyProcess1 ApplicatonID = 2 2018-09-18 19:34:00
	ApplicatonProcessProcedureID = 3
74	Procedure FWDemo.MonthlyProcessControllerMD - STARTING 2018-09-18 19:34:00
73	Procedure FWDemo.DailyProcessControllerMD - COMPLETED 2018-09-18 19:34:00
72	Procedure FWDemo.DailyProcess2 - COMPLETED 2018-09-18 19:34:00
71	Procedure FWDemo.DailyProcess2 - STARTING 2018-09-18 19:33:00
70	Executing FWDemo.DailyProcess2 ApplicatonID = 1 2018-09-18 19:33:00

```

                    ApplicationProcessProcedureID = 2
69                  Procedure FWDemo.DailyProcess1 - COMPLETED
2018-09-18 19:33:00
68                  Procedure FWDemo.DailyProcess1 - STARTING
2018-09-18 19:32:00
67                  Executing FWDemo.DailyProcess1 ApplicationID = 1
2018-09-18 19:32:00
                    ApplicationProcessProcedureID = 1
66                  Procedure FWDemo.DailyProcessControllerMD - STARTING
2018-09-18 19:32:00
65                  Procedure FWDemo.ProcessOrchestratorMD - STARTING
2018-09-18 19:32:00

```

We see in the messages that their ApplicationProcessProcedureID values are 1 and 2, respectively. Code like that in Listing 4-15 could be used to manage the restart and subsequent process reset.

Listing 4-15. Manage error restart for FWDemo.MonthlyProcess1

```
--DISABLE stored procedures from execution for RESTART

UPDATE FWDemo.ApplicationProcessProcedure
    SET Active = 0
WHERE ApplicationProcessProcedureID in (1, 2) -- or values from Process
Log
```

Listing 4-16 shows the log entries from the restarted application execution.

Listing 4-16. Process Log after restart

```
ProcessLogID      ProcessLogMessage
CreateDate
91                  Procedure FWDemo.ProcessOrchestratorMD - COMPLETED
2018-09-18 19:47:00
90                  Procedure FWDemo.MonthlyProcessControllerMD - COMPLETED
2018-09-18 19:47:00
89                  Procedure FWDemo.MonthlyProcess2 - COMPLETED
2018-09-18 19:47:00
```

```

88          Procedure FWDemo.MonthlyProcess2 - STARTING
2018-09-18 19:46:00
87          Executing FWDemo.MonthlyProcess2 ApplicationID = 2
2018-09-18 19:46:00
                        ApplicationProcessProcedureID = 4
86          Procedure FWDemo.MonthlyProcess1 - COMPLETED
2018-09-18 19:46:00
85          Procedure FWDemo.MonthlyProcess1 - STARTING
2018-09-18 19:45:00
84          Executing FWDemo.MonthlyProcess1 ApplicationID = 2
2018-09-18 19:45:00
                        ApplicationProcessProcedureID = 3
83          Procedure FWDemo.MonthlyProcessControllerMD - STARTING
2018-09-18 19:45:00
82          Procedure FWDemo.DailyProcessControllerMD - COMPLETED
2018-09-18 19:45:00
81          Procedure FWDemo.DailyProcessControllerMD - STARTING
2018-09-18 19:45:00
80          Procedure FWDemo.ProcessOrchestratorMD - STARTING
2018-09-18 19:45:00

```

Once completed, the code shown in Listing 4-17 is used to reset the metadata for the application back to its original setting.

Listing 4-17. Enable stored procedures from execution for next execution

```

UPDATE FWDemo.ApplicationProcessProcedure
    SET Active = 1
WHERE ApplicationProcessProcedureID in (1, 2) -- or values from Process Log

```

Initial setup of the metadata and deployment of new stored procedures to an application do, indeed, make the development and implementation of a process more complex. Hopefully, these metadata examples have demonstrated that management during execution, particularly in an error resolution situation, has become much simpler. Remember, setup and deployment may only occur once. Execution management will remain and occur over again as long as the process remains as part of the production activities. So, where do you want to make it easier?

This is the simplest use of metadata for our system and could be just the beginning. For a more advanced version, we could build a metadata layer on top of what we already have and let orchestrators execute using metadata for the controllers. But we will stop here because the idea has been established.

Summary

Metadata is used as the road map for the execution of the orchestrator and controllers. Most situations can now be handled by the manipulation of the metadata that controls the execution, and no manual deconstructions and executions are necessary. Any new stored procedures added to existing controllers or controllers added to existing orchestrators can be included simply by deploying the stored procedures to be compiled and adding the appropriate metadata. In most cases, if not all, no changes are required to existing stored procedures.

PART II

SSIS Frameworks

CHAPTER 5

A Simple, Custom, File-Based SSIS Framework

If you survey enterprises using SSIS for data integration/engineering, you will learn most enterprises do not use the SSIS catalog. Most enterprises execute SSIS from the file system. Why? Enterprises with a smaller number of SSIS packages outnumber enterprises with large numbers of SSIS packages. Executing a few dozen SSIS packages is different – really different – from managing the execution of thousands of SSIS packages. Don't take my word for it; ask any data engineer managing a larger enterprise.

Before the release of SSIS 2012 and the SSIS catalog, SSIS developers had to build their own SSIS frameworks. In this chapter, I share one way to build a custom file-based SSIS framework.

There are relatively few benefits to having 45 years' experience developing software. One benefit is living through several architecture pattern cycles, watching them wax and then wane in popularity. Old patterns get a fresh coat of virtual paint and become new again. (It reminds me of my older daughters asking me, in the 1990s, if I'd heard an awesome new band named Aerosmith.) We will see an example in a subsequent chapter when we deploy a remarkably similar SSIS framework to Azure Data Factory's Azure-SSIS integration runtime.

An SSIS Framework, Defined and Designed

We author SSIS frameworks according to a few principles:

1. Functionality
2. Empathy
3. Simplicity

Functionality

In my opinion, an SSIS framework must accomplish execution, logging, and configuration. Execution of SSIS packages includes *grouped* execution. Execution grouping is the ability to execute a collection of SSIS packages in a specified order. The goal of logging is to surface enough operational information about an SSIS package execution to allow an experienced operator or developer to troubleshoot any execution error. Configuration promotes code reuse and supports SSIS Design Patterns by allowing developers and operators to configure SSIS package properties, parameters, and variables at runtime.

Empathy

“What do you mean by ‘empathy,’ Andy?” I am glad you asked. Empathy in software development shows up in user experience, or UX, design considerations. In SSIS design and development, empathy manifests by considering the skills of the user. In this case, the users consist of SSIS developers and operators. SSIS frameworks should employ the KISS (“keep it simple, stupid”) principle. I assume the SSIS developer, or SSIS Developer Team, will be familiar with SSIS, so I build the SSIS framework *in SSIS*. I want the developer or team to manage and maintain their framework.

Simplicity

Software should be as complex as required to accomplish the goal, and not more complex than necessary.

—Andy Leonard, circa 2007

Managing complexity is hard for software developers. In the end, managing complexity is a balancing act between functionality, extensibility, and maintainability. An SSIS framework requires some amount of complexity to function, by its very nature. Supporting extensions via “extension points” always fails because framework authors cannot anticipate every use case in every enterprise. Designing any software – including an SSIS framework – should include thoughts about maintenance. Maintainability helps strike the balance between maximum functionality and minimum complexity.

Taken together, functionality, empathy, and simplicity are relatively straightforward concepts that are easy to comprehend but difficult to accomplish.

Building a File-Based SSIS Framework

In this portion of the book, over the next several chapters, we will build the components of an SSIS framework that will execute SSIS packages stored in the file system. The components of an SSIS framework are

- A database for execution and configuration metadata storage, log storage, and business logic
- An execution management engine built in SSIS

In this chapter, we focus on

- Building the metadata database
- Building a test SSIS project
- Adding metadata to the metadata database

Obtain the Code

To obtain the code for this book, visit the catalog page for this book on the Apress website (www.apress.com) or connect to the GitHub repository at

`github.com/aleonard763/FrameworksBook`

Save the code to a location you can readily access.

Metadata-Driven Execution Management

Metadata-driven execution management is not a complex topic even though it may appear complex at first blush. The example to follow builds a database named SSISConfig and the execution metadata tables shown in Figure 5-1.

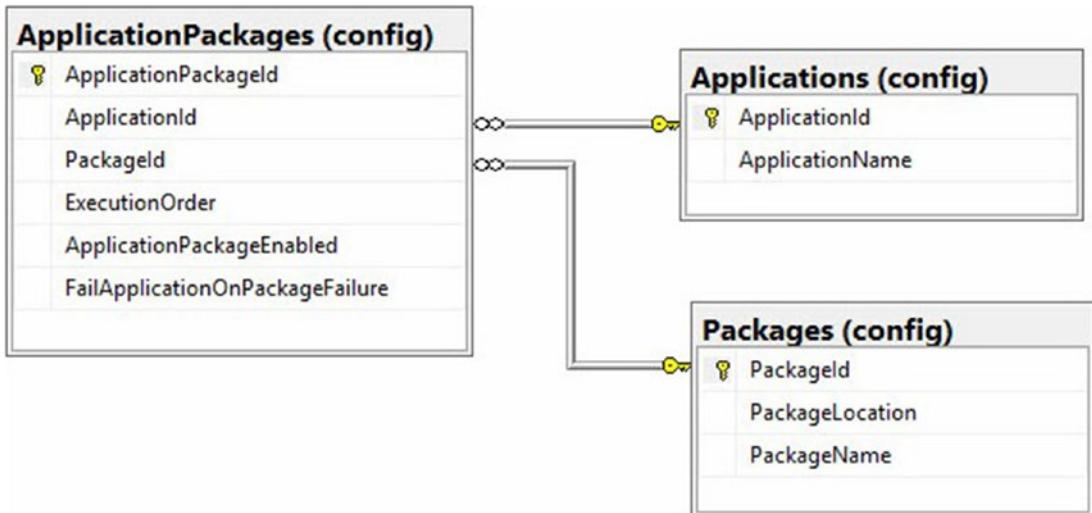


Figure 5-1. Metadata tables for execution management

The SSISConfig Database

The SSISConfig database is designed to contain metadata for the framework.

Create the SSISConfig database using the T-SQL shown in Listing 5-1.

Listing 5-1. Creating the SSISConfig database

```

Use [master]
go

print 'SSISConfig database'
If Not Exists(Select [databases].[name]
              From [sys].[databases]
              Where [databases].[name] = N'SSISConfig')
begin

```

```
print ' - Create SSISConfig database'
Create Database SSISConfig
print ' - SSISConfig database created'
end
Else
begin
    print ' - SSISConfig database already exists.'
end
print ''
go
```

SQL Server Management Studio (SSMS) Object Explorer should appear similar to Figure 5-2.

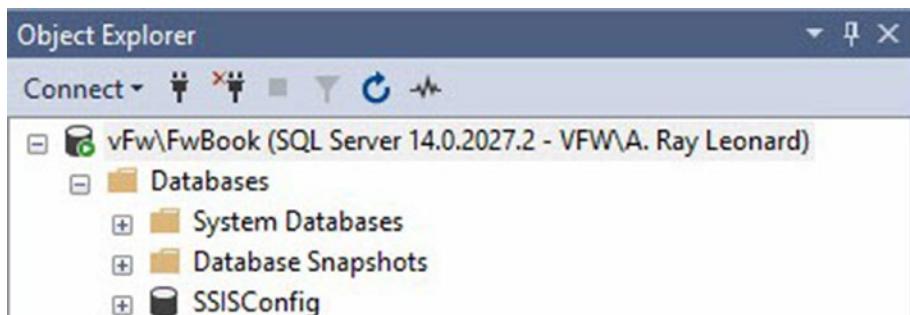


Figure 5-2. SSISConfig database created

Config Schema

At the top of the list of software development best practices is “separation of concerns.” One way to separate concerns is decoupling. One can argue database development is not software development. Others will argue some software development best practices may – perhaps *should* – be used in database development. I am in this second camp.

Add the first schema – config – to the SSISConfig database using the T-SQL shown in Listing 5-2.

Listing 5-2. Creating the config schema

```
use [SSISConfig]
go

print 'Config schema'
If Not Exists(Select [schemas].[name]
    From [sys].[schemas]
    Where [schemas].[name] = N'config')
begin
    print ' - Create config schema'
    declare @sql nvarchar(100) = N'Create Schema config'
    exec(@sql)
    print ' - Config schema created'
end
Else
begin
    print ' - Config schema already exists.'
end
print ''
go
```

Once the config schema is created, it can be viewed in the SSMS Object Explorer. Navigate to the SSISConfig ➤ Security ➤ Schemas node as shown in Figure 5-3.

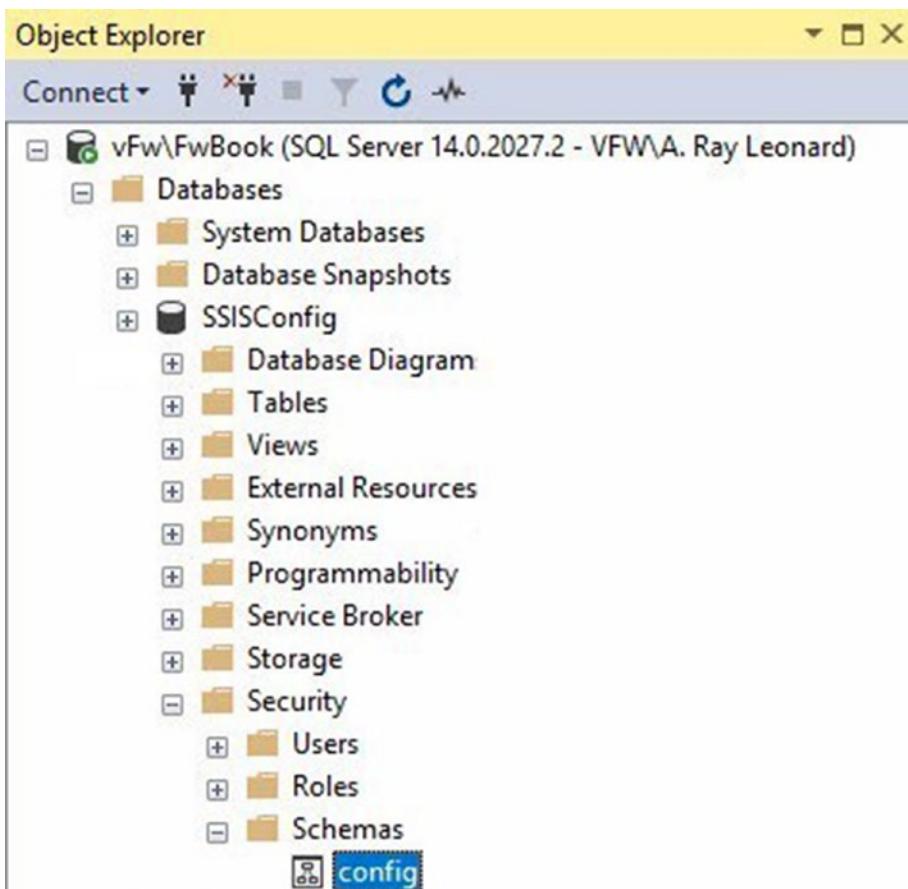


Figure 5-3. Viewing the config schema

The Config.Applications Table

In this framework, an *SSIS application* is defined as “a collection of SSIS packages configured to execute in a prescribed order.” Each application is identified by an ApplicationId (non-null int).

Each application has a name attribute, and the ApplicationName must be unique. Identifying the application by ApplicationId is important for referential integrity. Observing the T-SQL for tables in the config schema – especially the config.ApplicationPackages table – you will note the design is third normal form.

Execute the T-SQL in Listing 5-3 to add the Applications table to the config schema in the SSISConfig database.

Listing 5-3. Adding config.Applications

```

use [SSISConfig]
go

print 'Config.Applications table'
If Not Exists(Select [schemas].[name] + '.' + [tables].[name] As ▶
[Schema.Table]
    From [sys].[tables]
    Join [sys].[schemas]
        On [schemas].[schema_id] = [tables].[schema_id]
    Where [schemas].[name] = N'config'
        And [tables].[name] = N'Applications')

begin
    print ' - Create config.Applications table'
    Create Table [config].[Applications]
    (
        ApplicationId int identity(1, 1)
            Constraint PK_config_Applications Primary Key Clustered
        , ApplicationName nvarchar(255) Not NULL
            Constraint UQ_config_Applications_ApplicationName
                Unique
    )
    print ' - Config.Applications table created'
end

Else
begin
    print ' - Config.Applications table already exists.'
end
print ''
go

```

Once executed, the config.Applications table appears in Object Explorer as shown in Figure 5-4.

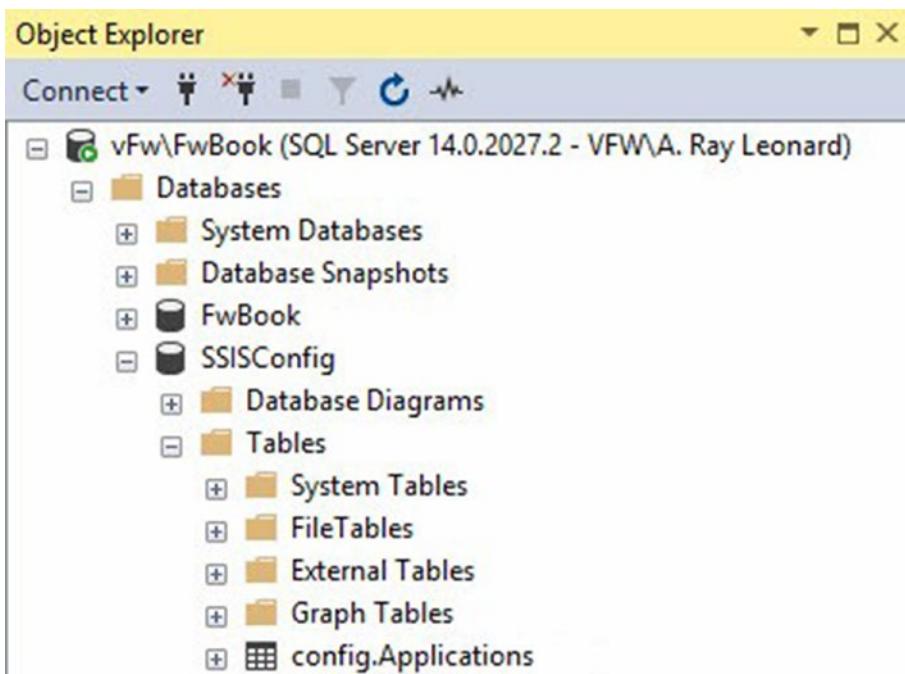


Figure 5-4. Config.Applications table

The Config.Packages Table

Package location metadata is stored in the config.Packages table. As we build out other versions of this SSIS framework, we will see this is the table that changes. The changes are based on where SSIS packages reside in the enterprise.

Create the config.Packages table by executing the T-SQL shown in Listing 5-4.

Listing 5-4. Creating the config.Packages table

```
use [SSISConfig]
go

print 'Config.Packages table'
If Not Exists(Select [schemas].[name] + '.' + [tables].[name] As >
[Schema.Table]
    From [sys].[tables]
    Join [sys].[schemas]
        On [schemas].[schema_id] = [tables].[schema_id]
```

CHAPTER 5 A SIMPLE, CUSTOM, FILE-BASED SSIS FRAMEWORK

```
Where [schemas].[name] = N'config'
      And [tables].[name] = N'Packages')

begin
    print ' - Create config.Packages table'
    Create Table [config].[Packages]
    (
        PackageId int identity(1, 1)
        Constraint PK_config_Packages Primary Key Clustered
        , PackageLocation nvarchar(255) Not NULL
        , PackageName nvarchar(255) Not NULL
        , Constraint UQ_config_Packages_PackageName
            Unique(PackageLocation, PackageName)
    )
    print ' - Config.Packages table created'
end
Else
begin
    print ' - Config.Packages table already exists.'
end
print ''
go
```

Once created, the config.Packages table appears in Object Explorer as shown in Figure 5-5.

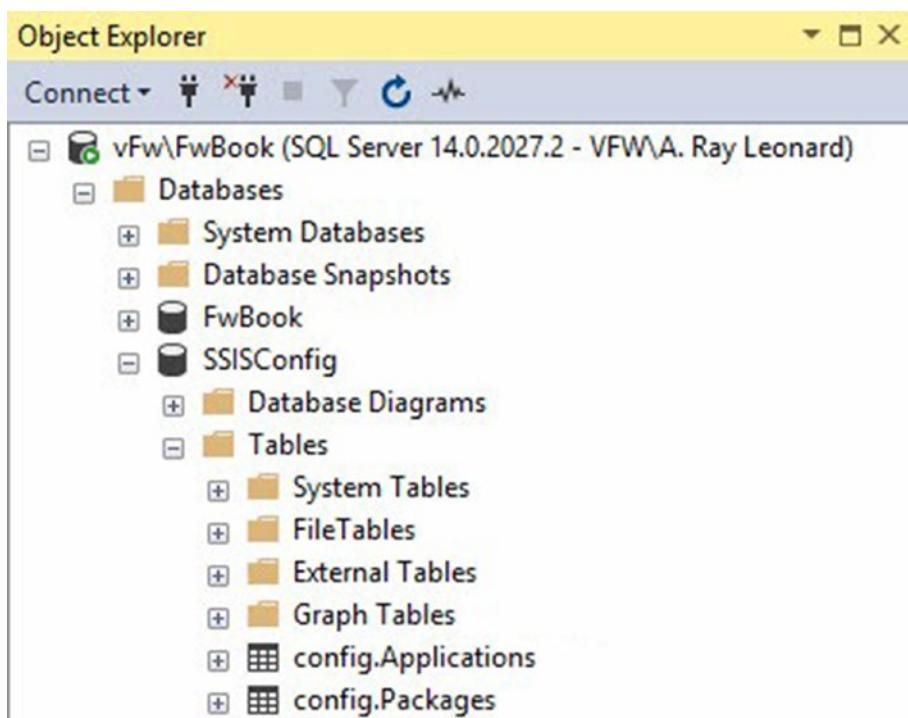


Figure 5-5. Creating config.Packages table

As mentioned earlier, an SSIS application in the framework is a collection of packages configured to execute in a specific order. Considering the cardinality of applications to packages, it's obvious the answer is one to many, at least in part.

Consider utility SSIS packages, such as a package that archives flat files after the data contained in them has been successfully loaded or staged in a database. Such a package – perhaps named ArchiveFile.dtsx – could be parameterized with parameters like SourceFilePath and DestinationLocation. ArchiveFile.dtsx could then be reused in several SSIS applications.

The cardinality of these applications to the ArchiveFile.dtsx package is many-to-one. Combining the cardinalities one-to-many and many-to-one, one gets many-to-many. Resolving many-to-many in a third normal form design requires an additional table: config.ApplicationPackages in this case. Create config.ApplicationPackages using the T-SQL shown in Listing 5-5.

Listing 5-5. Creating the config.ApplicationPackages table

```

use [SSISConfig]
go

print 'Config.ApplicationPackages table'
If Not Exists(Select [schemas].[name] + '.' + [tables].[name] As ▷
[Schema.Table]
    From [sys].[tables]
    Join [sys].[schemas]
        On [schemas].[schema_id] = [tables].[schema_id]
    Where [schemas].[name] = N'config'
        And [tables].[name] = N'ApplicationPackages')

begin
    print ' - Create config.ApplicationPackages table'
    Create Table [config].[ApplicationPackages]
    (
        ApplicationPackageId int identity(1, 1)
            Constraint PK_config_ApplicationPackages Primary Key Clustered
        , ApplicationId int Not NULL
            Constraint FK_config_ApplicationPackages_config_Applications
                Foreign Key References [config].[Applications](ApplicationId)
        , PackageId int Not NULL
            Constraint FK_config_ApplicationPackages_config_Packages
                Foreign Key References [config].[Packages](PackageId)
        , ExecutionOrder int Not NULL
            Constraint DF_config_ApplicationPackages_ExecutionOrder
                Default(10)
        , ApplicationPackageEnabled bit Not NULL
            Constraint DF_config_ApplicationPackages_ApplicationPackageEnabled
                Default(1)
        , FailApplicationOnPackageFailure bit Not NULL
            Constraint ▷ DF_config_ApplicationPackages_
                FailApplicationOnPackageFailure
                Default(1)
    )

```

```
, Constraint > UQ_config_ApplicationPackages_ApplicationId_
PackageId_ExecutionOrder
    Unique(ApplicationId, PackageId, ExecutionOrder)
)
print ' - Config.ApplicationPackages table created'
end
Else
begin
    print ' - Config.ApplicationPackages table already exists.'
end
print ''
go
```

When the config.ApplicationPackages is created, Object Explorer shows the SSISConfig tables as shown in Figure 5-6.

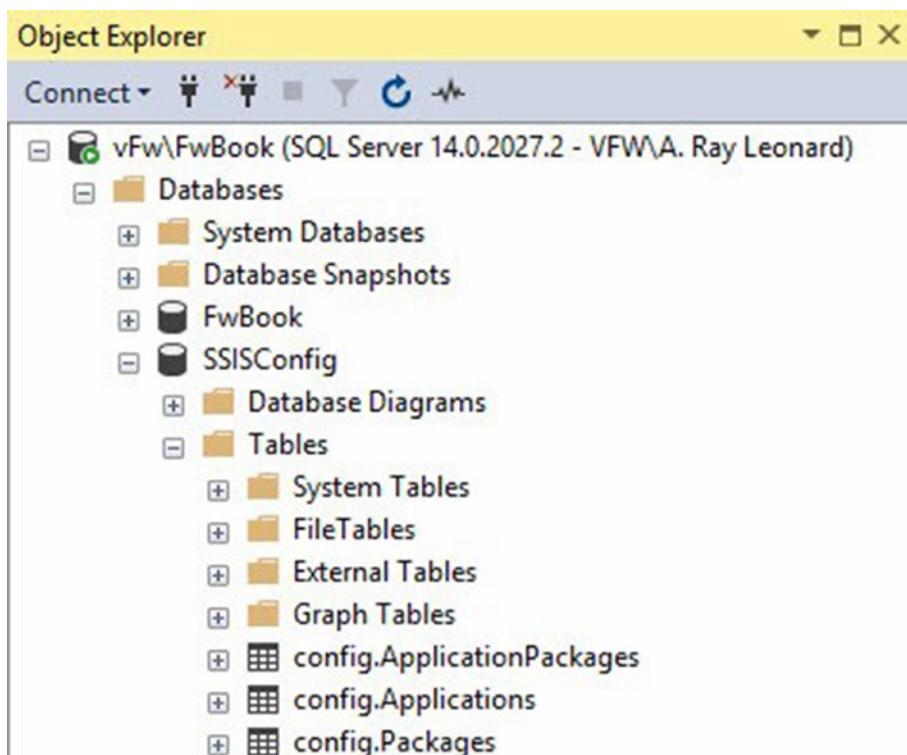


Figure 5-6. The config.ApplicationPackages table

Examine the config.ApplicationPackages table in Object Explorer by expanding the Columns, Keys, and Constraints virtual folders as shown in Figure 5-7.

The screenshot shows the Object Explorer in SQL Server Management Studio (SSMS) with the following tree structure:

- config.ApplicationPackages**
 - Columns**
 - ApplicationPackageId (PK, int, not null)
 - ApplicationId (FK, int, not null)
 - PackageId (FK, int, not null)
 - ExecutionOrder (int, not null)
 - ApplicationPackageEnabled (bit, not null)
 - FailApplicationOnPackageFailure (bit, not null)
 - Keys**
 - PK_config_ApplicationPackages
 - FK_config_ApplicationPackages_config_Applications
 - FK_config_ApplicationPackages_config_Packages
 - UQ_config_ApplicationPackages_ApplicationId_PackageId_ExecutionOrder
 - Constraints**
 - DF_config_ApplicationPackage_FailApplicationOnPackageFailure
 - DF_config_ApplicationPackages_ApplicationPackageEnabled
 - DF_config_ApplicationPackages_ExecutionOrder

Figure 5-7. Examining the config.ApplicationPackages table

The config.ApplicationPackages table columns are designed to resolve the many-to-many cardinality between the config.Applications and config.Packages tables. Each entry maps SSIS package metadata (via PackageId) to SSIS application metadata (via ApplicationId) to an identity column, ApplicationPackageId, thereby *resolving* (or *bridging*) the many-to-many relationship. Once resolved, this SSISConfig database design supports the execution of *Application Packages*, which are defined as an instance of an SSIS package mapped to an SSIS application. Application Packages are then, by design, the smallest unit of work in this SSIS framework.

Attributes of Execution

One implication of Application Packages being the smallest unit of work is Application Packages are where we define *attributes of execution*. The application package attributes of execution in this framework are

- ExecutionOrder
- ApplicationPackageEnabled
- FailApplicationOnPackageFailure

Execution Order is an integer that represents the “in a specified order” part of the SSIS application definition.

Application Package Enabled is a bit field that indicates whether the individual application package is configured to execute when the SSIS application is executed.

Fail Application On Package Failure is a bit that configures fault tolerance. Consider what happens when one executes an SSIS application: a number of SSIS packages are executed in a specified order. SSIS packages with a specified order within an SSIS application are application packages. What if we would like it very much if an SSIS package executed, completed, and succeeded, but if the package fails for some reason, it’s not vital for this particular SSIS application execution? What if we just don’t care so much if this particular package fails? The FailApplicationOnPackageFailure attribute will permit the package to fail *without* stopping the SSIS application execution.

Default Constraints

A default constraint manages the initial value of a column when a new record is added to a table and the value for field is not supplied an explicit value:

- DF_config_ApplicationPackages_FailApplicationOnPackageFailure
is a default constraint that sets the value of the FailApplicationOnPackageFailure bit to True (1) if no value is supplied when the config.ApplicationPackages record is first added to the table.
- DF_config_ApplicationPackages_ApplicationPackageEnabled
is a default constraint that sets the value of the ApplicationPackageEnabled bit to True (1) if no value is supplied when the config.ApplicationPackages record is first added to the table.

- DF_config_ApplicationPackages_ExecutionOrder is a default constraint that sets the value of the ExecutionOrder integer to 10 if no value is supplied when the config.ApplicationPackages record is first added to the table.

Relationships

Keys manage relationships between entities in the SSISConfig database and therefore in the SSIS framework:

- Config.ApplicationPackages is configured with PK_config_ApplicationPackages. Each table has a primary key for relationship management. In this framework, PK_config_ApplicationPackages serves to help identify each Application-Package-Execution Order combination.
- FK_config_ApplicationPackages_config_Applications is a foreign key between the config.ApplicationPackages.ApplicationId field and the config.Applications.ApplicationId field. Before a row is added to config.ApplicationPackages, a corresponding row *must* already exist in the config.Applications table. The Application is mapped to a Package that executes at this *position* – denoted by the value of ExecutionOrder – in the application.
- FK_config_ApplicationPackages_config_Packages is a foreign key between the config.ApplicationPackages.PackageId field and the config.Packages.PackageId field. Before a row is added to config.ApplicationPackages, a corresponding row *must* already exist in the config.Packages table. A Package that executes at a *position* – denoted by the value of ExecutionOrder – is mapped into an Application.
- UQ_config_ApplicationPackages_ApplicationId_PackageId_ExecutionOrder is a unique constraint that guarantees the distinctness – across the ApplicationId, PackageId, and ExecutionOrder fields – of each application package stored in the config.ApplicationPackages table. This means two (or more) records in the config.ApplicationPackages table *may not have identical* ApplicationId, PackageId, and ExecutionOrder field values.

Frequently Asked Questions

Frameworks are new to many data users, and there are several questions we hear regularly:

- Is it possible for two (or more) records in the config. ApplicationPackages table to have the same ApplicationId? Yes. In fact, the framework is designed for this use case. Several config. ApplicationPackages records with the same ApplicationId value is how the framework is used to configure an SSIS application.
- Is it possible for two (or more) records in the config. ApplicationPackages table to have the same PackageId? Yes. This supports code reuse, such as executing the ArchiveFile.dtsx SSIS package in several applications. This also supports an SSIS Design Pattern called “Range-Based Load.” In a Range-Based Load, the *same* SSIS package executes several times. New minimum and maximum range values, or other mathematical functions, are passed to the package for each execution. For example, the same SSIS package could execute ten times. The first execution could load records in which a numeric value begins with the numeral “0.” Subsequent executions could load data where the numeric value starts with the numerals “1,” “2,” and so on.

We will return to SSMS and development of the SSISConfig database in a bit. Let's next build a sample SSIS solution to use to test the framework.

A Sample SSIS Solution

Testing is hard. Testing requires thinking, differently, about how users will interact with an application. Developers are notoriously bad at testing their own code. Why? I believe it's because development focuses on delivering functionality that just works. For some (like me), it's difficult to change gears and begin thinking about all the permutations a user may enter – or *forget* to enter – when interacting with our code.

A first thought for building a sample SSIS solution for testing is, “I can do this with one SSIS package.” That's incorrect. Testing requires at least two SSIS packages, one package that succeeds and another that fails.

Configure Visual Studio 2019 for SSIS Development

Before you build an SSIS 2019 sample solution, you need to download and install Visual Studio 2019. To begin, browse to visualstudio.microsoft.com and select a version of Visual Studio 2019 to install, as shown in Figure 5-8.

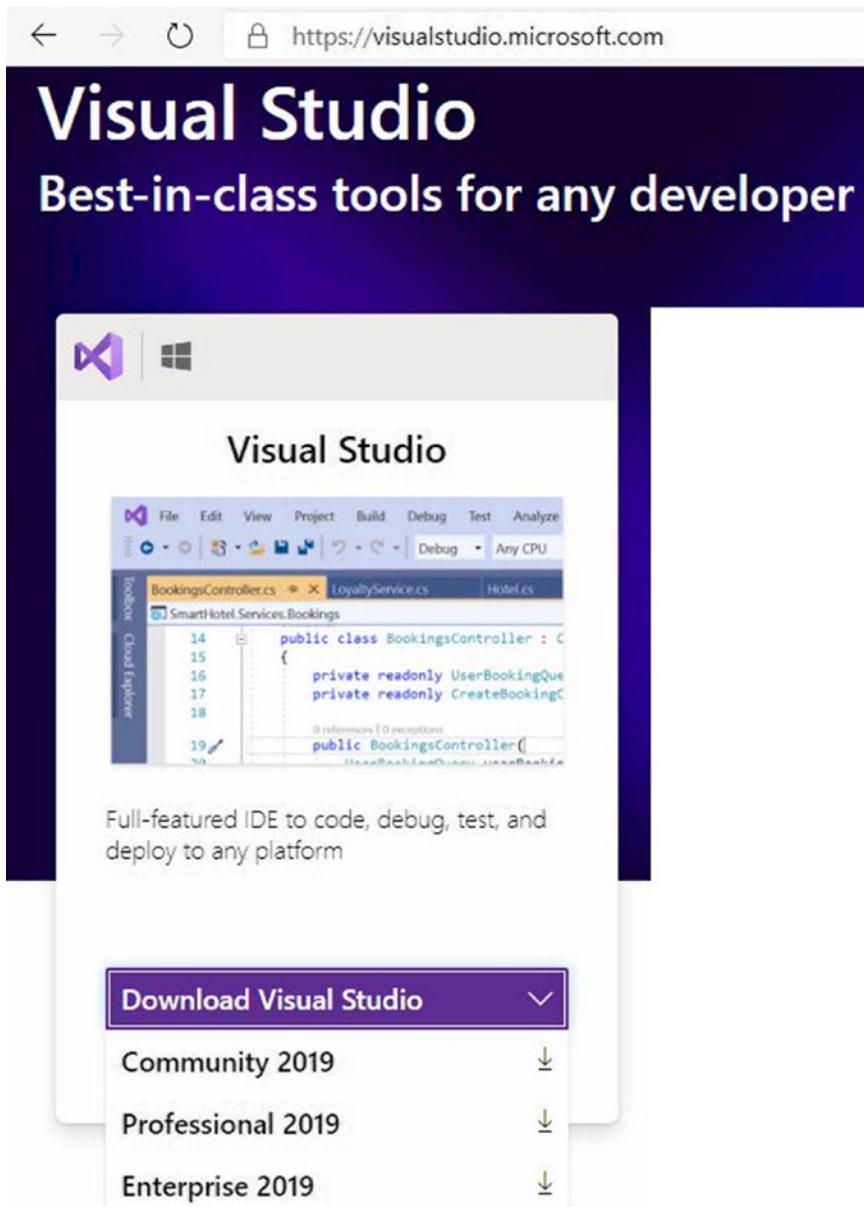


Figure 5-8. Preparing to download Visual Studio 2019

For the purposes of the example used in this book, the Community 2019 edition will suffice (plus, it's free).

The next step is to install the Integration Services extension. Visual Studio is an *integrated development environment*, or IDE. The IDE serves as a host, or shell, for various and sundry templates that enable software development experiences in different languages and platforms.

Before Visual Studio 2019, SSIS templates were installed by executing a SQL Server Data Tools, or SSDT, stand-alone installer. Starting with Visual Studio 2019, SSIS templates are managed in the Visual Studio Marketplace as another available Visual Studio extension.

To find and install the Integration Services extension in Visual Studio 2019, open Visual Studio and click Extensions ► Manage Extensions, as shown in Figure 5-9.

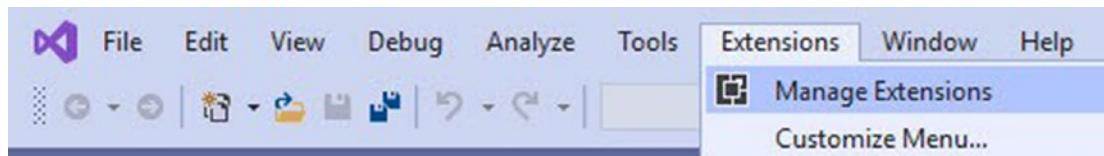


Figure 5-9. Preparing to install the Integration Services extension

When the Manage Extensions dialog displays, search for “Integration Services” in the search textbox. Select the “SQL Server Integration Services Projects” extension, and then click the Download button, as shown in Figure 5-10.

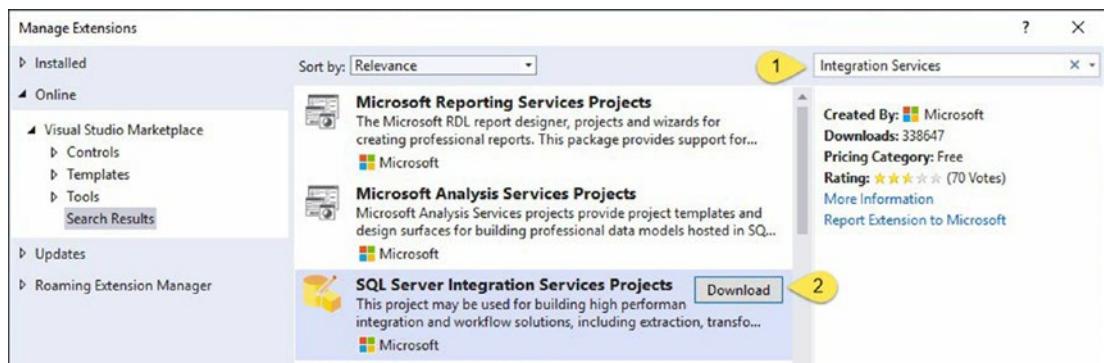


Figure 5-10. Downloading the “SQL Server Integration Services Projects” extension

When the new extension is downloaded, follow the instructions presented from the Visual Studio 2019 extension process. The extension installation process likely begins with closing the Visual Studio IDE. When the extension installation is complete, Visual Studio 2019 may be used to develop SSIS projects.

Create the Sample SSIS Solution

To begin, create a new SSIS solution named TestSSISSolution – containing a new SSIS project named TestSSISProject – as shown in Figure 5-11.

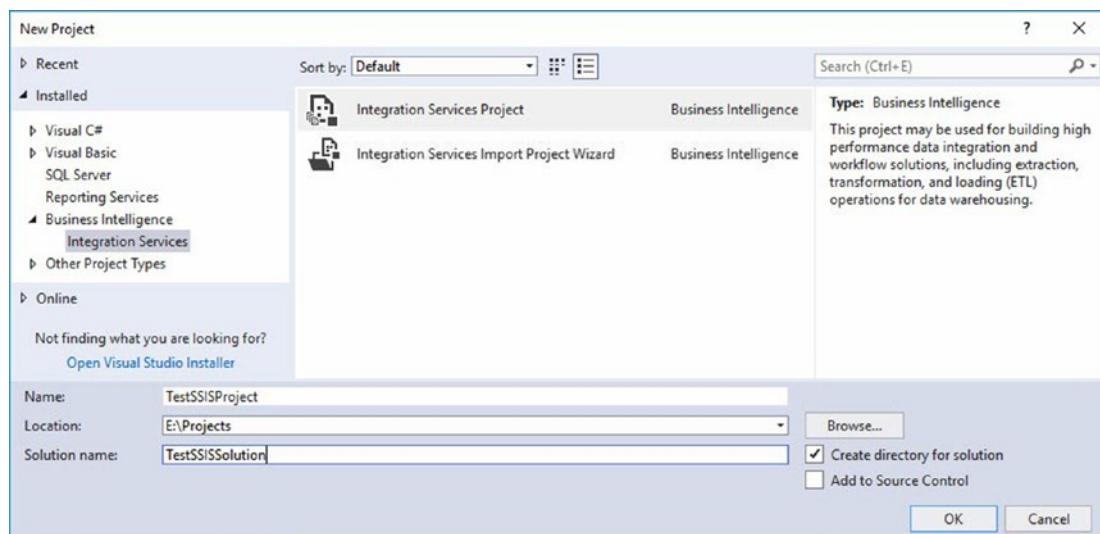


Figure 5-11. Creating the TestSSISProject in TestSSISSolution

When the project opens, rename the default SSIS package to “ReportAndSucceed.dtsx” as shown in Figure 5-12.

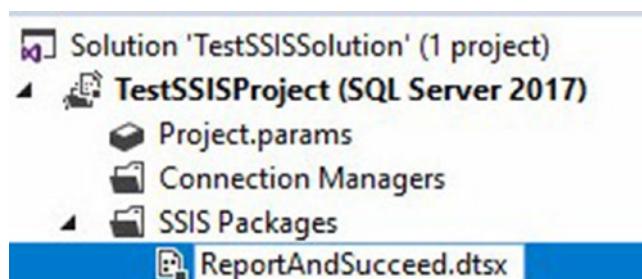


Figure 5-12. Renaming the package “ReportAndSucceed.dtsx”

Add a Script Task to the ReportAndSucceed's Control Flow by dragging a Script Task from the SSIS Toolbox onto the Control canvas, and then rename it to "SCR Log Values," as shown in Figure 5-13.

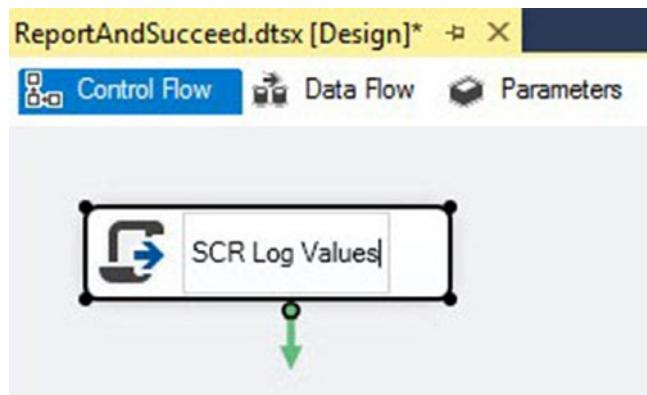


Figure 5-13. Adding and renaming “SCR Log Values” script task

Open the Script Task Editor by right-clicking the “SCR Log Values” script task and clicking “Edit.” When the “SCR Log Values” script task editor displays, add the following variables to the ReadOnlyVariables property list:

- System::PackageName
- System::TaskName

The Script Task Editor will appear as shown in Figure 5-14.

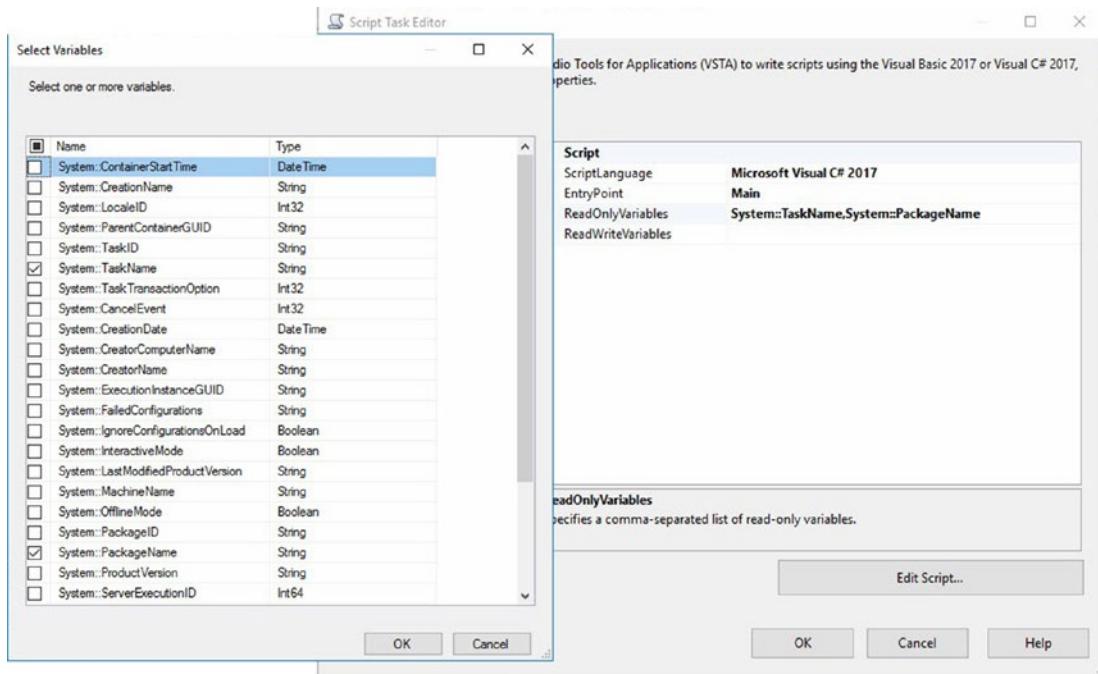


Figure 5-14. Adding System::PackageName and System::TaskName to ReadOnlyVariables

Click the Edit Script button, and enter the C# shown in Listing 5-6 into the public void Main() method.

Listing 5-6. Main() method C# to configure raising an Information event in the ReportAndSucceed SSIS package

```
public void Main()
{
    string packageName = ► Dts.Variables["System::PackageName"].Value.ToString();
    string taskName = ► Dts.Variables["System::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int informationCode = 1001;
    bool fireAgain = true;
    string description = "I am " + packageName;
```

```

    Dts.Events.FireInformation(informationCode, subComponent, >
        description, "", 0, ref fireAgain);

    Dts.TaskResult = (int)ScriptResults.Success;
}

```

When the code has been entered into the `Main()` method, it will appear as shown in Figure 5-15.

```

public void Main()
{
    string packageName = Dts.Variables["System::PackageName"].Value.ToString();
    string taskName = Dts.Variables["System::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int informationCode = 1001;
    bool fireAgain = true;
    string description = "I am " + packageName;

    Dts.Events.FireInformation(informationCode, subComponent, description, "", 0, ref fireAgain);

    Dts.TaskResult = (int)ScriptResults.Success;
}

```

Figure 5-15. *Main() method C# that raises an Information event*

Close the Script editor window, and click the OK button on the SCR Log Values script task. Click Debug ➤ Start Debugging as shown in Figure 5-16.

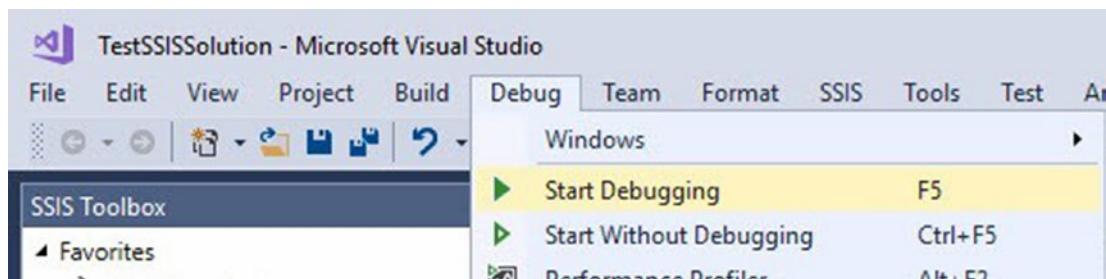


Figure 5-16. *Starting a debug execution of the ReportAndSucceed package*

When execution completes – and hopefully succeeds – click the Progress (or Execution Results, if you stop the debugger) to view the Information event we configure in the SCR Log Values script task as shown in Figure 5-17.

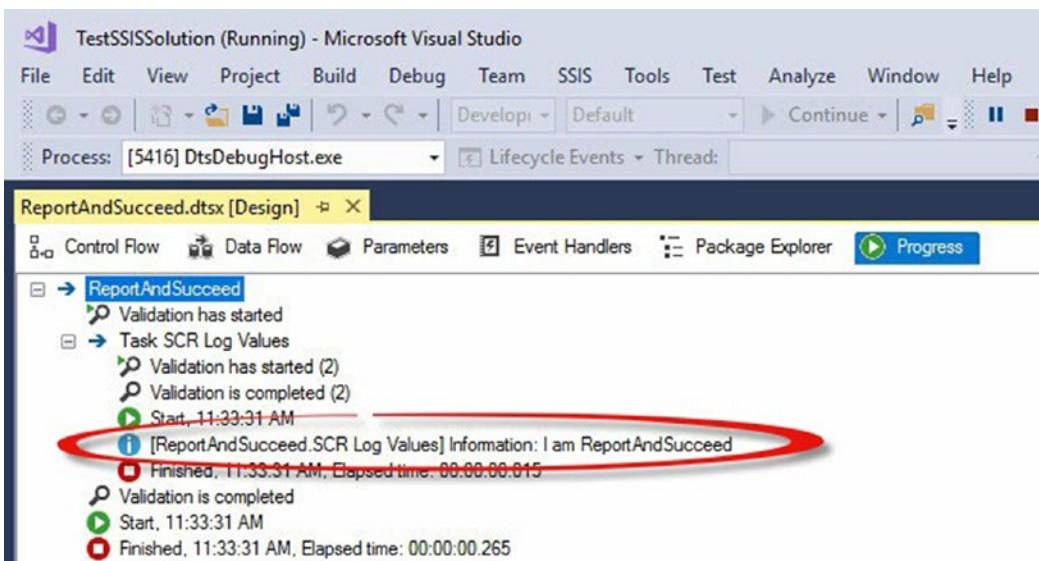


Figure 5-17. The Information event

The Information event reads: “[ReportAndSucceed.SCR Log Values] Information: I am ReportAndSucceed”.

The first part of the message is the subComponent and is enclosed in square brackets (“[]”). The subComponent is built from the Package name and Task name variables: “ReportAndSucceed.SCR Log Values”. The event name – Information – follows the subComponent. The description is next and reads “I am ReportAndSucceed”.

Report and succeed is exactly what happens when this package executes.

Add another SSIS package to the TestSSISProject SSIS project, and rename it to “ReportAndFail,” as shown in Figure 5-18.

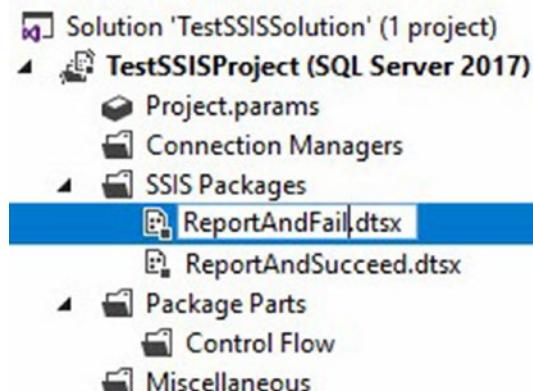


Figure 5-18. Adding the ReportAndFail SSIS package

As before, add a Script Task named “SCR Log Values” to the ReportAndFail package. Add the System::PackageName and System::TaskName variables to the ReadOnlyVariables property, and then click the Edit Script button. Add the C# in Listing 5-7 to the Main() method.

Listing 5-7. Main() method C# to configure and raise an Error event in the ReportAndFail SSIS package

```
public void Main()
{
    string packageName = ► Dts.Variables["System::PackageName"].Value.ToString();
    string taskName = ► Dts.Variables["System::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int errorCode = -1001;
    string description = packageName + " execution failed";

    Dts.Events.FireError(errorCode, subComponent, description,
    "", ► 0);

    Dts.TaskResult = (int)ScriptResults.Success;
}
```

When the code has been entered into the Main() method, it will appear as shown in Figure 5-19.

```
public void Main()
{
    string packageName = Dts.Variables["System::PackageName"].Value.ToString();
    string taskName = Dts.Variables["System::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int errorCode = -1001;
    string description = packageName + " execution failed";

    Dts.Events.FireError(errorCode, subComponent, description, "", 0);

    Dts.TaskResult = (int)ScriptResults.Success;
}
```

Figure 5-19. Main() method C# that raises an Error event

Close the Script editor window by clicking the “X” in the upper-right corner of the VstaProjects .Net code editor, or by clicking File ➤ Exit, and then click the OK button on the SCR Log Values script task. Click Debug ➤ Start Debugging to start debug execution of the ReportAndFail SSIS package.

When execution completes – and hopefully *fails* – click the Progress (or Execution Results, if you stop the debugger) to view the Information event we configure in the SCR Log Values script task as shown in Figure 5-20.

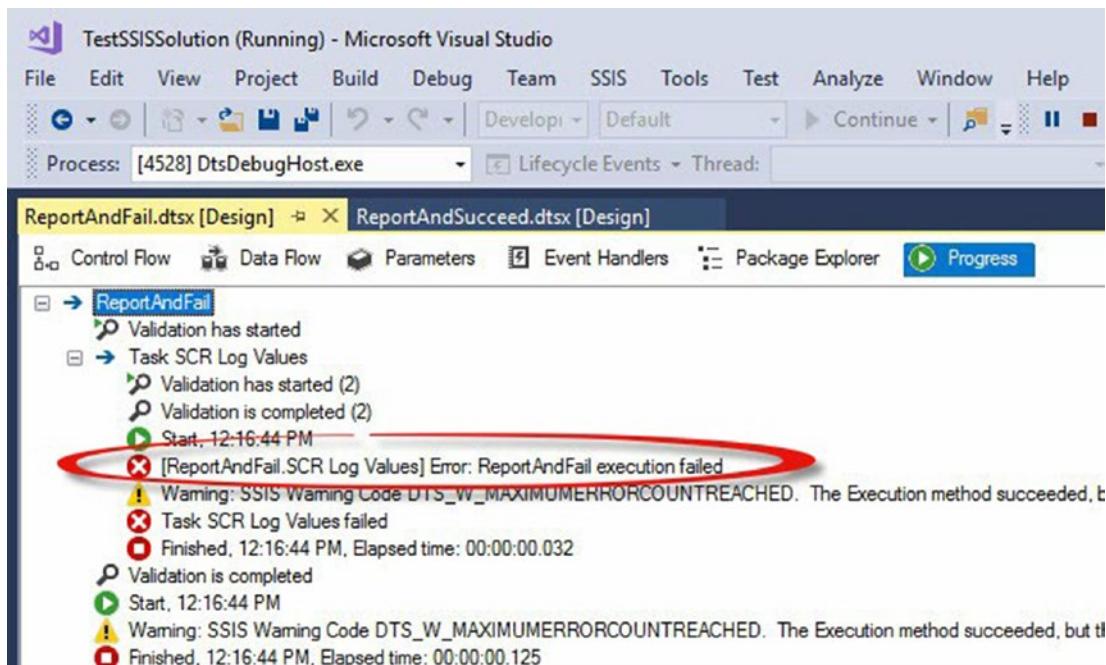


Figure 5-20. The Error event

The error event reads: “[ReportAndFail.SCR Log Values] Error: ReportAndFail execution failed”.

The first part of the message is the subComponent and is enclosed in square brackets (“[]”). The subComponent is built from the Package name and Task name variables: “ReportAndFail.SCR Log Values”. The event name – Error – follows the subComponent. The description is next and reads “ReportAndFail execution failed”.

Report and fail is exactly what happens when this package executes.

At a minimum, a test project should include components that succeed *and* fail.

SSIS Framework Metadata Management

In enterprise architecture or any flavor of development, solving one problem often creates a new problem, or problems. I call this “there’s no free lunch.” Growing through the “there’s no free lunch” phase is a rite of passage for developers, it’s a natural part of maturing, and it’s a *requirement* for technology architects.

For example, the SSIS framework described in this chapter will simplify the execution of several SSIS packages in a specified order by using metadata. There will be lots of metadata to manage.

There’s no free lunch.

All people who write T-SQL have preferences regarding capitalization and indentation. I am no different. I find my T-SQL capitalization and indentation helps me think about the problem I am trying to solve.

Asking yourself or your team, “What is the problem I am/we are trying to solve?” is potent. Try it and see.

Add an SSIS Application

Add an SSIS application to the SSIS framework by executing the T-SQL shown in Listing 5-8.

Listing 5-8. Adding an SSIS application

```
use [SSISConfig]
go

Set NoCount ON

declare @ApplicationName nvarchar(255) = N'Framework Test'

print @ApplicationName
declare @ApplicationId int = (Select [Applications].[ApplicationId]
                                From [config].[Applications]
                                Where [Applications].[ApplicationName] =
                                     > @ApplicationName)
```

```
If (@ApplicationId Is NULL)
begin
    print ' - Adding ' + @ApplicationName + ' application to ►
config.Applications table'

declare @AppTbl table(ApplicationId int)

Insert Into [config].[Applications]
(ApplicationName)
Output inserted.ApplicationId into @AppTbl
Values (@ApplicationName)

Set @ApplicationId = (Select ApplicationId
                      From @AppTbl)

print ' - ' + @ApplicationName + ' application added to ►
config.Applications table'
end
Else
begin
    print ' - ' + @ApplicationName + ' application already exists in ►
the config.Applications table.'
end

Select @ApplicationId As ApplicationId
print ''
```

The T-SQL in Listing 5-8 is *idempotent*. Idempotence is a mathematical term that means an operation may be applied multiple times and produce the same result. Idempotence is another way of saying “re-executable code.”

Applied to the T-SQL code in Listing 5-8, re-executable code produces the same result in the config.Applications table; metadata for a new SSIS application named “Framework Test” is added to the framework. The first execution of the T-SQL in Listing 5-8 returns the messages (to the Messages tab in SSMS) shown in Listing 5-9.

Listing 5-9. Messages returned when adding the Framework Test SSIS application for the first time

Framework Test

- Adding Framework Test application to config.Applications table
- Framework Test application added to config.Applications table

The result, which is the ApplicationId for the “Framework Test” SSIS application, is shown in Figure 5-21.

	ApplicationId
1	1

Figure 5-21. “Framework Test” ApplicationId

Re-executing the Add SSIS Application T-SQL shown in Listing 5-8 produces the messages shown in Figure 5-22.

Framework Test
- Framework Test application already exists in the config.Applications table.

Figure 5-22. Re-executing SSIS Application T-SQL

Re-executing the Add SSIS Application T-SQL shown in Listing 5-8 produces the same results, as shown in Figure 5-23.

	ApplicationId
1	1

Figure 5-23. Identical Framework Test ApplicationId

Subsequent re-executions of the Add SSIS Application T-SQL produce the same result.

Add SSIS Packages

Add SSIS packages in TestSSISProject to the SSIS framework by executing the T-SQL shown in Listing 5-10.

Listing 5-10. Adding SSIS packages from TestSSISProject

```
use [SSISConfig]
go
Set NoCount ON

declare @PackageLocation nvarchar(255) = ► N'E:\Projects\TestSSISolution\
TestSSISProject\' 
declare @PackageName nvarchar(255) = N'ReportAndSucceed.dtsx'

print @PackageLocation + @PackageName
declare @PackageId int = (Select [Packages].[PackageId]
                           From [config].[Packages]
                           Where [Packages].[PackageLocation] =
► @PackageLocation
                           And [Packages].[PackageName] = @PackageName)

If (@PackageId Is NULL)
begin
    print ' - Adding ' + @PackageName + ' package to config.Packages table'
    declare @PkgTbl table(PackageId int)

    Insert Into [config].[Packages]
    (PackageLocation, PackageName)
    Output inserted.PackageId into @PkgTbl
    Values (@PackageLocation, @PackageName)

    Set @PackageId = (Select PackageId
                      From @PkgTbl)

    print ' - ' + @PackageName + ' package added to config.Packages table'
end
Else
```

```
begin
    print ' - ' + @PackageName + ' package already exists in the ► config.
    Packages table.'
end

Select @PackageId As PackageId
print ''

set @PackageName = N'ReportAndFail.dtsx'

print @PackageLocation + @PackageName
set @PackageId = (Select [Packages].[PackageId]
                  From [config].[Packages]
                  Where [Packages].[PackageLocation] = @PackageLocation
                        And [Packages].[PackageName] = @PackageName)

If (@PackageId Is NULL)
begin
    print ' - Adding ' + @PackageName + ' application to config.Packages
    table'

    Delete @PkgTbl

    Insert Into [config].[Packages]
    (PackageLocation, PackageName)
    Output inserted.PackageId into @PkgTbl
    Values (@PackageLocation, @PackageName)

    Set @PackageId = (Select PackageId
                      From @PkgTbl)

    print ' - ' + @PackageName + ' package added to config.Packages table'
end
Else
begin
    print ' - ' + @PackageName + ' package already exists in the ► config.
    Packages table.'
end

Select @PackageId As PackageId
print ''
```

Like the T-SQL in Listing 5-8, the T-SQL in Listing 5-10 is also idempotent, so the T-SQL code in Listing 5-10 produces the same result in the config.Packages table; metadata for two new SSIS packages named “ReportAndSucceed.dtsx” and “ReportAndFail.dtsx” are added to the framework. The first execution of the T-SQL in Listing 5-10 returns the messages (to the Messages tab in SSMS) shown in Listing 5-11.

Listing 5-11. Messages returned when adding ReportAndSucceed.dtsx and ReportAndFail.dtsx SSIS packages for the first time

```
E:\Projects\TestSSISolution\TestSSISProject\ReportAndSucceed.dtsx
```

- Adding ReportAndSucceed.dtsx package to config.Packages table
- ReportAndSucceed.dtsx package added to config.Packages table

```
E:\Projects\TestSSISsolution\TestSSISProject\ReportAndFail.dtsx
```

- Adding ReportAndFail.dtsx application to config.Packages table
- ReportAndFail.dtsx package added to config.Packages table

The result, which is the PackageIds for the “ReportAndSucceed.dtsx” and “ReportAndFail.dtsx” SSIS packages, is shown in Figure 5-24.

	Packageld
1	1
1	2

Figure 5-24. “ReportAndSucceed.dtsx” and “ReportAndFail.dtsx” SSIS PackageIds

Re-executing the Add SSIS Packages T-SQL shown in Listing 5-10 produces the messages shown in Figure 5-25.

Results	Messages
E:\Projects\TestSSISsolution\TestSSISProject\ReportAndSucceed.dtsx	
	- ReportAndSucceed.dtsx package already exists in the config.Packages table.
E:\Projects\TestSSISsolution\TestSSISProject\ReportAndFail.dtsx	
	- ReportAndFail.dtsx package already exists in the config.Packages table.

Figure 5-25. Re-executing SSIS Packages T-SQL

Re-executing the Add SSIS Packages T-SQL shown in Listing 5-10 produces the *same* results, as shown in Figure 5-26.

		Results	Messages
		Packageld	
1		1	

		Results	Messages
		Packageld	
1		2	

Figure 5-26. Identical PackageIds for “ReportAndSucceed.dtsx” and “ReportAndFail.dtsx” SSIS packages

Subsequent re-executions of the Add SSIS Packages T-SQL produce the same result.

Assign SSIS Application Packages

Assign SSIS application packages in SSIS framework by executing the T-SQL shown in Listing 5-12.

Listing 5-12. Assigning SSIS application packages

```
use [SSISConfig]
go

Set NoCount ON

declare @ApplicationName nvarchar(255) = N'Framework Test'
declare @PackageLocation nvarchar(255) = ► N'E:\Projects\TestSSISolution\
TestSSISProject\
declare @PackageName nvarchar(255) = N'ReportAndSucceed.dtsx'
declare @ExecutionOrder int = 10

print @ApplicationName + ' - ' + @PackageLocation + @PackageName

declare @ApplicationId int = (Select [Applications].[ApplicationId]
                                From [config].[Applications]
                                Where [Applications].[ApplicationName] =
► @ApplicationName)
```

CHAPTER 5 A SIMPLE, CUSTOM, FILE-BASED SSIS FRAMEWORK

```
declare @PackageId int = (Select [Packages].[PackageId]
                           From [config].[Packages]
                           Where [Packages].[PackageLocation] =
                                ► @PackageLocation
                           And [Packages].[PackageName] = @PackageName)

declare @ApplicationPackageId int = (Select ApplicationPackageId
                                       From config.ApplicationPackages
                                       Where ApplicationId = @ApplicationId
                                         And PackageId = @PackageId
                                         And ExecutionOrder =
                                              @ExecutionOrder)

If (@ApplicationPackageId Is NULL)
begin
    print ' - Assigning ' + @PackageName + ' package to '
          + @ApplicationName + ' application'
          + ' in config.ApplicationPackages table'
          + ' at ExecutionOrder ' + Convert(varchar(9), @ExecutionOrder)

    Insert Into [config].[ApplicationPackages]
    (ApplicationId
     , PackageId
     , ExecutionOrder)
    Values (@ApplicationId
            , @PackageId
            , @ExecutionOrder)

    print ' - ' + @PackageName + ' package assigned to '
          + @ApplicationName + ' application'
          + ' in config.ApplicationPackages table'
          + ' at ExecutionOrder ' + Convert(varchar(9), @ExecutionOrder)
end
Else
begin
    print ' - ' + @PackageName + ' package already'
          + ' assigned to ' + @ApplicationName
```

```

+ ' application in config.ApplicationPackages table'
+ ' at ExecutionOrder ' + Convert(varchar(9), @ExecutionOrder)
+ '.'

end
print ''

set @PackageName = N'ReportAndFail.dtsx'
set @ExecutionOrder = 20

print @ApplicationName + ' - ' + @PackageLocation + @PackageName

set @ApplicationId = (Select [Applications].[ApplicationId]
                      From [config].[Applications]
                      Where [Applications].[ApplicationName] =
                            ► @ApplicationName)

set @PackageId = (Select [Packages].[PackageId]
                      From [config].[Packages]
                      Where [Packages].[PackageLocation] = @PackageLocation
                            And [Packages].[PackageName] = @PackageName)

set @ApplicationPackageId = (Select ApplicationPackageId
                                From config.ApplicationPackages
                                Where ApplicationId = @ApplicationId
                                      And PackageId = @PackageId
                                      And ExecutionOrder = @ExecutionOrder)

If (@ApplicationPackageId Is NULL)
begin
    print ' - Assigning ' + @PackageName + ' package to '
          + @ApplicationName + ' application'
          + ' in config.ApplicationPackages table'
          + ' at ExecutionOrder ' + Convert(varchar(9), @ExecutionOrder)

    Insert Into [config].[ApplicationPackages]
    (ApplicationId
    , PackageId
    , ExecutionOrder)
    Values (@ApplicationId

```

```

    , @PackageId
    , @ExecutionOrder)

print ' - ' + @PackageName + ' package assigned to '
    + @ApplicationName + ' application'
    + ' in config.ApplicationPackages table'
    + ' at ExecutionOrder ' + Convert(varchar(9), @ExecutionOrder)
end
Else
begin
    print ' - ' + @PackageName + ' package already'
        + ' assigned to ' + @ApplicationName
        + ' application in config.ApplicationPackages table'
        + ' at ExecutionOrder ' + Convert(varchar(9), @ExecutionOrder)
        + '.'
end
print ''

```

Like the T-SQL in Listing 5-8 and Listing 5-10, the T-SQL in Listing 5-12 is idempotent. The T-SQL code in Listing 5-12 produces the same result in the config.ApplicationPackages table; metadata for two new SSIS packages named “ReportAndSucceed.dtsx” and “ReportAndFail.dtsx” are assigned to the “Framework Test” application in the framework. The first execution of the T-SQL in Listing 5-12 returns the messages shown in Listing 5-13.

Listing 5-13. Messages returned when assigning ReportAndSucceed.dtsx and ReportAndFail.dtsx SSIS packages to the Framework Test application for the first time

```

Framework Test - ➤ E:\Projects\TestSSISolution\TestSSISProject\
ReportAndSucceed.dtsx
- Assigning ReportAndSucceed.dtsx package to Framework Test application in
➤ config.ApplicationPackages table at ExecutionOrder 10
- ReportAndSucceed.dtsx package assigned to Framework Test application in
➤ config.ApplicationPackages table at ExecutionOrder 10

```

- Framework Test - ➤ E:\Projects\TestSSISSolution\TestSSISProject\ReportAndFail.dtsx
- Assigning ReportAndFail.dtsx package to Framework Test application in ➤ config.ApplicationPackages table at ExecutionOrder 20
 - ReportAndFail.dtsx package assigned to Framework Test application in ➤ config.ApplicationPackages table at ExecutionOrder 20

Re-executing the Add SSIS Application Packages T-SQL shown in Listing 5-12 produces the messages shown in Figure 5-27.

```
Messages
Framework Test - E:\Projects\TestSSISSolution\TestSSISProject\ReportAndSucceed.dtsx
- ReportAndSucceed.dtsx package already assigned to Framework Test application in config.ApplicationPackages table at ExecutionOrder 10.

Framework Test - E:\Projects\TestSSISSolution\TestSSISProject\ReportAndFail.dtsx
- ReportAndFail.dtsx package already assigned to Framework Test application in config.ApplicationPackages table at ExecutionOrder 20.
```

Figure 5-27. Re-executing SSIS Application Packages T-SQL

Subsequent re-executions of the Add SSIS Application Packages T-SQL produce the same result.

The SSIS framework metadata database contains a schema named config. The config schema contains three tables:

1. Config.Applications
2. Config.Packages
3. Config.ApplicationPackages

The tables contain metadata for an application named Framework Test, which is related to two SSIS packages named ReportAndSucceed.dtsx and ReportAndFail.dtsx. It's possible to retrieve the metadata by executing the T-SQL shown in Listing 5-14.

Listing 5-14. Viewing the framework application contents

```
Use [SSISConfig]
go

declare @ApplicationName nvarchar(255) = N'Framework Test'

Select a.ApplicationName
    , p.PackageLocation + p.PackageName As PackagePath
    , ap.ExecutionOrder
    , ap.FailApplicationOnPackageFailure
```

CHAPTER 5 A SIMPLE, CUSTOM, FILE-BASED SSIS FRAMEWORK

```
From [config].[ApplicationPackages] ap
Join [config].[Applications] a
  On a.ApplicationId = ap.ApplicationId
Join [config].Packages p
  On p.PackageId = ap.PackageId
Where a.ApplicationName = @ApplicationName
  And ap.ApplicationPackageEnabled = 1
Order By ap.ExecutionOrder
```

The SSIS framework metadata returned from the T-SQL query shown in Listing 5-14 is shown in Figure 5-28.

	ApplicationName	PackagePath	ExecutionOrder	FailApplicationOnPackageFailure
1	Framework Test	E:\Projects\TestSSISSolution\TestSSISProject\ReportAndSucceed.dtsx	10	1
2	Framework Test	E:\Projects\TestSSISSolution\TestSSISProject\ReportAndFail.dtsx	20	1

Figure 5-28. Results from the SSIS framework metadata query

The metadata returned from the SSIS framework metadata is more than enough to execute in the SSIS framework execution engine. Our next step is to build the execution engine for the SSIS framework.

Conclusion

In this chapter, we focused on

- Building the metadata database
- Building a test SSIS project
- Adding metadata to the metadata database

The next step is building the execution engine, which we accomplish in the next chapter.

CHAPTER 6

Framework Execution Engine

Now that the metadata database and test application have been built in the previous chapter, it's time to focus on the execution engine. Recall that in the previous chapter, near the beginning, I emphasized the importance of empathy in software architecture and design. I wrote

I assume the SSIS developer, or SSIS Developer Team, will be familiar with SSIS, so I build the SSIS framework in SSIS. I want the developer or team to manage and maintain their framework.

Chapter 5 covered building out the framework and managing the metadata. Now it's time to think about how our SSIS packages will be executed.

Create a Parent SSIS Package

Create a new SSIS project named "SSISFrameworkProject" in a new SSIS solution named "SSISFrameworkSolution." Rename the default SSIS package to "Parent.dtsx" as shown in Figure 6-1.

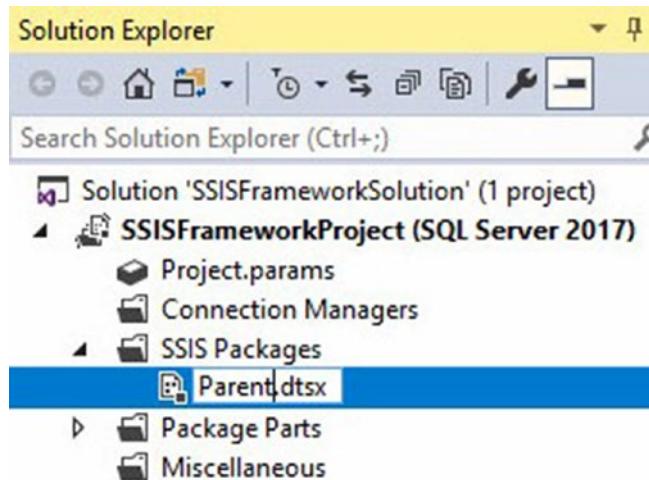


Figure 6-1. Renaming the Parent.dtsx SSIS package in the SSISFrameworkProject in SSISFrameworkSolution

Parent.dtsx will

- Log execution values
- Retrieve a list of SSIS application packages from the SSISConfig database – application packages that belong to an SSIS application
- Iterate the retrieved list of application packages
- Log metadata about each application package
- Execute each application package
- Log execution results

Log Execution Values

Begin by adding a String data type package parameter named “ApplicationName,” defaulted to “Framework Test,” as seen in Figure 6-2.

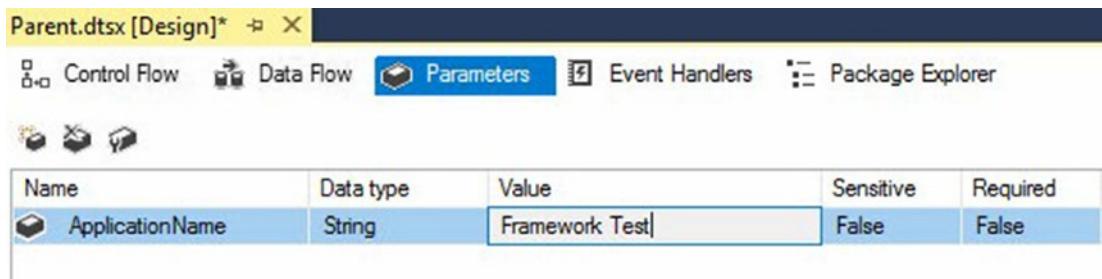


Figure 6-2. Adding the ApplicationName package parameter to Parent.dtsx

Leave Sensitive and Required parameter properties set to their default values: False. The ApplicationName parameter contains the name of the SSIS framework SSIS application to be executed when Parent.dtsx is executed. The value of the ApplicationName parameter will be overridden at runtime.

To log execution values, add a Script Task to the Parent.dtsx Control Flow. Rename the Script Task to “SCR Log Initial Values,” as shown in Figure 6-3.

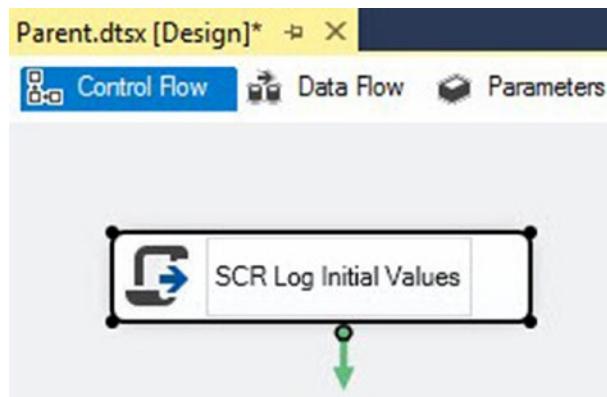


Figure 6-3. Adding and renaming “SCR Log Initial Values” script task

The SCR Log Initial Values script task will be used to capture *instrumentation* about the Parent.dtsx package’s initial state at execution start. Instrumentation is important in any engineering endeavor. This holds for data engineering, as well as ETL (extract, transform, and load) instrumentation. Capturing the initial state of settings at the start of SSIS package execution is *vital* for troubleshooting. As my friend Grant Fritchey says, “How do you know what’s wrong unless you know what *right* looks like?”

Open the SCR Log Initial Values script task editor, and add the following SSIS variables and parameter to the ReadOnlyVariables property, as shown in Figure 6-4 and Figure 6-5.

- System::PackageName
- System::TaskName
- \$Package::ApplicationName

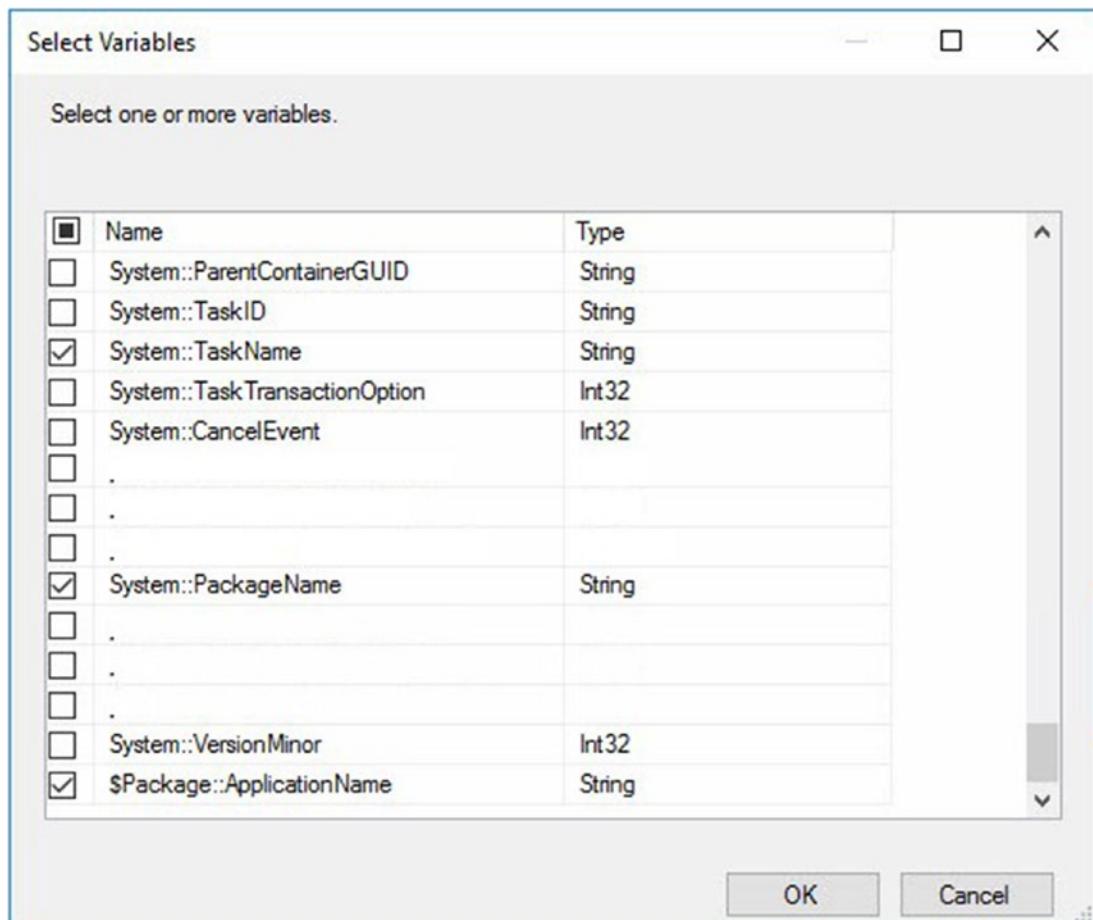


Figure 6-4. Selecting two SSIS variables and one parameter for the ReadOnlyVariables script task property

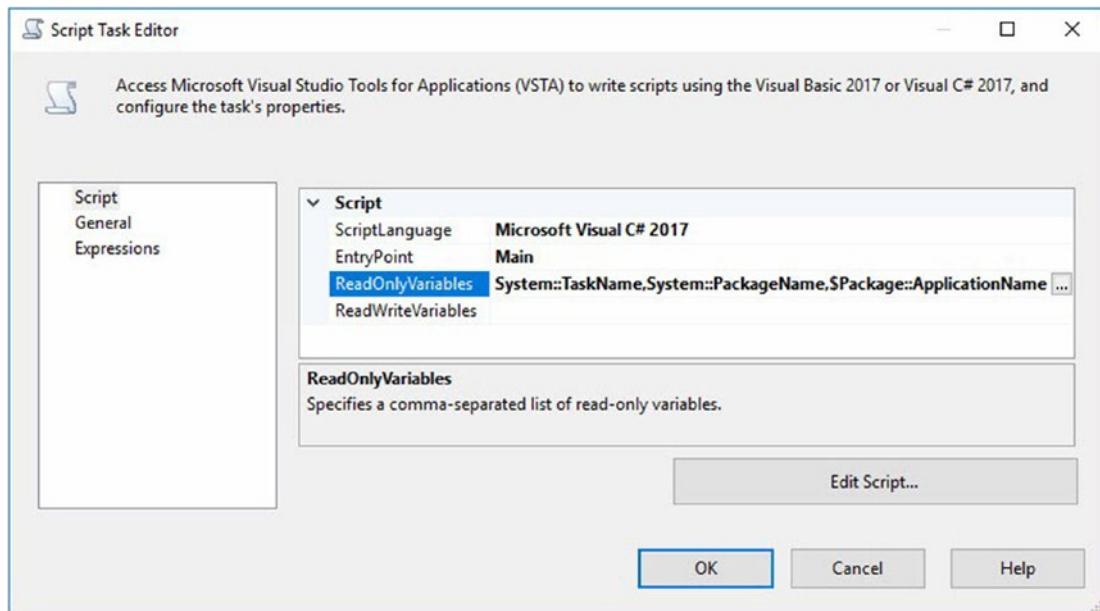


Figure 6-5. Two SSIS variables and one parameter added to the *ReadOnlyVariables* script task property

Copy the list of variables and parameters from the *ReadOnlyVariables* script task property. Click the *Edit Script* button to open the Visual Studio Tools for Applications (VSTA) .Net code editor window named *VstaProjects*. Navigate to `public void Main()`, and paste the SSIS variables and parameter in a comment, as shown in Figure 6-6.

```
public void Main()
{
    // Variables: System::TaskName, System::PackageName
    // Parameters: $Package::ApplicationName

    Dts.TaskResult = (int)ScriptResults.Success;
}
```

Figure 6-6. Copying the SSIS variables and parameter into *VstaProjects*

Add the .Net C# code shown in Listing 6-1 to log initial values or the *ApplicationName* parameter.

Listing 6-1. C# .Net code to log the initial value of the ApplicationName parameter

```
public void Main()
{
    // Variables: System::TaskName, System::PackageName
    // Parameters: $Package::ApplicationName

    string packageName = ► Dts.Variables["System::PackageName"].Value.ToString();
    string taskName = ► Dts.Variables["System::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int informationCode = 1001;
    bool fireAgain = true;

    string applicationName = ► Dts.Variables["$Package::ApplicationName"].Value.ToString();
    string description = "ApplicationName: " + applicationName;

    Dts.Events.FireInformation(informationCode, subComponent,►
        description, "", 0, ref fireAgain);

    Dts.TaskResult = (int)ScriptResults.Success;
}
```

Once the code in `public void Main()` matches Listing 6-1, the VstaProjects window should configure and raise an SSIS Information event, as shown in Figure 6-7.

```

public void Main()
{
    // Variables: System::TaskName, System::PackageName
    // Parameters: $Package::ApplicationName

    string packageName = Dts.Variables["System::PackageName"].Value.ToString();
    string taskName = Dts.Variables["System::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int informationCode = 1001;
    bool fireAgain = true;

    string applicationName = Dts.Variables["$Package::ApplicationName"].Value.ToString();
    string description = "ApplicationName: " + applicationName;

    Dts.Events.FireInformation(informationCode, subComponent, description, "", 0, ref fireAgain);

    Dts.TaskResult = (int)ScriptResults.Success;
}

```

Figure 6-7. C# .Net code that raises an SSIS Information event

Close the VstaProjects window, and click the OK button on the Script Task Editor. Execute the Parent.dtsx SSIS package in the debugger (F5). If all goes well, the SCR Log Initial Values script task should execute and succeed as shown in Figure 6-8.

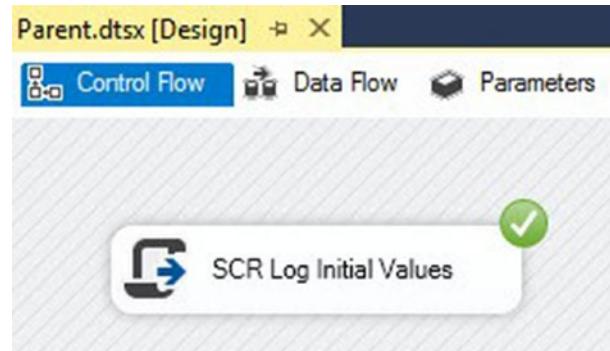


Figure 6-8. Successful debug execution of Parent.dtsx

Click the Progress tab (or Execution Results tab if you've already stopped the debugger), and view the Information event coded in the SCR Log Initial Values script task, as shown in Figure 6-9.

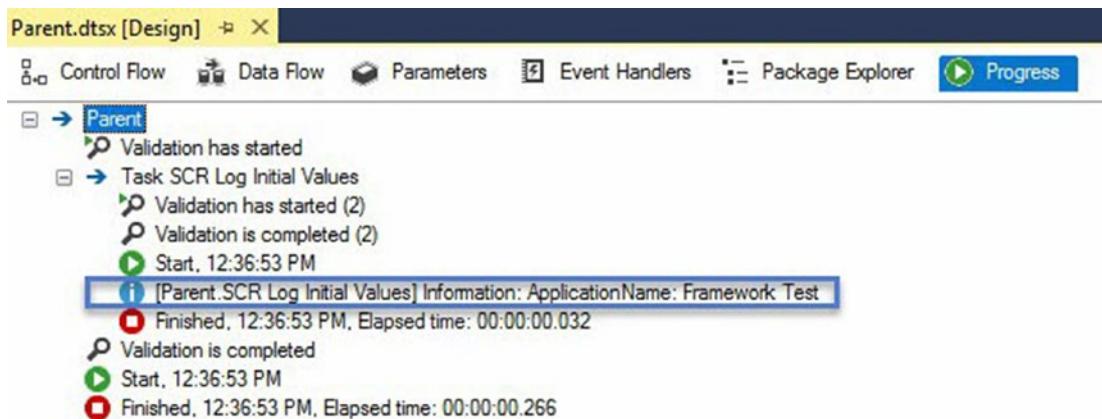


Figure 6-9. Viewing the Information event

The text of the Information event is shown in Listing 6-2.

Listing 6-2. Contents of the Information event

[Parent.SCR Log Initial Values] Information: ApplicationName: Framework Test

The value of the subComponent is listed first, enclosed in square brackets: [Parent. SCR Log Initial Values]. The event type follows the subComponent: Information. The description string is last: ApplicationName: Framework Test.

This message will appear in any logging configured for the Parent.dtsx SSIS package and will inform operators which SSIS Application Name was passed to the Parent.dtsx package.

Retrieve SSIS Application Packages from SSISConfig

The next step is to retrieve a list of packages to execute in the Parent.dtsx SSIS package. Add an Execute SQL Task to the Parent.dtsx control flow, and rename it to “SQL Get Application Packages,” as shown in Figure 6-10.



Figure 6-10. Adding the “SQL Get Application Packages” Execute SQL Task

Connect a precedence constraint from the SCR Log Initial Values script task to the SQL Get Application Packages execute SQL task, and then open the SQL Get Application Packages execute SQL task editor by right-clicking the task and then clicking “Edit,” as shown in Figure 6-11.

CHAPTER 6 FRAMEWORK EXECUTION ENGINE

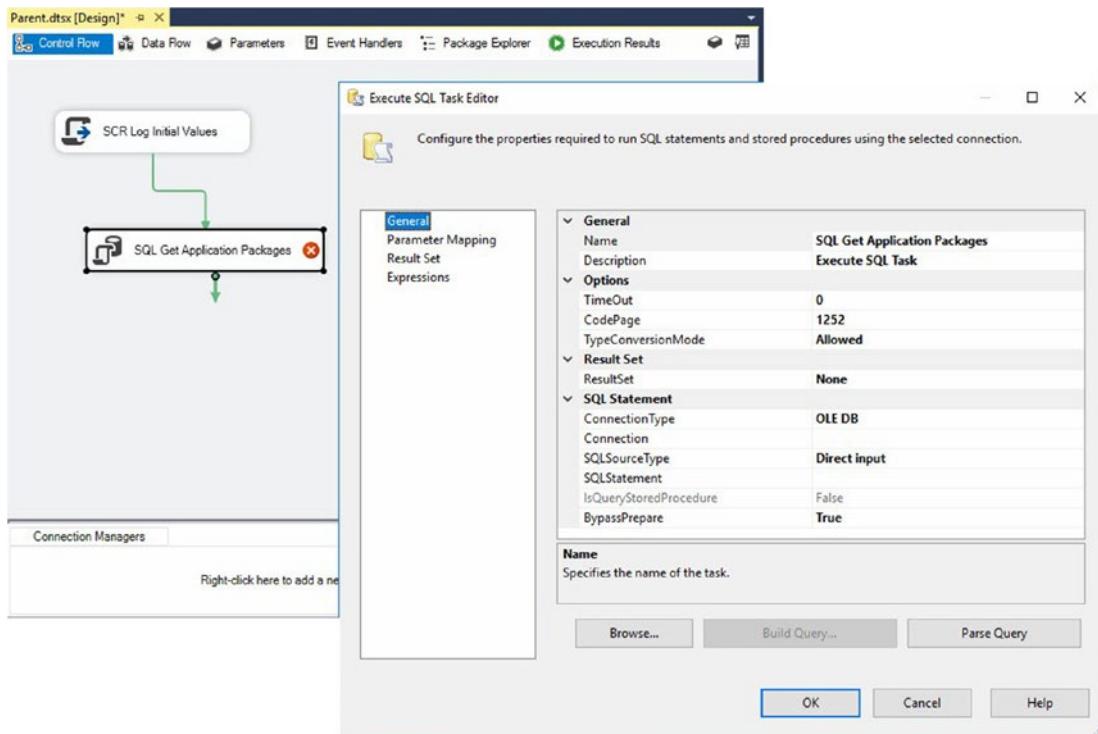


Figure 6-11. Connecting a precedence constraint and opening the SQL Get Application Packages execute SQL task editor

Set the SQL Get Application Packages execute SQL task's ConnectionType property to ADO.NET, as shown in Figure 6-12.

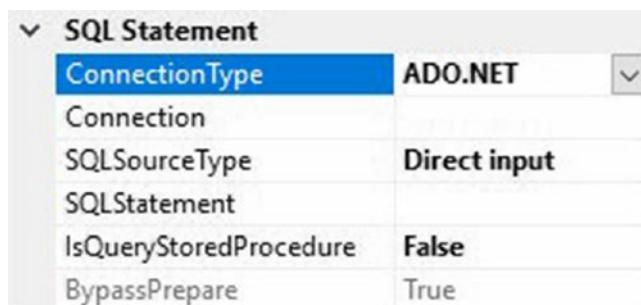


Figure 6-12. Setting the ConnectionType to ADO.NET

Click the drop-down in the Connection property, and click “<New Connection...>” as shown in Figure 6-13.

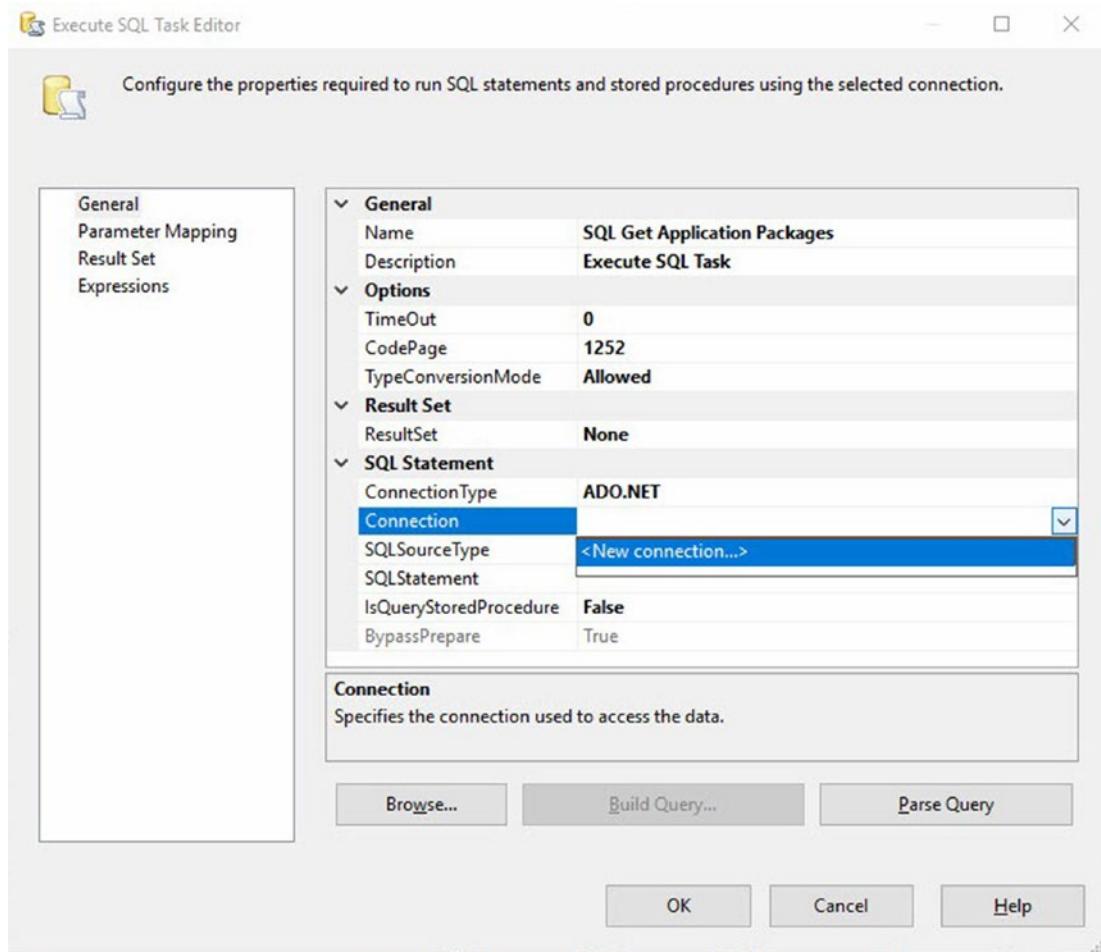


Figure 6-13. Selecting <New Connection...> from the Connection property

The Configure ADO.NET Connection Manager dialog displays, as shown in Figure 6-14.

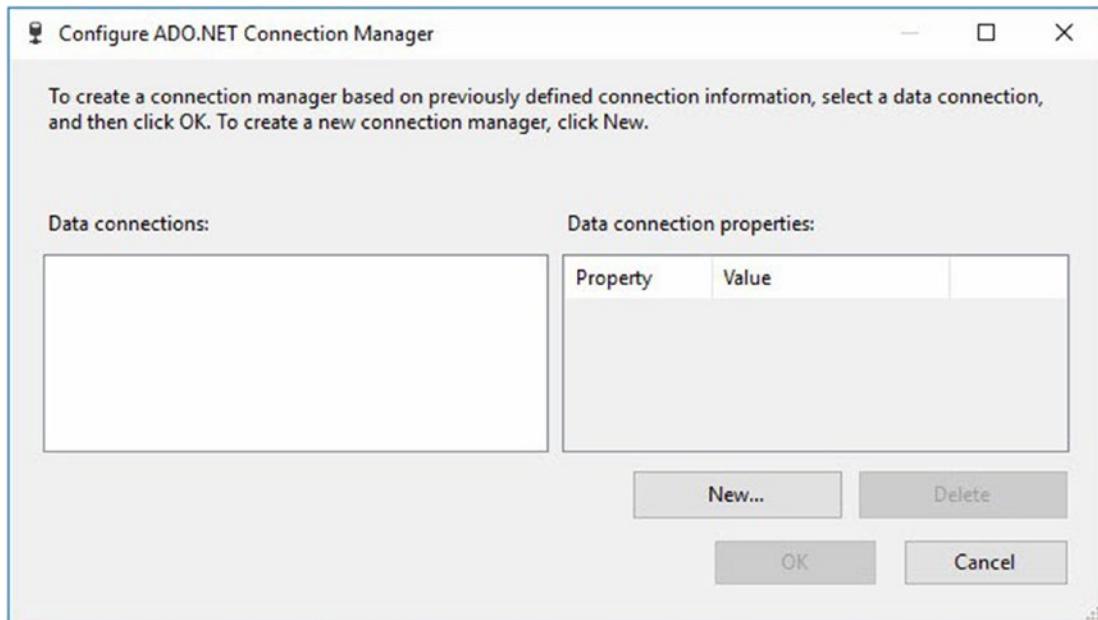


Figure 6-14. The Configure ADO.NET Connection Manager dialog

Click the New button to open the Connection Manager configuration dialog. Configure a new ADO.NET Connection Manager to connect to your instance of the SSISConfig database, as shown in Figure 6-15.

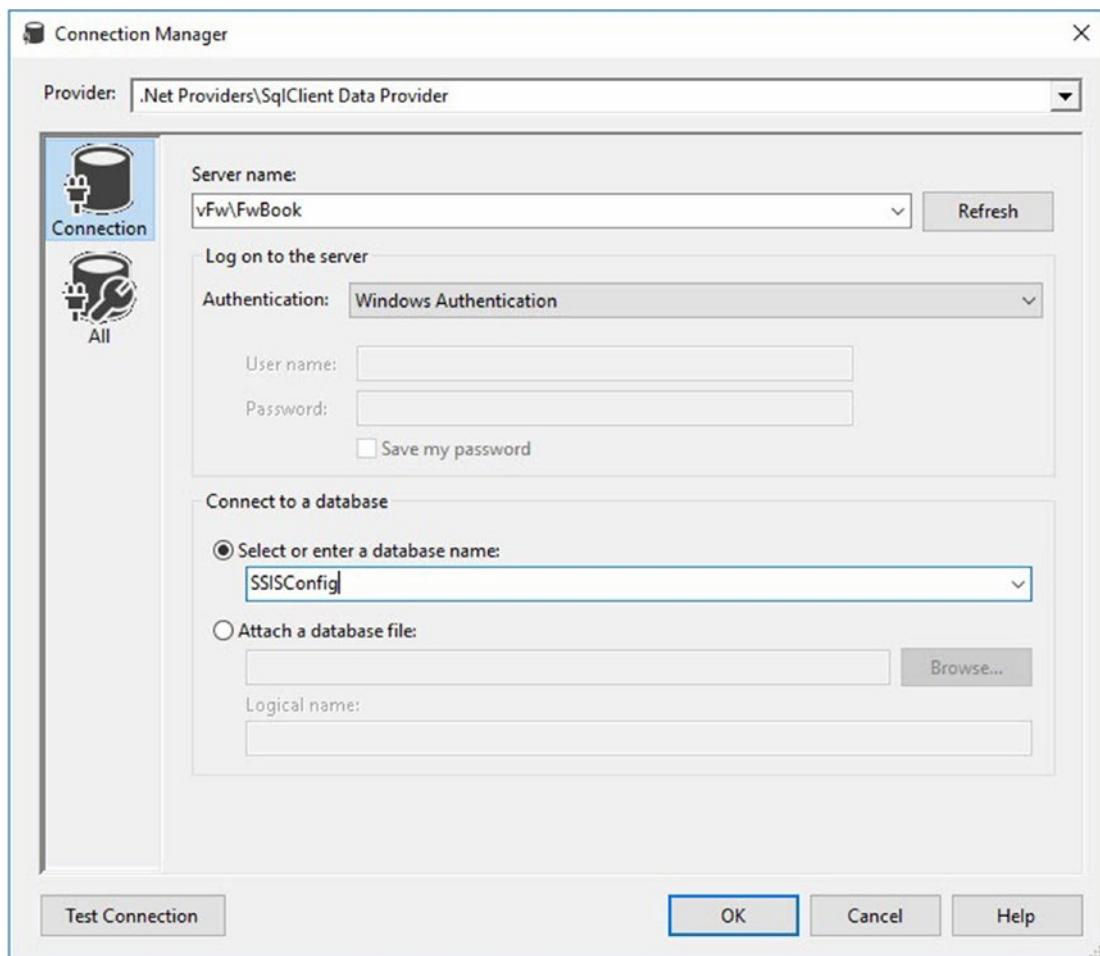


Figure 6-15. Configuring a new ADO.NET Connection Manager

Click the OK button to complete configuration of the ADO.NET Connection Manager. The new connection manager configuration contains connection information for your instance of the SSISConfig database and is now stored in the list of configured ADO.NET connection managers as shown in Figure 6-16.

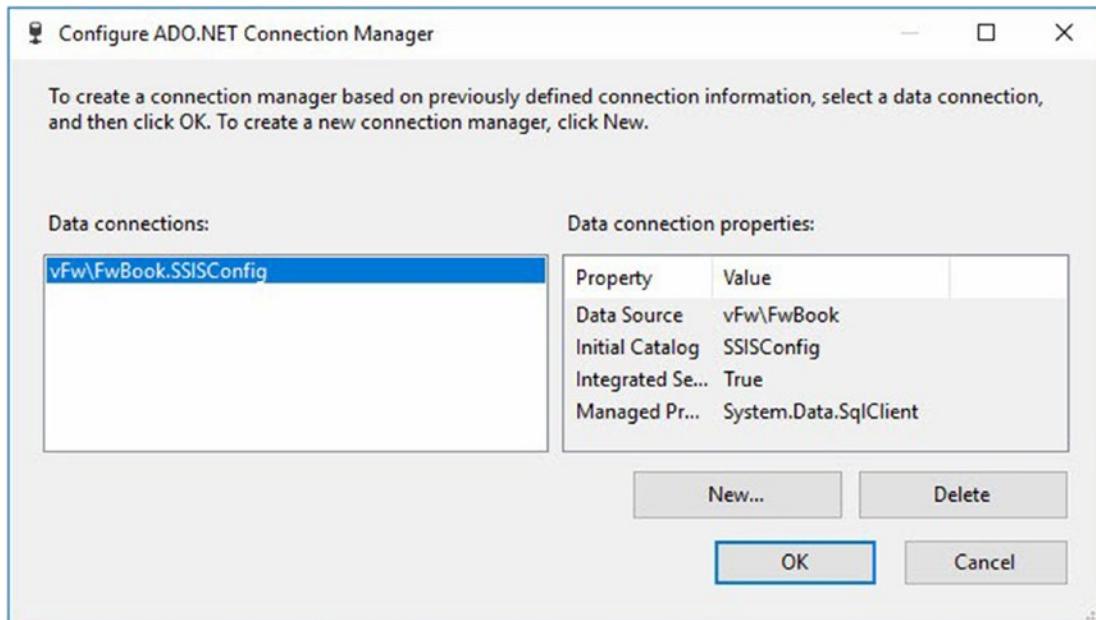


Figure 6-16. Your SSISConfig database configuration in the Configure ADO.NET Connection Manager

Click the OK button to complete the configuration of the ADO.NET Connection Manager. The Execute SQL Task Editor should now appear similar to Figure 6-17.

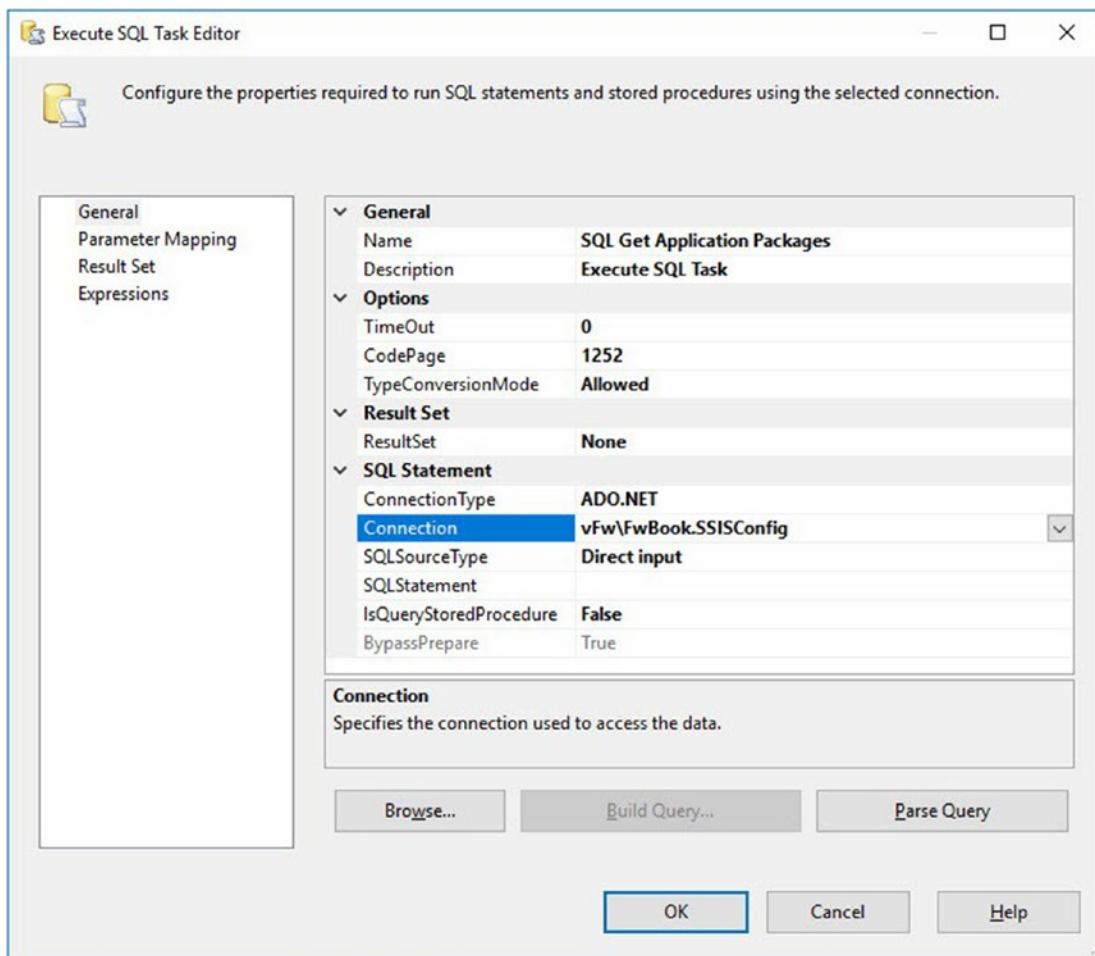


Figure 6-17. Connection configured for the SQL Get Application Packages execute SQL task

Listing 5-14 from the previous chapter contains the T-SQL required to return a list of SSIS application packages from the metadata stored in the SSISConfig database for a given SSIS application. Configure the SQL Get Application Packages execute SQL task's SQLStatement property with the T_SQL shown in Listing 6-3.

Listing 6-3. T-SQL to retrieve application packages for an SSIS application from the SSISConfig database

```
Select a.ApplicationName
    , p.PackageLocation + p.PackageName As PackagePath
    , ap.ExecutionOrder
```

CHAPTER 6 FRAMEWORK EXECUTION ENGINE

```
, ap.FailApplicationOnPackageFailure
From [config].[ApplicationPackages] ap
Join [config].[Applications] a
On a.ApplicationId = ap.ApplicationId
Join [config].Packages p
On p.PackageId = ap.PackageId
Where a.ApplicationName = @ApplicationName
And ap.ApplicationPackageEnabled = 1
Order By ap.ExecutionOrder
```

Add the T-SQL in Listing 6-3 to the SQLStatement property of the SQL Get Application Packages execute SQL task, as shown in Figure 6-18.

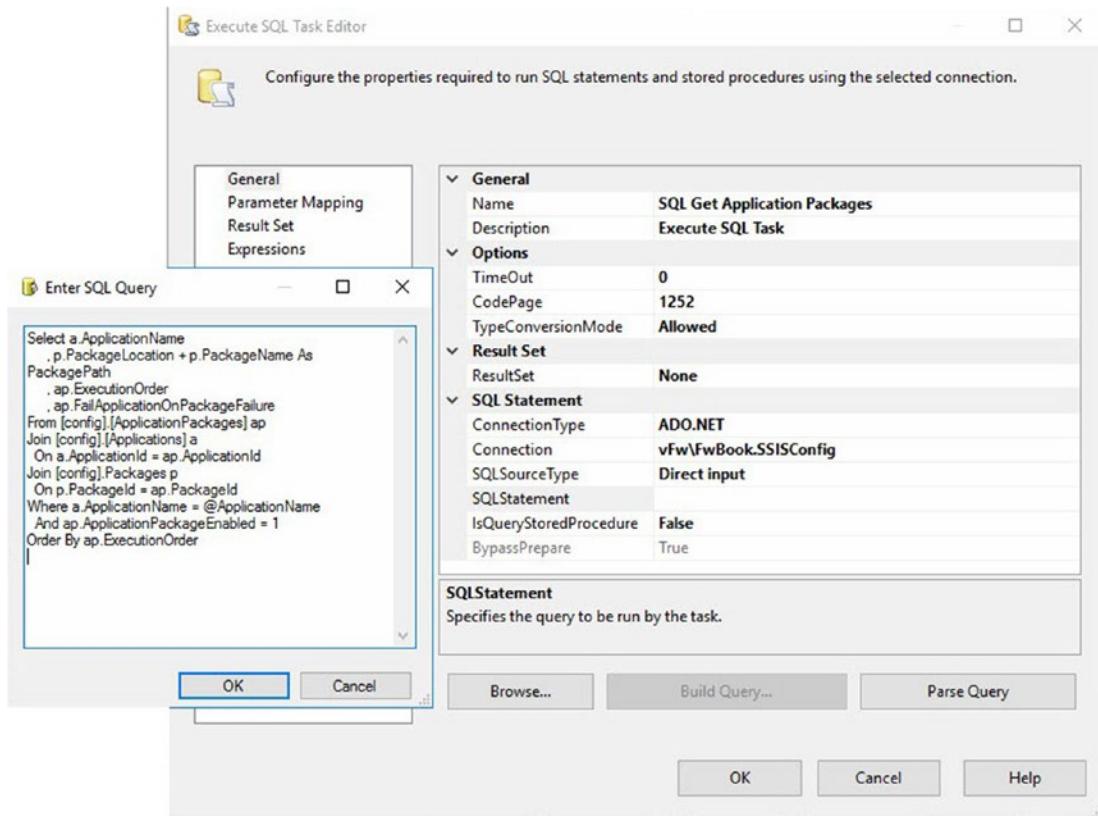


Figure 6-18. Adding the SQLStatement property's T-SQL

Click the OK button to close the “Enter SQL Query” dialog. The Execute SQL Task Editor should appear similar to that shown in Figure 6-19.

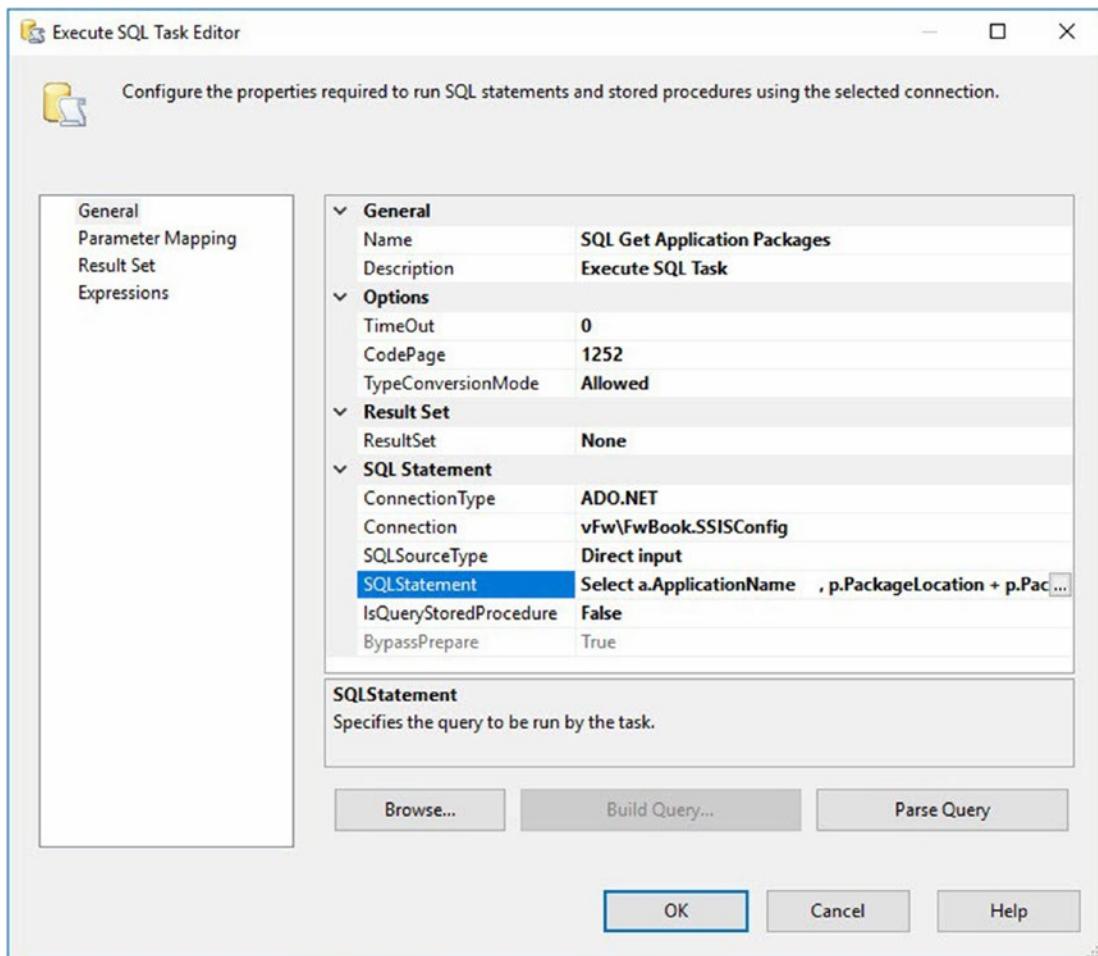


Figure 6-19. SQL Get Application Packages execute SQL task after configuring the SQLStatement property

In the list of Execute SQL Task Editor pages on the left side of the Execute SQL Task Editor, click the “Parameter Mapping” page, and then click the Add button to add a parameter mapping. Select the \$Package::ApplicationName parameter from the Variable Name column, select the String data type, and enter “ApplicationName” for the Parameter Name column as shown in Figure 6-20.

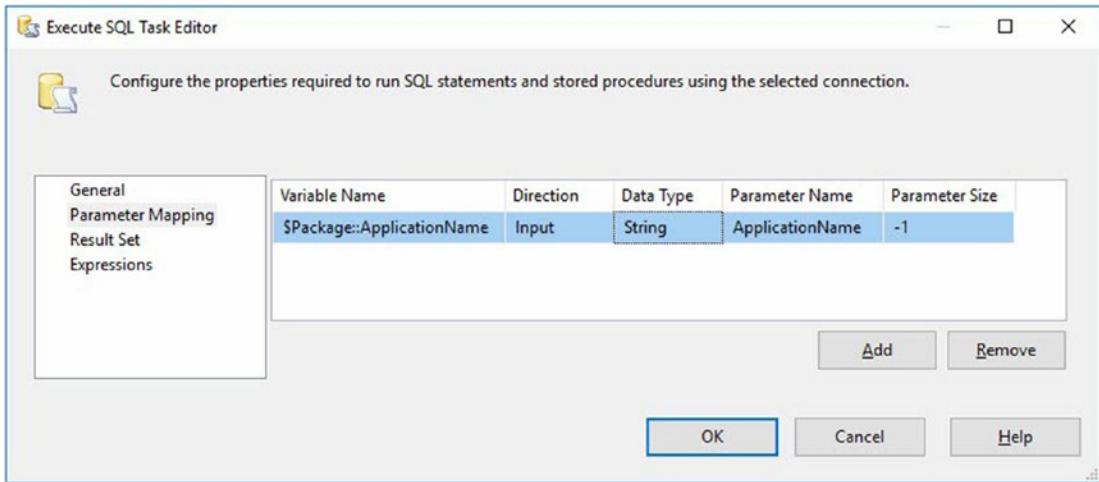


Figure 6-20. Adding the ApplicationName parameter

The ApplicationName parameter configured on the Parameter Mapping page will read the value from the SSIS package parameter named \$Package::ApplicationName (shown in the Variable Name column in Figure 6-20) and supply the value to the T-SQL query parameter configured in the SQL Get Application Packages execute SQL task's SQLStatement property found in Listing 6-17.

Return to the SQL Get Application Packages execute SQL task's General page. Click the drop-down in the ResultSet property, and select "Full result set," as shown in Figure 6-21.

General	
Name	SQL Get Application Packages
Description	Execute SQL Task
Options	
TimeOut	0
CodePage	1252
TypeConversionMode	Allowed
Result Set	
ResultSet	None
SQL Statement	
ConnectionType	None
Connection	Single row
SQLSourceType	Full result set
SQLStatement	XML
IsQueryStoredProcedure	False
BypassPrepare	True

Figure 6-21. Selecting Full result set

Click the Result Set page, and then click the Add button. Edit the Result Name value, and set it to “0”. In the Variable Name drop-down, select “<New variable...>” as shown in Figure 6-22.

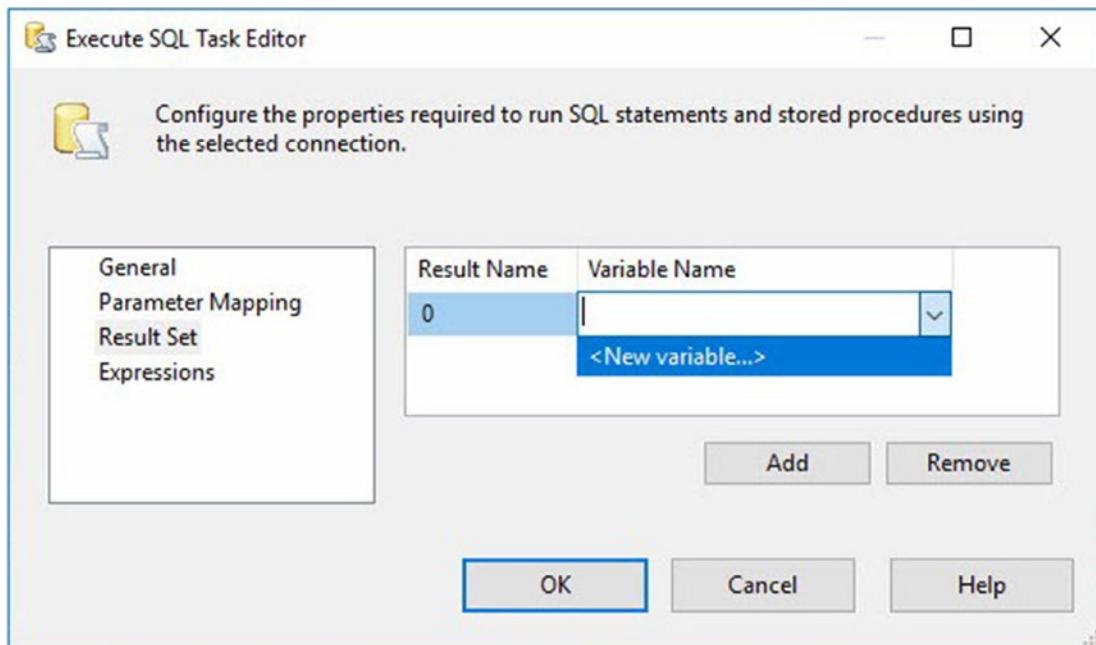


Figure 6-22. Selecting “<New variable...>” from the Result Set Variable Name drop-down

When the Add Variable dialog displays, make sure the Container is set to Parent as shown in Figure 6-23.

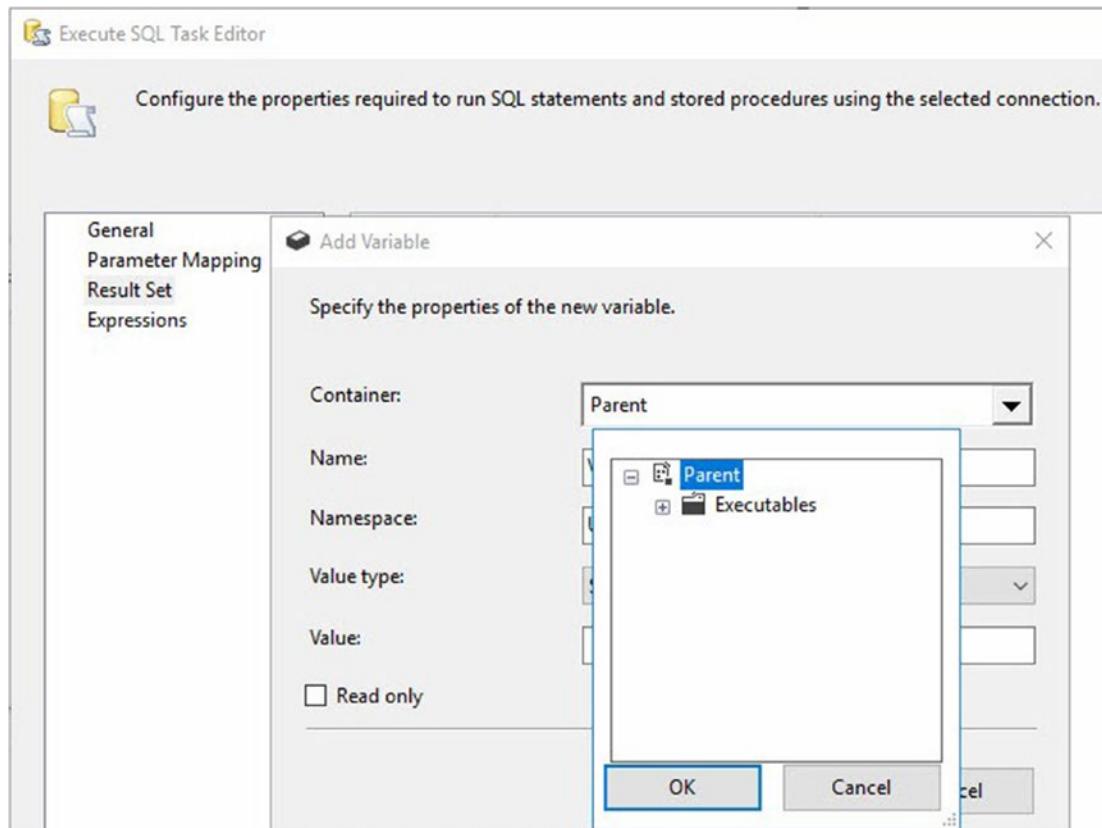


Figure 6-23. Verifying Parent Container is selected

Change the Name property to ApplicationPackages, leave the Namespace set to User, and set the Value type to Object as shown in Figure 6-24.

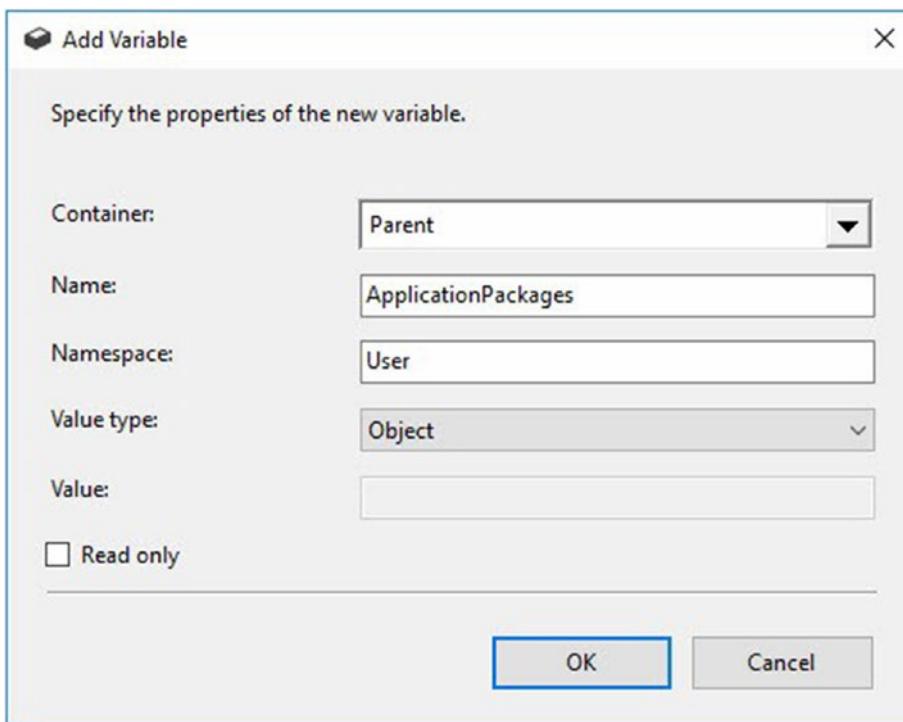


Figure 6-24. Configuring the Object variable ApplicationPackages

The User::ApplicationPackages variable will receive the ADO.Net dataset returned from the execution of the T-SQL query stored in the SQL Get Application Packages execute SQL task's SQLStatement property. Since the T-SQL query includes the ApplicationName parameter in the criteria (WHERE clause) and this parameter is set by the \$Package::ApplicationName SSIS parameter sent to the Parent.dtsx SSIS package, the ADO.Net dataset will contain the list of packages for which SSISConfig data has been configured. Since the \$Package::ApplicationName is defaulted to the SSIS framework application named “Framework Test,” the User::ApplicationPackages variable will contain the results shown in Figure 6-25.

	ApplicationName	PackagePath	ExecutionOrder	FailApplicationOnPackageFailure
1	Framework Test	E:\Projects\TestSSISolution\TestSSISProject\ReportAndSucceed.dtsx	10	1
2	Framework Test	E:\Projects\TestSSISolution\TestSSISProject\ReportAndFail.dtsx	20	1

Figure 6-25. Results sent to User::ApplicationPackages variable

Click the OK button to close the Add Variable dialog.

The General page of the SQL Get Application Packages execute SQL task editor should now appear similar to Figure 6-26.

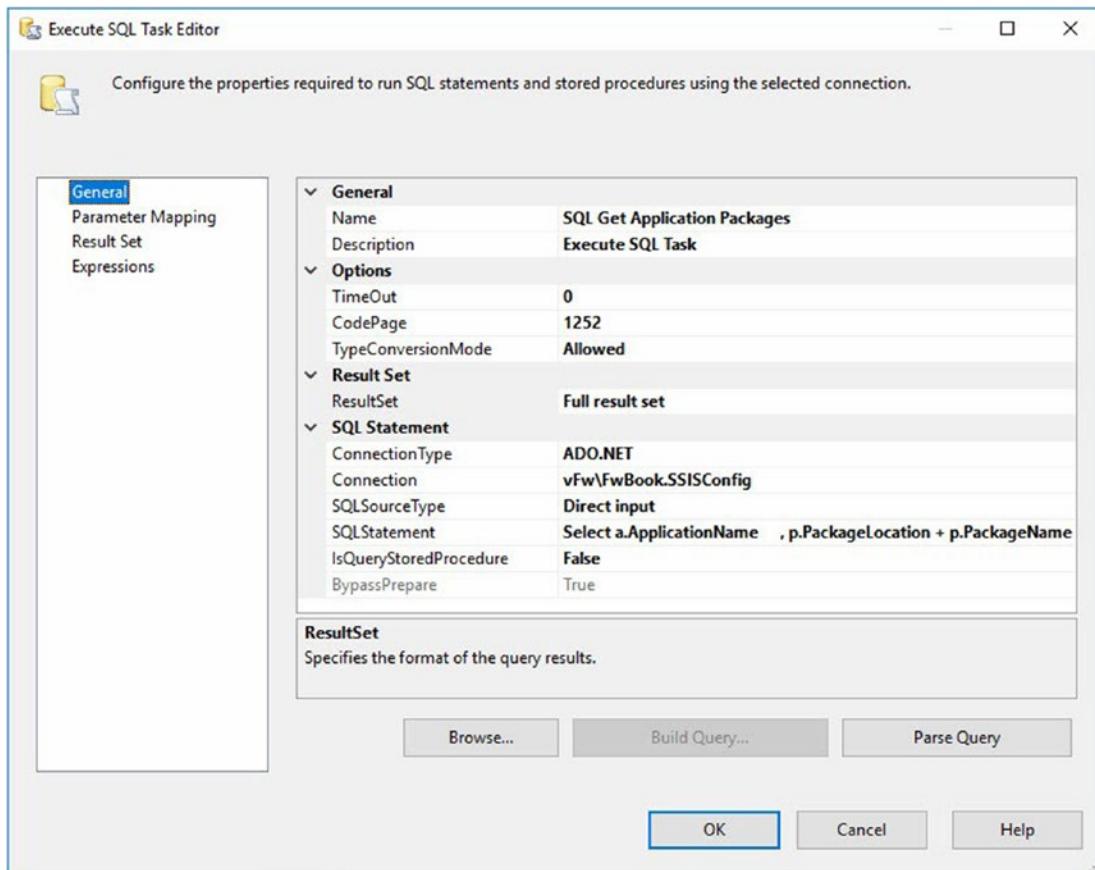


Figure 6-26. The SQL Get Application Packages execute SQL task's General page

Click the OK button to close the SQL Get Application Packages execute SQL task editor.

Iterate Application Packages

SSIS Foreach Loop Containers are used to iterate collections of many sorts. The User::ApplicationPackages SSIS variable is an object data type variable. Objects can contain scalars (or single, individual values), but they're not really built for scalars. Objects are designed to contain collections. Collections can be arrays or lists, as well as recordsets and datasets. Collections can also be, well, collections – a special list-like .Net variable type.

In the previous section, we added an Execute SQL Task. The configuration of the Execute SQL Task ConnectionType property is ADO.NET. We configured a Full result set and sent the results (of the result set) into an Object variable named User::ApplicationPackages. Because the Execute SQL Task ConnectionType property is set to ADO.NET, the type of object in User::ApplicationPackages is an ADO.Net dataset. ADO.Net datasets include a tables collection.

Had the Execute SQL Task been configured to use an OLE DB ConnectionType, the result set sent to the User::ApplicationPackages object variable would be an ADO recordset, which is a COM (Common Object Model) object. COM was introduced in the 32-bit (x86) era of Microsoft software, circa early 1990s.

The cool thing is that the SSIS Foreach Loop Container's Foreach ADO Iterator doesn't care whether you send it an ADO recordset or an ADO.Net dataset; it will iterate either.

Drag a Foreach Loop Container onto the control flow of Parent.dtsx and connect a precedence constraint from the SQL Get Application Package execute SQL task to the Foreach Loop Container. Rename the Foreach Loop Container to “FOREACH Application Package” as shown in Figure 6-27.

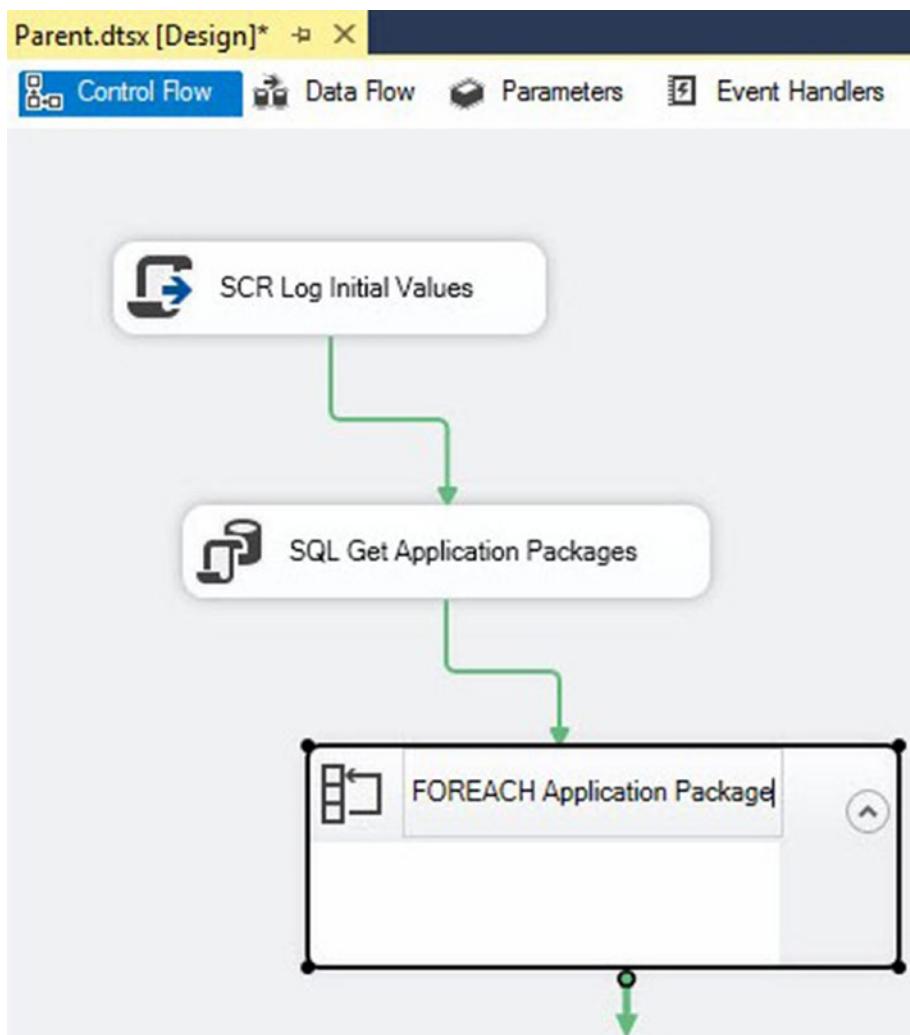


Figure 6-27. Renaming the Foreach Loop Container

Open the FOREACH Application Package's editor and navigate to the Collection page. Change the Enumerator to “Foreach ADO Enumerator” as shown in Figure 6-28.

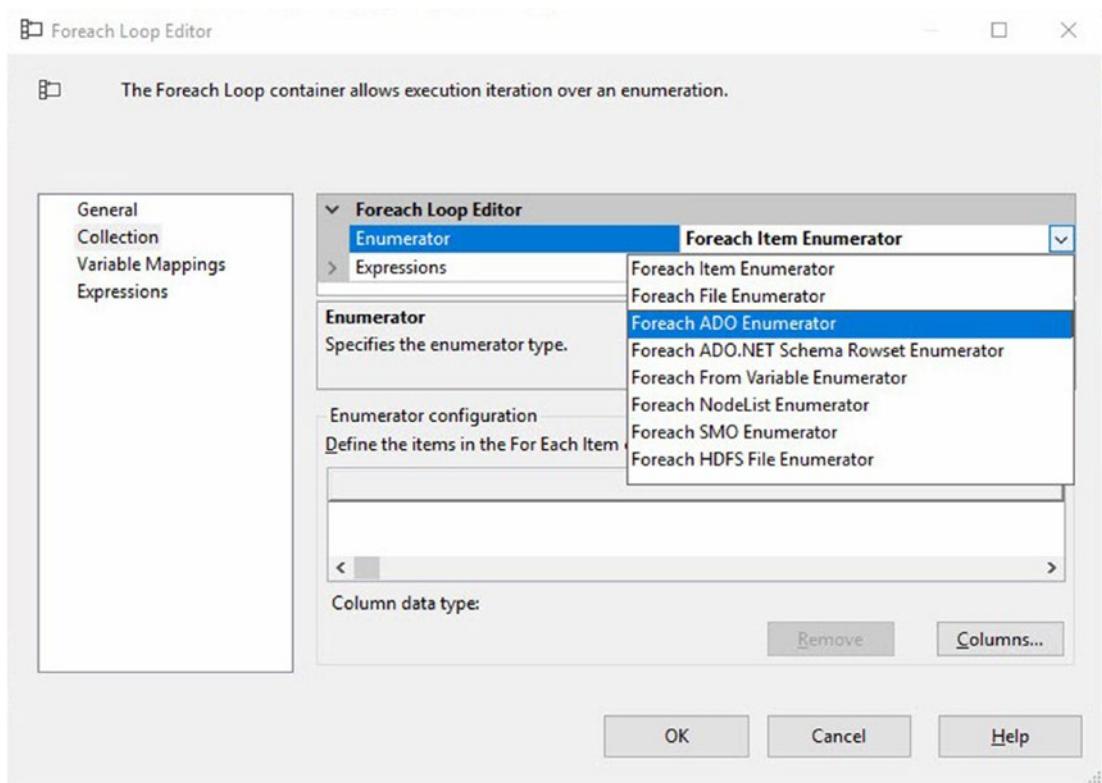


Figure 6-28. Changing the FOREACH Application Package Enumerator to Foreach ADO Enumerator

A Foreach Loop Container's Foreach ADO Enumerator "points" to one row each enumeration (iteration through the dataset). The values in each column may be mapped into SSIS variables by ordinal with each "pass" (enumeration/iteration). Most often, SSIS tasks inside the Foreach Loop Container are configured to *read* these values from the variables and act upon them in some way. More on Variable Mappings in a bit.

Each enumerator presents an enumerator-specific properties view. In the Enumerator configuration group box, click the "ADO object source variable" drop-down, and select the User::ApplicationPackages variable as shown in Figure 6-29.

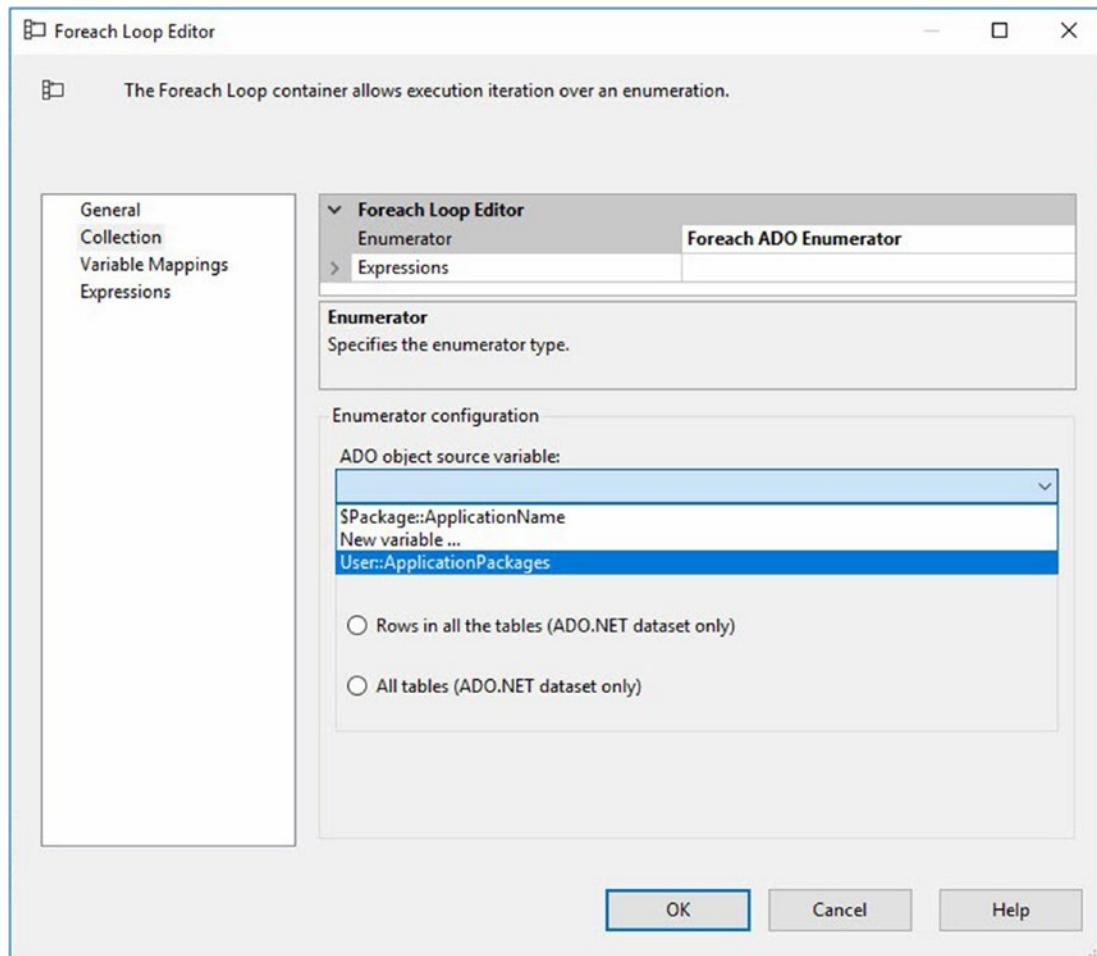


Figure 6-29. Selecting the ADO object source variable

Leave the Enumeration mode set to “Rows in the first table.” The Collection page is configured as shown in Figure 6-30.

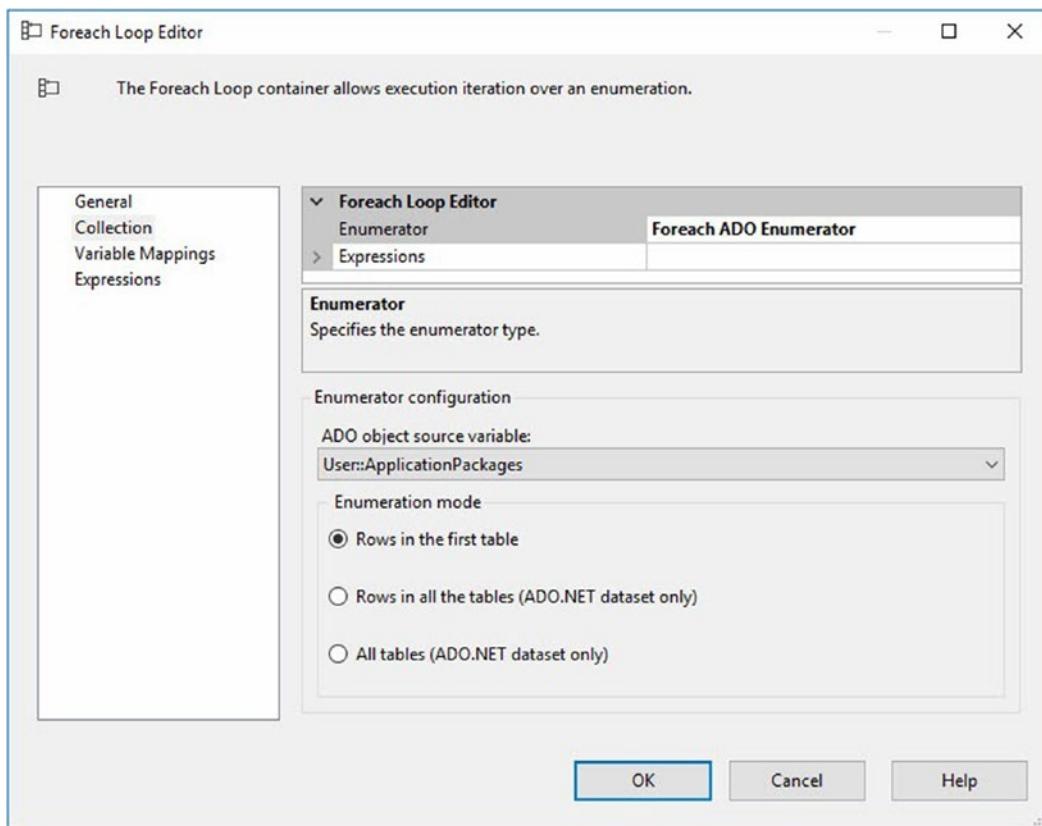


Figure 6-30. Collection page configured

Click the Variable Mappings page. Click inside the Variable column, click the drop-down, and then click “<New variable...>” as shown in Figure 6-31.

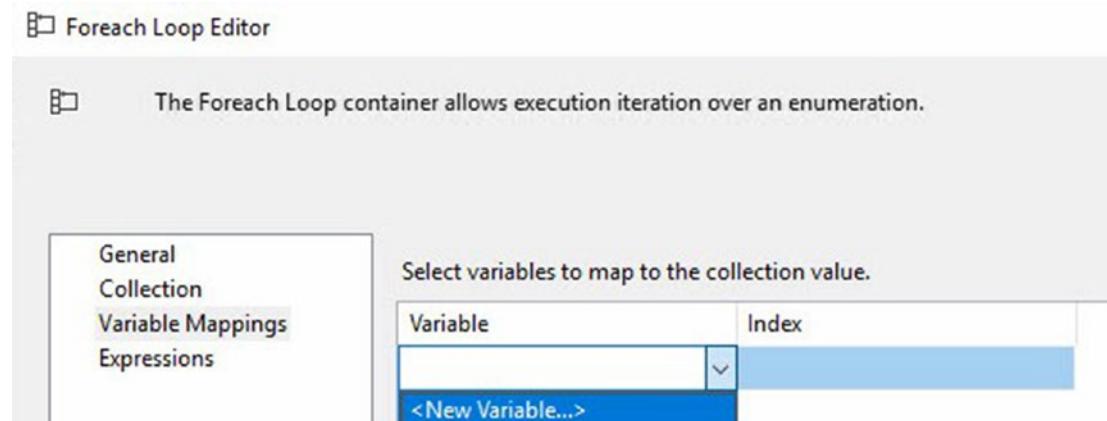


Figure 6-31. Clicking New Variable on the Variable Mappings page

As before, make sure the Container is set to the SSIS package scope (“Parent”) when the Add Variable dialog displays. Enter “PackagePath” as the variable name, and leave all other settings at their respective defaults, as shown in Figure 6-32.

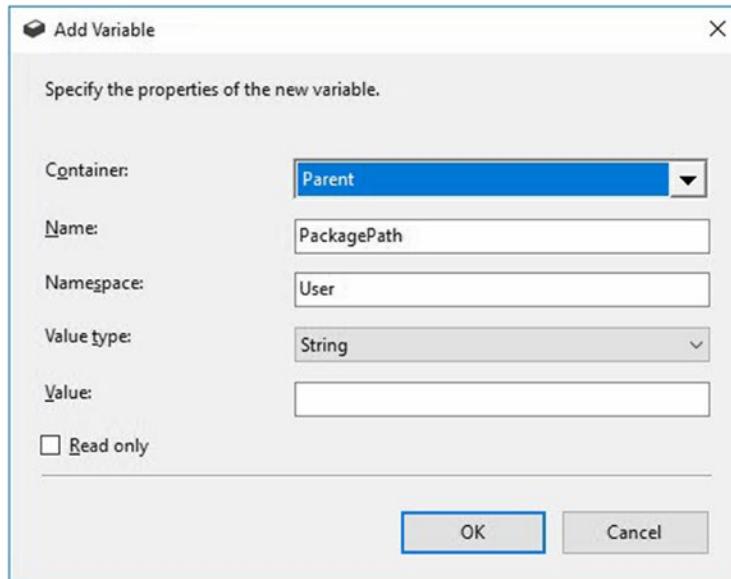


Figure 6-32. Adding the PackagePath variable

Click the OK button to close the Add Variable dialog. Note the Index column defaults to “0” for this first variable assignment as shown in Figure 6-33.

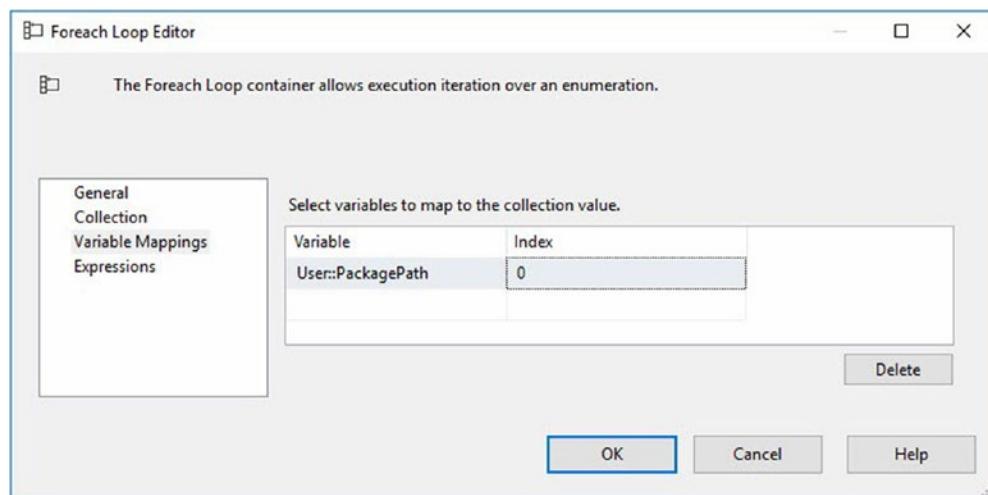


Figure 6-33. User::PackagePath assigned to ordinal “0”

What does “0” mean? I’m glad you asked. “0” indicates the ordinal of the column in the User::ApplicationPackages dataset. Please recall the contents of Tables(0) in the ADO.Net dataset as shown in Figure 6-34.

	ApplicationName	PackagePath	ExecutionOrder	FailApplicationOnPackageFailure
1	Framework Test	E:\Projects\TestSSISolution\TestSSISProject\ReportAndSucceed.dtsx	10	1
2	Framework Test	E:\Projects\TestSSISolution\TestSSISProject\ReportAndFail.dtsx	20	1

Figure 6-34. Column Ordinal “0” in the ADO.Net dataset in User::ApplicationPackages

Our package already *has* the value in this field – ApplicationName is passed to the Parent.dtsx SSIS package in the parameter named \$Package::ApplicationName. The SSIS variable we just created is named “User::PackagePath”. The value we want to assign here is in the second field, ordinal 1 (of this 0-based array).

Change the “0” to “1” in the Index column, as shown in Figure 6-35.

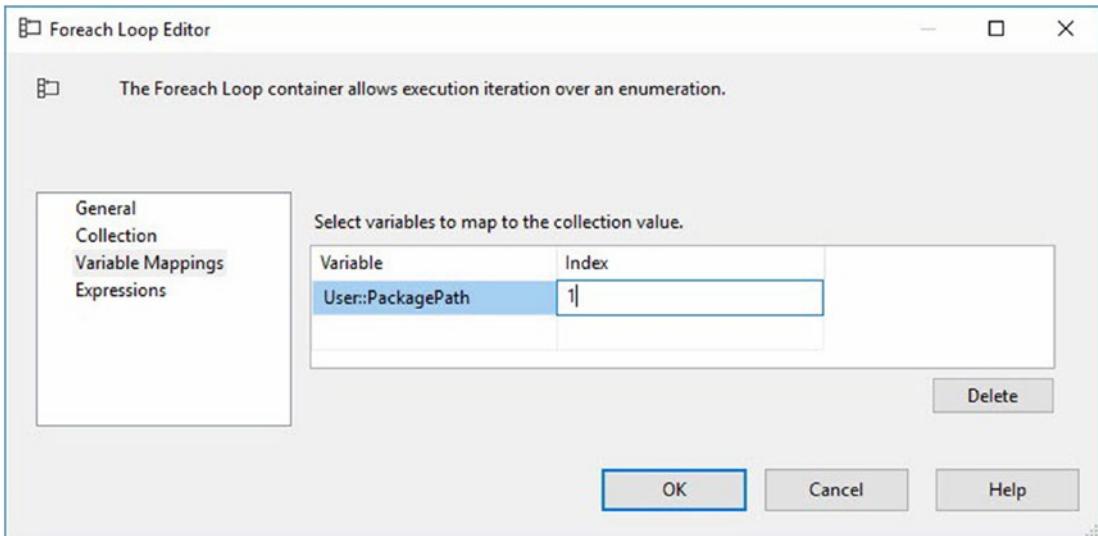


Figure 6-35. Changing the ordinal to “1”

The User::PackagePath will now be assigned the value of the second column – 1 in zero-based ordinals – in the dataset.

Create a new SSIS Int32 variable named User::ExecutionOrder to map the value of the ExecutionOrder from the dataset in the User::ApplicationPackages variable using the following settings on the Add Variable dialog, as shown in Figure 6-36.

- Container: Parent
- Name: ExecutionOrder
- Namespace: User
- Value type: Int32
- Value: -1

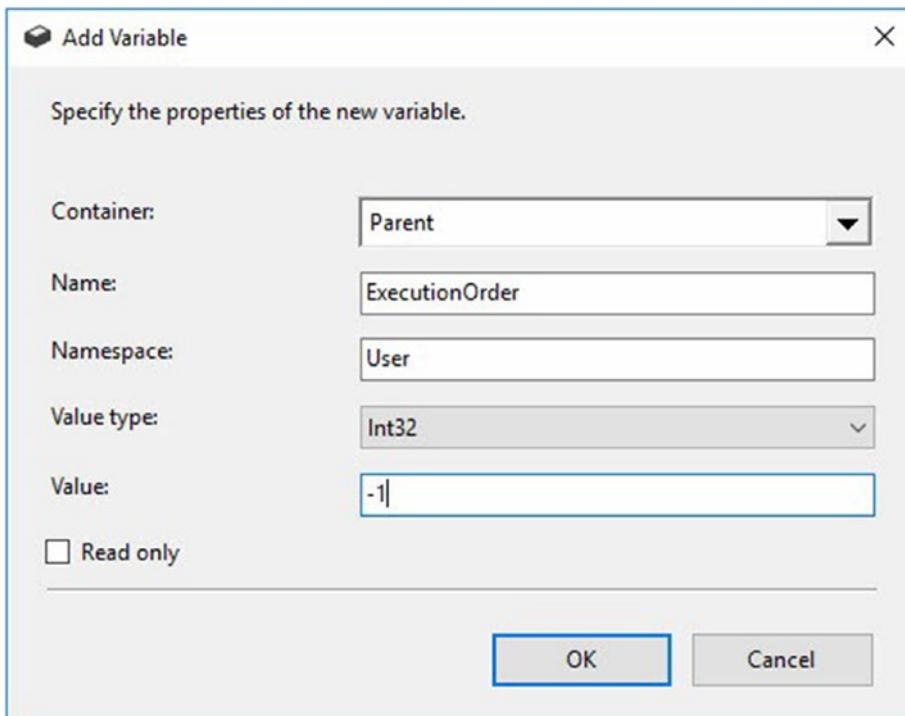


Figure 6-36. Mapping the ExecutionOrder SSIS variable

Update the Index column to 2 for the ExecutionOrder column. The FOREACH Application Package Foreach Loop Container's Variable Mappings page should appear similar to Figure 6-37.

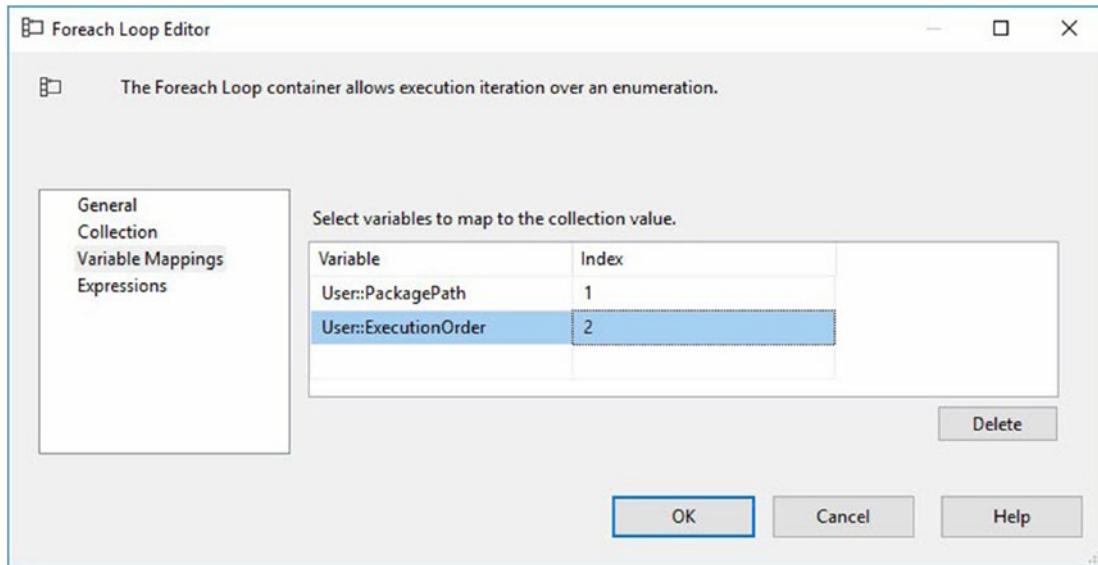


Figure 6-37. User::PackagePath and User::ExecutionOrder variables mapped

Create a new SSIS Boolean variable named User::FailApplicationOnPackageFailure to map the value of the FailApplicationOnPackageFailure from the dataset in the User::ApplicationPackages variable using the following settings on the Add Variable dialog, as shown in Figure 6-38.

- Container: Parent
- Name: FailApplicationOnPackageFailure
- Namespace: User
- Value type: Boolean
- Value: true

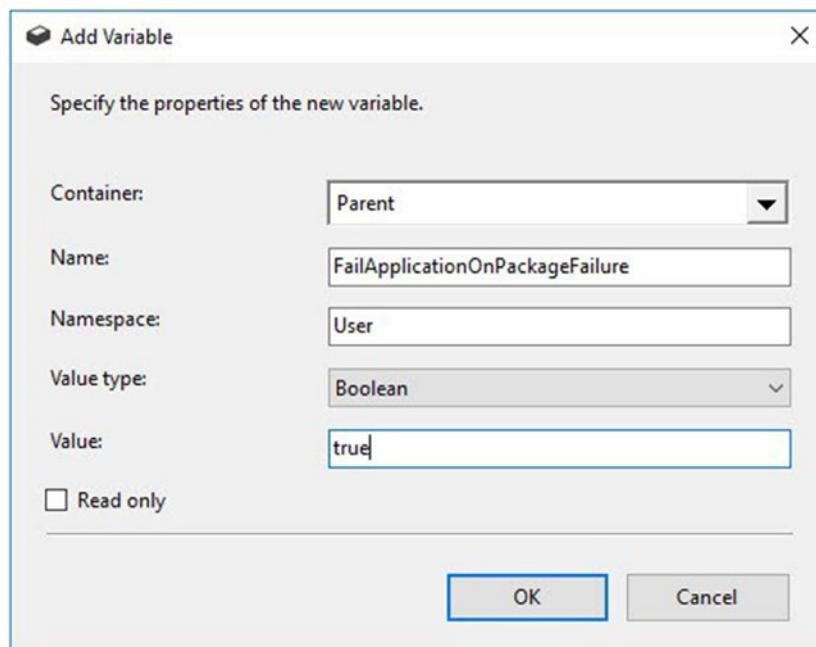


Figure 6-38. Mapping the FailApplicationOnPackageFailure SSIS variable

Update the Index column to 3 for the FailApplicationOnPackageFailure column. The FOREACH Application Package Foreach Loop Container's Variable Mappings page should appear similar to Figure 6-39.

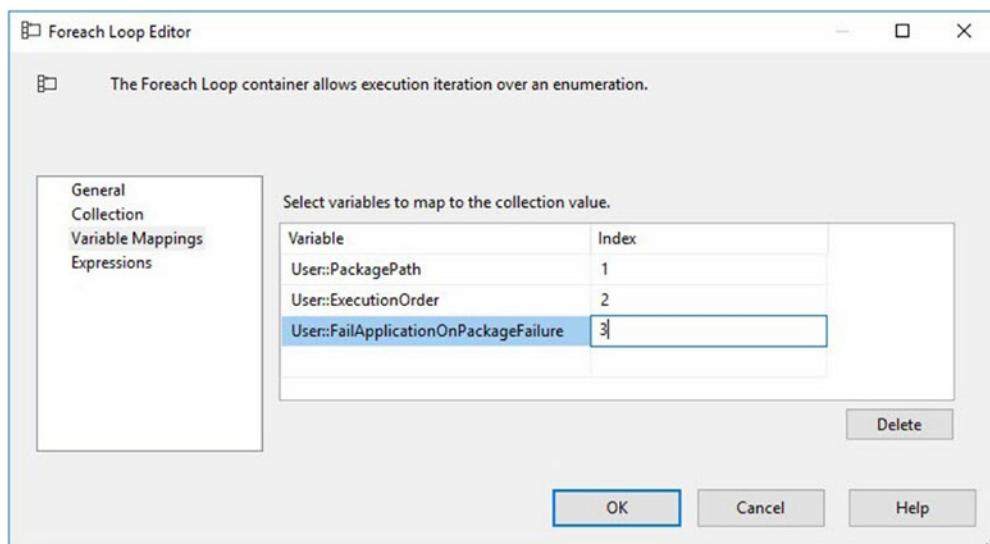


Figure 6-39. User:: FailApplicationOnPackageFailure variable mapped

Click the OK button to close the Foreach Loop Editor. Parent.dtsx should appear similar to Figure 6-40.

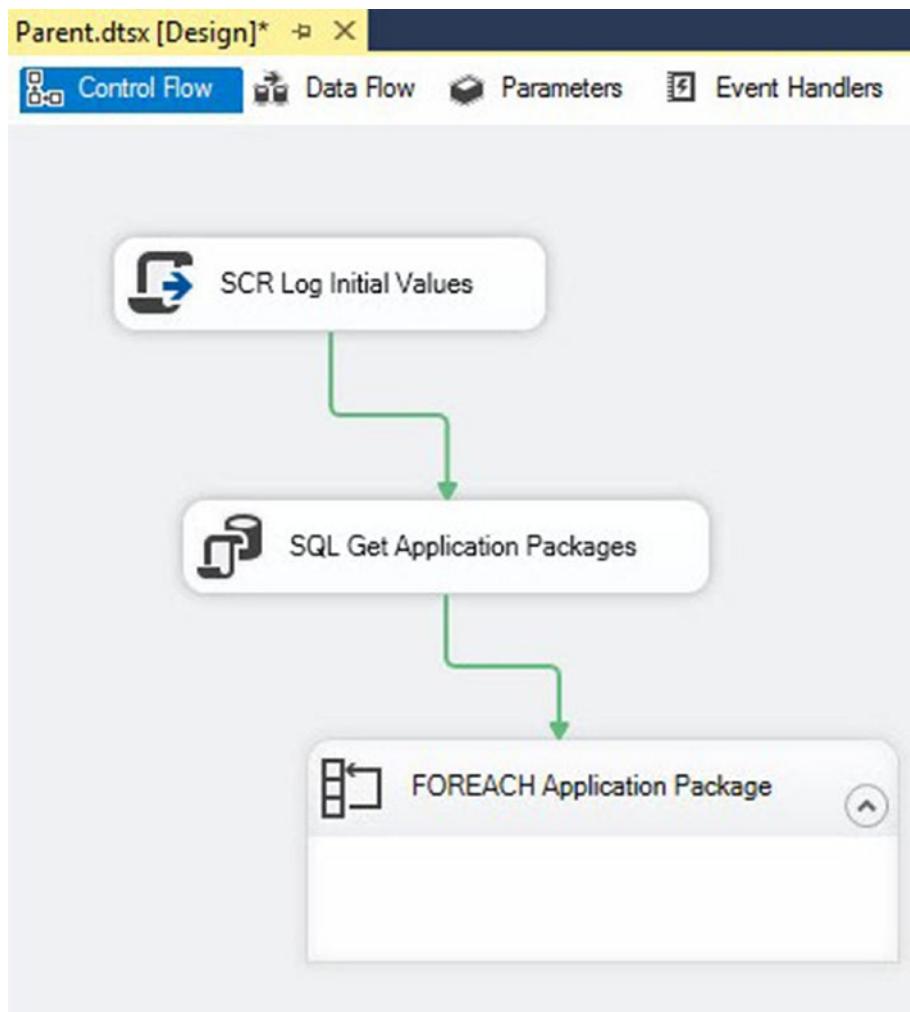


Figure 6-40. Parent.dtsx control flow after adding FOREACH Application Package

Log Application Package Values

As with logging initial values when the package starts, logging all package data before the package executes is helpful for troubleshooting purposes. Add a new Script Task to the Parent.dtsx package. Place the new Script Task inside the FOREACH Application Package Foreach Loop Container and rename the Script Task to “SCR Log Application Package Values” as shown in Figure 6-41.

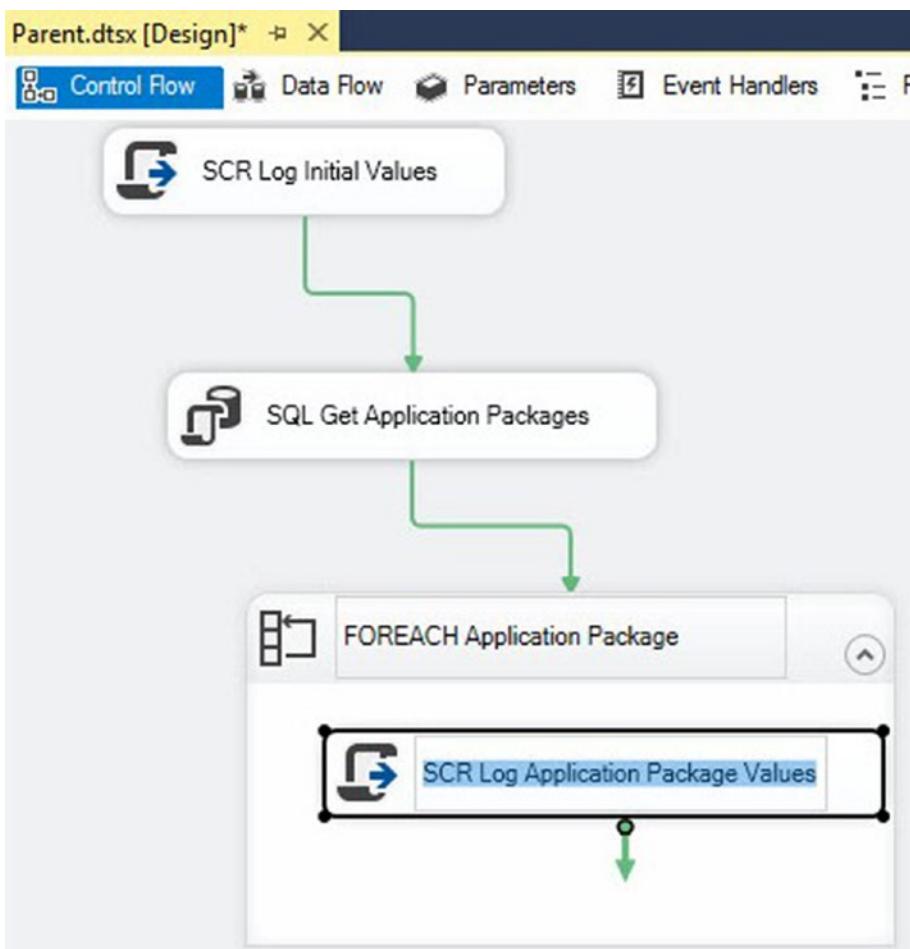


Figure 6-41. Adding the SCR Log Application Package Values Script Task

Open the Script Task Editor and add the following SSIS variables to the `ReadOnlyVariables` property:

- `System::PackageName`
- `System::TaskName`
- `User::PackagePath`
- `User::ExecutionOrder`
- `User::FailApplicationOnPackageFailure`

The Script Task's Script page should appear as shown in Figure 6-42.

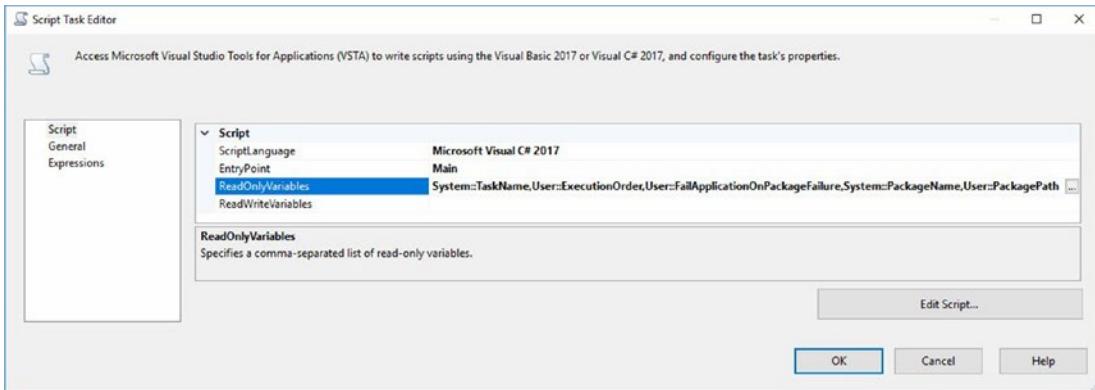


Figure 6-42. Script Task *ReadOnlyVariables* configured

Open the .Net code editor, and place the C# .Net code shown in Listing 6-4 into the `public void Main()` method.

Listing 6-4. Logging Application Package Values

```
public void Main()
{
    // System::PackageName, System::TaskName
    // User::PackagePath, User::ExecutionOrder, ► User::FailApplicationOnP
        ackageFailure

    string packageName = ► Dts.Variables["System::PackageName"].Value.ToString();
    string taskName = ► Dts.Variables["System::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int informationCode = 1001;
    bool fireAgain = true;

    string packagePath = ► Dts.Variables["User::PackagePath"].Value.ToString();
    string description = "PackagePath: " + packagePath;
    Dts.Events.FireInformation(informationCode, subComponent, ►
        description, "", 0, ref fireAgain);
}
```

```

int executionOrder = ► Convert.ToInt32(Dts.Variables["User::Ex
ecutionOrder"].Value);
description = "ExecutionOrder: " + executionOrder.ToString();
Dts.Events.FireInformation(informationCode, subComponent, ►
description, "", 0, ref fireAgain);

bool failApplicationOnPackageFailure = ► Convert.
.ToBoolean(Dts.Variables["User::FailApplicationOnPackageFailu
re"]►
.Value);

description = "FailApplicationOnPackageFailure: " + ►
failApplicationOnPackageFailure.ToString();
Dts.Events.FireInformation(informationCode, subComponent, ►
description, "", 0, ref fireAgain);

Dts.TaskResult = (int)ScriptResults.Success;
}

```

When the C# .Net code is added to the VstaProjects editor for the SCR Log Application Package Values script task, it should appear as shown in Figure 6-43.

```

public void Main()
{
    // System::PackageName, System::TaskName
    // User::PackagePath, User::ExecutionOrder, User::FailApplicationOnPackageFailure

    string packageName = Dts.Variables["System::PackageName"].Value.ToString();
    string taskName = Dts.Variables["System::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int informationCode = 1001;
    bool fireAgain = true;

    string packagePath = Dts.Variables["User::PackagePath"].Value.ToString();
    string description = "PackagePath: " + packagePath;
    Dts.Events.FireInformation(informationCode, subComponent, description, "", 0, ref fireAgain);

    int executionOrder = Convert.ToInt32(Dts.Variables["User::ExecutionOrder"].Value);
    description = "ExecutionOrder: " + executionOrder.ToString();
    Dts.Events.FireInformation(informationCode, subComponent, description, "", 0, ref fireAgain);

    bool failApplicationOnPackageFailure = Convert.ToBoolean(Dts.Variables["User::FailApplicationOnPackageFailure"].Value);
    description = "FailApplicationOnPackageFailure: " + failApplicationOnPackageFailure.ToString();
    Dts.Events.FireInformation(informationCode, subComponent, description, "", 0, ref fireAgain);

    Dts.TaskResult = (int)ScriptResults.Success;
}

```

Figure 6-43. SCR Log Application Package Values C# .Net code

Close the VstaProjects window and click the OK button to close the Script Task Editor. The control flow for Parent.dtsx should appear as shown in Figure 6-44.

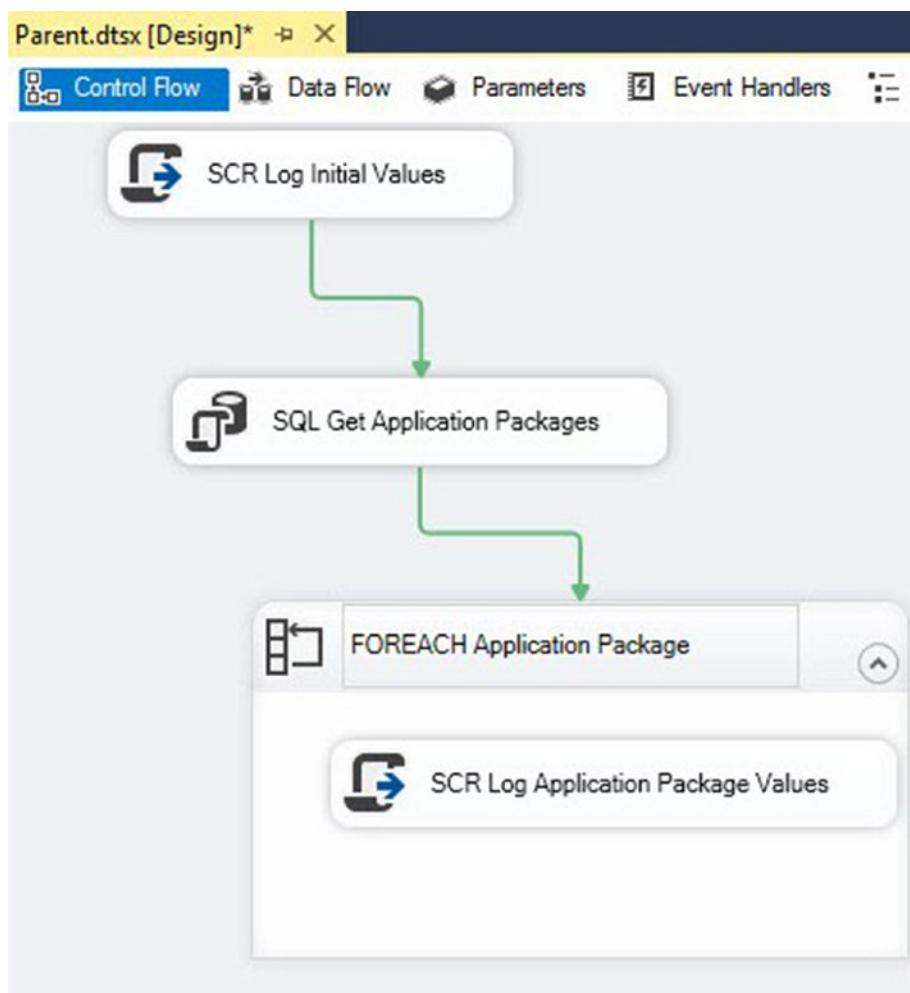


Figure 6-44. Parent.dtsx control flow after adding SCR Log Application Package Values

Test the new functionality by executing the Parent.dtsx package in the debugger. Press the F5 key, and then observe the Progress tab once debug execution completes, as shown in Figure 6-45.

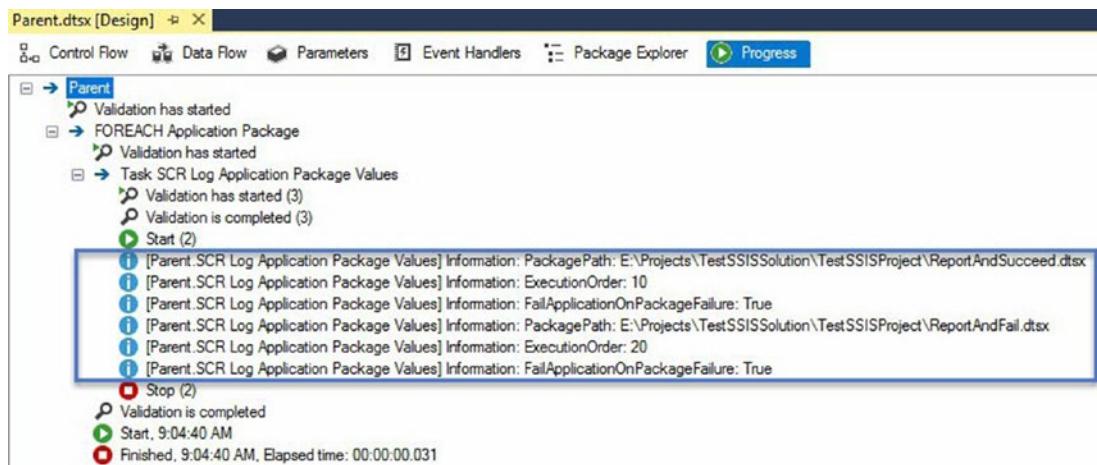


Figure 6-45. Viewing instrumentation messages raised by SCR Log Application Package Values

Instrumentation is nearly complete for the Parent.dtsx package, the SSIS framework execution engine. The next step is to execute the child packages for which metadata has been returned from the SSIS framework database.

Execute Application Packages

There are a number of mechanisms available for executing SSIS packages, including

- .Net (C# and Visual Basic)
- PowerShell
- DtExec command-line
- SQL Agent
- SSIS Execute Package Task

In this example, we will use the SSIS Execute Package Task. Drag an Execute Package Task into the FOREACH Application Package Foreach Loop Container. Connect a precedence constraint to the Execute Package Task from SCR Log Application Package Values, and then rename the Execute Package Task to “EPT Execute Child Package,” as shown in Figure 6-46.

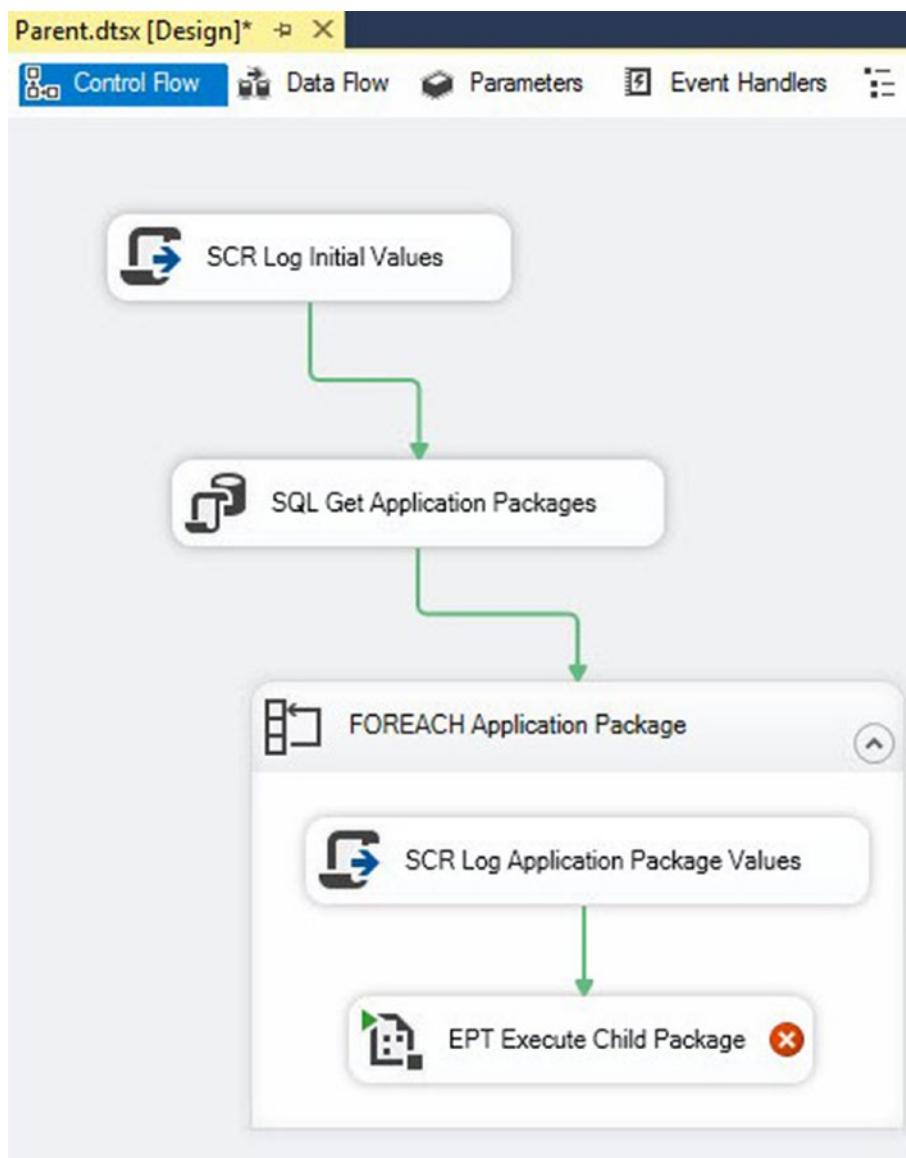


Figure 6-46. Adding the EPT Execute Child Package Execute Package Task

Open the Execute Package Task Editor and navigate to the Package page. Change the ReferenceType property to External Reference, as shown in Figure 6-47.

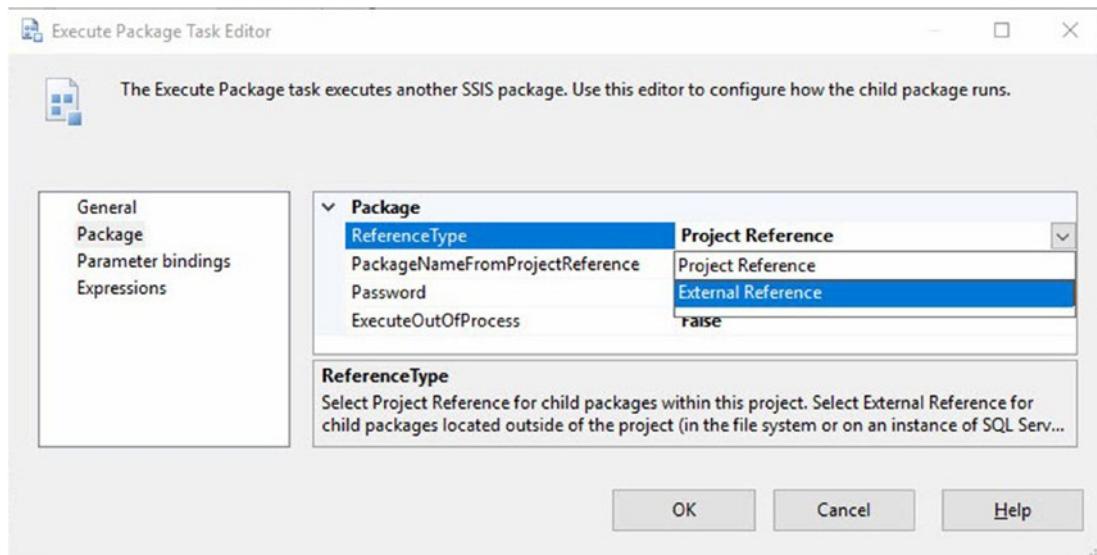


Figure 6-47. Changing the EPT Execute Child Package execute package task's ReferenceType property to External Reference

Change the EPT Execute Child Package execute package task's Location property to "File system," as shown in Figure 6-48.

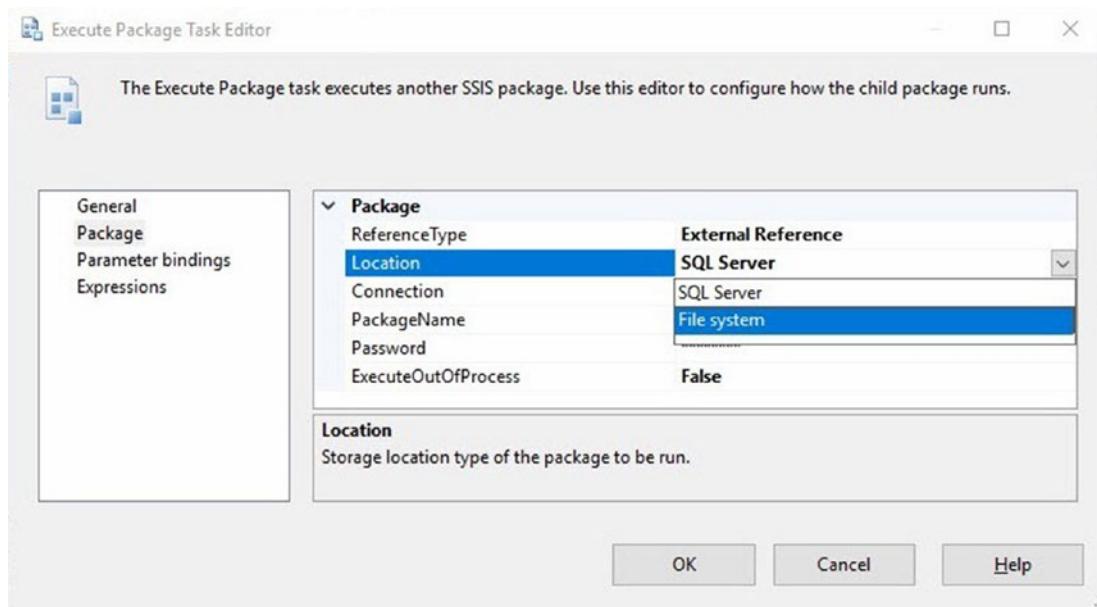


Figure 6-48. Setting the EPT Execute Child Package execute package task's Location property

CHAPTER 6 FRAMEWORK EXECUTION ENGINE

Once the EPT Execute Child Package execute package task's Location property is changed to File system, a new property – Connection – is surfaced in the editor. Begin configuring a new connection by clicking the drop-down and selecting “<New connection...>” as shown in Figure 6-49.

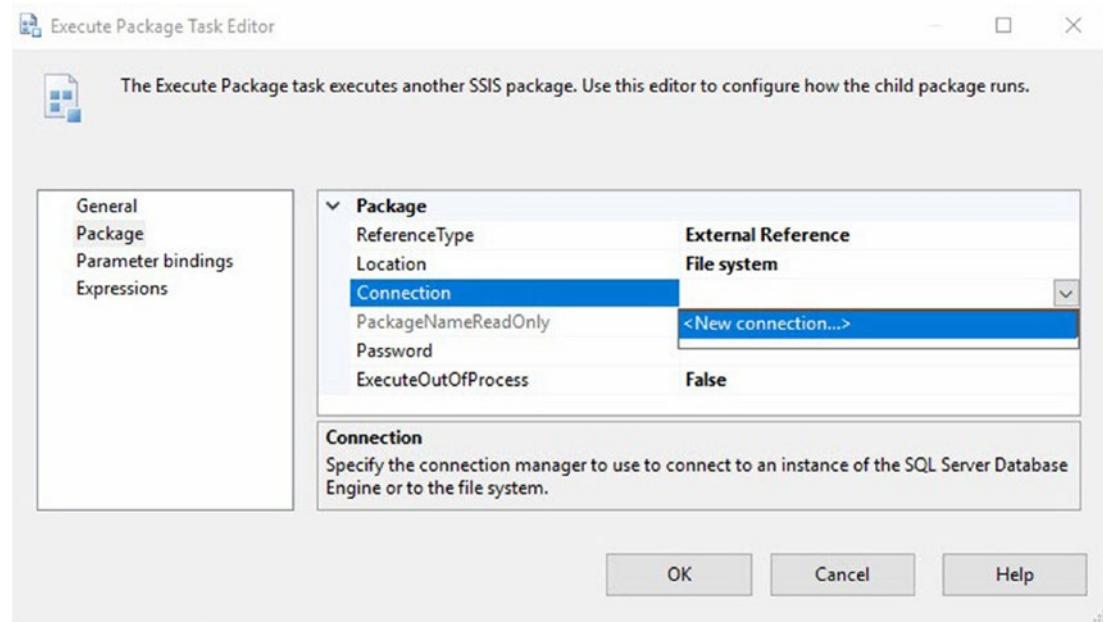


Figure 6-49. Select “<New connection...>” to create a new connection

Because the ReferenceType property is set to External Reference and the Location property is set to File system, the EPT Execute Child Package execute package task's Connection property expects to connect to a file – an *SSIS Package file* or dtsx file – so the New Connection operation creates a new File Connection Manager and opens its editor, as shown in Figure 6-50.

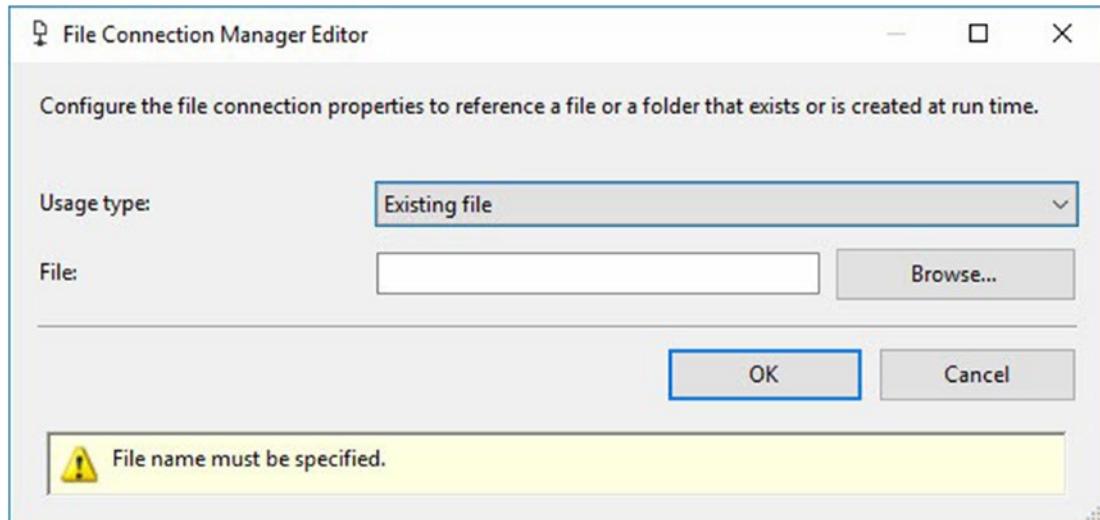


Figure 6-50. A File Connection Manager ready for configuration

Leave the “Usage type” file connection manager property set to “Existing file.” Click the Browse button and navigate to the ReportAndSucceed.dtsx file in the TestSSISSolution\TestSSISProject folder, as shown in Figure 6-51.

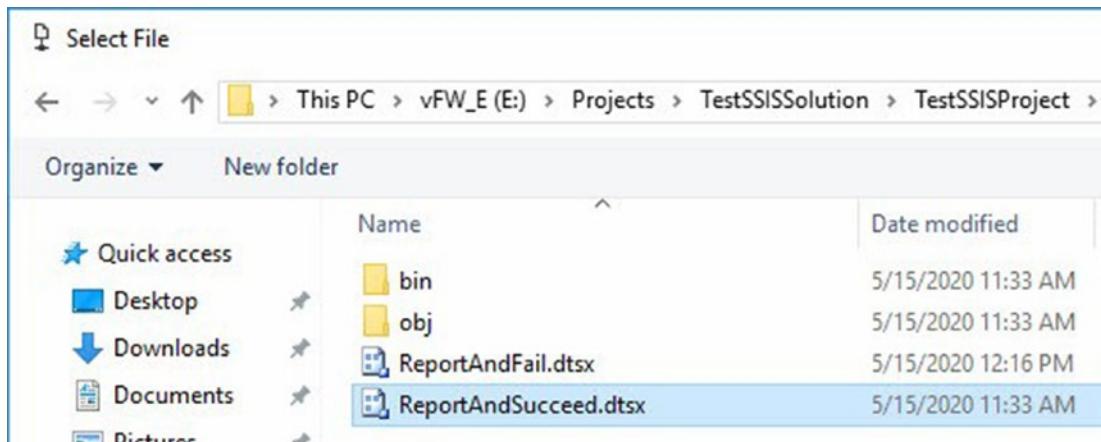


Figure 6-51. Browsing to the ReportAndSucceed.dtsx file in the TestSSISSolution\TestSSISProject folder

Once the TestSSISSolution\TestSSISProject\ReportAndSucceed.dtsx file is selected, the File Connection Manager Editor will appear similar to Figure 6-52.

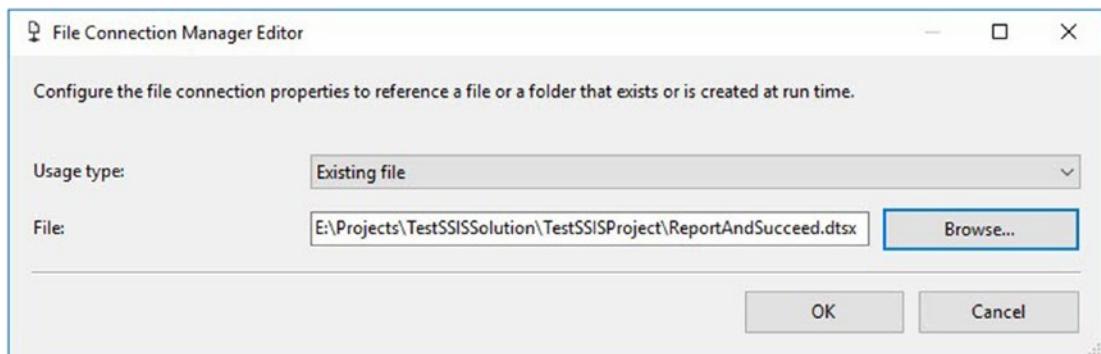


Figure 6-52. TestSSISSolution\TestSSISProject\ReportAndSucceed.dtsx file selected

Click the OK button on the File Connection Manager Editor to return to the EPT Execute Child Package execute package task editor, which appears similar to Figure 6-53.

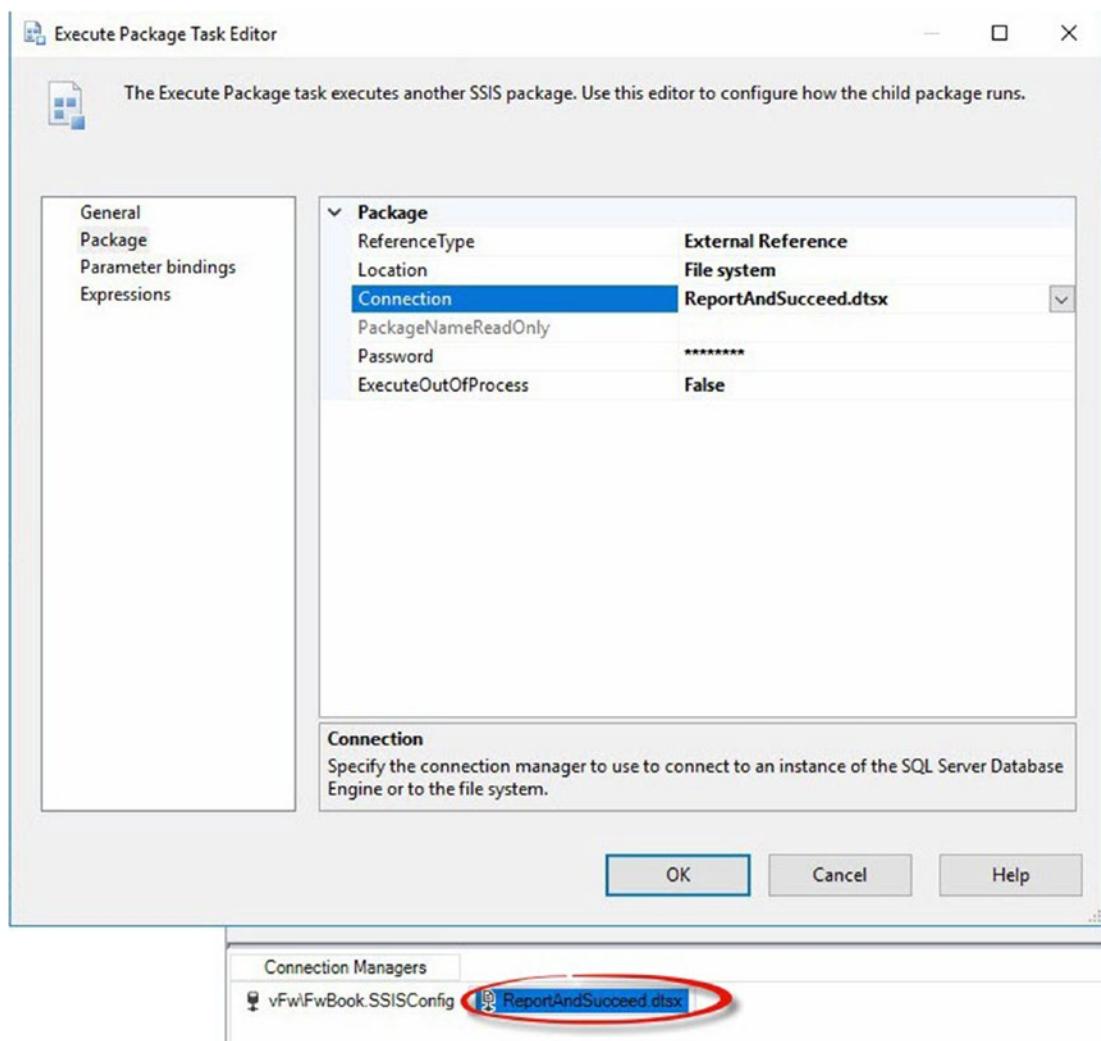


Figure 6-53. Configured EPT Execute Child Package and ReportAndSucceed.dtsx File Connection Manager

Click the OK button on the EPT Execute Child Package execute package task editor. Right-click the ReportAndSucceed.dtsx file connection manager, and then click Rename as shown in Figure 6-54.

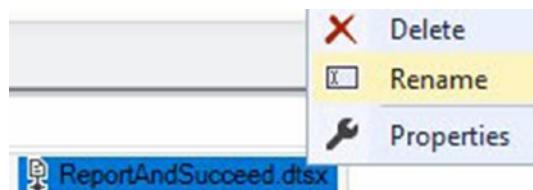


Figure 6-54. Renaming the ReportAndSucceed.dtsx File Connection Manager

Rename the File Connection Manager to “ChildPackage,” as shown in Figure 6-55.



Figure 6-55. Renaming the File Connection Manager “ChildPackage”

In order to execute the SSIS application package supplied from FOREACH Application Package (which gets metadata from querying the SSISConfig database), the ChildPackage file connection manager ConnectionString property needs to be dynamic. SSIS Expressions are designed to provide dynamic values to the package, containers, and tasks at runtime.

Click the ChildPackage file connection manager to select it, and then press the F4 key to open the ChildPackage file connection manager’s properties, as shown in Figure 6-56.

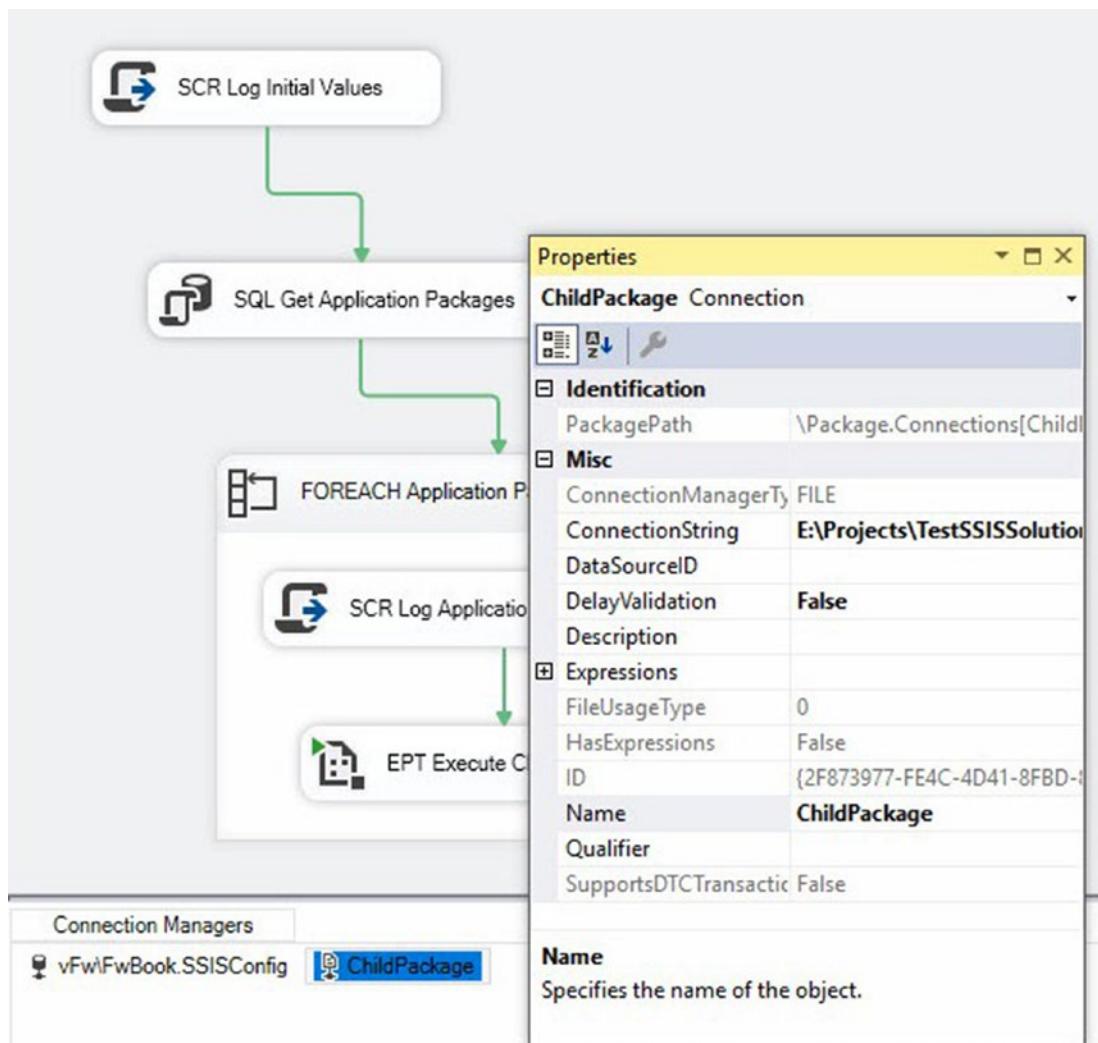
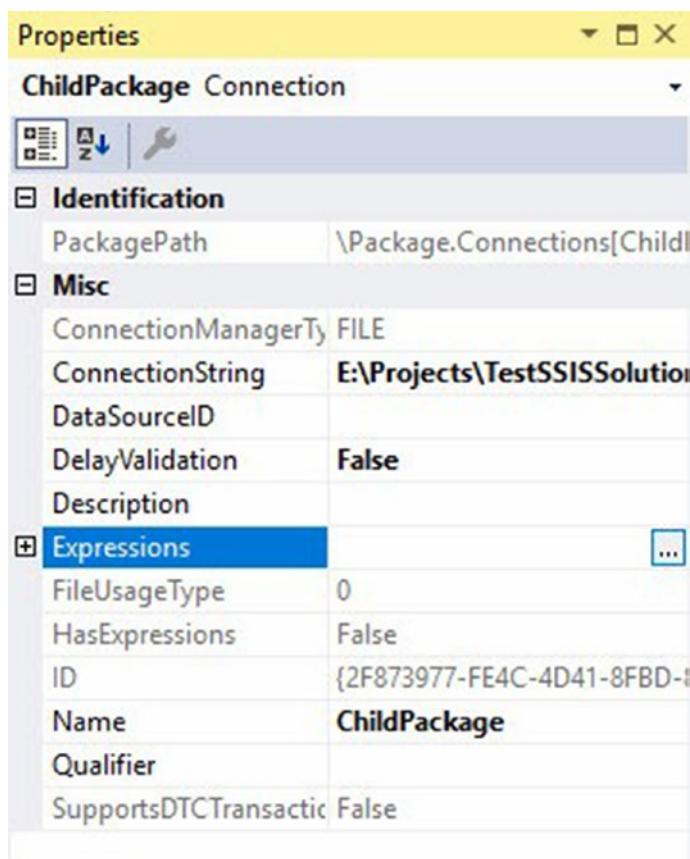


Figure 6-56. Opening Properties for the ChildPackage file connection manager

Click the Expressions property, and then click the ellipsis in the Expressions property value area, as shown in Figure 6-57.



Expressions

A collection of expressions. The evaluation result of each expression is assigned to a property and replac...

Figure 6-57. Clicking the Expressions ellipsis

Clicking the Expressions property ellipsis opens the Property Expressions Editor, as shown in Figure 6-58.

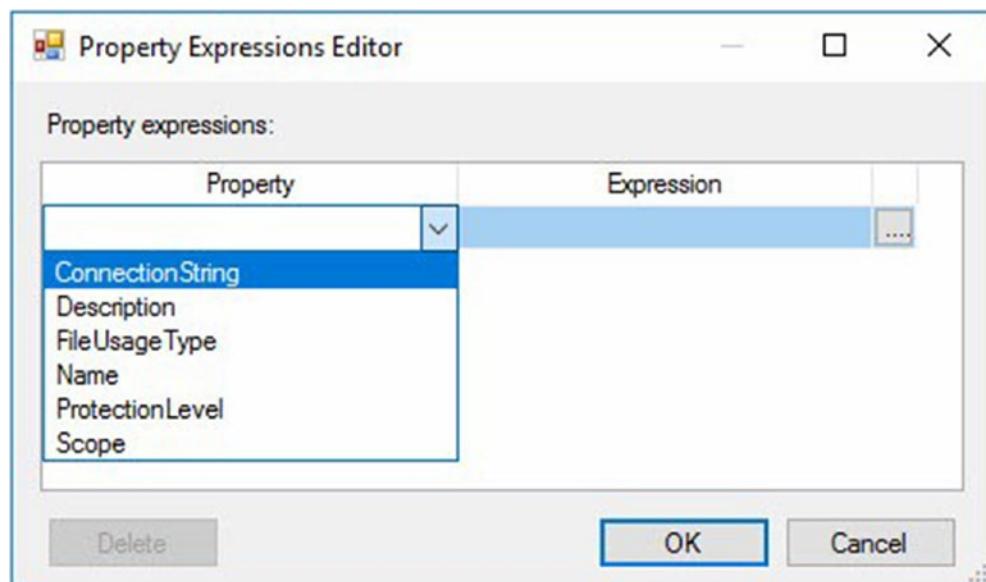


Figure 6-58. File Connection Manager Property Expressions Editor dialog properties

Select the `ConnectionString` property from the Property drop-down as shown in Figure 6-59.

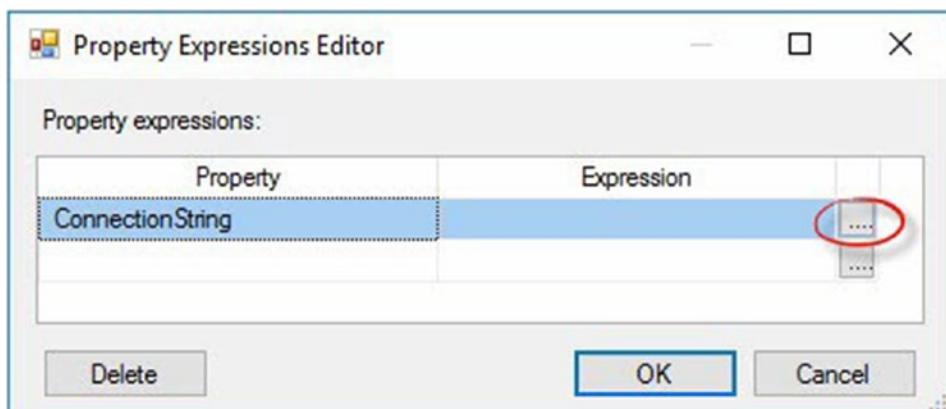


Figure 6-59. Selecting the `ConnectionString` property of the File Connection Manager

Click the ellipsis in the Expression column of the Property Expressions Editor (circled in Figure 6-59) to open the Expression Builder, as shown in Figure 6-60.

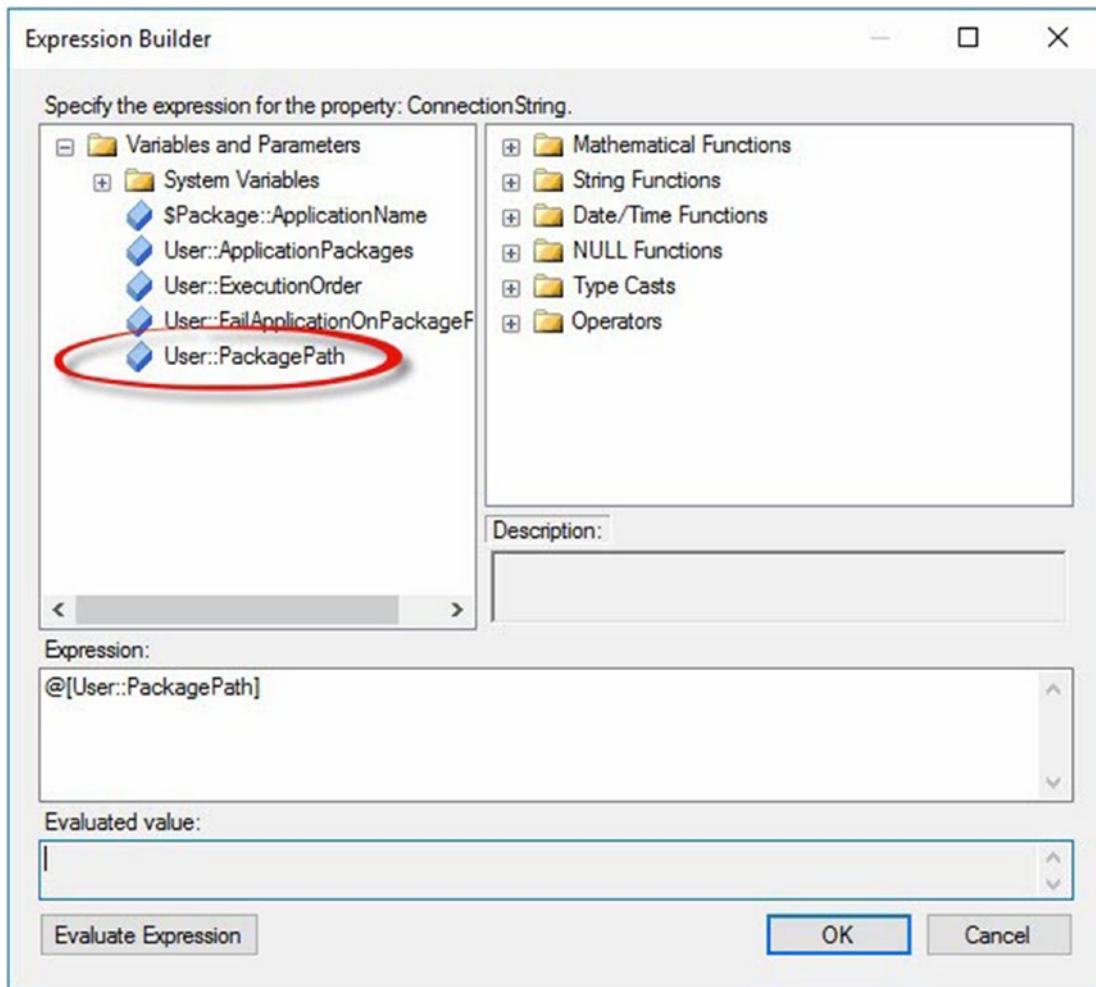


Figure 6-60. Expression Builder

In Expression Builder, expand the System Variables virtual folder, click and hold the User::PackagePath SSIS variable, and then drag the User::PackagePath SSIS variable into the “Expression” textbox as shown in Figure 6-60. Click the OK button on Expression Builder to return to the Property Expressions Editor, as shown in Figure 6-61.

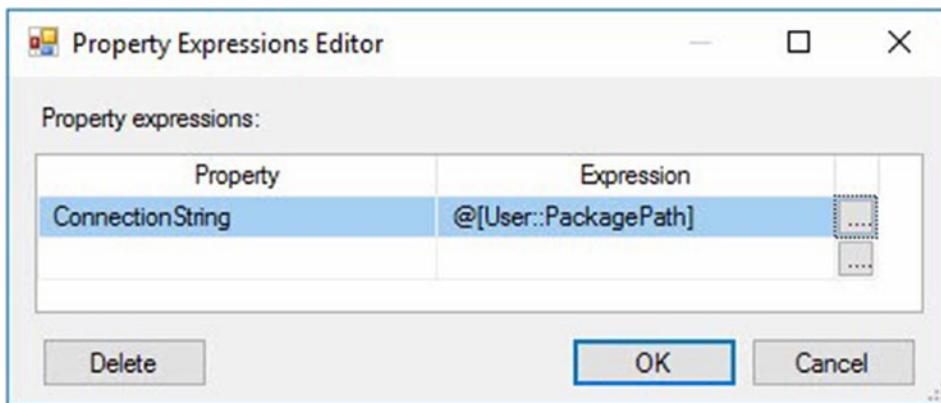


Figure 6-61. Property Expressions Editor with a configured ConnectionString property

Figure 6-61 displays the ChildPackage file connection manager's ConnectionString property that is now dynamic and driven by the value of the User::PackagePath SSIS variable. Remember, the User::PackagePath SSIS variable value is managed in the FOREACH Application Package foreach loop container, where the User::PackagePath SSIS variable is populated by each PackagePath value in the ADO.Net dataset returned from the query for metadata stored in the SSISConfig database. Thus, Parent.dtsx, as now configured (once the OK button is clicked), will execute *each and every* (enabled) SSIS framework application package associated with the SSIS application.

Click the OK button to close the Property Expressions Editor. Expand the Expressions property (collection) for the ChildPackage file connection manager, and note the ConnectionString property is now managed by the value of the User::PackagePath SSIS variable. Note also the ChildPackage file connection manager contains Expressions as indicated by the f(x) decoration, as shown in Figure 6-62.

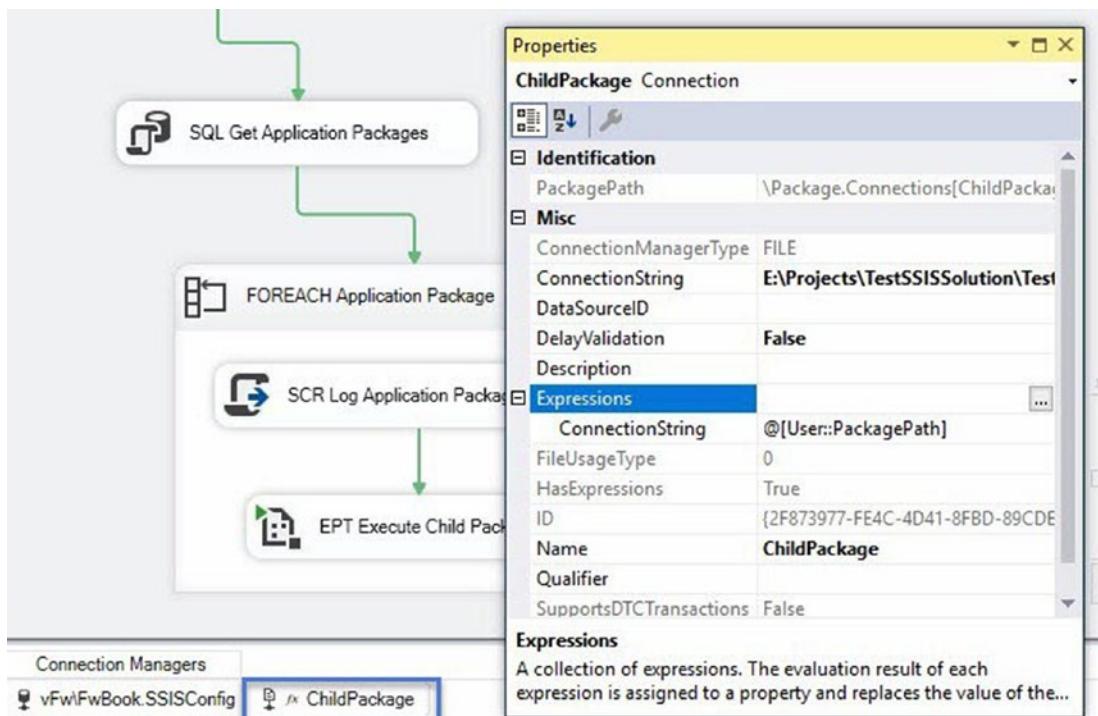


Figure 6-62. *ConnectionString Override and Expressions Decoration for ChildPackage file connection manager*

Test the operation of Parent.dtsx by starting the SSIS debugger (F5). At first glance, the results may appear askew, as shown in Figure 6-63.

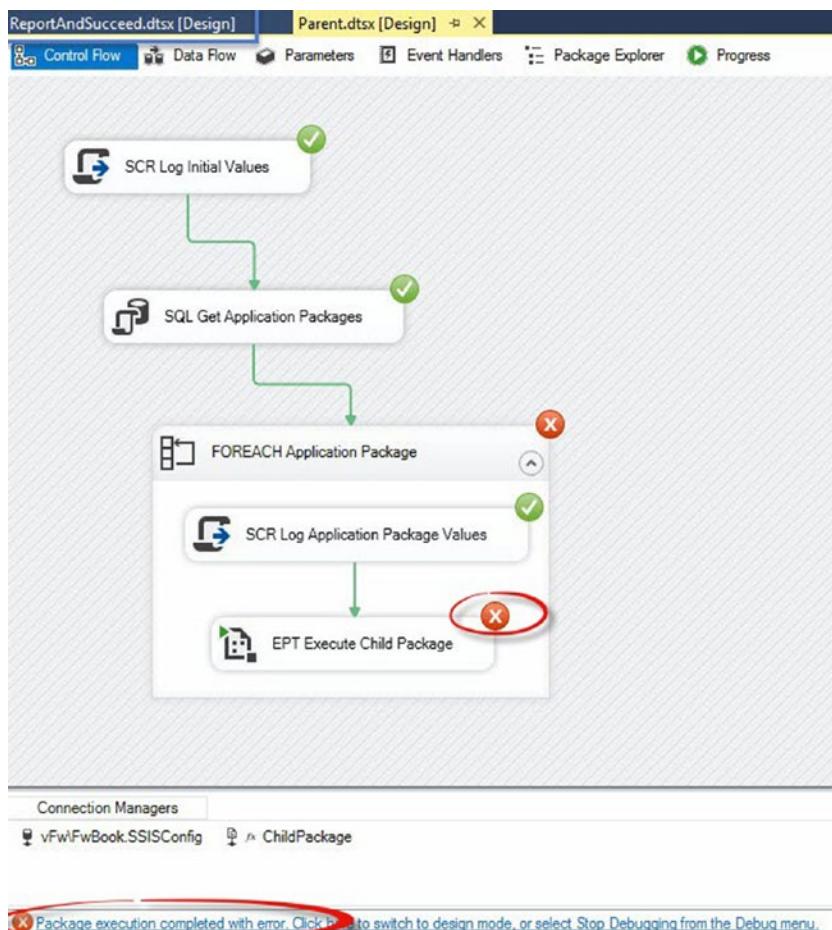


Figure 6-63. Test Execution Results

Note the test execution of Parent.dtsx failed. The lower package execution message affirms package execution failure, as does task-failed indications on the EPT Execute Child Package execute package task and the FOREACH Application Package foreach loop container.

What happened? Great question! Open the Progress/Execution Results tab on the Parent.dtsx package tab and view the messages recorded during execution, as shown in Figure 6-64.

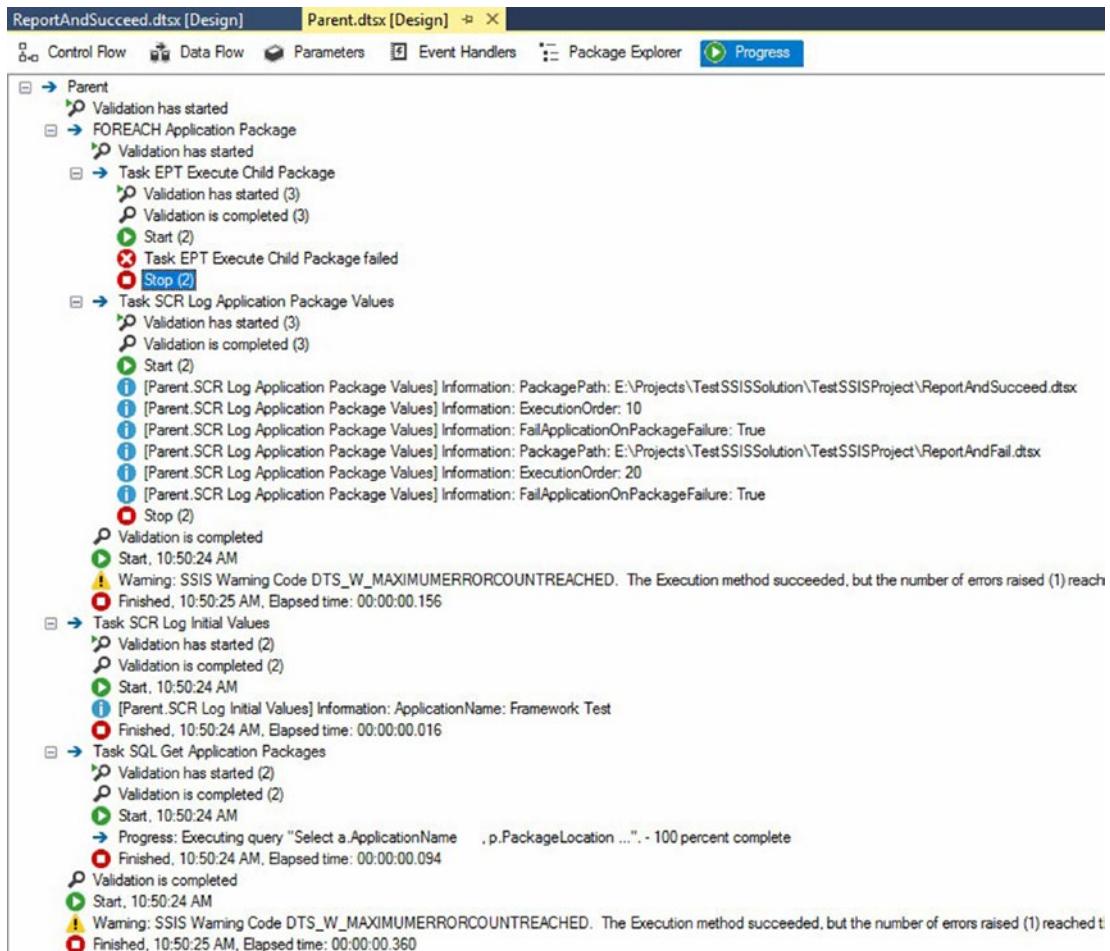


Figure 6-64. Viewing the Progress/Execution Results tab

Figure 6-64 shows the messages raised during the failed execution of Parent.dtsx. We see metadata from both files in the SSIS application – ReportAndSucceed.dtsx and ReportAndFail.dtsx – logged by the SCR Log Application Package Values script task. What's missing is an indication of whether each SSIS package execution succeeded or failed.

Log Execution Results

To begin logging execution results, drag a new Script Task into the FOREACH Application Package foreach loop container, and then rename the new Script Task to “SCR Log Package Execution Success,” as shown in Figure 6-65.

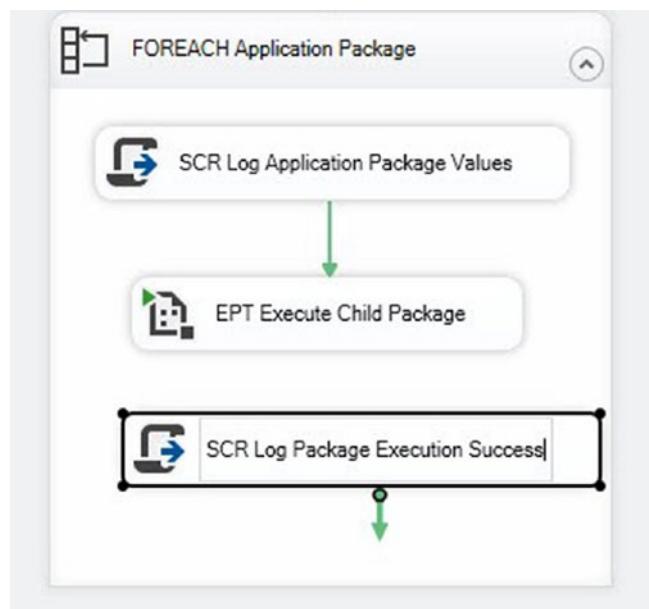


Figure 6-65. Adding the SCR Log Package Execution Success script task

Open the SCR Log Package Execution Success script task editor. As before, add System::PackageName, System::TaskName, and User::PackagePath SSIS variables to the ReadOnlyVariables property, as shown in Figure 6-66.

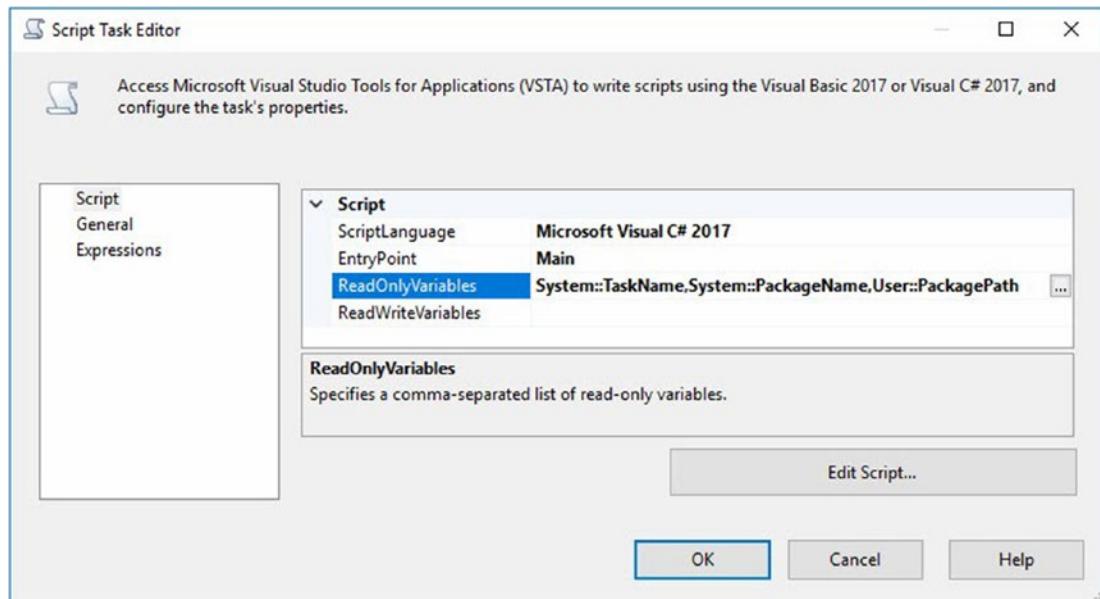


Figure 6-66. Populating the ReadOnlyVariables property

Click the Edit Script button to open the VstaProjects editor. Add the C# .Net code shown in Listing 6-5 to the public void Main() method to log the successful execution of an SSIS application package in the framework.

Listing 6-5. C# .Net code to log a successful SSIS application package execution

```
public void Main()
{
    // System:::PackageName, System:::TaskName
    // User:::PackagePath

    string packageName = ► Dts.Variables["System:::PackageName"].Value.
        ToString();
    string taskName = Dts.Variables["System:::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int informationCode = 1001;
    bool fireAgain = true;

    string packagePath = Dts.Variables["User:::PackagePath"].Value.
        ToString();
    string description = packagePath + " execution succeeded";
    Dts.Events.FireInformation(informationCode, subComponent, description,
        ► "", 0, ref fireAgain);

    Dts.TaskResult = (int)ScriptResults.Success;
}
```

The code in Listing 6-5 raises an Information event that builds a message informing operators and developers the SSIS framework application package execution has succeeded, as shown in Figure 6-67.

```

public void Main()
{
    // System::PackageName, System::TaskName
    // User::PackagePath

    string packageName = Dts.Variables["System::PackageName"].Value.ToString();
    string taskName = Dts.Variables["System::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int informationCode = 1001;
    bool fireAgain = true;

    string packagePath = Dts.Variables["User::PackagePath"].Value.ToString();
    string description = packagePath + " execution succeeded";
    Dts.Events.FireInformation(informationCode, subComponent, description, "", 0, ref fireAgain);

    Dts.TaskResult = (int)ScriptResults.Success;
}

```

Figure 6-67. Raising an Information event on successful SSIS framework application package execution

Close the VstaProjects window and then click the OK button on the SCR Log Package Execution Success script task editor. Connect an on-success precedence constraint from the EPT Execute Child Package execute package task to the SCR Log Package Execution Success script task, as shown in Figure 6-68.

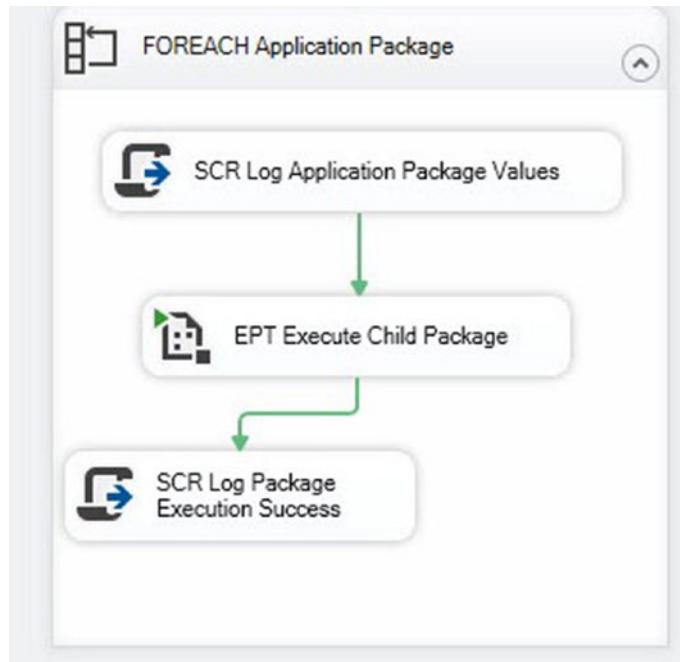


Figure 6-68. SCR Log Package Execution Success configured

Drag another Script Task into the FOREACH Application Package foreach loop container, and rename the Script Task to “SCR Log Package Execution Failure,” and connect an on-failure precedence constraint (red arrow) from the EPT Execute Child Package execute package task to the SCR Log Package Execution Failure script task, as shown in Figure 6-69.

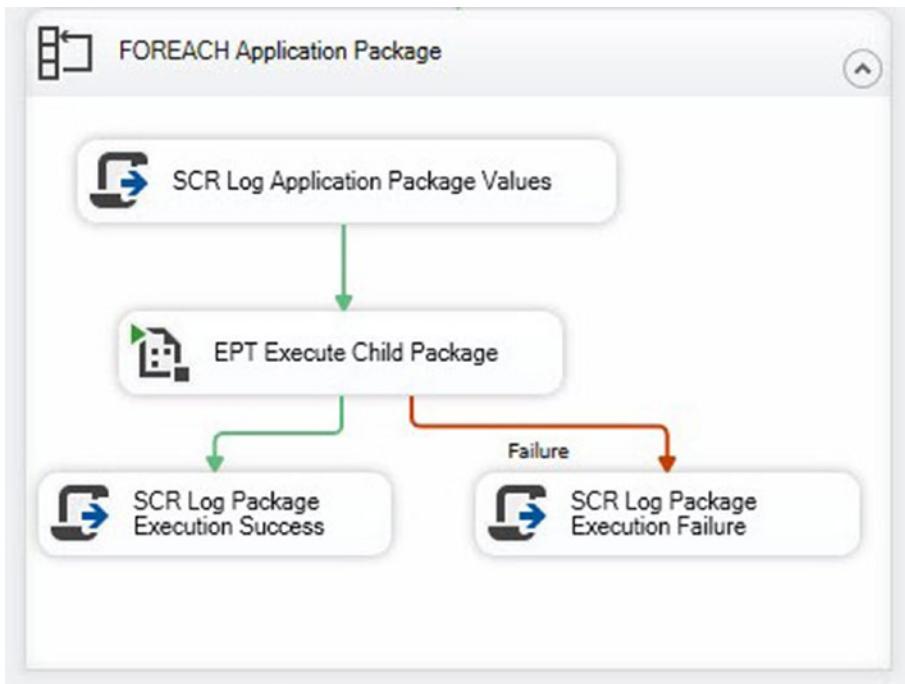


Figure 6-69. Adding SCR Log Package Execution Failure script task

Add the following SSIS variables to SCR Log Package Execution Failure script task's ReadOnlyVariables list:

- System::PackageName
- System::TaskName
- User::PackagePath
- User::FailApplicationOnPackageFailure

Once added, the SCR Log Package Execution Failure script task's ReadOnlyVariables property should appear as shown in Figure 6-70.

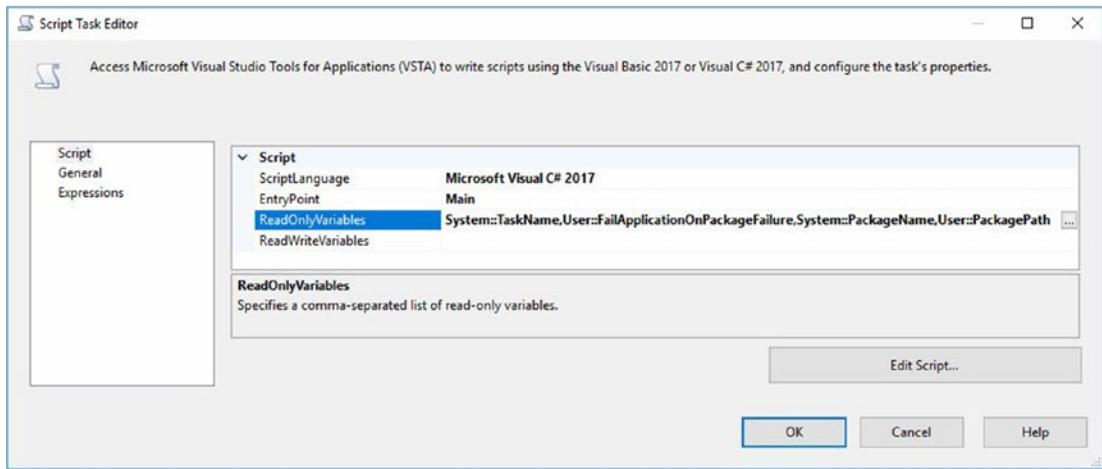


Figure 6-70. Viewing the SCR Log Package Execution Failure script task's *ReadOnlyVariables* property

Click the Edit Script button and add the C# .Net code shown in Listing 6-6 to the `public void Main()` method.

Listing 6-6. C# .Net code to log application package execution failure

```
public void Main()
{
    // System:::PackageName, System:::TaskName
    // User:::FailApplicationOnPackageFailure, User:::PackagePath

    string packageName = ► Dts.Variables["System:::PackageName"].Value.
        ToString();
    string taskName = Dts.Variables["System:::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int informationCode = 1001;
    int errorCode = -999;
    bool fireAgain = true;

    bool failApplicationOnPackageFailure = ► Convert.ToBoolean(Dts.Variables
        ["User:::FailApplicationOnPackageFailure"] ►
        .Value);

    string packagePath = Dts.Variables["User:::PackagePath"].Value.ToString();
    string description = String.Empty;
```

CHAPTER 6 FRAMEWORK EXECUTION ENGINE

```
if(failApplicationOnPackageFailure)
{
    description = packagePath + " execution failed and ➤
    FailApplicationOnPackageFailure is set (true)";
    Dts.Events.FireError(errorCode, subComponent, description, "", 0);
}
else
{
    description = packagePath + " execution failed and ➤
    FailApplicationOnPackageFailure is NOT set (false)";
    Dts.Events.FireInformation(informationCode, subComponent, ➤
        description, "", 0, ref fireAgain);
}

Dts.TaskResult = (int)ScriptResults.Success;
}
```

The C# .Net code to respond to SSIS framework application package failure should appear similar to Figure 6-71.

```
public void Main()
{
    // System::PackageName, System::TaskName
    // User::FailApplicationOnPackageFailure, User::PackagePath

    string packageName = Dts.Variables["System::PackageName"].Value.ToString();
    string taskName = Dts.Variables["System::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int informationCode = 1001;
    int errorCode = -999;
    bool fireAgain = true;

    bool failApplicationOnPackageFailure = Convert.ToBoolean(Dts.Variables["User::FailApplicationOnPackageFailure"].Value);
    string packagePath = Dts.Variables["User::PackagePath"].Value.ToString();
    string description = String.Empty;

    if(failApplicationOnPackageFailure)
    {
        description = packagePath + " execution failed and FailApplicationOnPackageFailure is set (true)";
        Dts.Events.FireError(errorCode, subComponent, description, "", 0);
    }
    else
    {
        description = packagePath + " execution failed and FailApplicationOnPackageFailure is NOT set (false)";
        Dts.Events.FireInformation(informationCode, subComponent, description, "", 0, ref fireAgain);
    }

    Dts.TaskResult = (int)ScriptResults.Success;
}
```

Figure 6-71. C# .Net code to log application package execution failure

Close the VstaProjects window and click the OK button on the SCR Log Package Execution Failure script task editor.

Let's test it! A test execution results (still) in a failure on the part of this Parent.dtsx SSIS package execution, as shown in Figure 6-72.

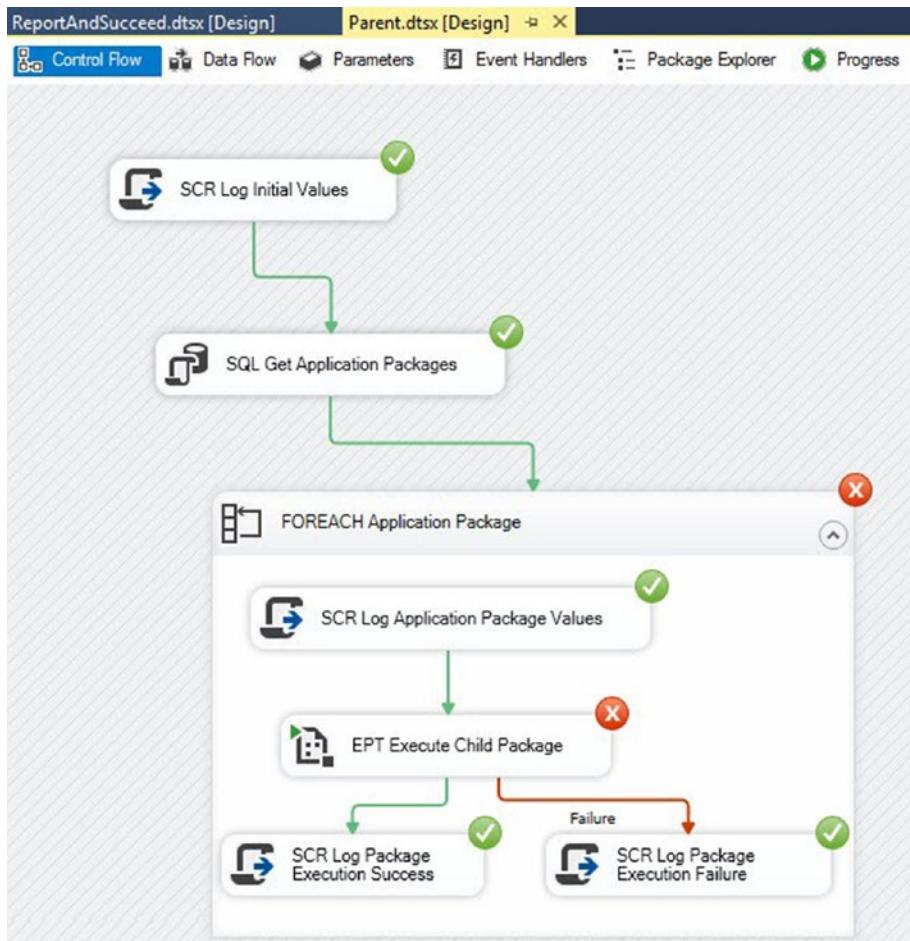


Figure 6-72. Parent.dtsx execution fails (again)

But this time, unlike previous executions, the information events instrumented in Parent.dtsx reveal what succeeded and what failed, as shown in Figure 6-73.

CHAPTER 6 FRAMEWORK EXECUTION ENGINE

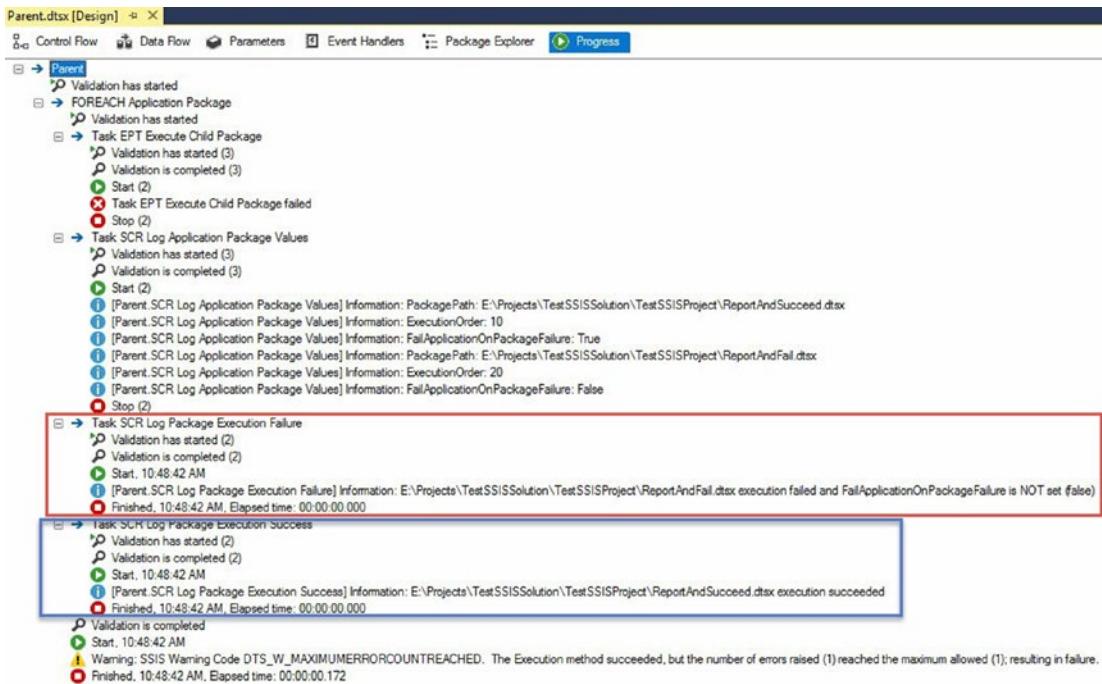


Figure 6-73. Viewing application package execution results

Instrumentation is good. Trust me. But instrumentation is almost useless if the messages are not persisted. I can hear you thinking, “Where might we capture these instrumentation messages, Andy?” I am glad you asked that excellent question. The answer is, “In SSISConfig, of course!”

Conclusion

This chapter was focused on building the execution engine for the file-based SSIS framework. In this chapter, we built the Parent.dtsx SSIS package and added instrumentation.

The next step is to record instrumentation messages in the SSISConfig database, which is covered in the next chapter.

CHAPTER 7

Framework Logging

In our journey to build a simple, custom, file-based SSIS framework, we have built a metadata database (named SSISDB), a test SSIS project, loaded some sample metadata (for the test SSIS project), created the execution engine (Parent.dtsx), and added instrumentation.

This chapter covers persisting instrumentation messages that are generated in the execution engine to the SSISConfig database.

Create a Log Schema

The information messages raised by the execution Parent.dtsx and displayed on the Progress/Execution Results tab are helpful, especially for troubleshooting. How may we store these messages for future access – especially when “something bad” has happened?

Return to your favorite T-SQL development editor (I am switching to Azure Data Studio version 1.17.1 for the remainder of this chapter), and connect to the SQL Server instance that hosts the SSISConfig database. Create a new schema named “log” in the SSISConfig database using the T-SQL in Listing 7-1.

Listing 7-1. Create the SSISConfig.log schema

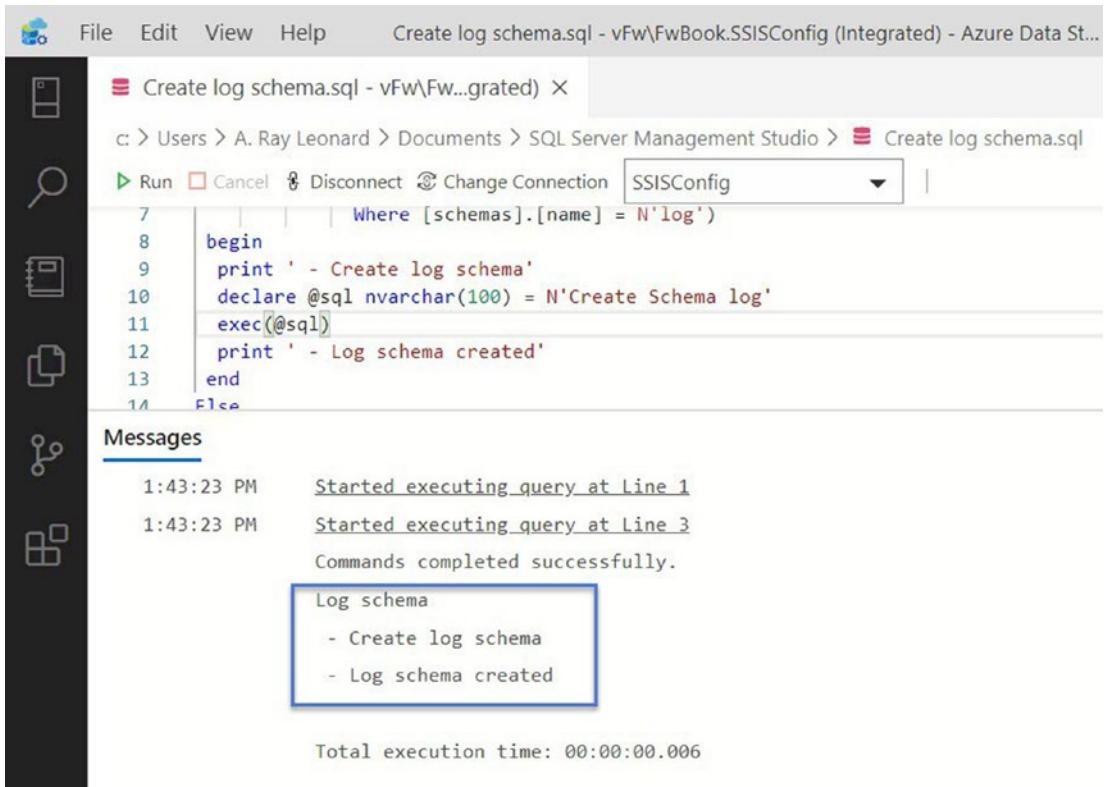
```
Use [SSISConfig]
go

print 'Log schema'
If Not Exists(Select [schemas].[name]
    From [sys].[schemas]
        Where [schemas].[name] = N'log')
begin
```

CHAPTER 7 FRAMEWORK LOGGING

```
print ' - Create log schema'
declare @sql nvarchar(100) = N'Create Schema log'
exec(@sql)
print ' - Log schema created'
end
Else
begin
print ' - Log schema already exists.'
end
print ''
go
```

Messages appear a little different in Azure Data Studio, as shown in Figure 7-1.



The screenshot shows the Azure Data Studio interface with a query window titled "Create log schema.sql". The code in the window is:

```
7 begin
8     print ' - Create log schema'
9     declare @sql nvarchar(100) = N'Create Schema log'
10    exec(@sql)
11    print ' - Log schema created'
12 end
13 Else
14
```

Below the code, the "Messages" pane displays the following output:

- 1:43:23 PM Started executing query at Line 1
- 1:43:23 PM Started executing query at Line 3
- Commands completed successfully.
- Log schema
 - Create log schema
 - Log schema created

Total execution time: 00:00:00.006

Figure 7-1. Viewing messages from creating a log schema

Messages we would expect in SQL Server Management Studio (SSMS) are highlighted in the “box” in Figure 7-1.

Before adding a table to capture “instances of application package execution” in the log schema, add a table to capture “instances of application execution” in the log schema. The first new table is named log.ApplicationInstance, and the T-SQL code to build this table is found in Listing 7-2.

Listing 7-2. Create the log.ApplicationInstance table in the SSISConfig database

```
use [SSISConfig]
go

print 'Log.ApplicationInstance table'
If Not Exists(Select [schemas].[name]
    + '.' + [tables].[name] As [Schema.Table]
    From [sys].[tables]
    Join [sys].[schemas]
        On [schemas].[schema_id] = [tables].[schema_id]
    Where [schemas].[name] = N'log'
        And [tables].[name] = N'ApplicationInstance')
begin
    print ' - Create log.ApplicationInstance table'
    Create Table [log].[ApplicationInstance]
    (
        ApplicationInstanceId int identity(1, 1)
            Constraint PK_log_ApplicationInstance Primary Key Clustered
        , ApplicationId int Not NULL
            Constraint FK_log_ApplicationInstance_config_Applications
                Foreign Key References [config].[Applications](ApplicationId)
        , ApplicationStartTime datetimeoffset(7) Not NULL
            Constraint DF_log_ApplicationInstance_ApplicationStartTime
                Default(sysdatetimeoffset())
        , ApplicationEndTime datetimeoffset(7) NULL
        , ApplicationStatus nvarchar(25) Not NULL
            Constraint DF_log_ApplicationInstance_ApplicationStatus
                Default(N'Running')
    )
    print ' - Log.ApplicationInstance table created'
end
```

```

Else
begin
    print ' - Log.ApplicationInstance table already exists.'
end
print ''
go

```

When the T-SQL in Listing 7-2 is executed in Azure Data Factory, the messages shown in Figure 7-2 are displayed.

<u>Messages</u>	
2:43:02 PM	<u>Started executing query at Line 1</u>
2:43:02 PM	<u>Started executing query at Line 3</u>
	Commands completed successfully.
	Log.ApplicationInstance table
	- Create log.ApplicationInstance table
	- Log.ApplicationInstance table created

Figure 7-2. Log.ApplicationInstance table creation messages

A new record will be inserted into the log.ApplicationInstance table when an SSIS framework application starts executing in the Parent.dtsx package. When the SSIS framework application completes executing, the record will be updated. But wait, there's more!

The second new table is named log.ApplicationPackageInstance, and the T-SQL code to build this table is found in Listing 7-3.

Listing 7-3. Create the log.ApplicationPackageInstance table in the SSISConfig database

```

use [SSISConfig]
go

print 'Log.ApplicationPackageInstance table'
If Not Exists(Select [schemas].[name]
    + '.' + [tables].[name] As [Schema.Table]
    From [sys].[tables]

```

```

Join [sys].[schemas]
  On [schemas].[schema_id] = [tables].[schema_id]
Where [schemas].[name] = N'log'
  And [tables].[name] = N'ApplicationPackageInstance')

begin
print ' - Create log.ApplicationPackageInstance table'
Create Table [log].[ApplicationPackageInstance]
(
  ApplicationPackageInstanceId int identity(1, 1)
    Constraint PK_log_ApplicationPackageInstance Primary Key Clustered
  , ApplicationInstanceId int Not NULL
    Constraint FK_log_ApplicationPackageInstance_log_
    ApplicationInstance
      Foreign Key References [log].[ApplicationInstance]
        (ApplicationInstanceId)
  , ApplicationPackageId int Not NULL
    Constraint FK_log_ApplicationPackageInstance_config_
    ApplicationPackages
      Foreign Key References [config].[ApplicationPackages]
        (ApplicationPackageId)
  , ApplicationPackageStartTime datetimeoffset(7) Not NULL
    Constraint DF_log_ApplicationPackageInstance_ApplicationPackageStartTime
    Default(sysdatetimeoffset())
  , ApplicationPackageEndTime datetimeoffset(7) NULL
  , ApplicationPackageStatus nvarchar(25) Not NULL
    Constraint DF_log_ApplicationPackageInstance_
    ApplicationPackageStatus
      Default(N'Running')
)
print ' - Log.ApplicationPackageInstance table created'
end

```

```

Else
begin
    print ' - Log.ApplicationPackageInstance table already exists.'
end
print ''
go

```

When the T-SQL in Listing 7-3 is executed in Azure Data Factory, the messages shown in Figure 7-3 are displayed.

<u>Messages</u>	
2:56:44 PM	<u>Started executing query at Line 1</u>
2:56:44 PM	<u>Started executing query at Line 3</u>
	Commands completed successfully.
	Log.ApplicationPackageInstance table
	- Create log.ApplicationPackageInstance table
	- Log.ApplicationPackageInstance table created

Figure 7-3. Log.ApplicationPackageInstance table creation messages

A new record will be inserted into the log.ApplicationPackageInstance table when an SSIS framework application package starts executing in the EPT Execute Application Package in the Parent.dtsx package's FOREACH Application Package foreach loop container. When the SSIS framework application completes execution, the record will be updated.

The “EPT” prefix stands for “Execute Package Task.” Naming conventions in SSIS are important because the event model does not include a field that identifies “Executable Type.” You can find an SSIS naming convention at ssis.tips/pages/naming.html.

Add Application Instance Logging to Parent.dtsx

Return to the Parent.dtsx SSIS package. Drag a new Execute SQL Task onto the control flow. Delete the precedence constraint between the SCR Log Initial Values script task and the SQL Get Application Packages execute SQL task. Connect new precedence constraints – the first between SCR Log initial Values script task and the new execute SQL

task, and the second between the new Execute SQL Task and the SQL Get Application Packages execute SQL task. Rename the new Execute SQL Task “SQL Log Application Instance Start,” as shown in Figure 7-4.

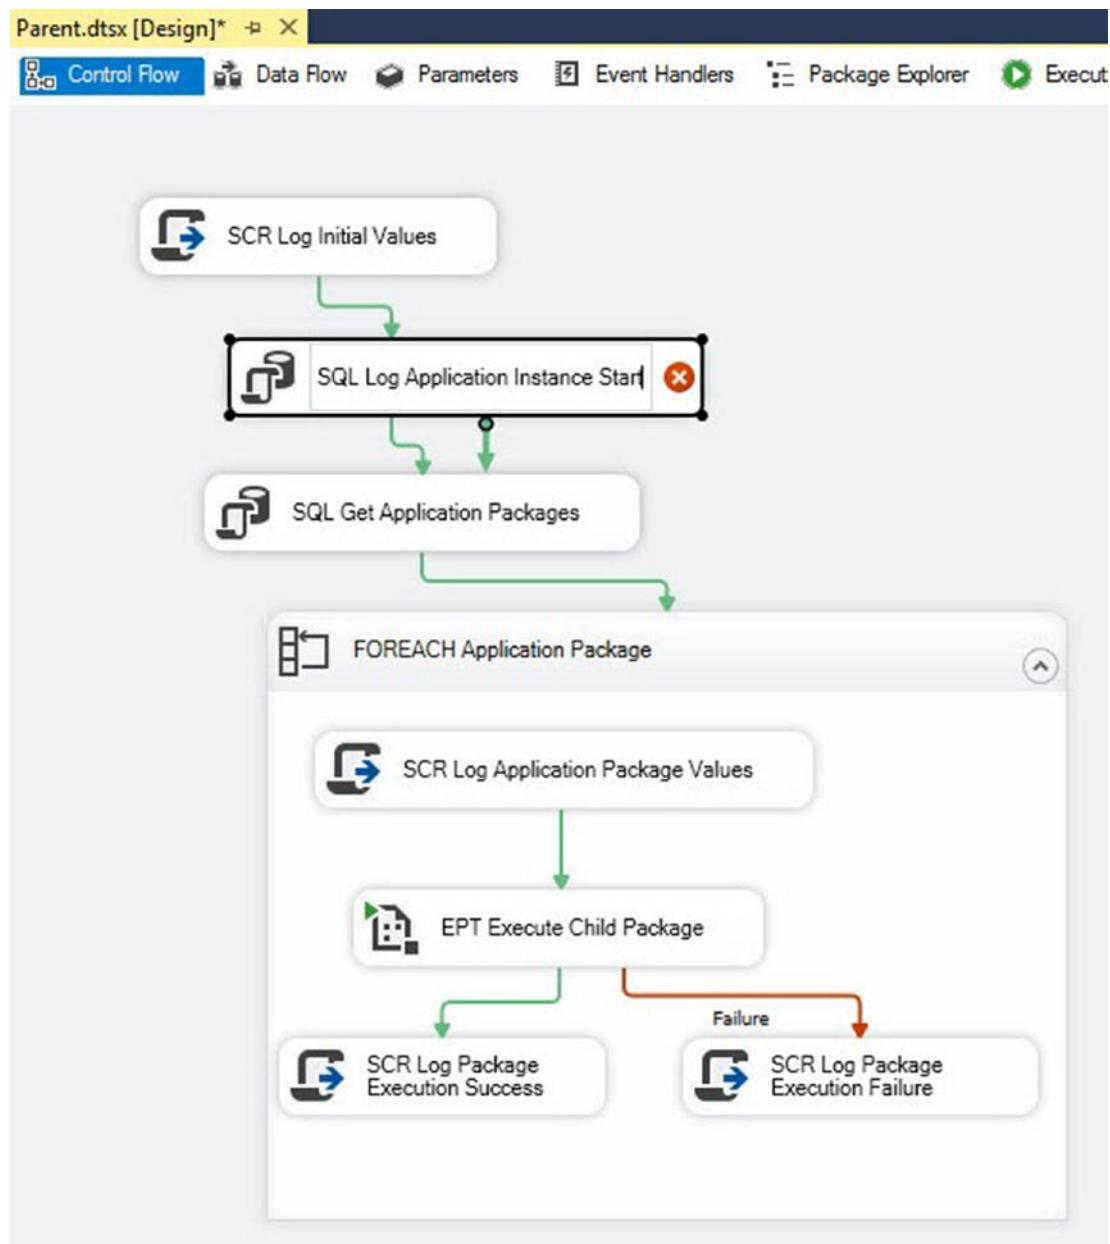


Figure 7-4. Adding the SQL Log Application Instance Start execute SQL task

Open the SQL Log Application Instance Start execute SQL task editor, and set the ConnectionType property to ADO.NET. Set the Connection property to the ADO.NET connection manager aimed at the SSISConfig database (I've renamed my connection manager “SSISConfig”). Set the ResultSet property to “Single row,” as shown in Figure 7-5.

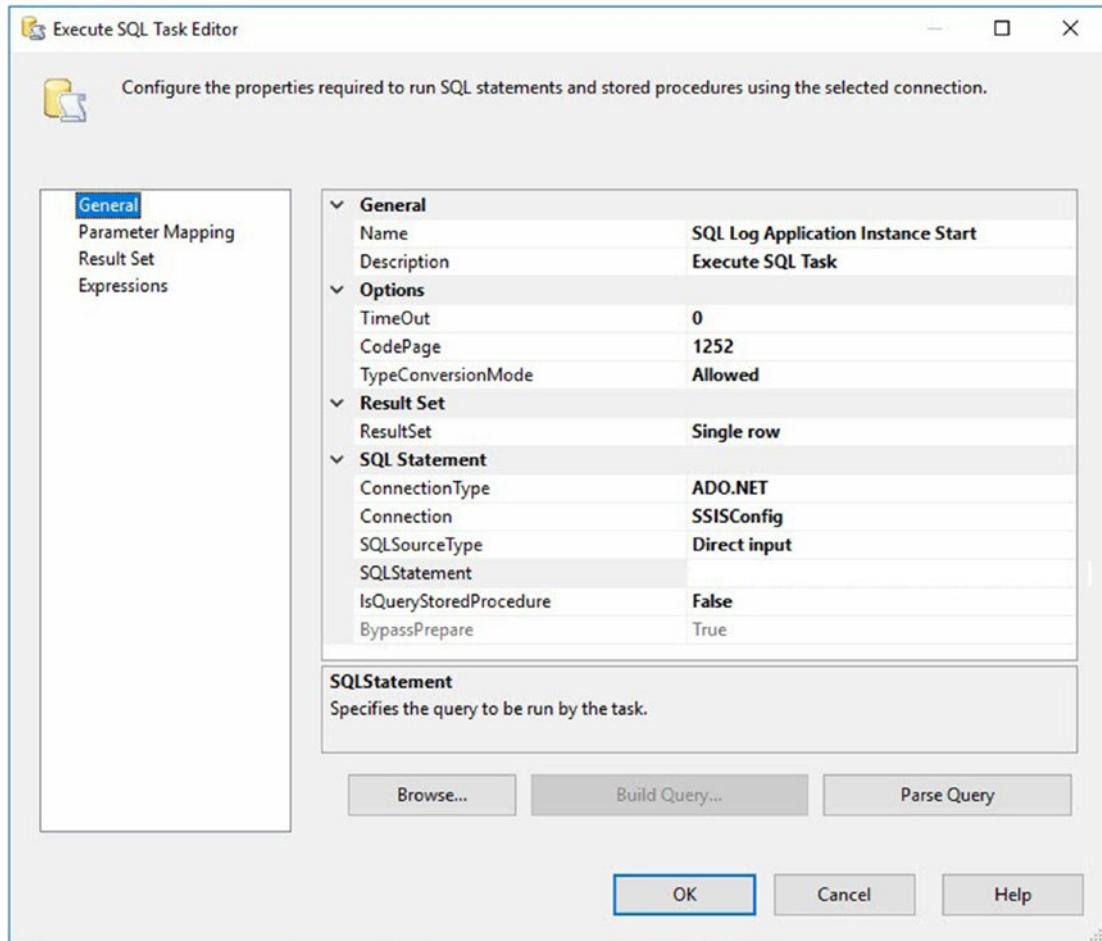


Figure 7-5. Configuring the SQL Log Application Instance Start execute SQL task’s ResultSet, ConnectionType, and Connection properties

Click the ellipsis in the SQLStatement property and paste the T-SQL from Listing 7-4 into the Enter SQL Query dialog.

Listing 7-4. Inserting a row into the [log].[ApplicationInstance] table

```
declare @ApplicationId int = (Select ApplicationId  
From config.Applications  
Where ApplicationName = @ApplicationName)  
  
Insert Into [log].ApplicationInstance (ApplicationId)  
Output inserted.ApplicationInstanceId  
Values (@ApplicationId)
```

The Enter SQL Query dialog should appear as shown in Figure 7-6.

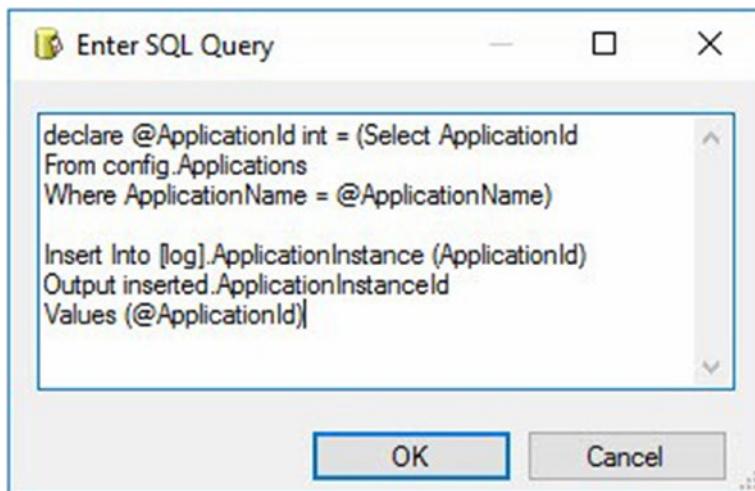


Figure 7-6. Inserting a row into the [log].[ApplicationInstance] table

Click the OK button to close the Enter SQL Query dialog.

Click the Parameter Mapping page and map the parameter

\$Package::ApplicationName to the ApplicationName parameter in the SQLStatement property as shown in Figure 7-7.

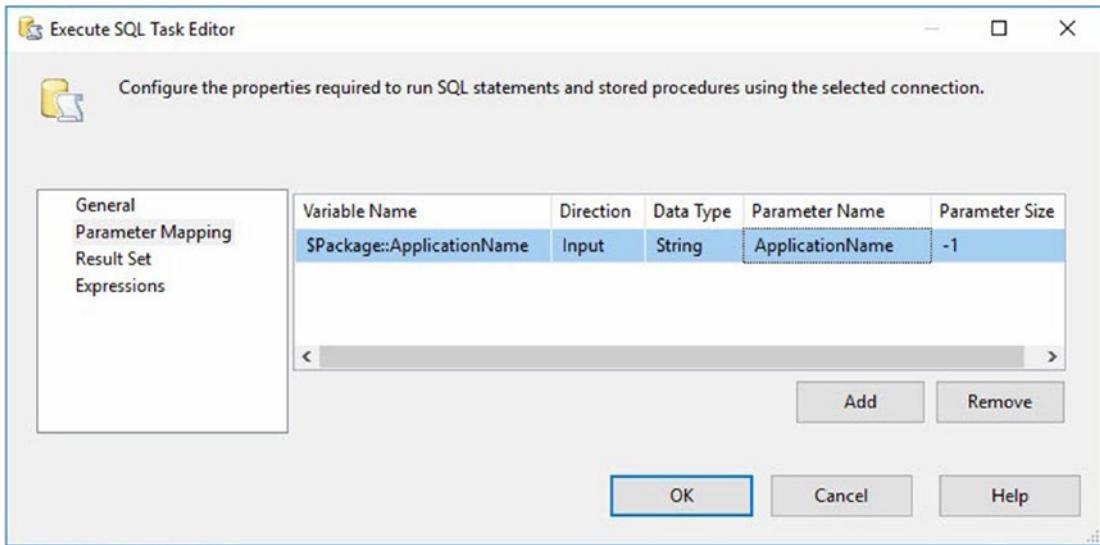


Figure 7-7. Mapping \$Package::ApplicationName to the ApplicationName query parameter

Click the Result Set page, and click the Add button. Set the Result Name to “0” and select “<New variable...>” from the Variable Name drop-down, as shown in Figure 7-8.

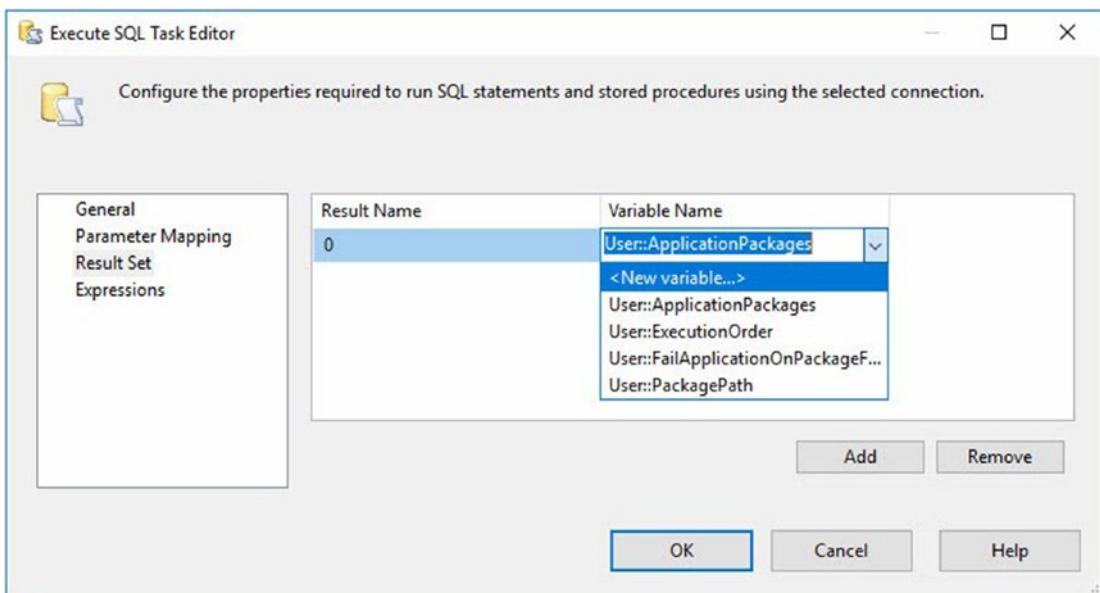


Figure 7-8. Creating a new SSIS variable for the single-row result set

When the Add Variable dialog displays, configure the new SSIS variable properties as shown in Figure 7-9:

- Container: Parent
- Name: ApplicationInstanceID
- Namespace: User
- Value type: Int32
- Value: -1

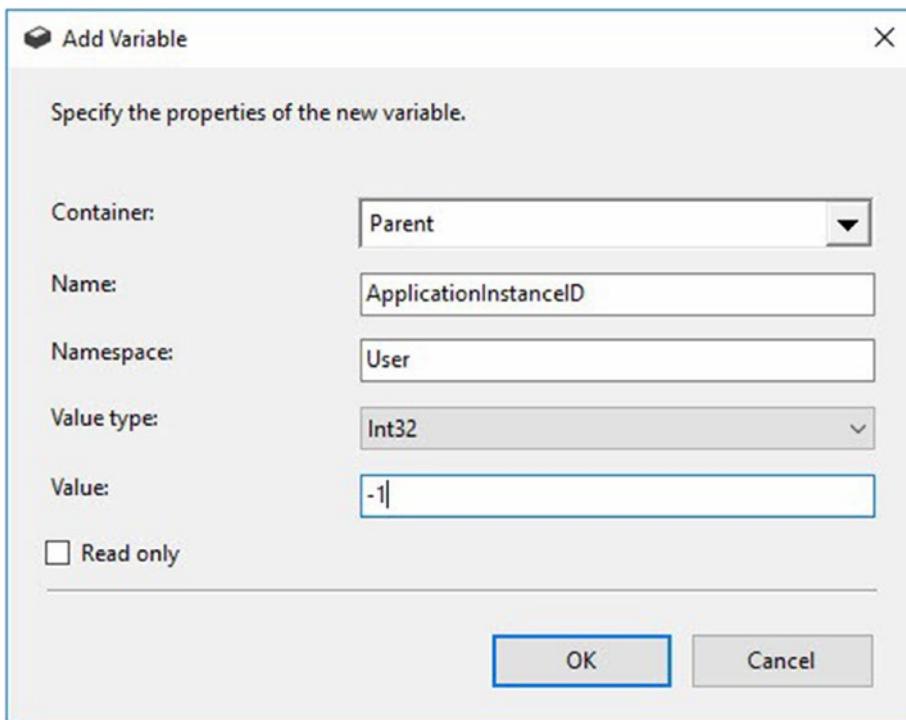


Figure 7-9. Configuring the User::ApplicationInstanceID SSIS variable

Click the OK button to close the Execute SQL Task Editor. The Parent.dtsx control flow should now appear as shown in Figure 7-10.

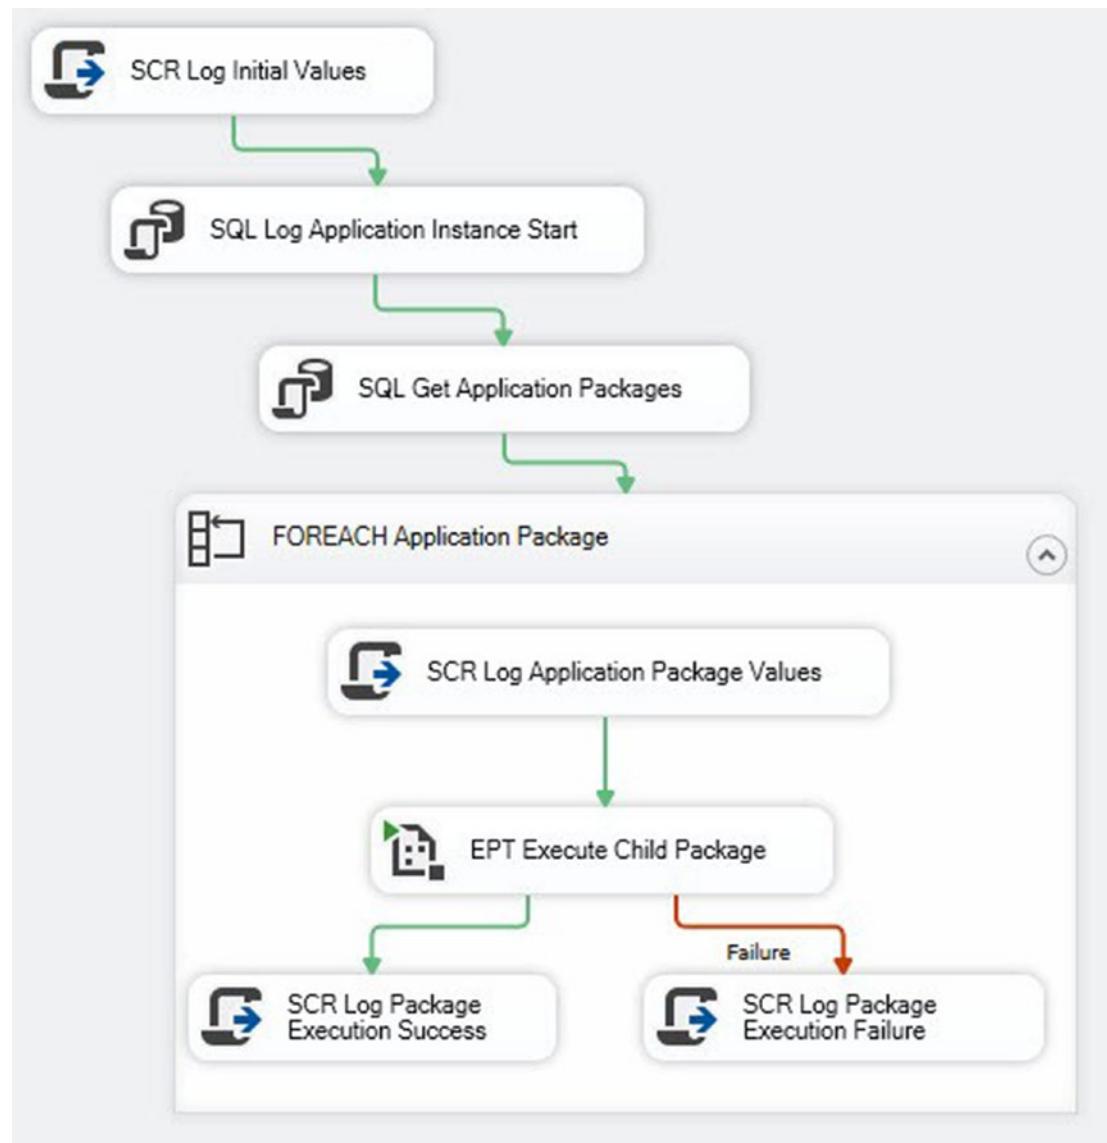


Figure 7-10. SQL Log Application Instance Start execute SQL task configured

Drag an Execute SQL Task onto the Parent.dtsx control flow below the FOREACH Application Package foreach loop container, and rename the Execute SQL Task to “SQL Log Application Instance Success,” as shown in Figure 7-11.

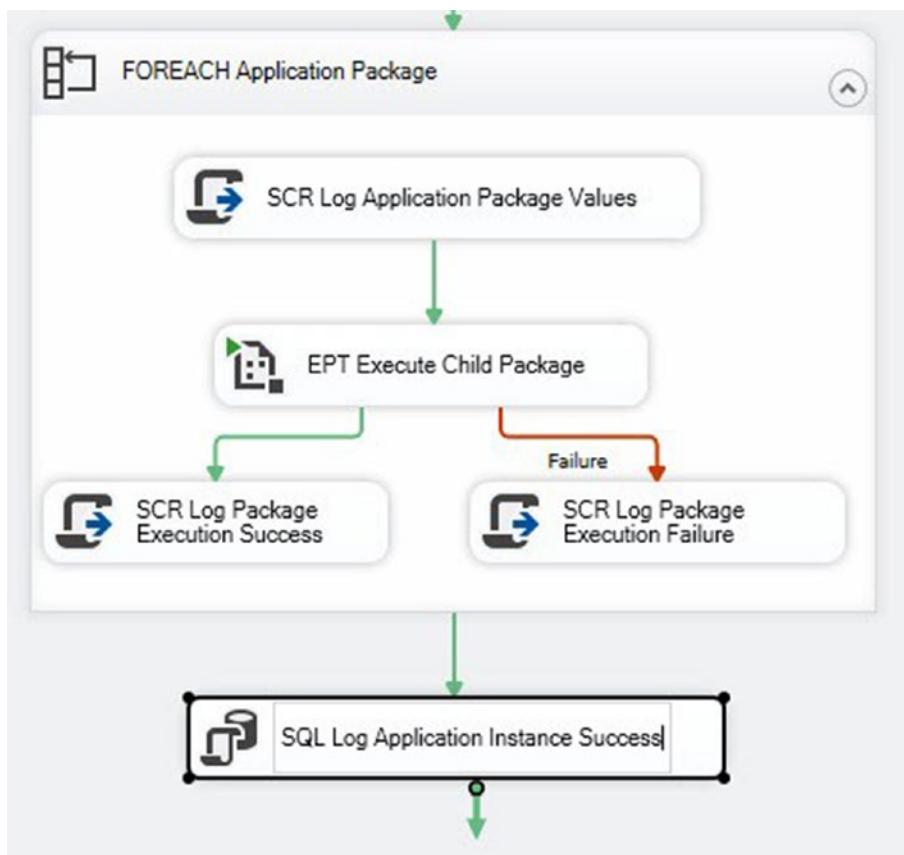


Figure 7-11. Adding the SQL Log Application Instance Success execute SQL task

Open the SQL Log Application Instance Success execute SQL task editor. Change the ConnectionType property to ADO.NET, and select the SSISConfig Connection. Click the ellipsis in the SQLStatement property. When the Enter SQL Query dialog displays, enter the T-SQL shown in Listing 7-5.

Listing 7-5. T-SQL to update the [log].[ApplicationInstance] table when the SSIS application succeeds

```

Update [log].ApplicationInstance
Set ApplicationEndTime = sysdatetimeoffset()
  , ApplicationStatus = 'Succeeded'
Where ApplicationInstanceId = @ApplicationInstanceId
  
```

The Enter SQL Query dialog should appear as shown in Figure 7-12.

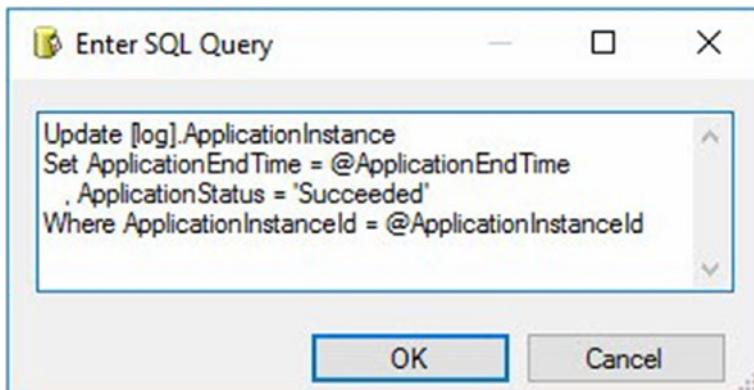


Figure 7-12. Adding T-SQL to update the [log].[ApplicationInstance] table

Click the OK button to close the Enter SQL Query dialog.

Click the Parameter Mapping page, and then click the Add button. Map the Variable Name “User::ApplicationInstanceId” to the Parameter Name “ApplicationInstanceId,” as shown in Figure 7-13.

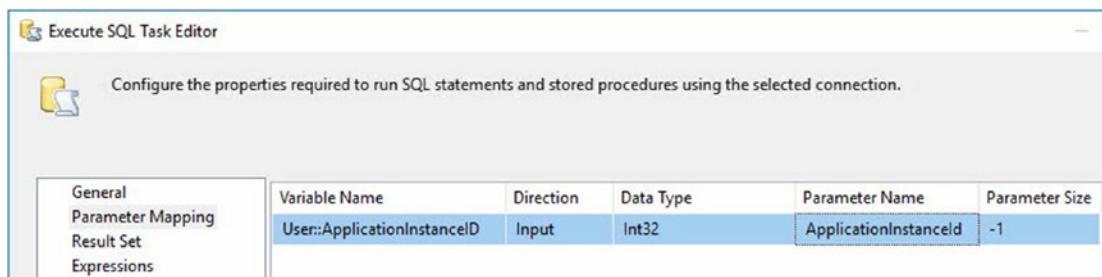


Figure 7-13. Mapping the User:ApplicationInstanceId variable to the ApplicationInstanceId parameter

Click the OK button to close the Execute SQL Task Editor.

The SQL Log Application Instance Success execute SQL task is now configured to update the row in the [log].[ApplicationInstance] table that was inserted when the SQL Log Application Instance Start execute SQL task executed earlier in the Parent.dtsx SSIS package, updating the row’s ApplicationEndTime and ApplicationStatus fields to indicate the SSIS framework application instance completed successfully at such and such a time.

The next step is to update the ApplicationEndTime and ApplicationStatus fields in the row in [log].[ApplicationInstance] when the SSIS framework application instance *fails*.

Click the Event Handlers tab in the Parent.dtsx SSIS package, as shown in Figure 7-14.

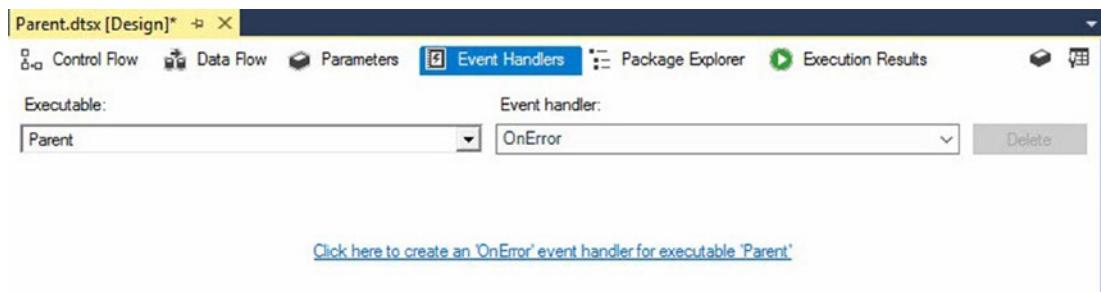


Figure 7-14. The Parent.dtsx OnError Event handler

If the Executable is not set to Parent, click the Executable drop-down, and select Parent as shown in Figure 7-15.

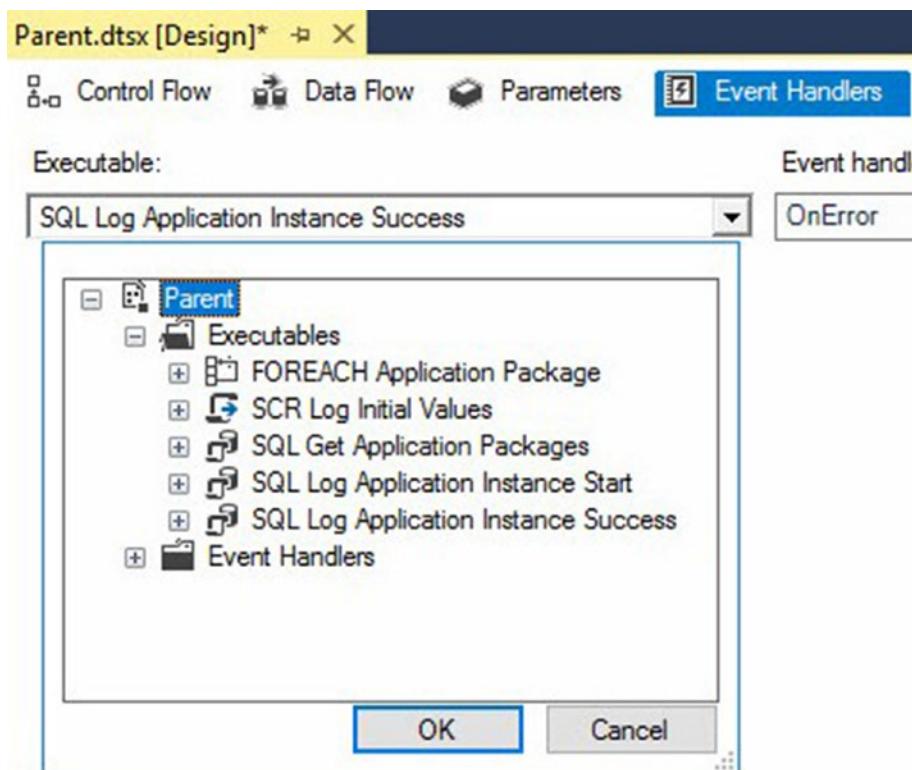


Figure 7-15. Selecting the Parent Executable

Click the link in the OnError Event handler labeled “Click here to create an ‘OnError’ event handler for the executable ‘Parent’.” Drag an Execute SQL Task onto the Event handler surface, and rename it “SQL Log Application Instance Failure,” as shown in Figure 7-16.

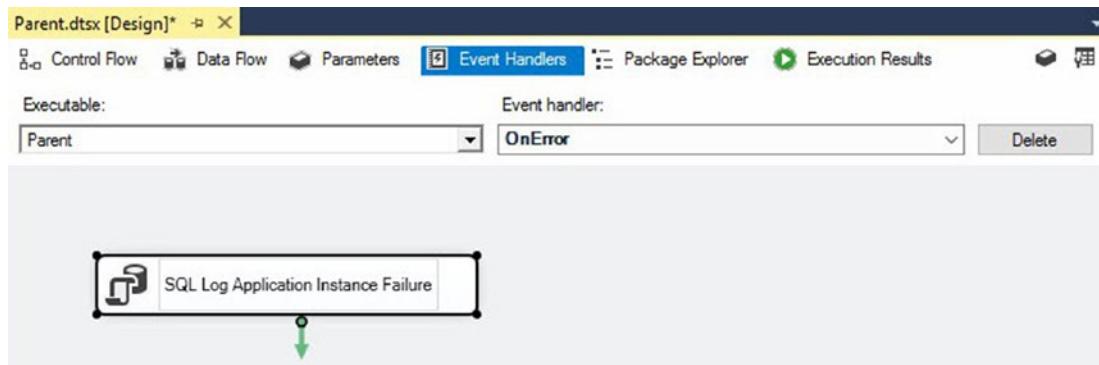


Figure 7-16. Adding the SQL Log Application Instance Failure execute SQL task to the Parent OnError Event handler

Set the ConnectionType property to ADO.NET, and select the SSISConfig Connection, as shown in Figure 7-17.

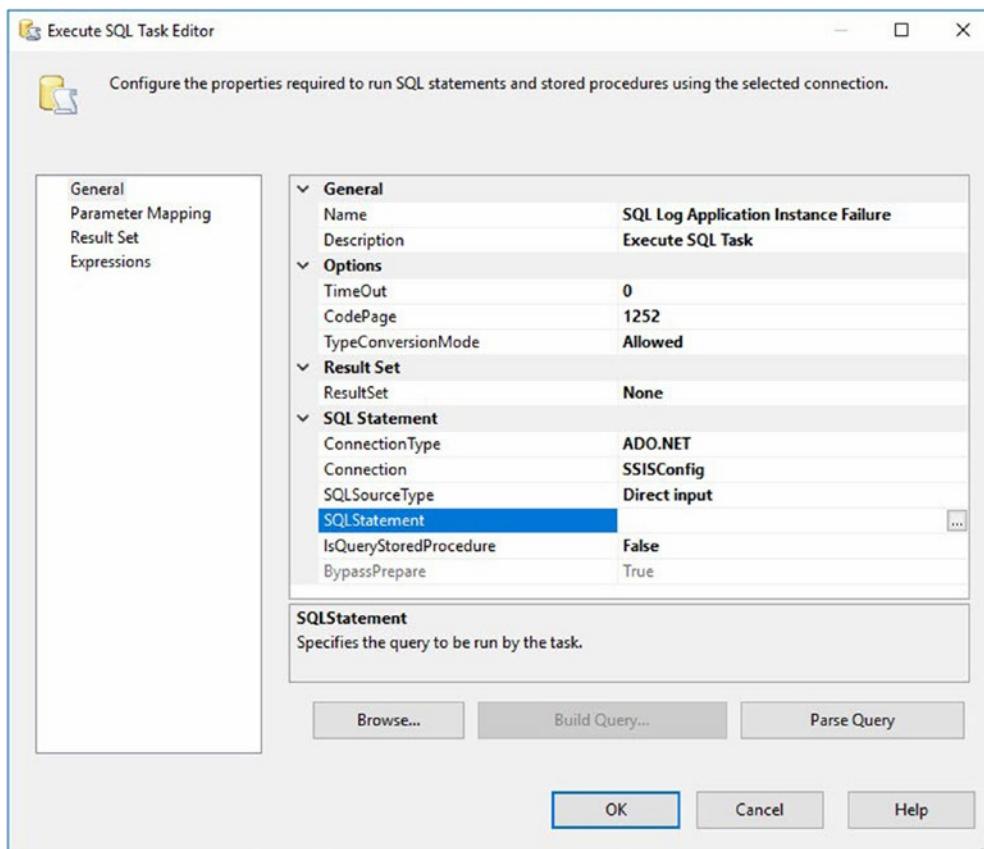


Figure 7-17. Configuring the ConnectionType and Connection properties

Click the ellipsis in the SQLStatement property to open the Enter SQL Query dialog, and enter the T-SQL shown in Listing 7-6 to update the [log].[ApplicationInstance] row when the SSIS application instance fails.

Listing 7-6. T-SQL to update [log].[ApplicationInstance] when the SSIS application fails

```
Update [log].ApplicationInstance  
Set ApplicationEndTime = sysdatetimeoffset()  
    , ApplicationStatus = 'Failed'  
Where ApplicationInstanceId = @ApplicationInstanceId
```

The Enter SQL Query dialog should appear as shown in Figure 7-18.

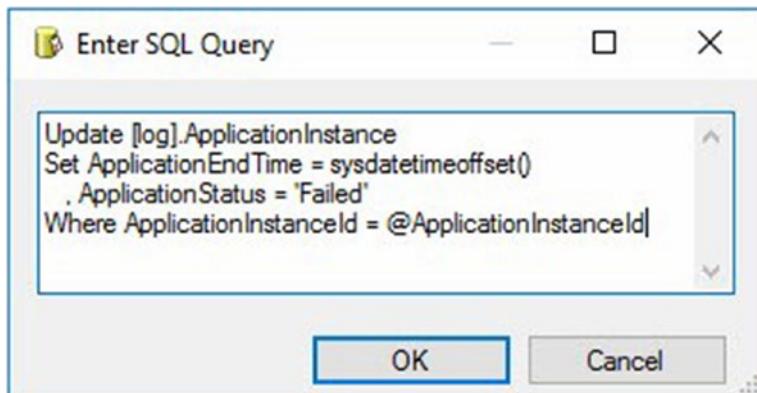


Figure 7-18. Updating the [log].[ApplicationInstance] row

Click the OK button to close the Enter SQL Query dialog.

Click the Parameter Mapping page, and then click the Add button. Map the Variable Name “User::ApplicationInstanceId” to the Parameter Name “ApplicationInstanceId,” as shown in Figure 7-19.

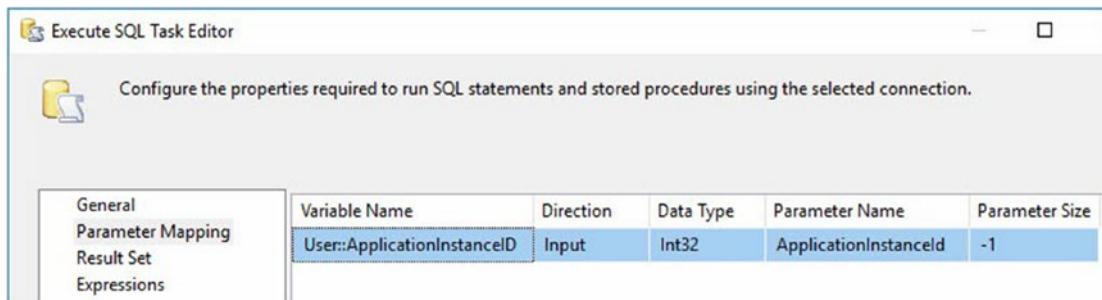


Figure 7-19. Mapping the User:ApplicationInstanceId variable to the ApplicationInstanceId parameter

Click the OK button to close the Execute SQL Task Editor.

The SQL Log Application Instance Failure execute SQL task is now configured to update the row in the [log].[ApplicationInstance] table that was inserted when the SQL Log Application Instance Start execute SQL task executed earlier on the Parent.dtsx SSIS package's control flow, updating the row's ApplicationEndTime and ApplicationStatus fields to indicate the SSIS framework application instance execution failed at such and such a time.

Add Application Package Instance Logging to Parent.dtsx

Drag a new Execute SQL Task into the FOREACH Application Package foreach loop container. Delete the precedence constraint between the SCR Log Application Package Values script task and the EPT Execute Child Package execute package task. Connect new precedence constraints – the first between SCR Log Application Package Values script task and the new execute SQL task, and the second between the new Execute SQL Task and the EPT Execute Child Package execute package task. Rename the new Execute SQL Task “SQL Log Application Package Instance Start,” as shown in Figure 7-20.

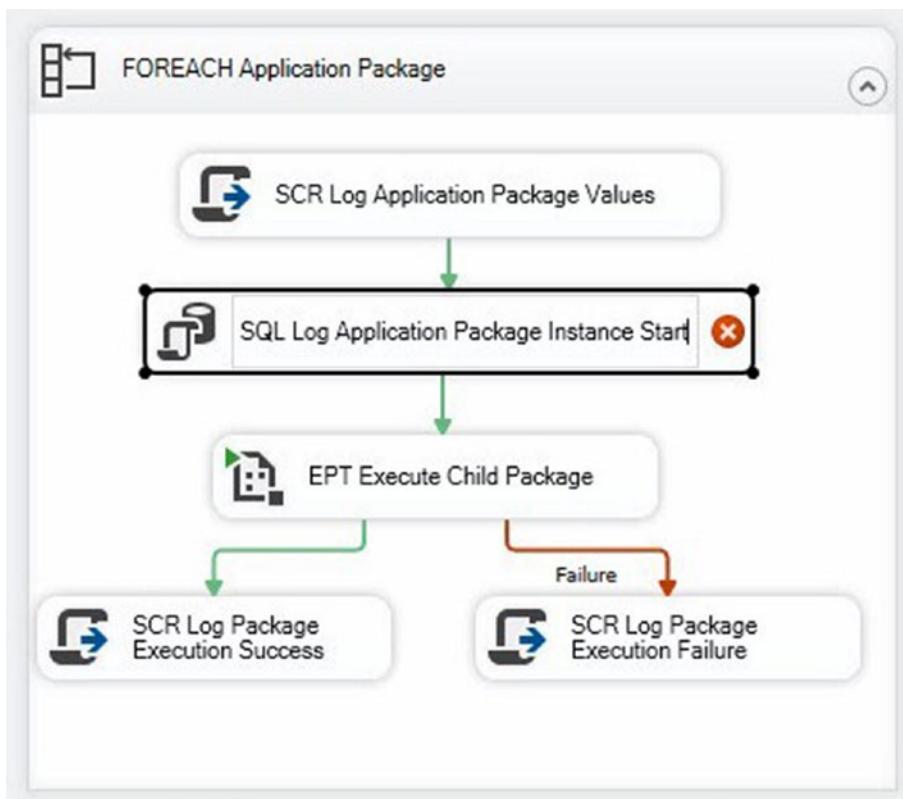


Figure 7-20. Adding the SQL Log Application Package Instance Start execute SQL task

Open the SQL Log Application Package Instance Start execute SQL task editor, and set the ConnectionType property to ADO.NET. Set the Connection property to the SSISConfig ADO.Net connection manager. Set the ResultSet property to “Single row,” as shown in Figure 7-21.

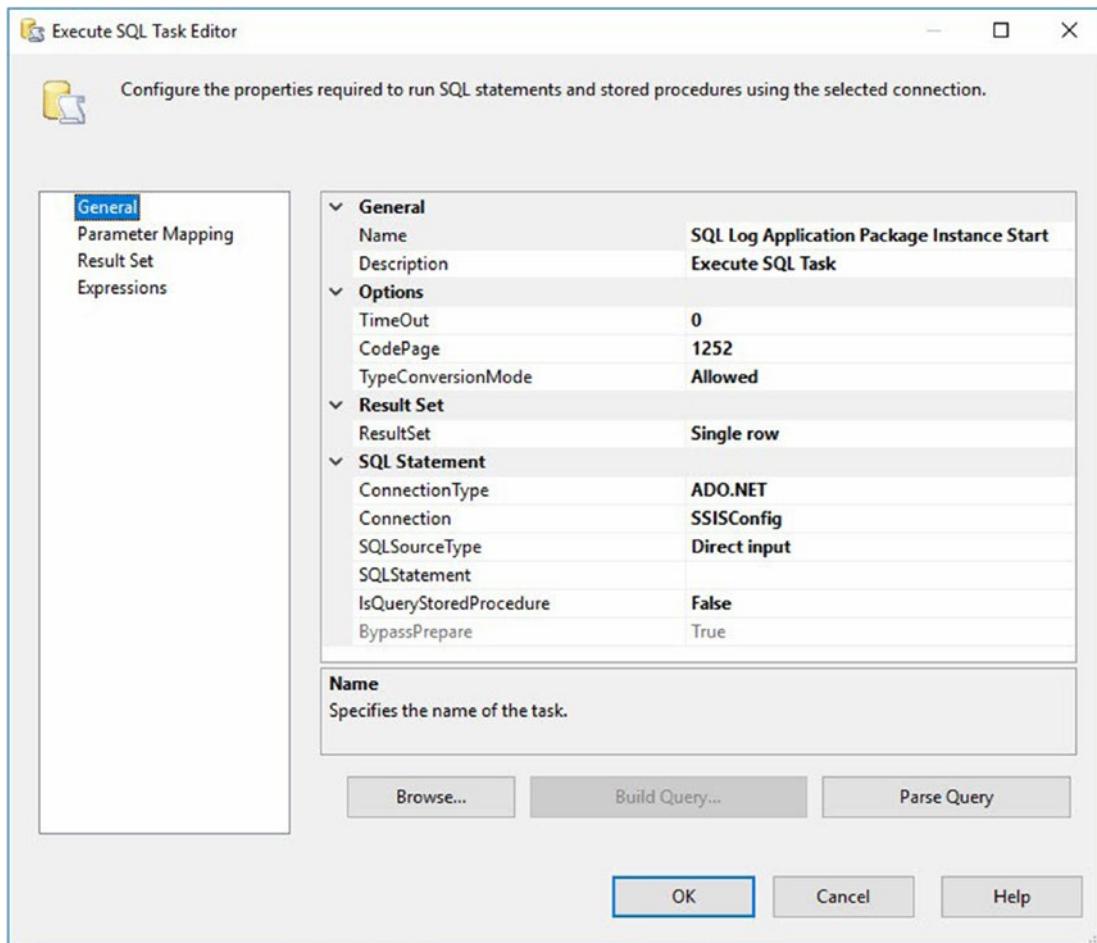


Figure 7-21. Configuring the SQL Log Application Package Instance Start execute SQL task's *ResultSet*, *ConnectionType*, and *Connection* properties

Click the ellipsis in the *SQLStatement* property, and paste the T-SQL from Listing 7-7 into the Enter SQL Query dialog.

Listing 7-7. Inserting a row into the [log].[ApplicationPackageInstance] table

```
declare @ApplicationPackageId int =
Select ap.ApplicationPackageId
From config.ApplicationPackages ap
Join config.Applications a
On a.ApplicationId = ap.ApplicationId
Join config.Packages p
```

```
On p.PackageId = ap.PackageId
Where ApplicationName = @ApplicationName
And p.PackageLocation + p.PackageName = @PackagePath
)

Insert Into [log].ApplicationPackageInstance
(ApplicationInstanceId, ApplicationPackageId)
Output inserted.ApplicationPackageInstanceId
Values (@ApplicationInstanceId, @ApplicationPackageId)
```

The Enter SQL Query dialog should appear as shown in Figure 7-22.

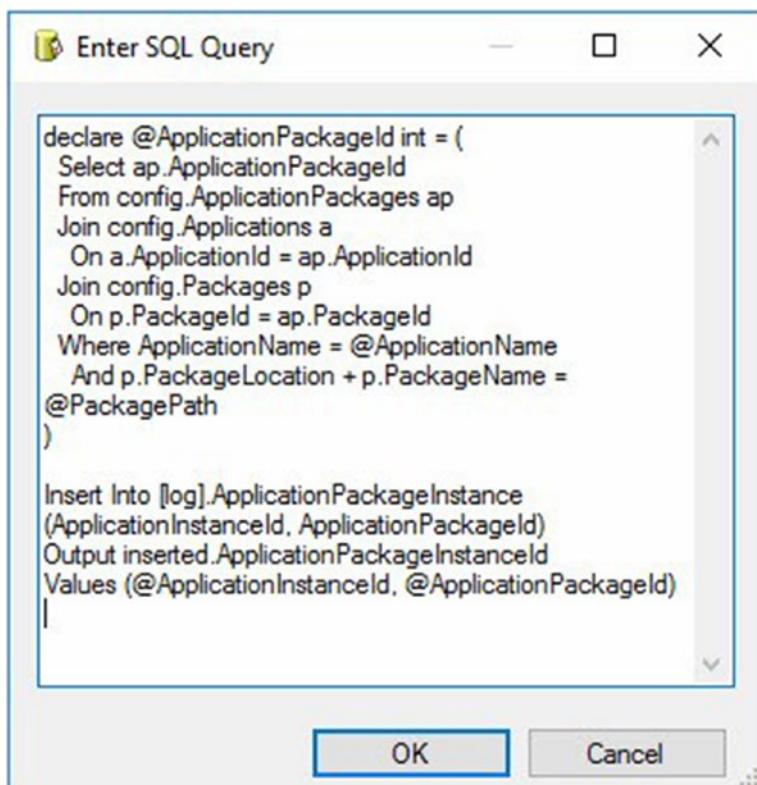


Figure 7-22. Inserting a row into the [log].[ApplicationPackageInstance] table

Click the OK button to close the Enter SQL Query dialog.

Click the Parameter Mapping page, and map the following parameters (Variable Name, Data Type, Parameter Name):

- \$Package::ApplicationName, String, ApplicationName
- User::PackagePath, String, PackagePath
- User::ApplicationInstanceId, Int32, ApplicationInstanceId

When complete, the Parameter Mapping page will appear as shown in Figure 7-23.

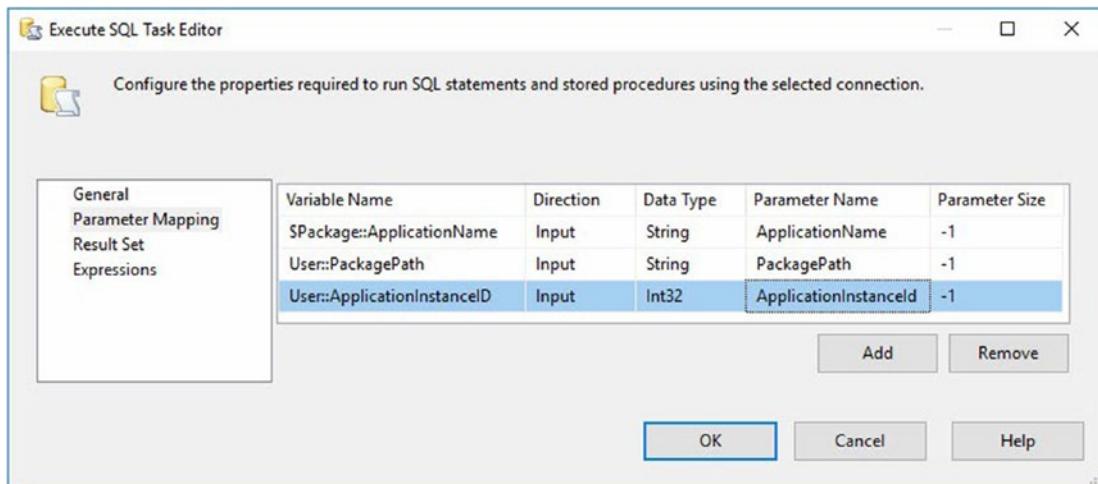


Figure 7-23. Mapping ApplicationPackageInstanceId parameters

Click the Result Set page, and click the Add button. Set the Result Name to “0” and select “<New variable...>” from the Variable Name drop-down, as shown in Figure 7-24.

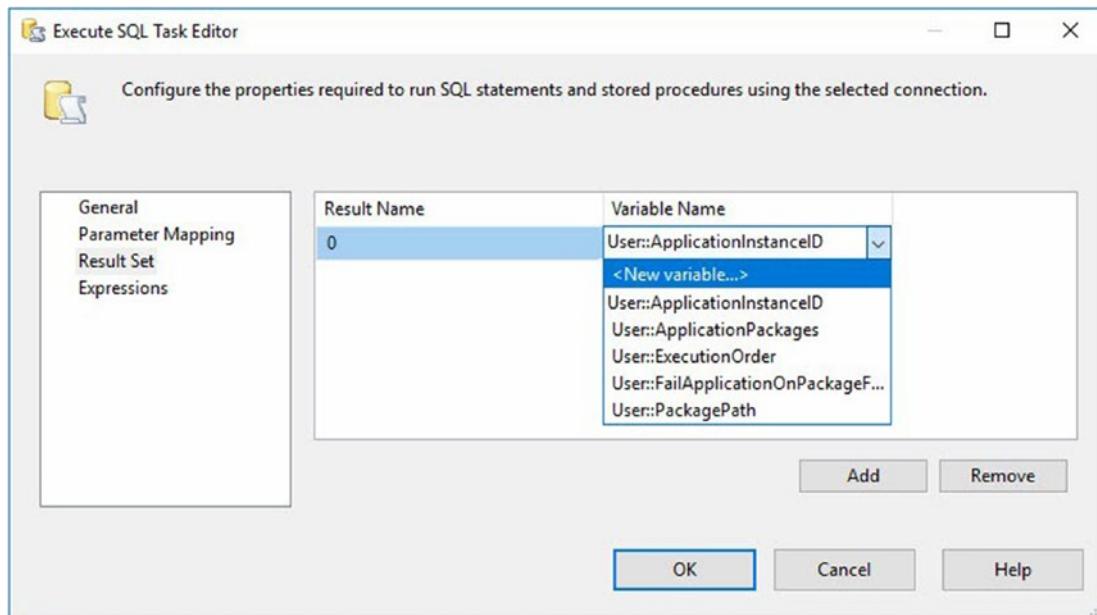


Figure 7-24. Creating a new SSIS variable for the single-row result set

When the Add Variable dialog displays, configure the new SSIS variable properties as shown in Figure 7-25.

- Container: Parent
- Name: ApplicationPackageInstanceId
- Namespace: User
- Value type: Int32
- Value: -1

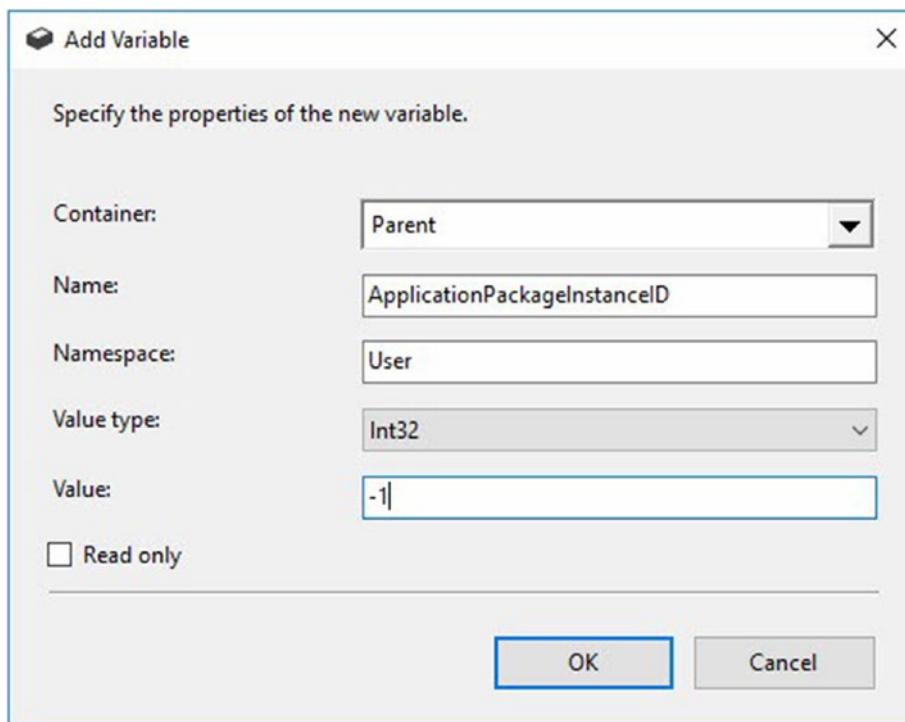


Figure 7-25. Configuring the User::ApplicationPackageInstanceId SSIS variable

Click the OK button to close the Execute SQL Task Editor. The FOREACH Application Package foreach loop container should now appear as shown in Figure 7-26.

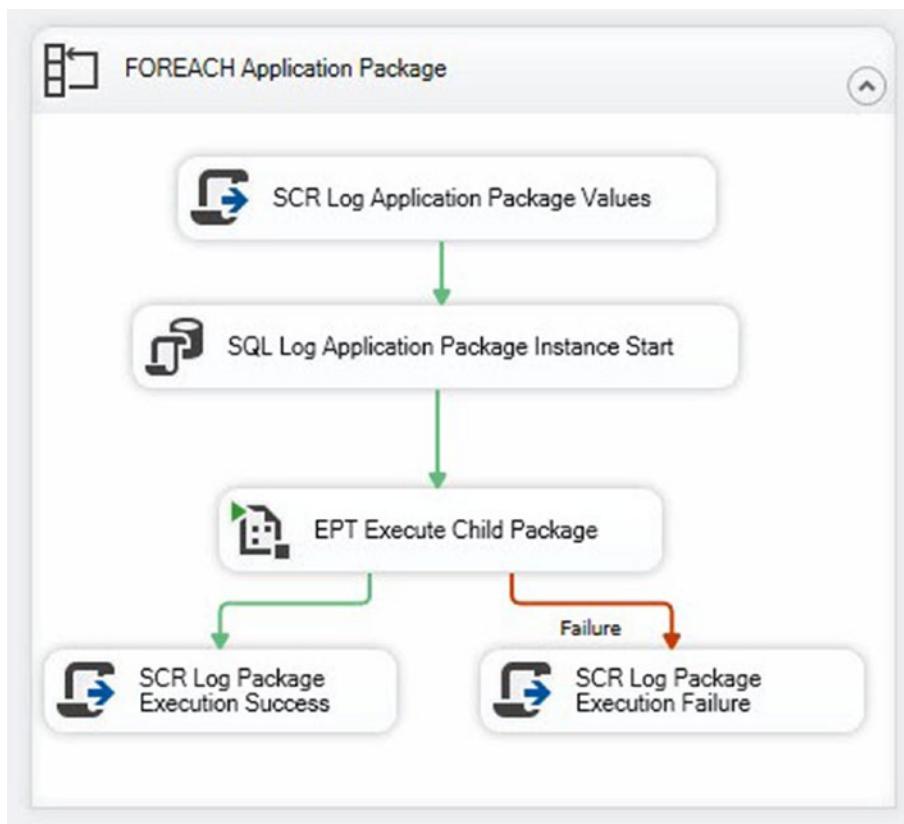


Figure 7-26. SQL Log Application Package Instance Start execute SQL task configured

Drag an Execute SQL Task onto the Parent.dtsx control flow below the SCR Log Package Execution Success script task in the FOREACH Application Package foreach loop container. Rename the Execute SQL Task “SQL Log Application Package Instance Success,” as shown in Figure 7-27.

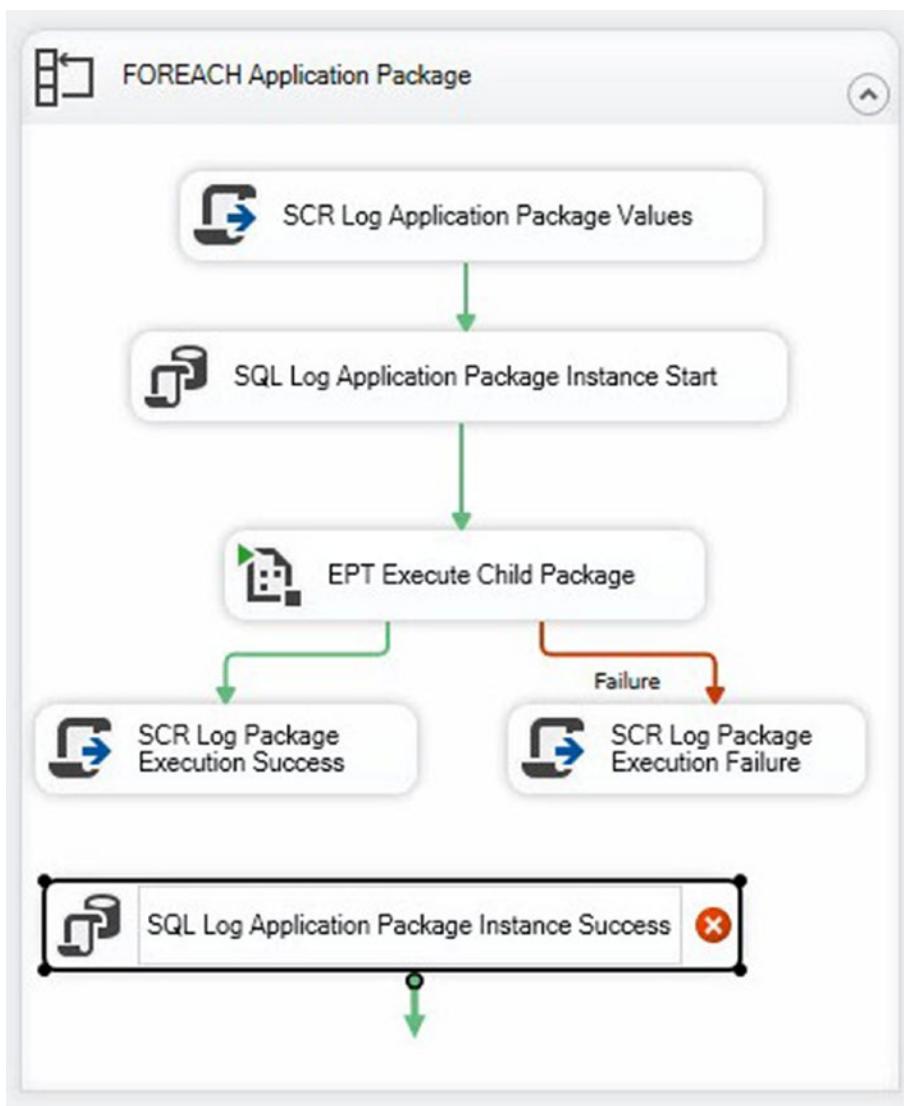


Figure 7-27. Adding the SQL Log Application Package Instance Success execute SQL task

Connect a precedence constraint from the SCR Log Package Execution Success script task to the SQL Log Application Package Instance Success execute SQL task. Open the SQL Log Application Package Instance Success execute SQL task editor. Change the ConnectionType property to ADO.NET, and select the SSISConfig Connection. Click the ellipsis in the SQLStatement property. When the Enter SQL Query dialog displays, enter the T-SQL shown in Listing 7-8.

Listing 7-8. T-SQL to update the [log].[ApplicationPackageInstance] table when the SSIS application package succeeds

```
Update [log].ApplicationPackageInstance
Set ApplicationPackageEndTime = sysdatetimeoffset()
, ApplicationPackageStatus = 'Succeeded'
Where ApplicationPackageInstanceId = @ApplicationPackageInstanceId
```

The Enter SQL Query dialog should appear as shown in Figure 7-28.

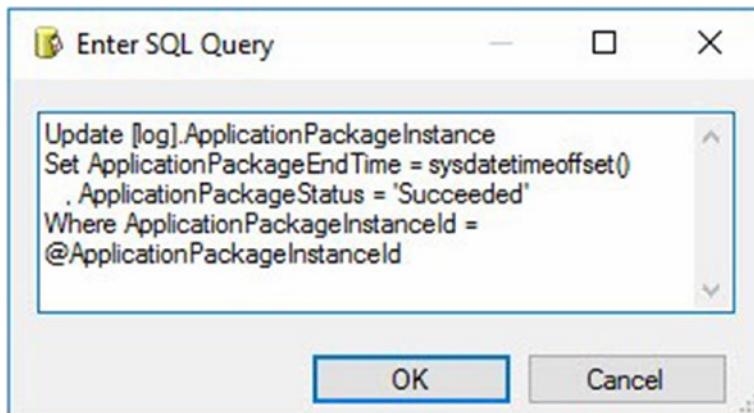


Figure 7-28. Adding T-SQL to update the [log].[ApplicationPackageInstance] table

Click the OK button to close the Enter SQL Query dialog.

Click the Parameter Mapping page, and then click the Add button. Map the Variable Name “User::ApplicationPackageInstanceId” to the Parameter Name “ApplicationPackageInstanceId,” as shown in Figure 7-29.

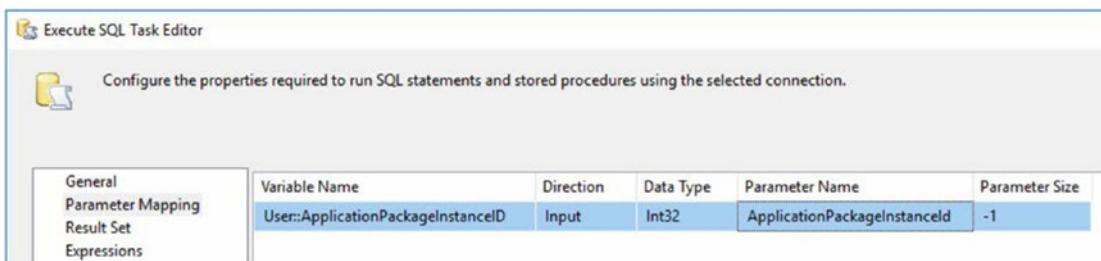


Figure 7-29. Mapping the User:ApplicationPackageInstanceId variable to the ApplicationPackageInstanceId parameter

Click the OK button to close the Execute SQL Task Editor.

The SQL Log Application Package Instance Success execute SQL task is now configured to update the row in the [log].[ApplicationPackageInstance] table that was inserted when the SQL Log Application Package Instance Start execute SQL task executed earlier in the FOREACH Application Package foreach loop container, updating the row's ApplicationPackageEndTime and ApplicationPackageStatus fields to indicate the SSIS framework application package instance completed successfully at such and such a time.

The next step is to update the ApplicationPackageEndTime and ApplicationPackageStatus fields in the row in [log].[ApplicationPackageInstance] when the SSIS framework application package instance *fails*.

Drag an Execute SQL Task onto the Parent.dtsx control flow below the SCR Log Package Execution Failure script task in the FOREACH Application Package foreach loop container. Rename the Execute SQL Task “SQL Log Application Package Instance Failure,” as shown in Figure [7-30](#).

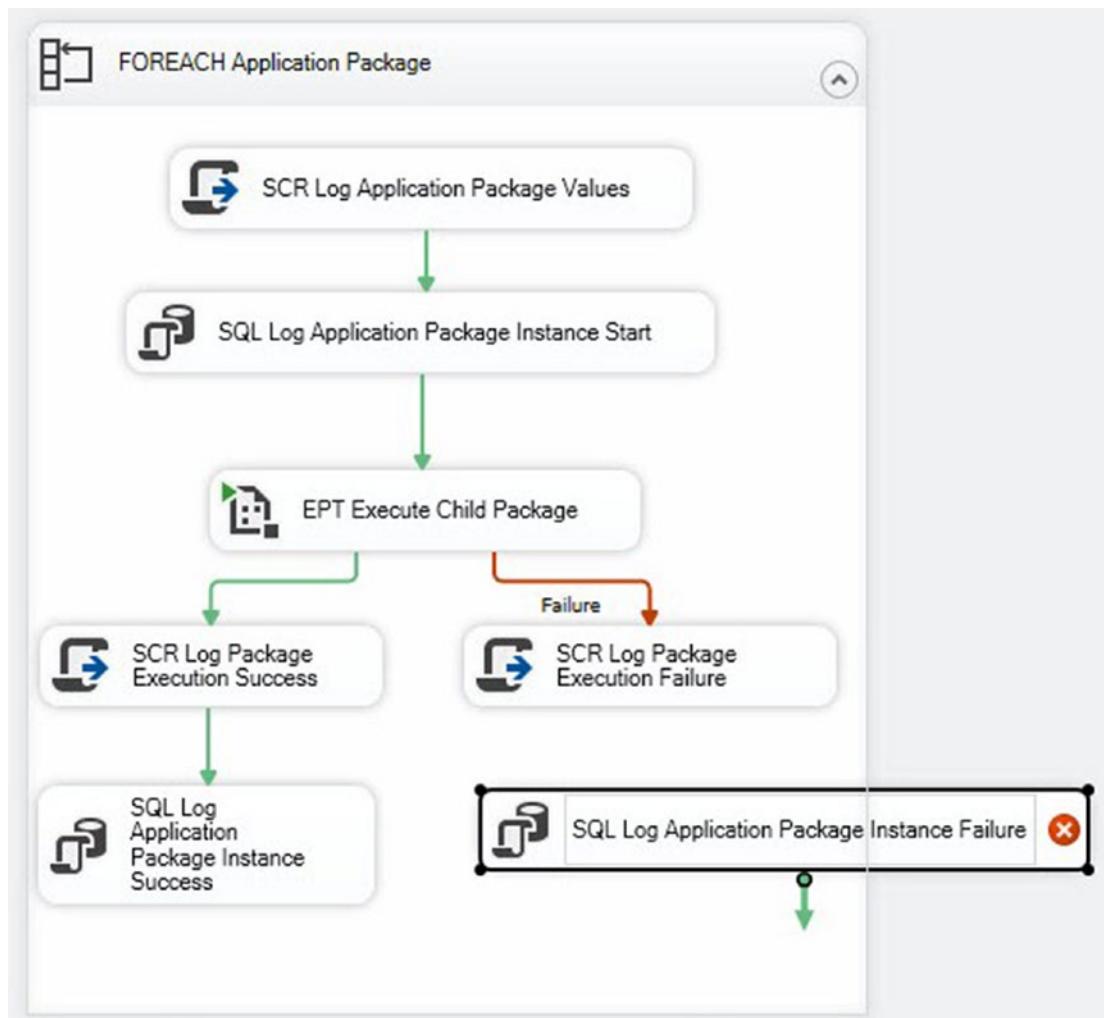


Figure 7-30. Adding the SQL Log Application Package Instance Failure execute SQL task

Connect a precedence constraint from the SCR Log Package Execution Failure script task to the SQL Log Application Package Instance Failure execute SQL task. Open the SQL Log Application Package Instance Failure execute SQL task editor. Change the ConnectionType property to ADO.NET, and select the SSISConfig Connection. Click the ellipsis in the SQLStatement property. When the Enter SQL Query dialog displays, enter the T-SQL shown in Listing 7-9.

Listing 7-9. T-SQL to update [log].[ApplicationPackageInstance] when the SSIS application package fails

```
Update [log].ApplicationPackageInstance
Set ApplicationPackageEndTime = sysdatetimeoffset()
    , ApplicationPackageStatus = 'Failed'
Where ApplicationPackageInstanceId = @ApplicationPackageInstanceId
```

The Enter SQL Query dialog should appear as shown in Figure 7-31.

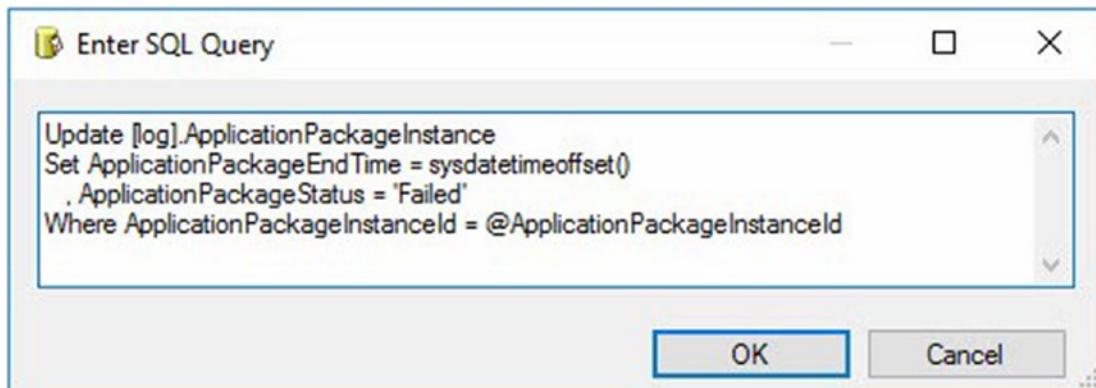


Figure 7-31. Updating the [log].[ApplicationPackageInstance] row

Click the OK button to close the Enter SQL Query dialog.

Click the Parameter Mapping page, and then click the Add button. Map the Variable Name “User::ApplicationPackageInstanceId” to the Parameter Name “ApplicationPackageInstanceId,” as shown in Figure 7-32.

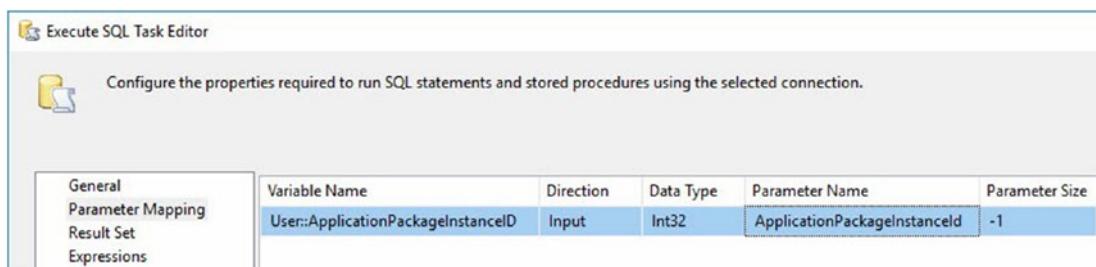


Figure 7-32. Mapping the User:ApplicationPackageInstanceId variable to the ApplicationPackageInstanceId parameter

Click the OK button to close the Execute SQL Task Editor.

The SQL Log Application Package Instance Failure execute SQL task is now configured to update the row in the [log].[ApplicationPackageInstance] table that was inserted when the SQL Log Application Package Instance Start execute SQL task executed earlier in the FOREACH Application Package foreach loop container, updating the row's ApplicationPackageEndTime and ApplicationPackageStatus fields to indicate the SSIS framework application package instance execution failed at such and such a time.

The Parent.dtsx SSIS package's FOREACH Application Package foreach loop container should appear similar to Figure 7-33.

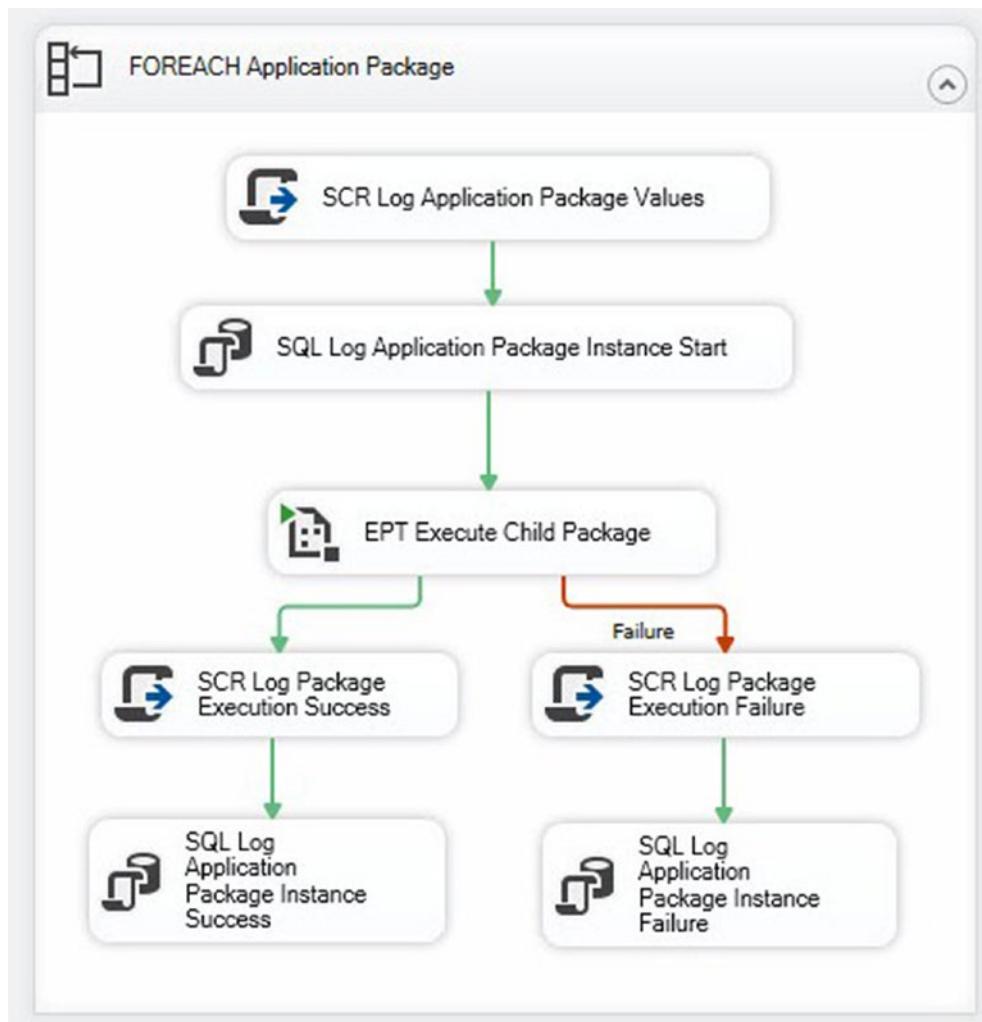


Figure 7-33. The FOREACH Application Package foreach loop container configured to capture SSIS framework application package instrumentation

Let's test it! Start the Parent.dtsx SSIS package in the debugger (press the F5 key). As currently configured, the SSIS framework application should execute and fail as shown in Figure 7-34.

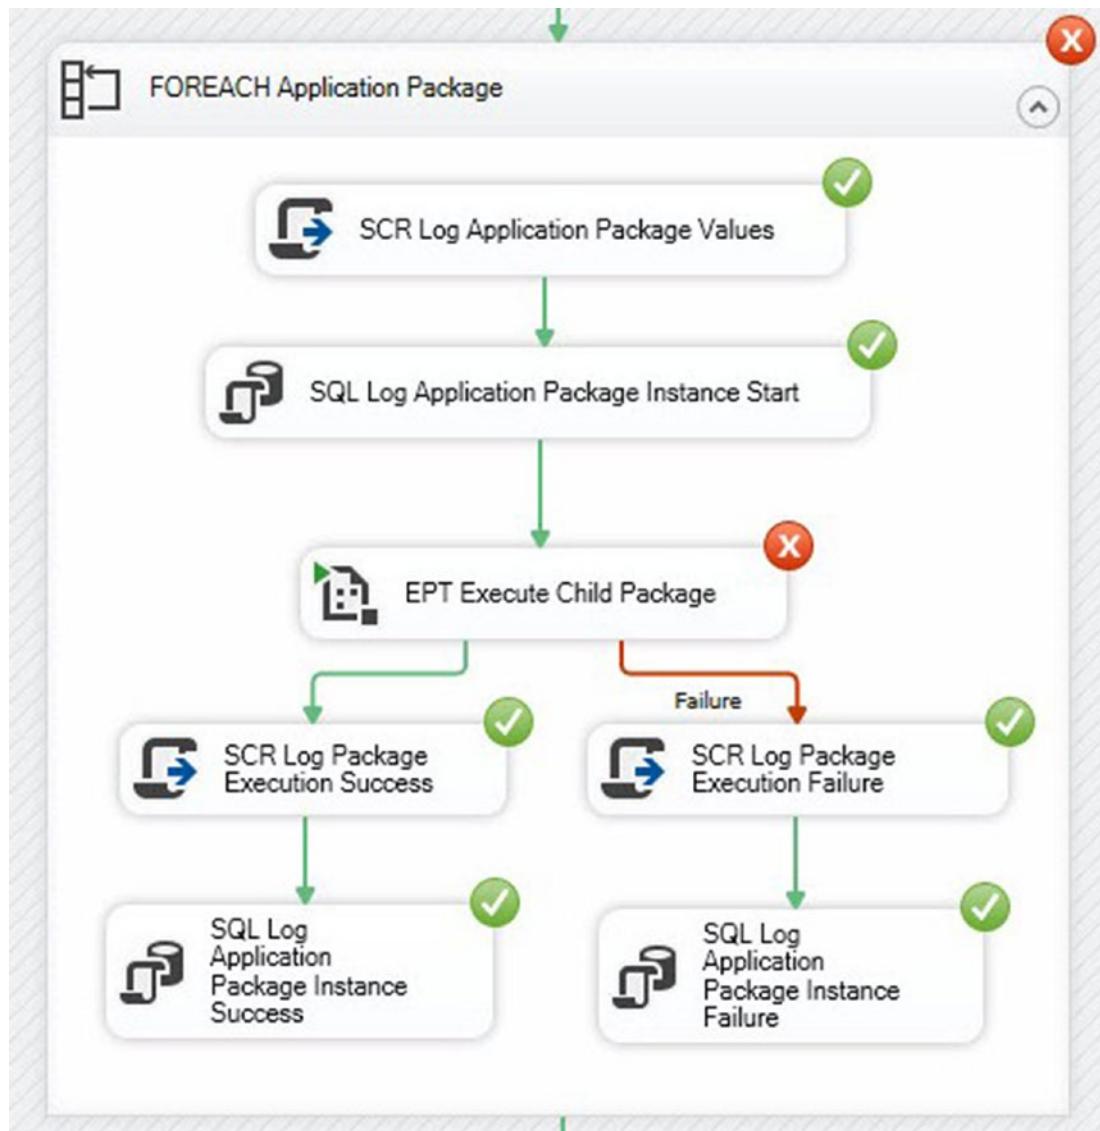


Figure 7-34. Viewing the execution of the FOREACH Application Package foreach loop container

The EPT Execute Child Package execute package task displayed in Figure 7-34 indicates failure. Two very important questions are

1. Did the execution test succeed?
2. Did the test execution succeed?

The answer to both questions lies in the log tables.

Viewing Execution Report Data

SSIS framework logging is stored in two tables name [log].[ApplicationInstance] and [log].[ApplicationPackageInstance]. Each row in the [log].[ApplicationInstance] table contains data about a single execution of an SSIS framework application or an *instance* of an application execution. View application instance data using the T-SQL query shown in Listing 7-10.

Listing 7-10. Viewing application instance log data

```
Select a.ApplicationName
      , ai.ApplicationStartTime
      , DateDiff(ms, ai.ApplicationStartTime, ai.ApplicationEndTime) As
ApplicationRunMilliSeconds
      , ai.ApplicationStatus
From [SSISConfig].[log].[ApplicationInstance] ai
Join [SSISConfig].[config].[Applications] a
  On a.ApplicationId = ai.ApplicationId
Order By ai.ApplicationStartTime Desc
```

Results from the application instance query will appear similar to those shown in Figure 7-35.

Results	Messages			
	ApplicationName	ApplicationStartTime	ApplicationRunMilliSeconds	ApplicationStatus
1	Framework Test	2020-05-22 08:50:06.9532116 ...	351	Failed
2	Framework Test	2020-05-21 09:42:47.1302552 ...	317	Failed
3	Framework Test	2020-05-19 16:32:07.6675319 ...	209	Failed

Figure 7-35. Viewing application instance results

In answer to our first question – “Did the execution test succeed?” – the results of the SSIS application query indicate three test application instance executions are recorded, which means, “Yes, the execution tests succeeded three times.”

Examining the application package instance data is next. Execute the T-SQL shown in Listing 7-11 to view application package instance data.

Listing 7-11. Viewing application package instance log data

```
Select a.ApplicationName
    , p.PackageName
    , api.ApplicationPackageStartTime
    , ai.ApplicationStatus
    , api.ApplicationPackageStatus
    , api.ApplicationPackageStartTime
    , ap.FailApplicationOnPackageFailure
    , DateDiff(ms, api.ApplicationPackageStartTime, ► api.
ApplicationPackageEndTime) As ApplicationPackageRunMilliSeconds
    , ai.ApplicationStartTime
    , DateDiff(ms, ai.ApplicationStartTime, ai.ApplicationEndTime) As ►
ApplicationRunMilliSeconds
From [SSISConfig].[log].[ApplicationPackageInstance] api
Join [SSISConfig].[log].[ApplicationInstance] ai
    On ai.ApplicationInstanceId = api.ApplicationInstanceId
Join [SSISConfig].[config].[ApplicationPackages] ap
    On ap.ApplicationPackageId = api.ApplicationPackageId
Join [SSISConfig].[config].[Applications] a
    On a.ApplicationId = ap.ApplicationId
Join [SSISConfig].[config].[Packages] p
    On p.PackageId = ap.PackageId
Order By api.ApplicationPackageStartTime Desc
```

Results from the application package instance query will appear similar to those shown in Figure 7-36.

Results		Messages					
ApplicationName	PackageName	ApplicationPackageStartTime	ApplicationStatus	ApplicationPackageStatus	ApplicationPackageEndTime	FailApplicationOnPackageFailure	ApplicationPackageRunTime
1 Framework Test	ReportAndFail.dtsx	2020-05-22 08:50:07.1402...	Failed	Failed	2020-05-22 08:50:07.140275...	1	216
2 Framework Test	ReportAndSucceed.dtsx	2020-05-22 08:50:07.0141...	Failed	Succeeded	2020-05-22 08:50:07.014190...	1	182
3 Framework Test	ReportAndFail.dtsx	2020-05-21 09:42:47.3443...	Failed	Failed	2020-05-21 09:42:47.344349...	1	183
4 Framework Test	ReportAndSucceed.dtsx	2020-05-21 09:42:47.1952...	Failed	Succeeded	2020-05-21 09:42:47.195273...	1	121

Figure 7-36. Viewing application package instance results

In answer to our second question – “Did the test execution succeed?” – the results of the SSIS application package query indicate two test application instance executions have failed, and each test execution of the ReportAndFail.dtsx application package failed, which means, “Yes, the test executions succeeded two times.” Wait a minute. Why are two failures considered a success? Two failures are considered a success because the ReportAndFail.dtsx application package is designed to fail and the FailApplicationOnPackageFailure bit is set to true (1).

Test the assertion by setting the FailApplicationOnPackageFailure bit for the ReportAndFail.dtsx application package to False (0) using the T-SQL in Listing 7-12.

Listing 7-12. Resetting the FailApplicationOnPackageFailure bit for ReportAndFail.dtsx

```
Select p.PackageName
    , ap.FailApplicationOnPackageFailure
From [SSISConfig].[config].[ApplicationPackages] ap
Join [SSISConfig].[config].[Packages] p
    On p.PackageId = ap.PackageId
Where p.PackageName = N'ReportAndFail.dtsx'

Update ap
Set ap.FailApplicationOnPackageFailure = 0
From [SSISConfig].[config].[ApplicationPackages] ap
Join [SSISConfig].[config].[Packages] p
    On p.PackageId = ap.PackageId
Where p.PackageName = N'ReportAndFail.dtsx'

Select p.PackageName
    , ap.FailApplicationOnPackageFailure
From [SSISConfig].[config].[ApplicationPackages] ap
Join [SSISConfig].[config].[Packages] p
    On p.PackageId = ap.PackageId
Where p.PackageName = N'ReportAndFail.dtsx'
```

Note the Select statements surface the values of the FailApplicationOnPackageFailure bit values before and after the Update statement. Strive to communicate intent in T-SQL statements; “Commands completed successfully” is *not* enough feedback.

Results from the application package update will appear similar to those shown in Figure 7-37.

Results		Messages
	PackageName	FailApplicationOnPackageFailure
1	ReportAndFail.dtsx	1
<hr/>		
	PackageName	FailApplicationOnPackageFailure
1	ReportAndFail.dtsx	0

Figure 7-37. Viewing application package update results

Executing Parent.dtsx in the SSIS debugger reveals *no change* in SSIS execution results, as shown in Figure 7-38.

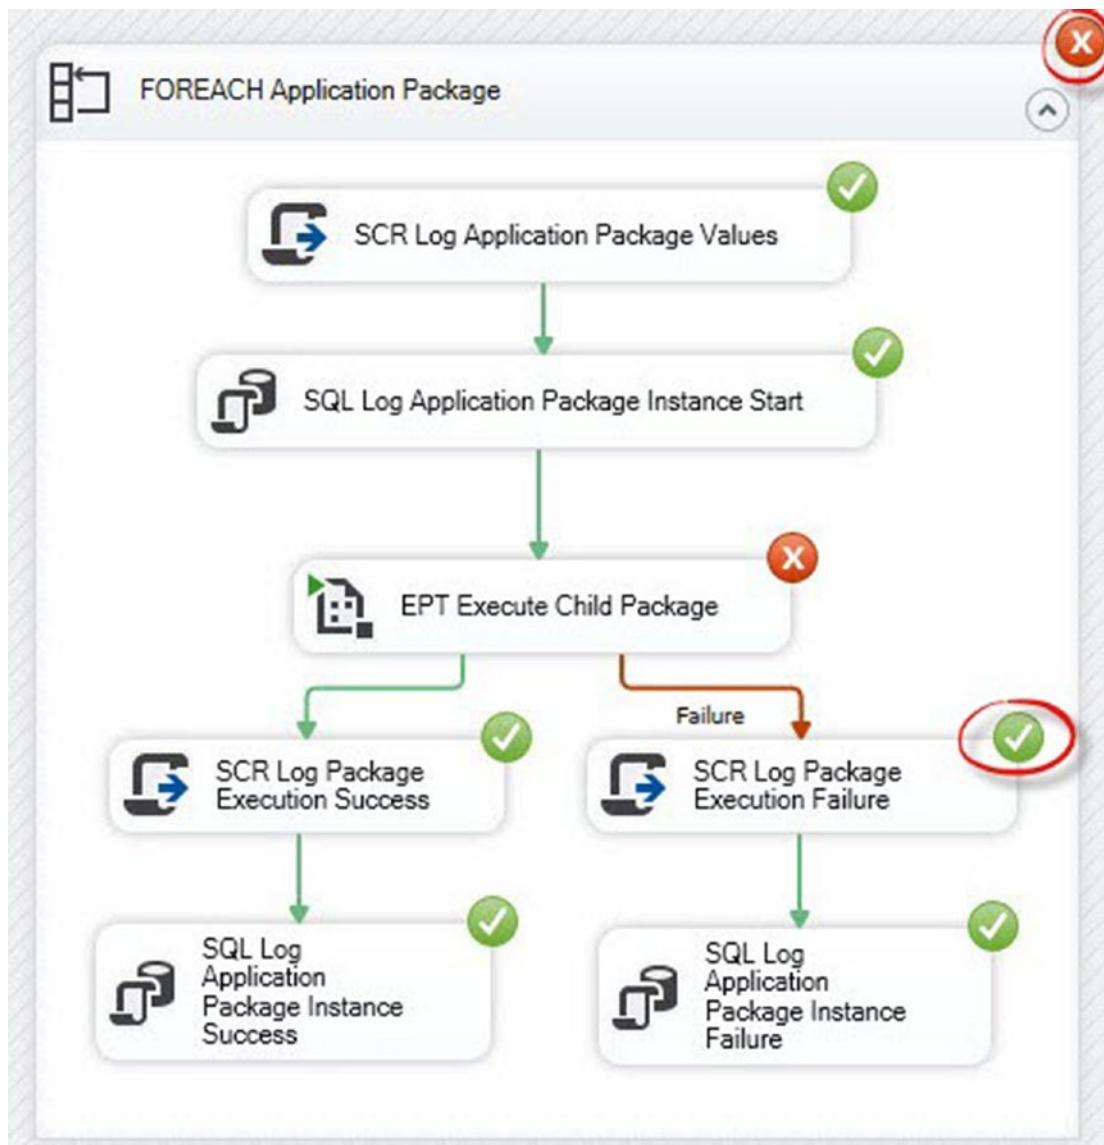


Figure 7-38. Parent.dtsx execution fails, again s

What gives? Fault tolerance for Parent.dtsx is not properly configured. Before we properly configure the fault tolerance for Parent.dtsx SSIS package, please note the FOREACH Application Package foreach loop container failed, indicating fault tolerance is not properly configured. Note also the SCR Log Package Execution Failure script task succeeded, indicating the FailApplicationOnPackageFailure logic – shown in Listing 7-13 – is working.

Listing 7-13. FailApplicationOnPackageFailure logic

```
public void Main()
{
    // System::PackageName, System::TaskName
    // User::FailApplicationOnPackageFailure, User::PackagePath

    string packageName = ► Dts.Variables["System::PackageName"].Value.
        ToString();
    string taskName = Dts.Variables["System::TaskName"].Value.ToString();
    string subComponent = packageName + "." + taskName;
    int informationCode = 1001;
    int errorCode = -999;
    bool fireAgain = true;

    bool failApplicationOnPackageFailure = ► Convert.ToBoolean(Dts.Variabl
        es["User::FailApplicationOnPackageFailure"] ►
    .Value);

    string packagePath = Dts.Variables["User::PackagePath"].Value.
        ToString();
    string description = String.Empty;

    if(failApplicationOnPackageFailure)
    {
        description = packagePath + " execution failed and ►
            FailApplicationOnPackageFailure is set (true)";
        Dts.Events.FireError(errorCode, subComponent, description, "", 0);
    }
    else
    {
        description = packagePath + " execution failed and ►
            FailApplicationOnPackageFailure is NOT set (false)";
        Dts.Events.FireInformation(informationCode, subComponent, ►
            description, "", 0, ref fireAgain);
    }

    Dts.TaskResult = (int)ScriptResults.Success;
}
```

The Progress tab message from the SCR Log Package Execution Failure script task confirms FailApplicationOnPackageFailure functionality is working as designed, as shown in Figure 7-39.

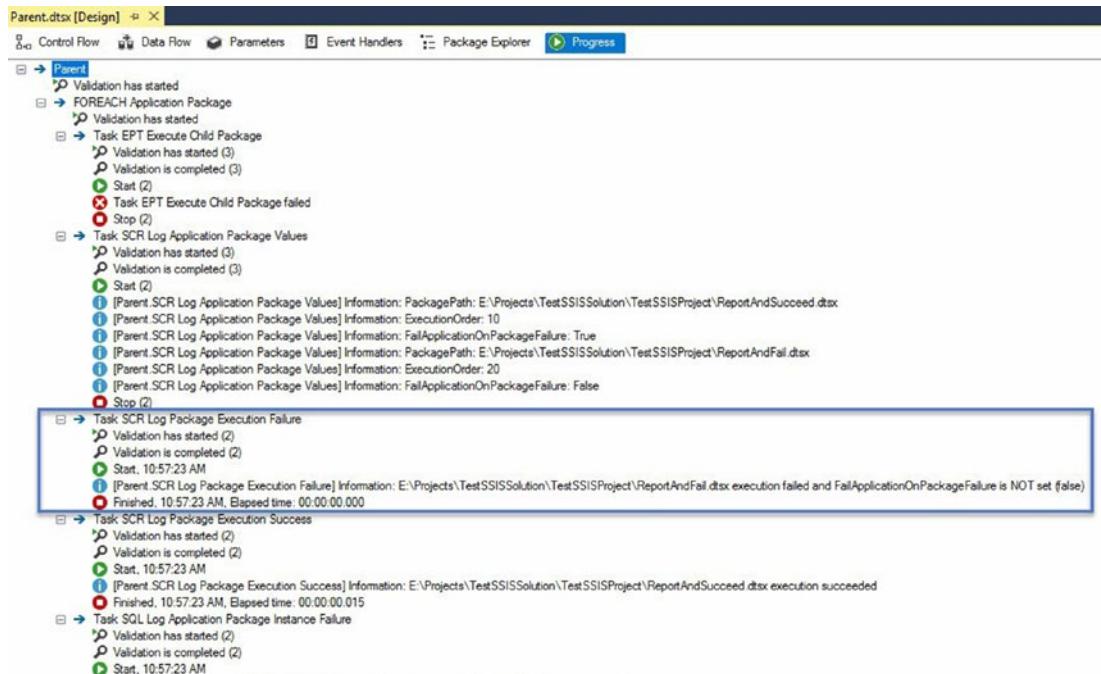


Figure 7-39. FailApplicationOnPackageFailure is working

The next step is to configure fault tolerance for the Parent.dtsx SSIS package.

Configure Fault Tolerance in SSIS Framework Application Metadata

Begin configuring fault tolerance for the Parent.dtsx SSIS package so that execution fails if and only if two conditions are met:

1. An application package execution fails.
2. The FailApplicationOnPackageFailure bit is set to true for the application package.

If one or both of these conditions are false, the desired behavior is that Parent.dtsx continues executing. To configure the *always-keep-running* functionality, click the FOREACH Application Package foreach loop container, and press the F4 key to display its properties, as shown in Figure 7-40.

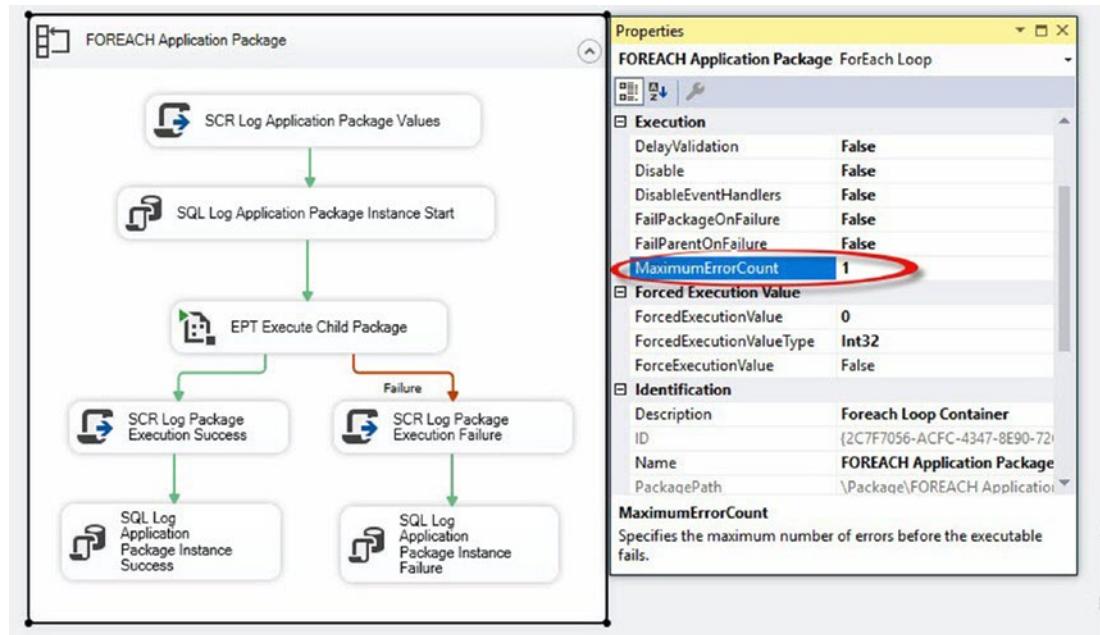


Figure 7-40. Displaying properties of the FOREACH Application Package foreach loop container

In Figure 7-40, the MaximumErrorCount property is highlighted. The net effect of the default MaximumErrorCount property value of 1 is as follows: if one error occurs within the FOREACH Application Package foreach loop container at execution time, the FOREACH Application Package foreach loop container execution will fail.

The FOREACH Application Package foreach loop container's MaximumErrorCount property may be increased but, if so, to what new value? Suppose 10, 20, or 300 application packages are configured to execute as part of an SSIS framework application. While there is actually a way to manage this very scenario, there exists a more elegant solution.

A more elegant solution is to ignore *all* the execution errors. Change the FOREACH Application Package foreach loop container's MaximumErrorCount property value to 0 to configure the FOREACH Application Package foreach loop container to ignore all errors during execution.

Credit is due to Julie Smith (@juliechix on Twitter) for teaching me about setting MaximumErrorCount to 0.

Because events *bubble* in SSIS, any error that occurs inside the FOREACH Application Package foreach loop container will “bubble up” to the Parent.dtsx SSIS package, which is the FOREACH Application Package foreach loop container’s container. Click anywhere in the whitespace of the Parent.dtsx’s control flow to display the Parent.dtsx package’s properties, and change the MaximumErrorCount property to 0, as shown in Figure 7-41.

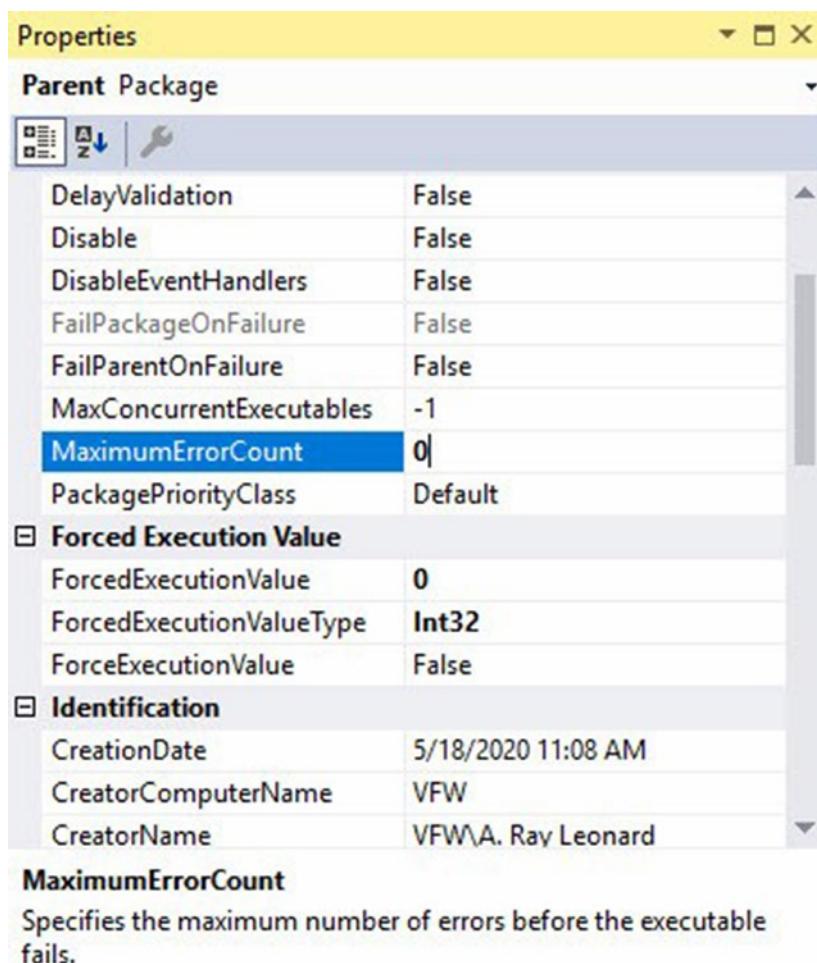


Figure 7-41. Changing the Parent.dtsx package’s MaximumErrorCount property

As *currently* configured, SSIS framework application packages will continue to execute until the last application package is executed, regardless of whether the application package instance fails *and* regardless of the setting of the application package's FailApplicationOnPackageFailure bit.

The next step is to configure the Parent.dtsx SSIS package to *stop* executing if and only if two conditions are met:

1. An application package execution fails.
2. The FailApplicationOnPackageFailure bit is set to true for the application package.

Click the SCR Log Package Execution Failure script task. In properties, change the SCR Log Package Execution Failure script task's FailPackageOnFailure property to True, as shown in Figure 7-42.

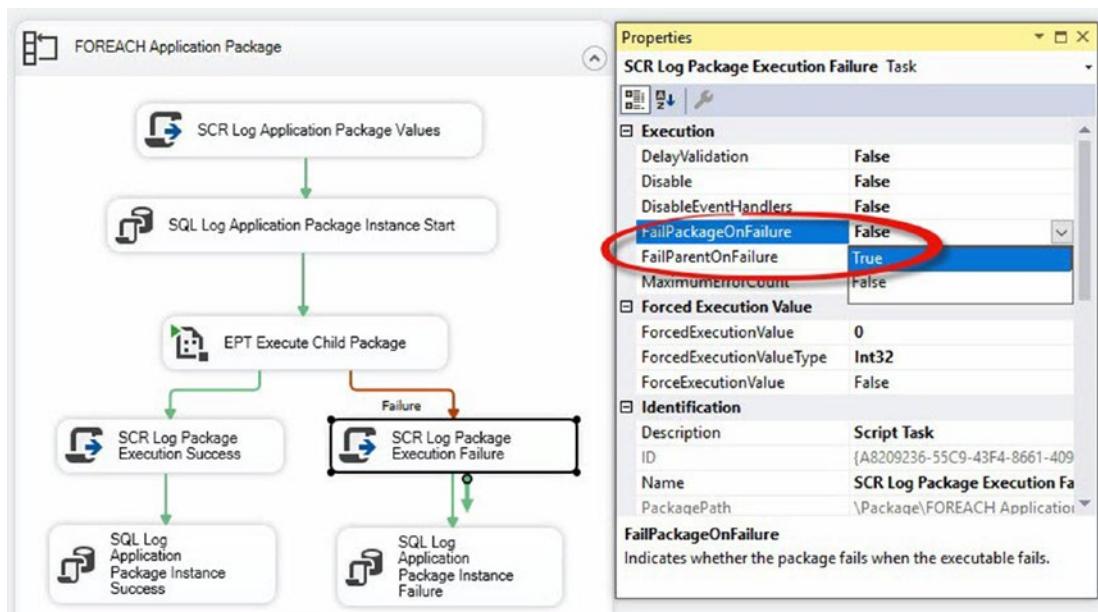


Figure 7-42. Changing the SCR Log Package Execution Failure script task's FailPackageOnFailure property to True

Testing execution now reveals the desired functionality; SSIS framework application package execution instances behave as configured, and fault tolerance is working as designed, as shown in Figure 7-43.

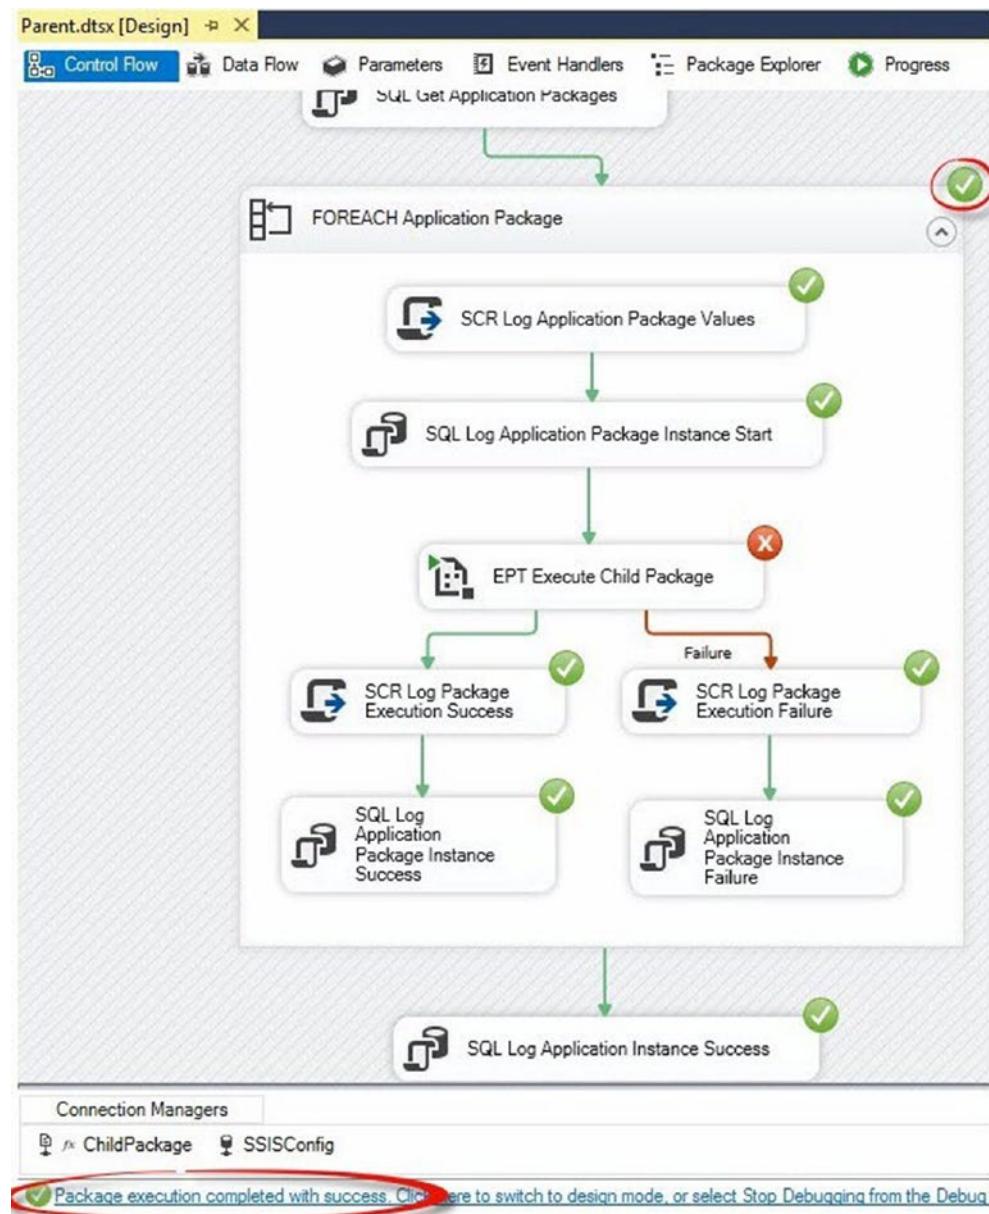


Figure 7-43. SSIS application package fault tolerance working as designed

Remember, the ReportAndFail.dtsx application package execution *should* fail, but the failure of a ReportAndFail.dtsx application package execution instance should *not* cause the application to fail (in its current configuration).

Update the FailApplicationOnPackageFailure bit configuration for the ReportAndFail.dtsx application package using the T-SQL in Listing 7-14.

Listing 7-14. Updating the FailApplicationOnPackageFailure bit for ReportAndFail.dtsx

```

Select p.PackageName
      , ap.FailApplicationOnPackageFailure
From [SSISConfig].[config].[ApplicationPackages] ap
Join [SSISConfig].[config].[Packages] p
  On p.PackageId = ap.PackageId
Where p.PackageName = N'ReportAndFail.dtsx'

Update ap
Set ap.FailApplicationOnPackageFailure = 1
From [SSISConfig].[config].[ApplicationPackages] ap
Join [SSISConfig].[config].[Packages] p
  On p.PackageId = ap.PackageId
Where p.PackageName = N'ReportAndFail.dtsx'

Select p.PackageName
      , ap.FailApplicationOnPackageFailure
From [SSISConfig].[config].[ApplicationPackages] ap
Join [SSISConfig].[config].[Packages] p
  On p.PackageId = ap.PackageId
Where p.PackageName = N'ReportAndFail.dtsx'

```

Updating the SSIS framework application package metadata for the ReportAndFail.dtsx application package results appears as shown in Figure 7-44.

Results		Messages
	PackageName	FailApplicationOnPackageFailure
1	ReportAndFail.dtsx	0
	PackageName	FailApplicationOnPackageFailure
1	ReportAndFail.dtsx	1

Figure 7-44. Update ReportAndFail.dtsx application package metadata results

Re-executing Parent.dtsx confirms we can fail if we want to, as shown in Figure 7-45.

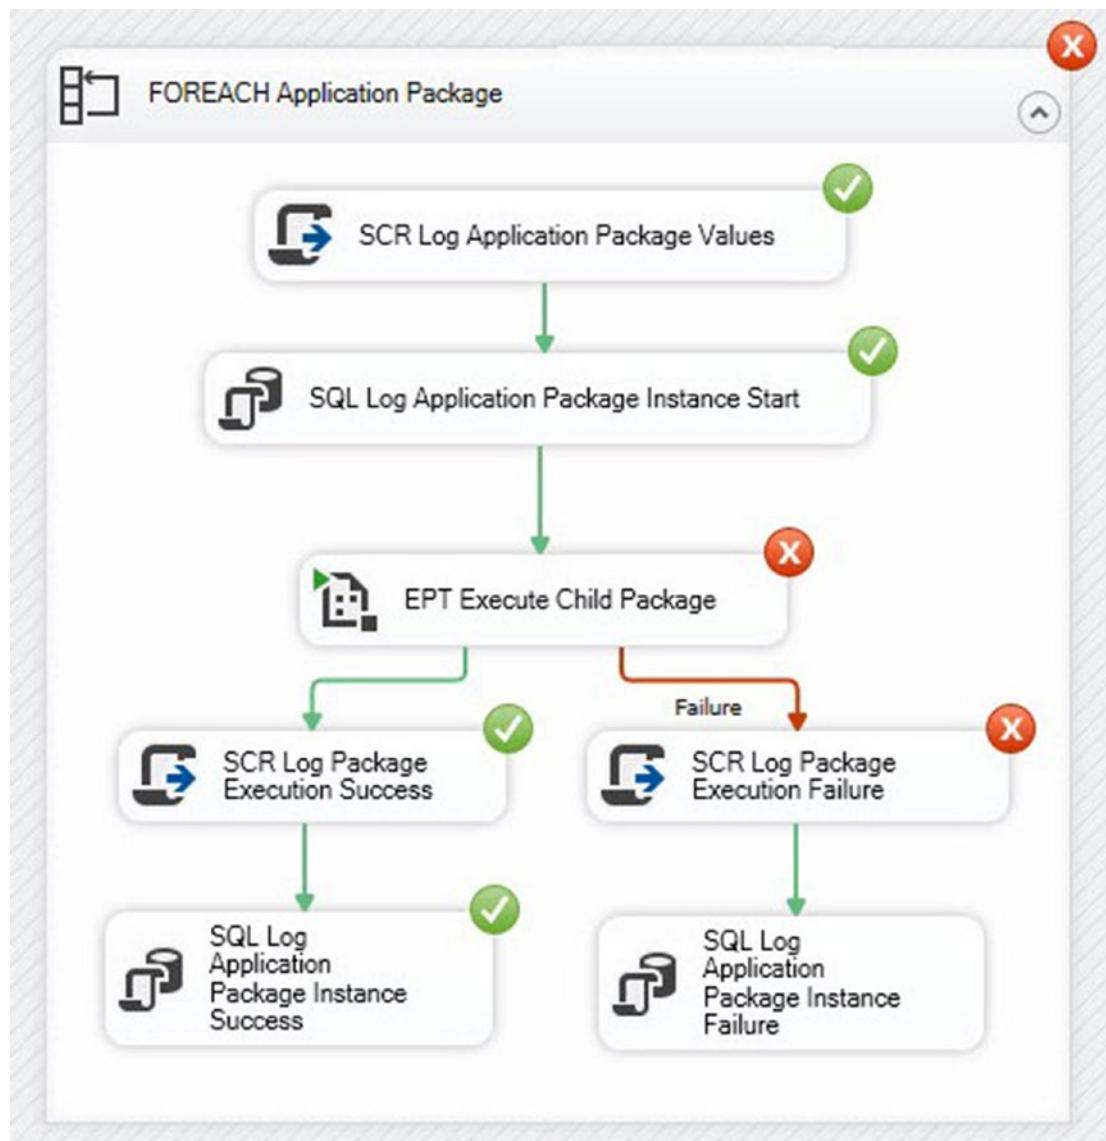


Figure 7-45. Application Package Failure managed by SSIS framework metadata

The current state of the Parent.dtsx SSIS package represents a fairly robust execution engine for the SSIS framework.

One issue remains, however: re-executing the query shown in Listing 7-11 returns the ApplicationPackageStatus result “Running,” similar to the results shown in Figure 7-46.

ApplicationName	PackageName	ApplicationPackageStartTime	ApplicationStatus	ApplicationPackageStatus
Framework Test	ReportAndFail.dtsx	2020-05-22 11:55:16.8872218 ...	Failed	Running

Figure 7-46. ApplicationPaackageStatus is “Running”

Why is the ApplicationPackageStatus “Running”? View the screenshot in Figure 7-45, and note the “SQL Log Application Package Instance Failure” execute SQL task *did not execute* because the precedence constraint connecting the “SCR Log Package Execution Failure” script task to the “SQL Log Application Package Instance Failure” execute SQL task is configured to evaluate on success *only*. Remember, looking at Listing 7-13, an application package failure may result in success *or* failure of the “SQL Log Application Package Instance Failure” execute SQL task, depending on the application package’s FailApplicationOnPackageFailure bit value in the [config].[ApplicationPackages] table.

Listing 7-14 and Figure 7-44 show we set the FailApplicationOnPackageFailure bit to 1 (true), which informed the “SCR Log Package Execution Failure” script task. The “SCR Log Package Execution Failure” script task then raised an error, as prescribed by the .Net C# code in Listing 7-13.

Fortunately, the fix is simple and straightforward – delete the precedence constraint connecting the “SCR Log Package Execution Failure” script task to the “SQL Log Application Package Instance Failure” execute SQL task. Connect a new *failure* precedence constraint from the “EPT Execute Child Package” execute package task to the “SQL Log Application Package Instance Failure” execute SQL task. Re-executing Parent.dtsx in the debugger results in execution of the “SQL Log Application Package Instance Failure” execute SQL task, as shown in Figure 7-47.

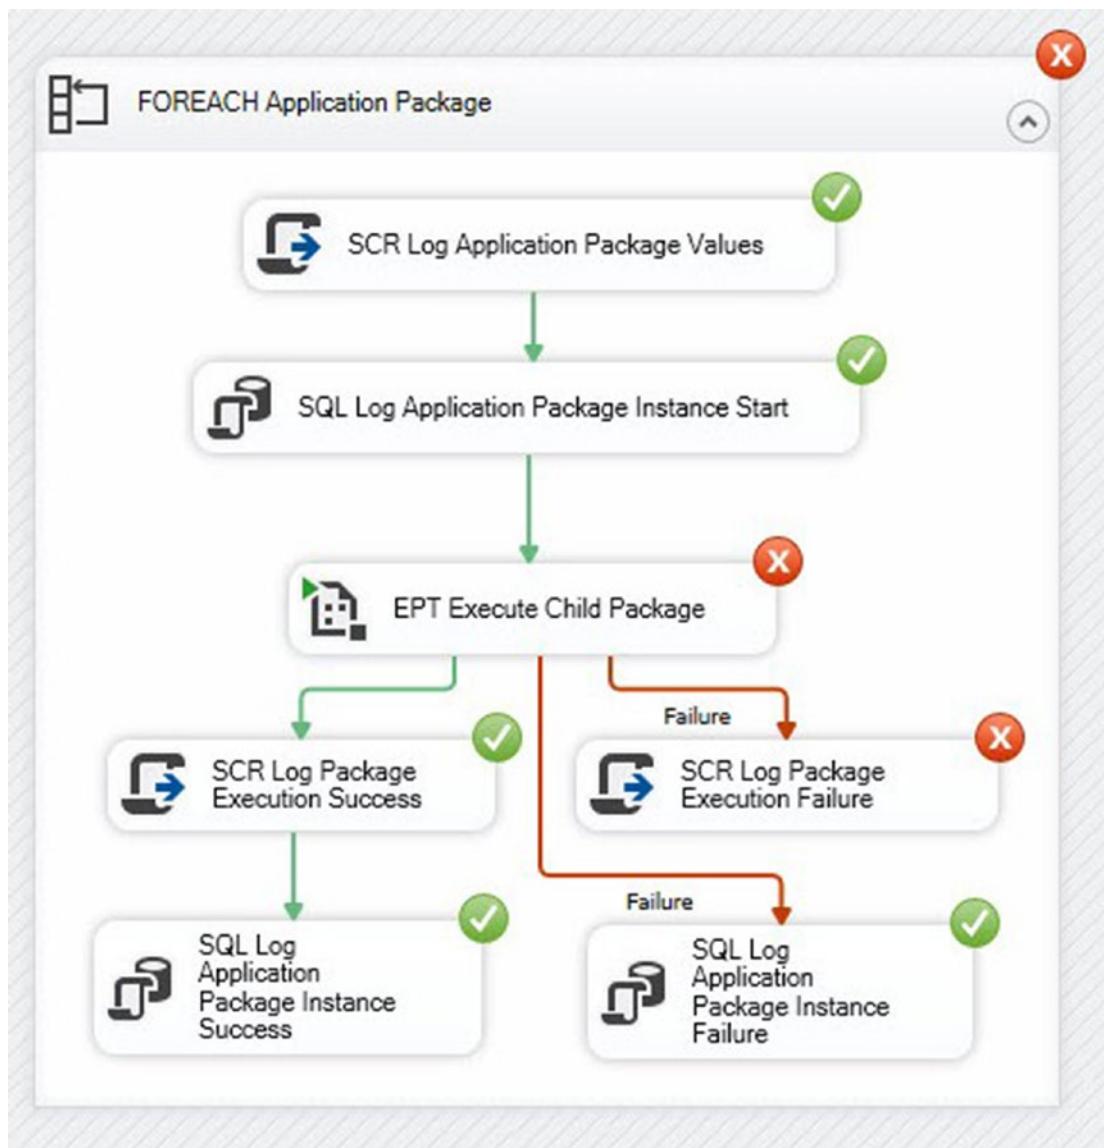


Figure 7-47. The “SQL Log Application Package Instance Failure” execute SQL task executes

Re-executing the application package status query in Listing 7-11 returns the ApplicationPackageStatus result “Failed,” similar to the results shown in Figure 7-48.

CHAPTER 7 FRAMEWORK LOGGING

ApplicationName	PackageName	ApplicationPackageStartTime	ApplicationStatus	ApplicationPackageStatus
Framework Test	ReportAndFail.dtsx	2020-06-01 13:05:16.7096071 ...	Failed	Failed

Figure 7-48. ApplicationPackageStatus is “Failed”

Our SSIS framework execution engine is now ready for deployment elsewhere.

Conclusion

This chapter covered SSIS framework logging that extended the instrumentation we built in Chapter 6. The next few chapters cover migrating – and adapting – this simple SSIS framework to additional scenarios, including Azure Data Factory.

CHAPTER 8

Azure-SSIS Integration Runtime

The Microsoft Azure-SSIS team has been hard at work reducing the friction between executing SSIS packages on-premises and executing SSIS packages in Azure. Many enterprises execute SSIS packages stored in the file system on-premises. In the past, migrating SSIS packages executed in an on-premises file system meant changing enterprise SSIS storage and execution patterns. With the advent of Azure-SSIS File Share-based execution, executing SSIS packages from files in Azure is no longer an issue.

In fact, the SSIS framework from the previous chapter works well in Azure-SSIS. In this chapter, we will discuss and demonstrate

- Getting Started with Azure
- Provisioning an Azure Data Factory
- Provisioning Azure Storage
- Provisioning Azure-SSIS

Getting Started with Azure

Before creating and interacting with resources in Azure, one must first create an Azure account at azure.com, as shown in Figure 8-1.

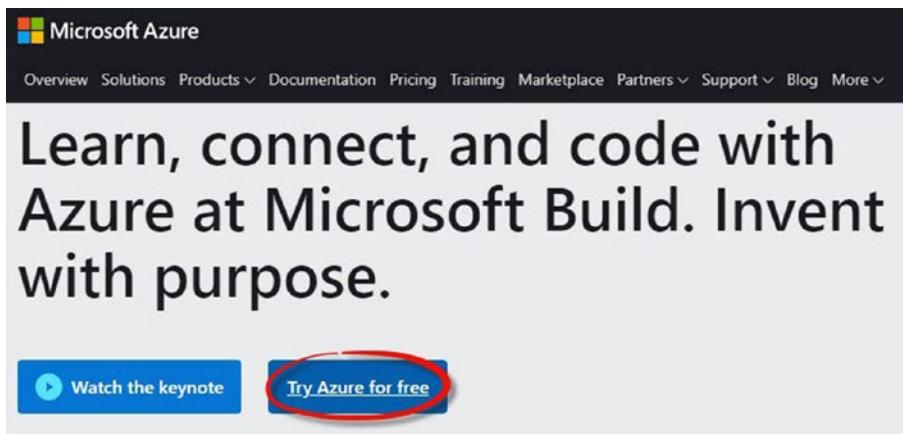


Figure 8-1. Azure.com

Please note the text on the link button in Figure 8-1 reads, “Try Azure for free.” To date, the good people at Microsoft have consistently offered free introductory access to some Azure services. Free introductory offers lie solely with Microsoft and are subject to changes at Microsoft’s discretion.

After an account has been created, the next step is to provision an Azure Data Factory.

Azure changes daily. Some of the screenshots and procedures in this book will be out of date before publication.

Provisioning an Azure Data Factory

Azure hosts lots of services. In general, Azure services may be categorized as

1. Infrastructure as a Service, or IaaS, in which Azure hosts enterprise infrastructure such as virtual machines (VM’s).
2. Platform as a Service, or PaaS, in which Azure surfaces enterprise platforms such as SQL Server or Azure Data Factory.

Azure users begin by *provisioning* – or creating instances of – Azure service offerings. Begin provisioning an instance of Azure Data Factory – or ADF – by clicking “Create a resource” in the left menu, as shown in Figure 8-2.

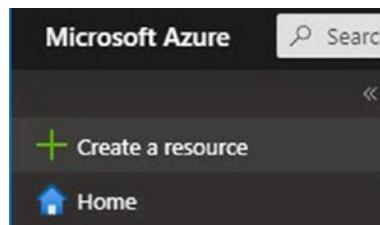


Figure 8-2. Create a resource in Azure's left menu

When the New page displays, search for “Data Factory,” as shown in Figure 8-3.

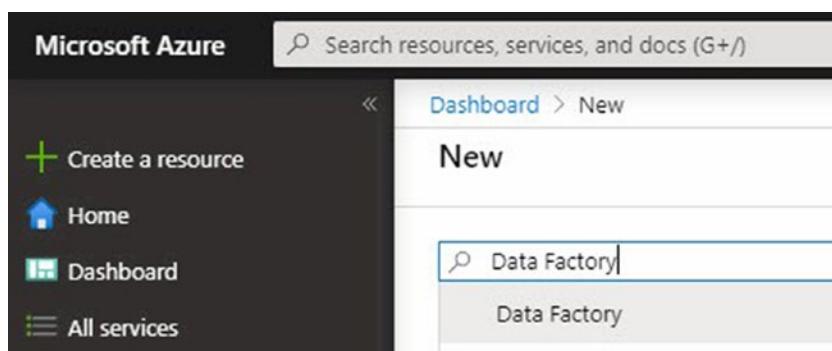


Figure 8-3. Searching for Data Factory on the New page

On the New page, click “Data Factory” beneath the search textbox on the New page to open the Data Factory page, as shown in Figure 8-4.

CHAPTER 8 AZURE-SSIS INTEGRATION RUNTIME

The screenshot shows the Microsoft Azure Data Factory page. At the top, there's a navigation bar with 'Dashboard > New > Data Factory'. Below it, the title 'Data Factory' is displayed with the Microsoft logo. A large blue button labeled 'Create' is prominent. To the left of the main content area is a sidebar with 'Overview' and 'Plans' tabs, where 'Overview' is selected. The main content area contains a brief description of Azure Data Factory and a bulleted list of features. At the bottom left, there's a 'Useful Links' section with links to 'Documentation', 'Service overview', and 'Pricing details'.

Microsoft Azure Data Factory is a cloud-based data integration service that automates the movement and transformation of data. You can quickly create, deploy, schedule, and monitor highly-available, fault tolerant data flow pipelines. Move and transform data of all shapes and sizes, and deliver the results to a range of destination storage services. Monitor all of your data pipelines and service health at a glance with a rich visual experience. Easily consume the data produced with BI, analytics tools, and other applications to drive key business insights and decisions.

- Compose data storage, movement and processing services into data flow pipelines
- Enhanced HDInsight integration including HCAT and on-demand cluster management
- Schedule data pipelines with fine-tuned control
- New data connectors for on-premises and cloud data sources
- Integration with Azure Machine Learning and Azure Batch
- Globally deployed data movement as a service
- Create, edit and deploy data pipelines with a Visual Studio plug-in

Useful Links
Documentation
Service overview
Pricing details

Figure 8-4. The Data Factory page

Click the Create button on the Data Factory page to open the New data factory page, as shown in Figure 8-5.

New data factory

The screenshot shows the 'New data factory' configuration page. It includes fields for Name (adfFrameworks), Version (V2), Subscription (Enterprise Data & Analytics), Resource Group (rgFrameworks), Location ((US) East US 2), and Enable GIT (unchecked). A 'Create' button is at the bottom.

Name *

adfFrameworks

Version ⓘ

V2

Subscription *

Enterprise Data & Analytics

Resource Group *

rgFrameworks

Create new

Location * ⓘ

(US) East US 2

Enable GIT ⓘ

Create

Figure 8-5. Adding a new instance of an Azure Data Factory

Enter values for the required fields on the New data factory page.

Azure Data Factory, Resource Group, and Storage Account names are globally unique and case insensitive, according to Azure Data Factory – naming rules (docs.microsoft.com/en-us/azure/data-factory/naming-rules).

Resource groups are one way to group related Azure resources. When using Azure resources to learn Azure, resource groups allow students to *delete* the resource group and all constituents from a single screen.

You may select an existing resource group from the Resource Group drop-down, or create a new resource group by clicking the “Create new” link beneath the Resource Group drop-down, as shown in Figure 8-6.

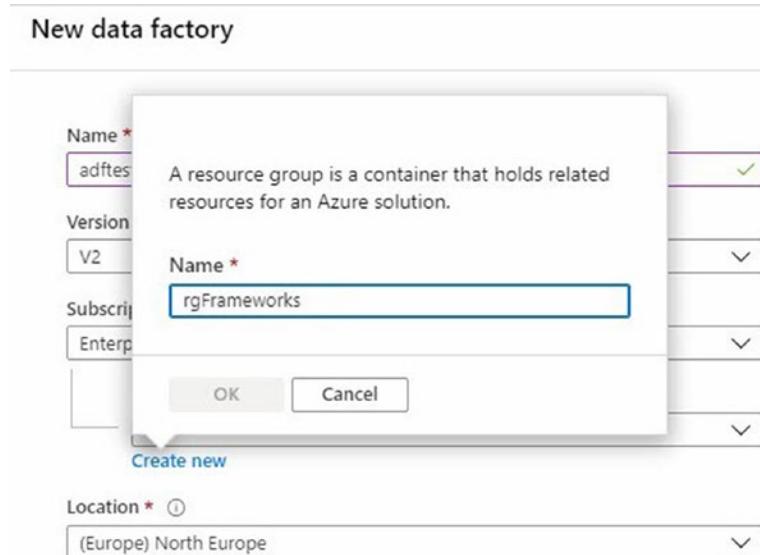


Figure 8-6. Creating a new Resource Group

At the time of this writing, Microsoft maintains more than 20 data center locations around the globe. The locations exist to help get your applications and data closer to clients. Multiple locations also facilitate backups and redundancy.

For best response, select a location near you from the Locations drop-down, as shown in Figure 8-7.

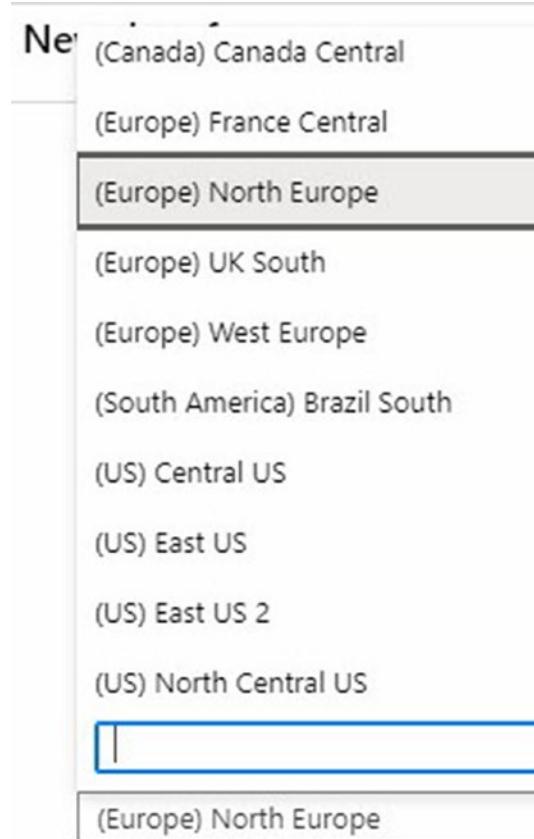


Figure 8-7. Selecting a Location

The “Enable GIT” is checked by default and this is awesome, and shown in Figure 8-8.

The screenshot shows a configuration form for creating an Azure Data Factory instance. At the top, there is a dropdown menu labeled "Location * ⓘ" with "(US) East US 2" selected. Below it is a checkbox labeled "Enable GIT ⓘ" which is checked. This checkbox is highlighted with a blue rectangular border. The form also includes fields for "GIT URL * ⓘ", "Repo name * ⓘ", "Branch Name * ⓘ", and "Root folder * ⓘ". At the bottom of the form is a large blue "Create" button.

Figure 8-8. The Enable GIT checkbox

This chapter does not cover source control, so please uncheck the Enable GIT checkbox.

Click the Create button to provision an instance of ADF configured as specified.

After a few minutes, the Azure Data Factory is provisioned. You may visit the resource, or pin it to a dashboard using the buttons in the notification, as shown in Figure 8-9.

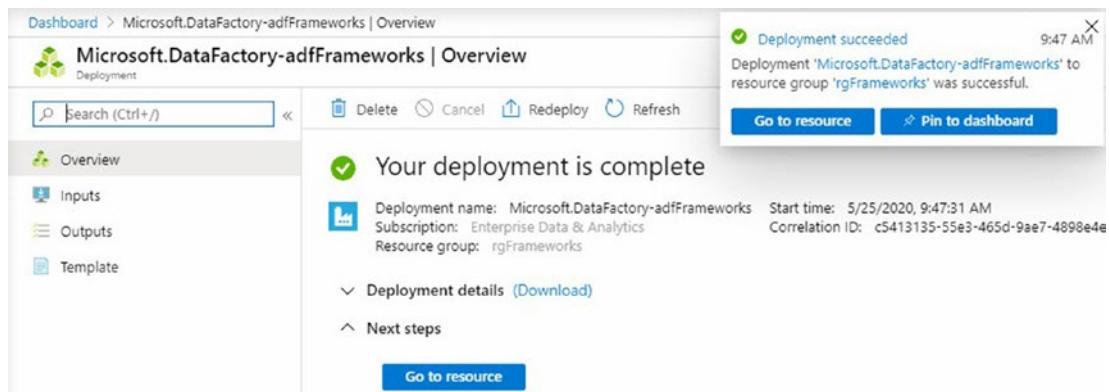


Figure 8-9. Azure Data Factory, provisioned

If you click the “Pin to dashboard” button from the notification “toast,” a Data Factory tile appears on the last Azure dashboard you visited, and appears as shown in Figure 8-10.

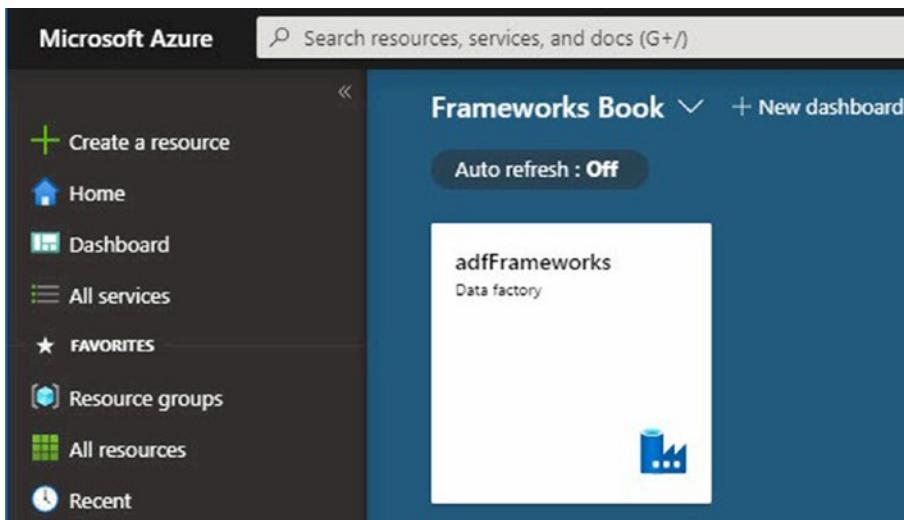


Figure 8-10. An Azure Data Factory instance pinned to an Azure dashboard

Like resource groups, Azure dashboards offer a way to group Azure resources. Azure dashboards allow visual grouping of Azure resources from many resource groups – and even subscriptions for those who manage more than one Azure subscription.

Click the tile to visit the page for the Azure Data Factory. The Overview displays by default, as shown in Figure 8-11.

CHAPTER 8 AZURE-SSIS INTEGRATION RUNTIME

The screenshot shows the Azure Data Factory interface for the 'adfFrameworks' instance. The left sidebar contains navigation links: Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Settings, Locks, General, Properties, Getting Started, and Quick start. The main panel displays basic information about the resource group ('rgFrameworks'), including its status (Succeeded), location (East US 2), and subscription (Enterprise Data & Analytics). It also includes links for Documentation and Author & Monitor.

Figure 8-11. The Azure Data Factory page

The “pushpin” in the upper right corner of the Azure Data Factory page allows you to pin a tile for this ADF instance to an Azure Dashboard. Remember, you may configure more than one Azure dashboard, and the ADF instance tile will pin to the last dashboard visited.

Click the Author & Monitor link button to visit adf.azure.com for this instance of Azure Data Factory, as shown in Figure 8-12.

The screenshot shows the 'Let's get started' section of the Data Factory Overview page. It features six circular icons with corresponding text labels: Create pipeline, Create data flow, Create pipeline from template, Copy data, Configure SSIS Integration, and Set up code repository. The background has a blue grid pattern and some faint charts.

Figure 8-12. The Data Factory Overview page

At the time of this writing, the Data Factory Overview page includes shortcut links for the following operations:

- Create pipeline
- Create data flow
- Create pipeline from template
- Copy data
- Configure SSIS integration
- Set up code repository

Now that an Azure Data Factory is provisioned, the next step is to add Azure Storage.

Provisioning Azure Storage

Azure Blob Storage is a file system available to Azure resources and services. Like the file system on servers and laptops, Azure Blob Storage plays a vital role in all things Azure. The goal of this section is to describe provisioning a simple Azure Storage account. Later, we will use an Azure File Share in this account to store SSIS packages and ISPAC files. To learn more about Azure file share, please see docs.microsoft.com/en-us/azure/storage/files/storage-how-to-create-file-share.

As when provisioning an ADF instance, click Create a resource in the Azure left menu. Click Storage, and then click “Storage account – blob, file, table, queue,” as shown in Figure 8-13.

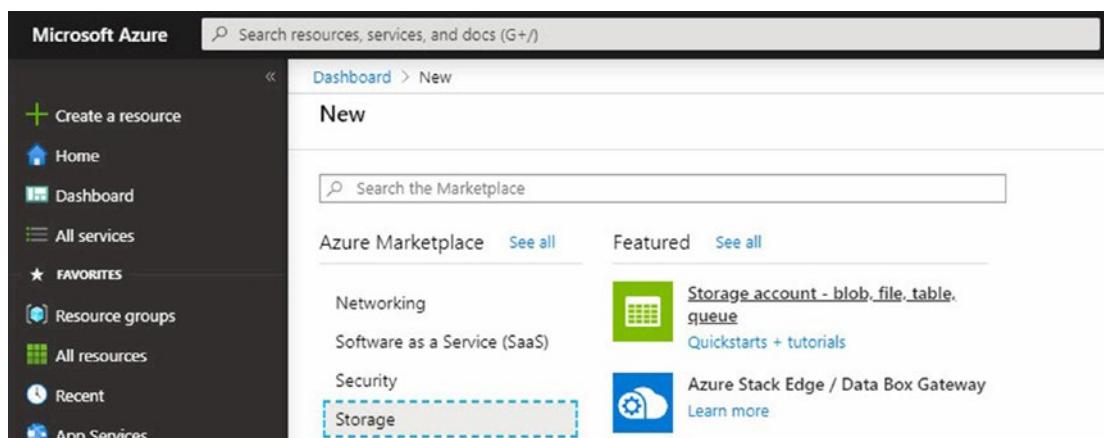


Figure 8-13. Provisioning Azure Storage

When the “Create storage account” page displays, configure the subscription and resource group in the “Project details” section, as shown in Figure 8-14.

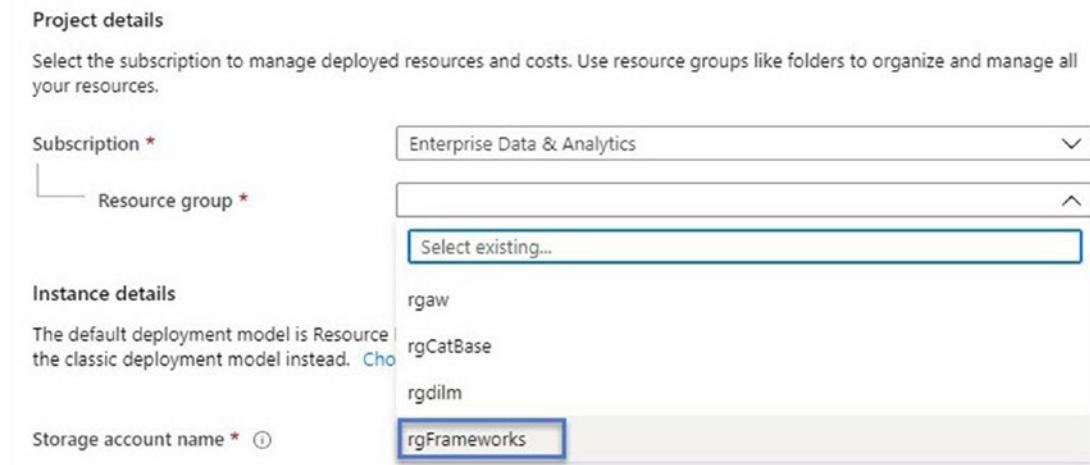


Figure 8-14. Configuring the storage account subscription and resource group

In the “Instance details” section, configure the “Storage account name” property, as shown in Figure 8-15.



Figure 8-15. Configuring the Storage account name property

The next step is to configure the Location property in the “Instance details” section, as shown in Figure 8-16.

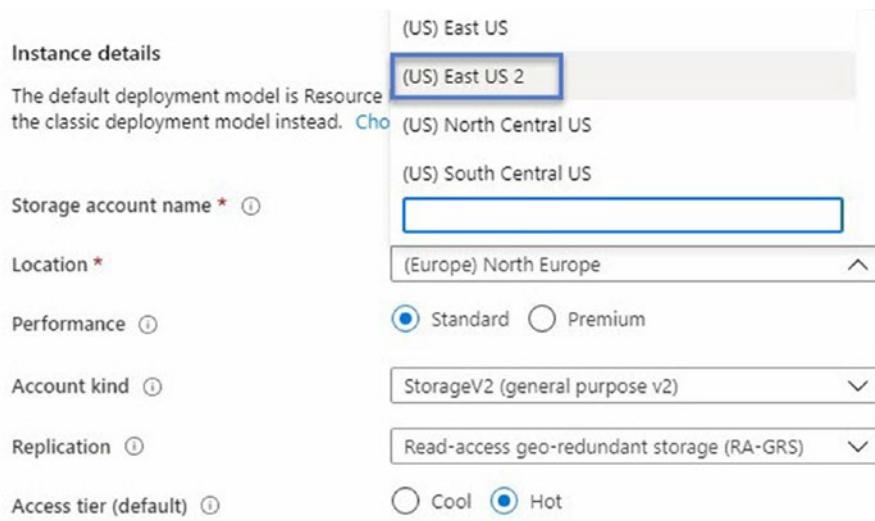


Figure 8-16. Configuring the Location property

For the purposes of this example, “Standard” Performance will suffice, as shown in Figure 8-17.



Figure 8-17. Configuring the Performance property

Configure the “Account kind” property next. Set Account kind to “StorageV2 (general purpose v2),” as shown in Figure 8-18.

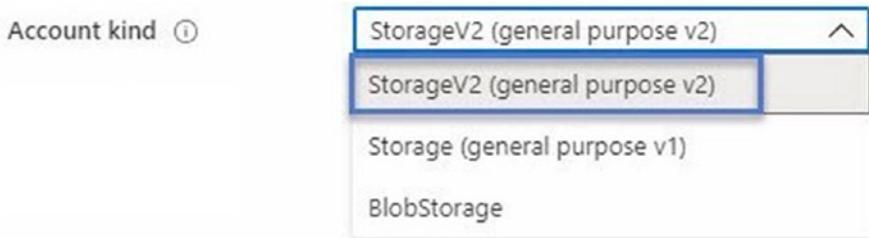


Figure 8-18. Configuring the Account kind property

For the purposes of the example, setting the Replication property to “Locally-redundant storage (LRS)” will suffice, as shown in Figure 8-19.

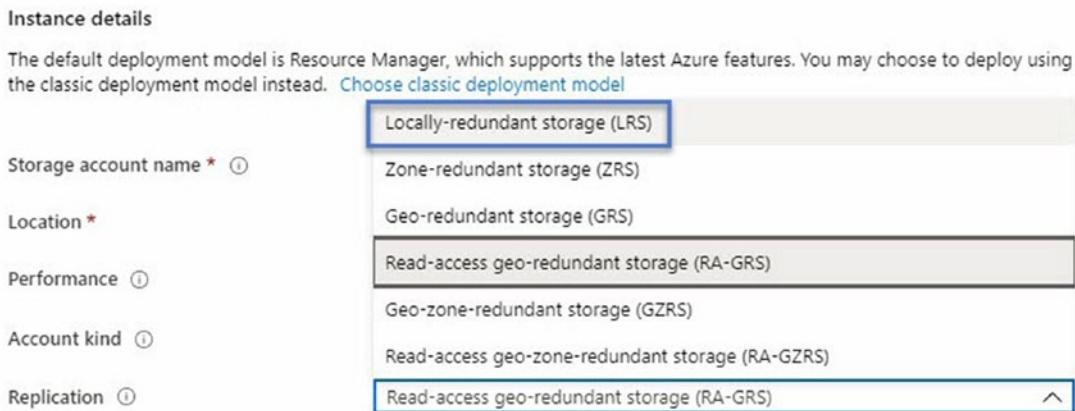


Figure 8-19. Configuring the Replication property of the Storage account

Accept the default “Access tier” property setting – “Hot” – as shown in Figure 8-20.



Figure 8-20. Configuring the Access tier property

The Basics tab of the “Create storage account” page is configured as shown in Figure 8-21.

Create storage account

Basics Networking Advanced Tags Review + create

Azure Storage is a Microsoft-managed service providing cloud storage that is highly available, secure, durable, scalable, and redundant. Azure Storage includes Azure Blobs (objects), Azure Data Lake Storage Gen2, Azure Files, Azure Queues, and Azure Tables. The cost of your storage account depends on the usage and the options you choose below.
[Learn more about Azure storage accounts](#)

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription *	Enterprise Data & Analytics	▼
Resource group *	rgFrameworks	▼
	Create new	

Instance details

The default deployment model is Resource Manager, which supports the latest Azure features. You may choose to deploy using the classic deployment model instead. [Choose classic deployment model](#)

Storage account name * ⓘ	stframeworks	✓
Location *	(US) East US 2	▼
Performance ⓘ	<input checked="" type="radio"/> Standard <input type="radio"/> Premium	
Account kind ⓘ	StorageV2 (general purpose v2)	▼
Replication ⓘ	Locally-redundant storage (LRS)	▼
Access tier (default) ⓘ	<input type="radio"/> Cool <input checked="" type="radio"/> Hot	

[Review + create](#) [< Previous](#) [Next : Networking >](#)

Figure 8-21. The Basics tab, configured

For the purposes of this example, leave settings on the Networking, Advanced, and Tags tabs set to default. At the bottom of the Basics tab, click the “Review + create” button (shown in Figure 8-21).

Storage account settings are validated, and the results are displayed, as shown in Figure 8-22.

Create storage account

Validation passed

Basics Networking Advanced Tags Review + create

Basics

Subscription	Enterprise Data & Analytics
Resource group	rgFrameworks
Location	East US 2
Storage account name	stframeworks
Deployment model	Resource manager
Account kind	StorageV2 (general purpose v2)
Replication	Locally-redundant storage (LRS)
Performance	Standard
Access tier (default)	Hot

Networking

Connectivity method	Public endpoint (all networks)
---------------------	--------------------------------

Advanced

Secure transfer required	Enabled
Large file shares	Disabled
Blob soft delete	Disabled
Blob change feed	Disabled
Versioning	Disabled
Hierarchical namespace	Disabled
NFS v3	Disabled

Buttons: Create, < Previous, Next >

Figure 8-22. Storage account configuration validation results

Note also the text on the button at the bottom of the Create storage account page changes to “Create.” Click the Create button to create the new Storage account. Once deployment is complete, the Storage account page and notification will display as shown in Figure 8-23.

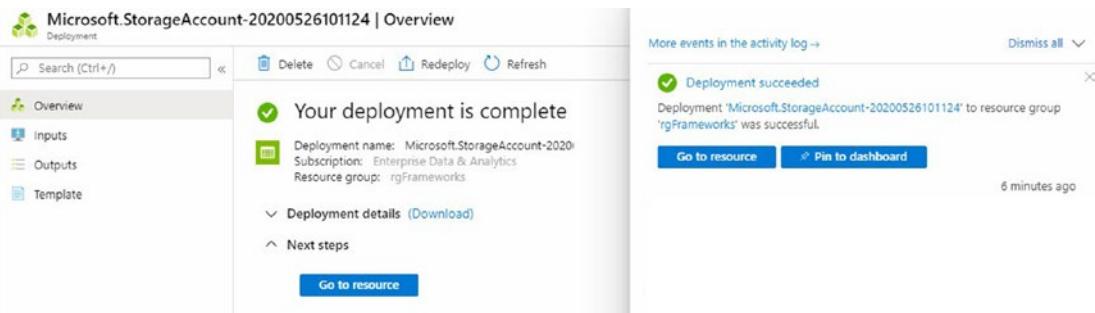


Figure 8-23. Storage account, created

Now that an Azure Data Factory and Storage account are provisioned, the next step is to add an Azure-SSIS integration runtime.

Provisioning an Azure-SSIS Integration Runtime for SSIS Package Files

On 30 June 2019, Microsoft released a new version of the ADF Azure-SSIS integration runtime. Enterprises could execute SSIS packages in the cloud previously using an SSIS Catalog hosted by an Azure database; the new release allowed SSIS package execution from an Azure file share.

To begin creating an Azure-SSIS integration runtime, navigate to an instance of Azure Data Factory’s Manage page, and click “Integration runtimes” in the Connections category, as shown in Figure 8-24.

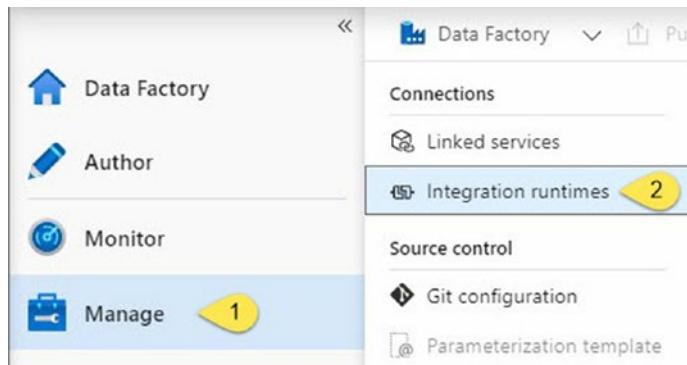


Figure 8-24. Opening ADF Connections

When the Connections page displays, click the “Integration runtimes” tab as shown in Figure 8-25.

Integration runtimes

The integration runtime (IR) is the compute infrastructure to provide the following data integration capabilities across different network environment. [Learn more](#)

This screenshot shows the 'Integration runtimes' tab in the ADF portal. At the top, there is a header with a 'New' button (circled in red), a 'Refresh' button, and a search bar. Below the header, it says 'Showing 1 - 1 of 1 items'. There is a table with columns: NAME, TYPE, SUB-TYPE, STATUS, and REGION. The single entry listed is 'AutoResolveIntegr...', which is of type 'Azure', has a 'Public' sub-type, is 'Running', and is located in the 'Auto Resolve' region.

NAME ↑↓	TYPE ↑↓	SUB-TYPE ↑↓	STATUS ↑↓	REGION ↑↓
AutoResolveIntegr...	Azure	Public	Running	Auto Resolve

Figure 8-25. ADF Connections “Integration runtimes” tab

Click the “New” button – circled in Figure 8-25 – to open the “Integration runtime setup” blade and select the “Azure-SSIS” option as shown in Figure 8-26.

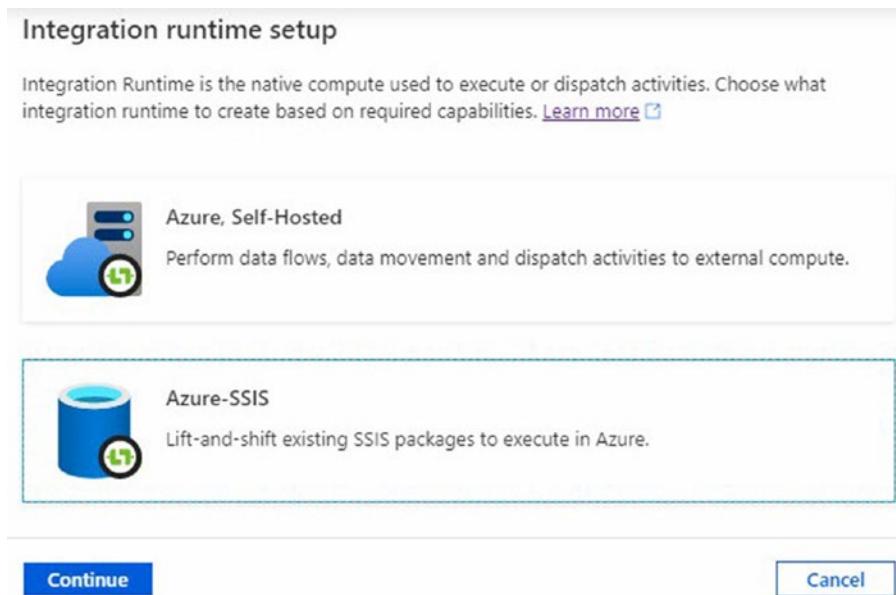


Figure 8-26. Selecting Azure-SSIS in the Integration runtime setup blade

The other option displayed in Figure 8-26 is “Azure, Self-Hosted.” You may learn more about *all* Azure Data Factory integration runtimes by visiting docs.microsoft.com/en-us/azure/data-factory/concepts-integration-runtime#self-hosted-integration-runtime. For the purposes of this example, we will configure a “stand-alone” Azure-SSIS Integration Runtime – or IR – an IR that requires minimal security configuration. More information is available about joining an Azure-SSIS IR to a virtual network at docs.microsoft.com/en-us/azure/data-factory/join-azure-ssis-integration-runtime-virtual-network.

Click the Continue button to open the General settings page on the Integration runtime setup blade. Configure the name and optional description of the Azure-SSIS IR, and then set the Location property. For this example, set the Location property to the same location as the Azure Data Factory (which is the default), as shown in Figure 8-27.

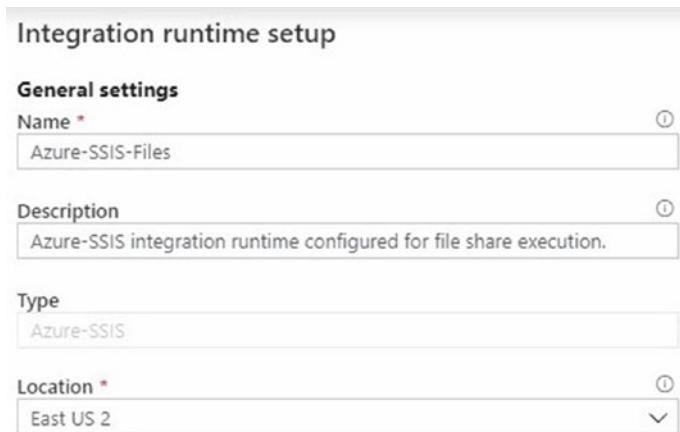


Figure 8-27. Configuring Name, Description, and Location properties

The next two steps configure virtual machines (VMs) upon which Azure-SSIS will execute SSIS packages. “Node size” represents the configuration of *each node* configured, and there are some beefy options, such as those listed in Figure 8-28.

D32_v3 (32 Core(s), 131072 MB)

D64_v3 (64 Core(s), 262144 MB)

E2_v3 (2 Core(s), 16384 MB)

E4_v3 (4 Core(s), 32768 MB)

E8_v3 (8 Core(s), 65536 MB)

E16_v3 (16 Core(s), 131072 MB)

E32_v3 (32 Core(s), 262144 MB)

E64_v3 (64 Core(s), 442368 MB)

Figure 8-28. Some Node size options for Azure-SSIS

The first letter of the VM series indicates the general class and purpose of the virtual machine. A-Series VMs are specifically designed to save money during the testing and development phases of a project. To learn more about Azure virtual machine series and costs, please visit azure.microsoft.com/en-us/pricing/details/virtual-machines/series/.

Select the “A8_v2 (8 Core(s), 16384 MB)” Node size from the drop-down, as shown in Figure 8-29.

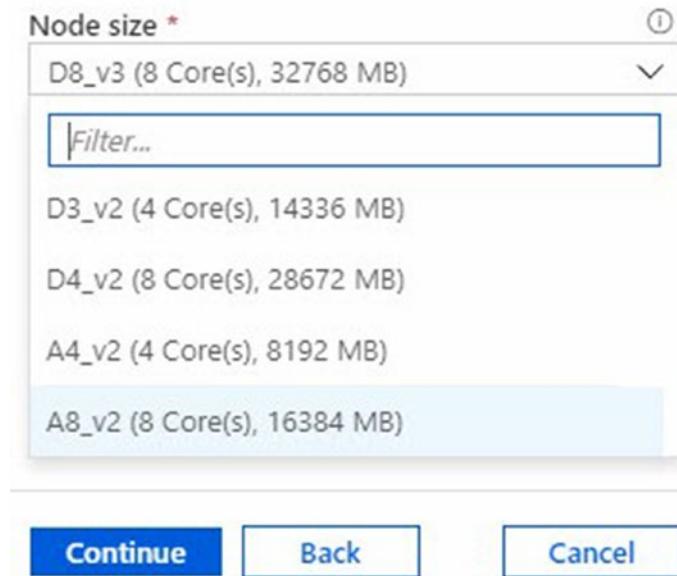


Figure 8-29. Selecting a Node size for Azure-SSIS

The “Node number” slider indicates the number of virtual machines available to Azure-SSIS. Each node will be configured according to the Node size selected in the previous drop-down. Node number is the number of nodes, configured in the Node size property, available to the Azure-SSIS IR. For this example, leave the Node number property set at 2 (the default), as shown in Figure 8-30.



Figure 8-30. Configuring the “Node number” property to 2

The next step is to configure the “Edition/license” property. Select “Standard” for now, as shown in Figure 8-31.

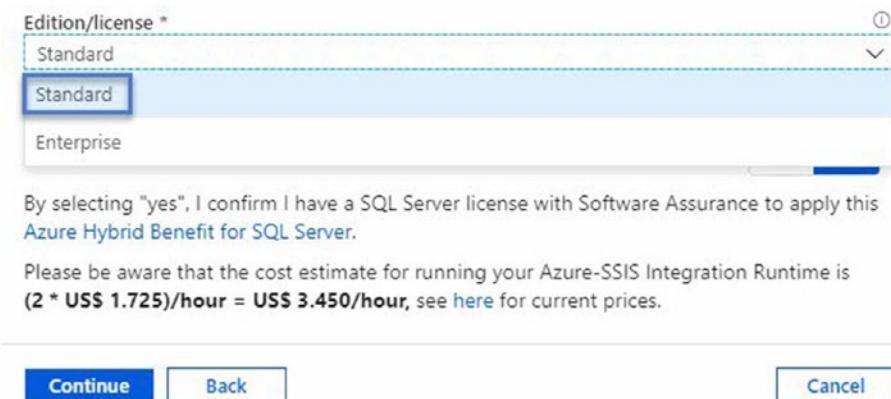


Figure 8-31. Setting the Edition/license property

A feature known as Bring Your Own License, or “BYOL,” is another way to reduce spending n Azure. If you or your enterprise already own a SQL Server license, click the “Yes” to reduce the cost of running the VMs behind Azure-SSIS, as shown in Figure 8-32.

Node size *

Node number *

Edition/license *

Save money

Save with a license you already own. Already have a SQL Server license? Yes No

By selecting "yes", I confirm I have a SQL Server license with Software Assurance to apply this [Azure Hybrid Benefit for SQL Server](#).

Please be aware that the cost estimate for running your Azure-SSIS Integration Runtime is **(2 * US\$ 0.945)/hour = US\$ 1.890/hour**, see [here](#) for current prices.

Figure 8-32. Reduce Azure costs with BYOL

Note changing the Edition/license property does *not* impact the Azure-SSIS cost, as shown in Figure 8-33.

Edition/license *

Save money

Save with a license you already own. Already have a SQL Server license? Yes No

By selecting "yes", I confirm I have a SQL Server license with Software Assurance to apply this [Azure Hybrid Benefit for SQL Server](#).

Please be aware that the cost estimate for running your Azure-SSIS Integration Runtime is **(2 * US\$ 0.945)/hour = US\$ 1.890/hour**, see [here](#) for current prices.

Figure 8-33. Select Enterprise Edition does not impact Azure-SSIS cost

Read azure.microsoft.com/en-us/pricing/details/data-factory/ssis/ to learn more about Azure-SSIS costs. You may also wish to try different variations to manage performance and costs, such as the “D2_v3 (2 Core(s), 8192 MB)” Enterprise configuration shown in Figure 8-34.

The screenshot shows a configuration form for an Azure Integration Runtime. At the top, there's a dropdown for 'Node size *' set to 'D2_v3 (2 Core(s), 8192 MB)'. Below it is a 'Node number *' field with a slider set to '2'. Under 'Edition/license *', it says 'Enterprise'. A note below reads 'Save money' followed by 'Save with a license you already own. Already have a SQL Server license?'. There are 'Yes' and 'No' buttons; 'Yes' is highlighted in blue. A note below states: 'By selecting "yes", I confirm I have a SQL Server license with Software Assurance to apply this Azure Hybrid Benefit for SQL Server.' Another note below says: 'Please be aware that the cost estimate for running your Azure-SSIS Integration Runtime is (2 * US\$ 0.290)/hour = US\$ 0.580/hour, see [here](#) for current prices.'

Figure 8-34. Testing other combinations of Node size

\$0.58 is around 31 percent of \$1.89.

Please visit azure.microsoft.com/en-us/pricing/details/data-factory/ssis/ to learn more about the cost of Azure-SSIS instances. One way to manage Azure-SSIS costs is to schedule an Azure Data Factory pipeline to shut down each Azure-SSIS instance nightly, in case you forget or are called away while learning. andyleonard.blog/2020/05/stop-an-azure-ssis-files-integration-runtime-safely/ is one detailed blog post about how to stop an Azure-SSIS IR.

The author and/or publisher is not responsible for any costs incurred.

Once General settings are configured, click the Continue button to open the “Deployment settings” page in the Integration runtime setup blade, as shown in Figure 8-35.

Integration runtime setup

Deployment settings

Create SSIS catalog (SSISDB) hosted by Azure SQL Database server/Managed Instance to store your projects/packages/environments/execution logs (See more info [here](#))

Subscription *
Enterprise Data & Analytics

Location
Select all

Catalog database server endpoint *

Use AAD authentication with the managed identity for your Data Factory (See how to enable it [here](#))

Admin username *
Admin username is empty

Admin password *
Admin password is empty

Catalog database service tier *
S1

Create package stores to manage your packages that are deployed into file system/Azure Files/SQL Server database (MSDB) hosted by Azure SQL Database Managed Instance (See more info [here](#))

Test connection **Back** **Cancel**

Figure 8-35. The “Deployment settings” page in the Integration runtime setup blade

I can hear you thinking, “Gosh, Andy! That’s a lot of stuff to configure. Plus, you wrote *nothing* about creating an Azure database. What gives?” Simmer down. We are creating an instance of Azure-SSIS for *files*, remember? Uncheck the “Create SSIS catalog (SSISDB) hosted by Azure SQL Database server/Managed Instance to store your projects/packages/environments/execution logs” checkbox – circled in Figure 8-35. The “Deployment settings” page in the Integration runtime setup blade now appears similar to that shown in Figure 8-36.

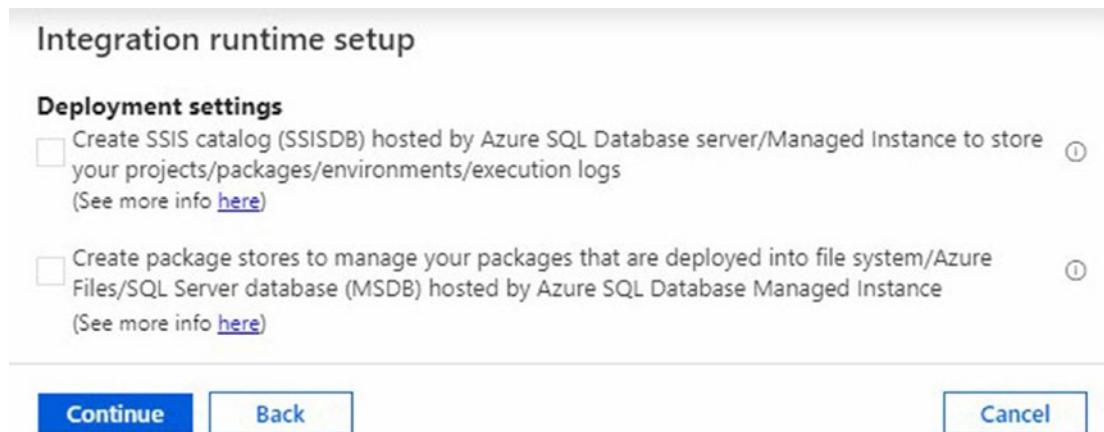


Figure 8-36. The “Deployment settings” page in the Integration runtime setup blade after unchecking the “Create SSIS catalog...” checkbox

Click the Continue button to open the Integration runtime setup blade’s Advanced settings page. Four options are available on the Advanced settings page:

1. Maximum parallel executions per node
2. Azure-SSIS integration runtime customization
3. VNet
4. Self-hosted integration runtime

The “Maximum parallel executions per node” drop-down does what it says. Each Azure-SSIS node (VM) may execute the number of SSIS packages configured by the Maximum parallel executions per node property.

Azure-SSIS integration runtimes may include custom configurations and/or third-party tools such as SentryOne’s Task Factory (sentryone.com/products/task-factory) and SSIS+ Components Suite by COZYROC (cozyroc.com/products).

Most enterprises using the cloud store some, but not all, of their data in the cloud in a *hybrid* architecture. VNet is one way to access on-premises enterprise data. At the time of this writing, classic Azure virtual network is being deprecated and replaced with VNet. VNet is recommended for enterprises using or desiring to use

- Classic Azure virtual network
- Public IP addresses with Azure-SSIS IR
- Customized Azure-SSIS

Learn more by visiting the article “Join an Azure-SSIS integration runtime to a virtual network” at docs.microsoft.com/en-us/azure/data-factory/join-azure-ssis-integration-runtime-virtual-network.

A self-hosted Azure-SSIS integration runtime permits access to on-premises enterprise data *without* requiring VNet.

For the purposes of this example, accept the defaults – no checkboxes checked – as shown in Figure 8-37.

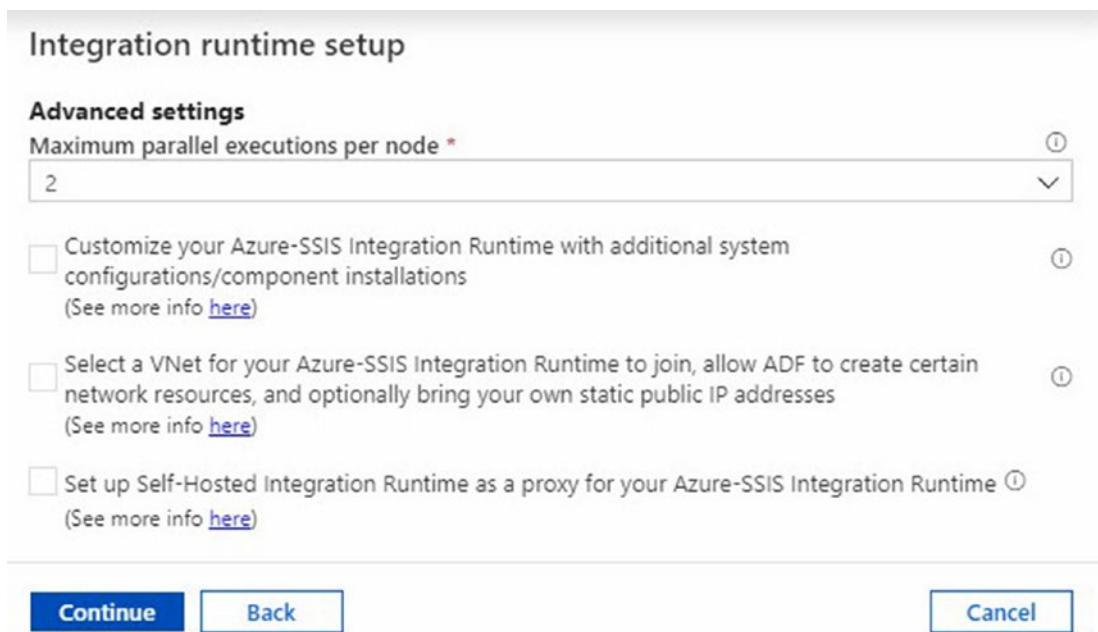


Figure 8-37. Configuring Azure-SSIS Advanced settings

Click the Continue button to proceed to the Integration runtime setup’s Summary page. The Summary page displays Azure-SSIS configuration choices, as shown in Figure 8-38.

Integration runtime setup

Summary

Your Azure-SSIS Integration Runtime (IR) is created with the following settings:

Azure Data Factory Settings

- **Subscription:** 78ff08f6-334c-4e53-b737-8b5feaf74ecc
- **Resource group:** rgFrameworks
- **Name:** adfFrameworks
- **Location:** eastus2

General settings

- **Name:** Azure-SSIS-Files
- **Description:** Azure-SSIS integration runtime configured for file share execution.
- **Location:** East US 2
- **Node size:** Standard_D2_v3
- **Node number:** 2
- **Edition:** Enterprise
- **Azure Hybrid Benefit:** BasePrice

Advanced settings

- **Maximum parallel executions per node:** 2
- If you need to access data on premises, click **Previous** to do any of the followings:
 - Join your Azure-SSIS IR to a VNet connected to your on-premises network OR
 - Set up Self-Hosted Integration Runtime as a proxy for your Azure-SSIS Integration Runtime

If you want to change any of the above settings, click **Previous** to do so.

Once your Azure-SSIS IR is running, you can execute your packages on it after [deploying](#) them into your file system/Azure Files.

Please be aware that the cost estimate for running your Azure-SSIS Integration Runtime is **(2 * US\$ 0.290)/hour = US\$ 0.580/hour**, see [here](#) for current prices.

To manage the running cost of your Azure-SSIS IR, you can [stop & restart](#) it whenever convenient or [schedule](#) it just in time.

[Create](#)[Previous](#)[Cancel](#)

Figure 8-38. Azure-SSIS configuration choices

Click the Create button to provision the Azure-SSIS integration runtime. Connections ➤ Integration runtimes displays the new Azure-SSIS integration runtime in the “Starting” status, as shown in Figure 8-39.

Integration runtimes

The integration runtime (IR) is the compute infrastructure to provide the following data integration capabilities across

NAME ↑↓	TYPE ↑↓	SUB-TYPE ↑↓	STATUS ↑↓
AutoResolveIntegrationRuntime	Azure	Public	Running
Azure-SSIS-Files	Azure-SSIS	---	Starting

Figure 8-39. Azure-SSIS-Files is starting

After some time – usually 3–5 minutes (maximum) at the time of this writing – the Azure-SSIS IR is started and Connections ➤ Integration runtimes displays “Running” status, as shown in Figure 8-40.

Integration runtimes

The integration runtime (IR) is the compute infrastructure to provide the following data integration capabilities across di environment. [Learn more](#)

NAME ↑↓	TYPE ↑↓	SUB-TYPE ↑↓	STATUS ↑↓
AutoResolveIntegrationRuntime	Azure	Public	Running
Azure-SSIS-Files	Azure-SSIS	---	Running

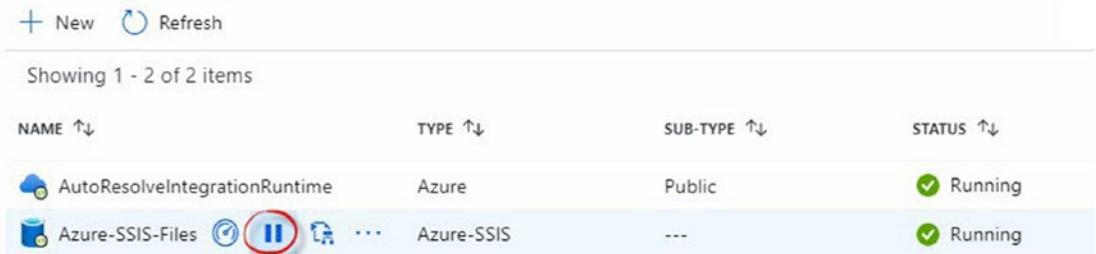
Figure 8-40. Azure-SSIS-Files has started

Stopping the Azure-SSIS Integration Runtime

One way to stop an Azure-SSIS IR is using the “Stop” button (the Stop button displays a Pause icon) found on Connections ➤ Integration runtimes, as shown in Figure 8-41.

Integration runtimes

The integration runtime (IR) is the compute infrastructure to provide the following data integration capabilities across different environment. [Learn more](#)



NAME ↑	TYPE ↑↓	SUB-TYPE ↑↓	STATUS ↑↓
AutoResolveIntegrationRuntime	Azure	Public	✓ Running
Azure-SSIS-Files	● II	Azure-SSIS	✓ Running

Figure 8-41. Stopping the Azure-SSIS IR

Clicking the Stop button triggers a confirmation dialog, as shown in Figure 8-42.

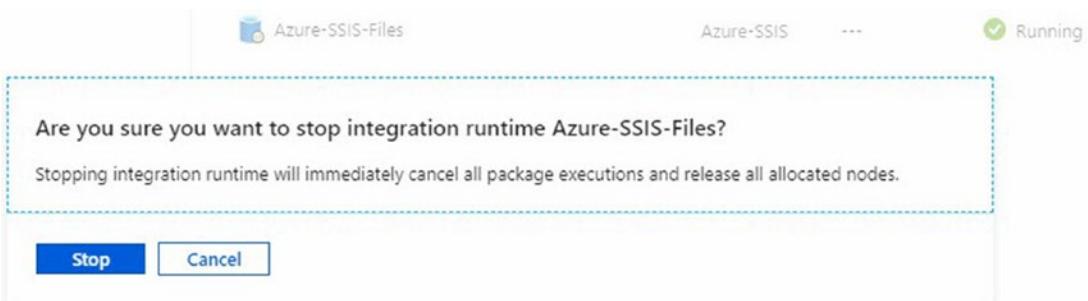


Figure 8-42. Confirming the stop command

At the time of this writing, clicking the Stop button on the confirmation dialog triggers the survey form shown in Figure 8-43.

Thanks for using our product! Please tell us why you choose to stop SSIS Integration Runtime. Your feedback helps us to create a better experience for you.

The reason for the stop (Multiple Selection):

- It is a part of regular routine. Don't ask me again.
- Everything is good. I am just looking and may come back later.
- The product performance is poor.
- The core functions of the product don't achieve the desired effect.
- The required functions/features are not available.
- The product is overpriced.

Other reasons

Would you like to provide us with your email address so that we can contact you later in case of any additional questions? (Optional)

Your feedback is collected by Microsoft and used to improve your experience.

[Privacy and cookies](#)

[Submit](#) [Cancel](#)

Figure 8-43. A survey dialog triggered by the confirmation dialog

Once the survey form is submitted (or cancelled), the Azure-SSIS integration runtime enters a “Stopping” state, as shown in Figure 8-44.

NAME ↑↓	TYPE ↑↓	SUB-TYPE ↑↓	STATUS ↑↓	REGION ↑↓
 AutoResolveIntegrationRuntime	Azure	Public	 Running	Auto Resolve
 Azure-SSIS-Files	Azure-SSIS	---	 Stopping	East US 2

Figure 8-44. Azure-SSIS is stopping

Another way to stop the Azure-SSIS IR is from the Monitor ➤ Integration runtimes page, as shown in Figure 8-45.

CHAPTER 8 AZURE-SSIS INTEGRATION RUNTIME

The screenshot shows the Azure portal's left sidebar with 'Monitor' selected. In the main area, the 'Integration runtimes' section is highlighted. A table lists two runtimes: 'AutoResolveIntegrationRuntime' (Azure, Public, Running) and 'Azure-SSIS-Files' (Azure-SSIS, ..., Running). The table has columns for NAME, TYPE, SUB-TYPE, and STATUS.

NAME ↑	TYPE ↑↓	SUB-TYPE ↑↓	STATUS ↑↓
AutoResolveIntegrationRuntime	Azure	Public	Running
Azure-SSIS-Files	Azure-SSIS	...	Running

Figure 8-45. The Monitor ➤ Integration runtimes page

Click the name of the Azure-SSIS IR – Azure-SSIS-Files in this case – to open the details dashboard, as shown in Figure 8-46.

The screenshot shows the 'Resource monitor (details)' page for the 'Azure-SSIS-Files' runtime. It displays various configuration settings:

STATUS	Running	TYPE	Azure-SSIS	REGION	East US 2
NODE SIZE	Standard_D2_v3 2 Core(s). 8192MB	RUNNING / REQUESTED NODE(S)	1/2		

Figure 8-46. Azure-SSIS details dashboard

Click the Status (“Running”) to open the Azure-SSIS status dialog shown in Figure 8-47.

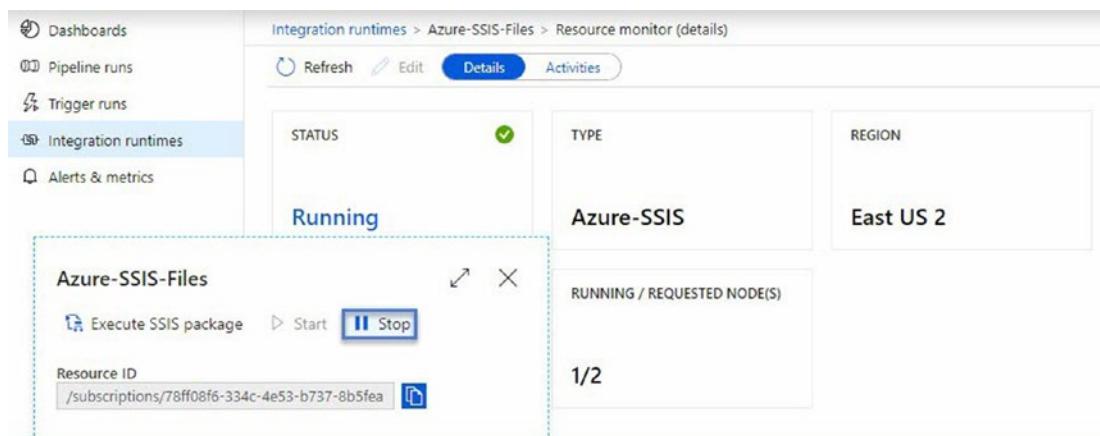


Figure 8-47. Azure-SSIS status dialog

Click the Stop button to stop the Azure-SSIS integration runtime, as shown in Figure 8-48.

STATUS	TYPE	REGION
Stopping	Azure-SSIS	East US 2
NODE SIZE	RUNNING / REQUESTED NODE(S)	
Standard_D2_v3 2 Core(s). 8192MB	1/2	

Figure 8-48. Azure-SSIS integration runtime stopping

Please note, clicking the Stop button on the Azure-SSIS status dialog triggers the “Are you sure?/Survey” process.

Restart the Azure-SSIS integration runtime to continue.

Now that an Azure-SSIS IR has been configured for execution from Azure File Shares, the next step is to provision an Azure SQL database.

Conclusion

The focus of this chapter has been provisioning an instance of Azure-SSIS and prerequisites. This chapter walked through the following steps:

- Getting started with Azure
- Provisioning an Azure Data Factory
- Provisioning Azure Storage
- Provisioning Azure-SSIS

The next step is to provision an Azure SQL Database.

CHAPTER 9

Deploy a Simple, Custom, File-Based Azure-SSIS Framework

The Microsoft Azure-SSIS team has been hard at work reducing the friction between executing SSIS packages on-premises and executing SSIS packages in Azure. Many enterprises execute SSIS packages stored in the file system on-premises. In the past, migrating SSIS packages executed in an on-premises file system meant changing enterprise SSIS storage and execution patterns. With the advent of Azure-SSIS File Share-based execution, executing SSIS packages from files in Azure is no longer an issue.

In fact, the SSIS framework in previous chapters works well in Azure-SSIS. In this chapter, we will discuss and demonstrate the following.

Provisioning the SSISConfig Database

The focus of this section is provisioning the SSISConfig database to an instance of Azure SQL Database upon which the SSISConfig database will reside.

Azure SQL can be a somewhat loose term. It can be taken to mean Azure SQL Database, Azure SQL Managed Instance, and some even use it to mean SQL Server running on an Azure VM. In this chapter, we are working in Azure SQL Database, which is the Software as a Service offering as described at docs.microsoft.com/en-us/azure/sql/azure-sql-iaas-vs-paas-what-is-overview.

Navigate to the Azure portal to begin the provisioning process. In the left Azure menu, hover over “SQL databases” until the SQL databases hover card displays. When the SQL databases hover card displays, click “+ Create” similar to Figure 9-1.

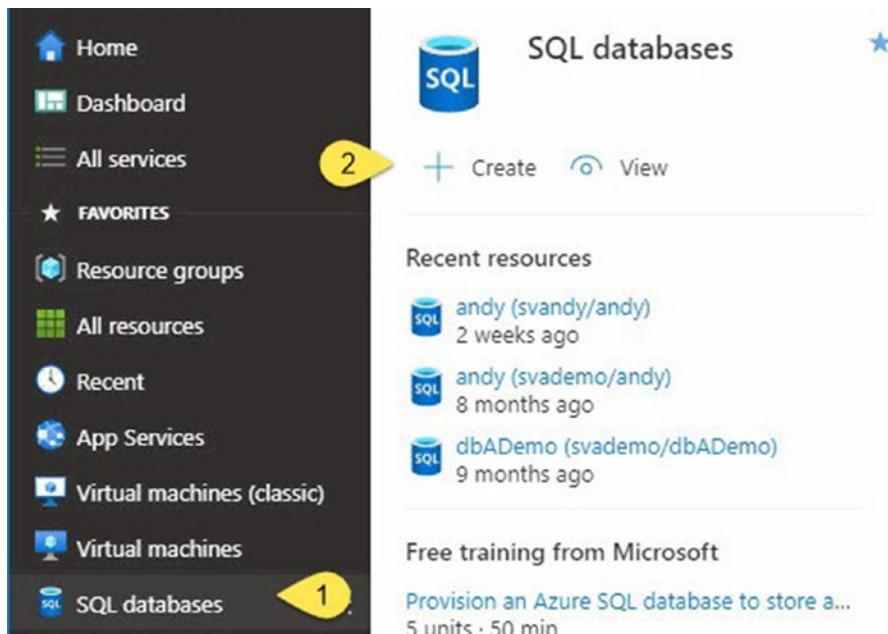


Figure 9-1. Starting the Azure SQL Database provisioning process

Clicking “+ Create” opens the Create SQL Database blade. Begin by configuring the Subscription property to your subscription name and the Resource group property to your resource group name, as shown in Figure 9-2.

[Dashboard >](#)

Create SQL Database

Microsoft

[Basics](#) [Networking](#) [Additional settings](#) [Tags](#) [Review + create](#)

Create a SQL database with your preferred configurations. Complete the Basics tab then go to Review + Create to provision with smart defaults, or visit each tab to customize. [Learn more](#) 

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription *	Enterprise Data & Analytics 
Resource group *	rgFrameworks  Create new

Figure 9-2. Azure SQL Database Subscriptions and Resource group properties configuration

The example uses the resource group named “rgFrameworks.” Azure resource group names are globally unique (and case insensitive), so you will need to use a different name for your resource group.

Next, configure the Data details section by entering “SSISConfig” for the “Database name” property. Please note server names are globally unique, but database names are only unique to the server, as shown in Figure 9-3.

Database details

Enter required settings for this database, including picking a logical server and configuring the compute and storage resources

Database name *	SSISConfig	✓
Server *	(new) svssis (East US 2)	▼
	Create new	

Figure 9-3. Configuring the Database name and Server properties

For the purposes of this example, click the “No” option to answer the question, “Want to use SQL elastic pool?” The “Compute + storage” property defaults to “General purpose Gen5, 2 vCores, 32 GB storage,” as shown in Figure 9-4.

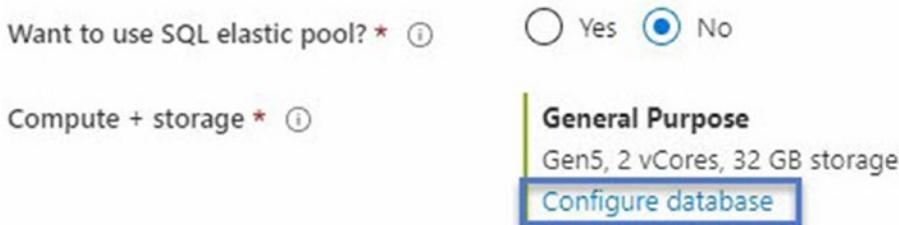


Figure 9-4. SQL elastic pool and “Compute + storage” property configuration

The default database configuration is a little heavy for the purposes of this example. Click the “Configure database” link shown in Figure 9-4 to open the Configure blade, which will allow changes to database capacity and performance options, as shown in Figure 9-5.

Configure

The screenshot shows the Azure portal interface for configuring an Azure SQL database. At the top left is a 'Feedback' button. To the right, under 'General Purpose', it says 'Scalable compute and storage options' with performance metrics '500 - 20,000 IOPS' and '2-10 ms latency'. A blue box highlights the link 'Looking for basic, standard, premium?'. Below this is a large 'SQL' icon. The 'Cost summary' section details the configuration and cost:

Cost summary	
Gen5 - General Purpose (GP_Gen5_2)	
Cost per vCore (in USD)	113.25
vCores selected	x 2
Cost per GB (in USD)	0.12
Max storage selected (in GB)	x 41.6
ESTIMATED COST / MONTH	231.28 USD

Figure 9-5. Configuring the Azure SQL database

As you can tell from Figure 9-5, the default database configuration is also a little *expensive* for the purposes of this example. Click the “Looking for basic, standard, premium?” link to navigate to, and select, the “Basic” option, as shown in Figure 9-6.

Configure

Feedback

Basic
For less demanding workloads

Standard
For workloads with typical performance requirements

DTUs [What is a DTU?](#)

5 (Basic)

Data max size

100 MB

2 GB

2 GB



Cost summary

Cost per DTU (in USD)	1.00
DTUs selected	x 5
ESTIMATED COST / MONTH	4.99 USD

Apply

Figure 9-6. Reconfiguring the Azure SQL database

Two GB and five DTUs is plenty for the example. Click the Apply button to proceed. Database configuration now displays out updated settings, as shown in Figure 9-7.

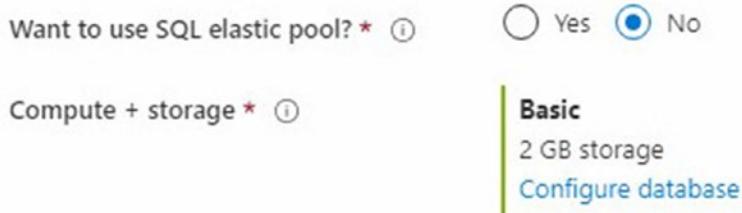


Figure 9-7. Updated Compute + storage settings

Before creating the new SSISConfig database, double-check the “Additional settings” tab’s “Enable advanced data security” property, as shown in Figure 9-8.

Advanced data security

Protect your data using advanced data security, a unified security package including data classification, vulnerability assessment and advanced threat protection for your server. [Learn more](#)

Get started with a 30 day free trial period, and then 15 USD/server/month.

Enable advanced data security *

[Start free trial](#)

Not now

Figure 9-8. Configuring the “Enable advanced data security” property

If the “Start free trial” option is selected, change the setting to “Not now” for the purposes of this example.

On the Networking and Tags tabs, leave the default configuration settings.

The “Review + create” tab should appear similar to Figure 9-9.

Create SQL Database

Microsoft

Basics Networking Additional settings Tags **Review + create**

Product details

SQL database by Microsoft Terms of use Privacy policy	Estimated cost per month 4.99 USD View pricing details
---	---

Terms

By clicking "Create", I (a) agree to the legal terms and privacy statement(s) associated with the fees associated with the offering(s), with the same billing frequency as my subscription, with the provider(s) of the offering(s) for support, billing and other transactional issues. [Marketplace Terms.](#) 

Basics

Subscription	Enterprise Data & Analytics
Resource group	rgFrameworks
Region	East US 2
Database name	SSISConfig
Server	(new) svssis
Compute + storage	Basic: 2 GB storage

Networking

Allow Azure services and resources to access this server	No
Private endpoint	None

Additional settings

Use existing data	Blank
Collation	SQL_Latin1_General_CI_AS
Advanced data security	Not now

Tags

Figure 9-9. Azure SQL Database Review + create page

Click the Create button to create the Azure SQL Database. Provisioning an Azure SQL Database takes a few minutes. Once complete, the portal should appear similar to Figure 9-10.

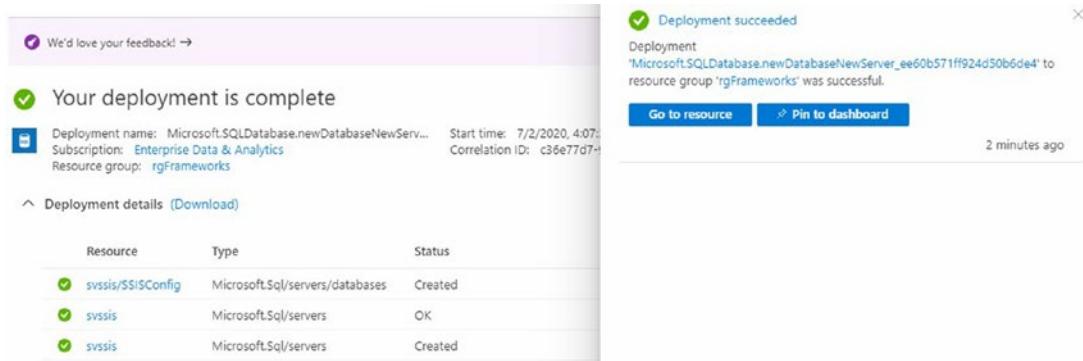


Figure 9-10. Azure SQL Database, provisioned

The next step is adding SSISConfig database artifacts to the new Azure SQL Database.

Deploy the Simple, Custom, File-Based Azure-SSIS Framework

The focus of this section is deploying the SSISConfig database designed in Chapters 5 and 7 to an Azure SQL SSISConfig database instance.

To begin, open Azure Data Studio (or SQL Server Management Studio) and connect to the recently provisioned Azure SQL Database instance, as shown in Figure 9-11.

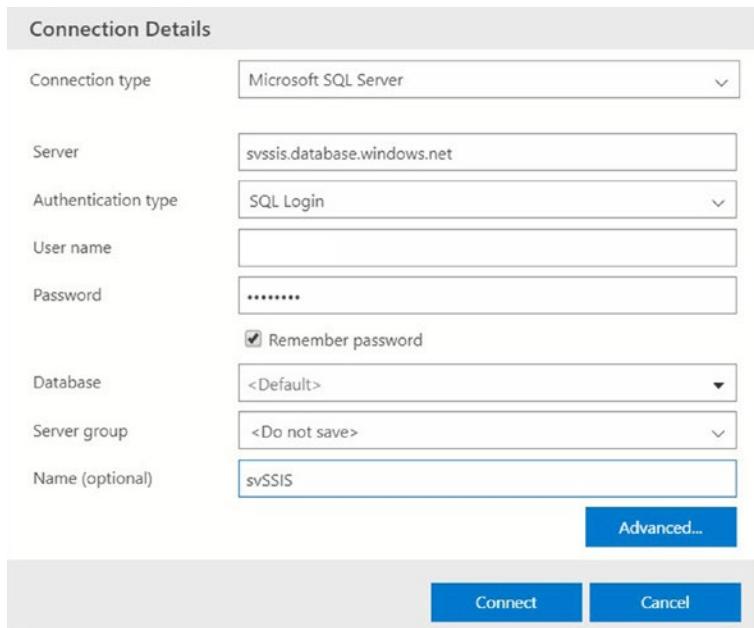


Figure 9-11. Connecting to the new Azure SQL Database

The SSISConfig database is created in the previous step. A script to create the SSISConfig database is similar to the scripts used in Chapters 5 and 7, as shown in Listings 9-1.

Listing 9-1. Creating SSISConfig in the Azure SQL database

```
print 'SSISConfig database'
If Not Exists(Select [databases].[name]
              From [sys].[databases]
              Where [databases].[name] = N'SSISConfig')
begin
    print ' - Create SSISConfig database'
    Create Database SSISConfig
    print ' - SSISConfig database created'
end
Else
```

```
begin
    print ' - SSISConfig database already exists.'
end
print ''
go
```

When execution completes, Azure Data Studio should display messages similar to the messages shown in Figure 9-12.

Messages	
7:44:14 AM	<u>Started executing query at Line 1</u> Commands completed successfully.
7:44:14 AM	<u>Started executing query at Line 3</u> SSISConfig database - Create SSISConfig database - SSISConfig database created

Figure 9-12. SSISConfig database created

At the time of this writing, the T-SQL in Azure SQL databases does not support the Use command. If you attempt to use the Use command, the error shown in Figure 9-13 is generated.

Instead of the Use command, use the database selector drop-down highlighted in Figure 9-13.

E: > FrameworksBook > ch09 > listing0902 Build SSISConfig database.sql

Run Cancel Disconnect Change Connection

```
1 use [SSISConfig]
2
```

master
SSISConfig
dbSSIS
master

Messages

7:54:39 AM Started executing query at Line 1

Msg 40508, Level 16, State 1, Line 1
USE statement is not supported to switch between databases.
Total execution time: 00:00:00.043

Figure 9-13. It's no "use"...

Listing 9-2 combines the remaining SSISConfig DDL (data definition language) statements found in Chapters 5 and 7 to create the SSISConfig database artifacts and metadata.

Listing 9-2. Creating SSISConfig database artifacts

```
print 'Config schema'
If Not Exists(Select [schemas].[name]
    From [sys].[schemas]
    Where [schemas].[name] = N'config')
begin
    print ' - Create config schema'
    declare @sql nvarchar(100) = N'Create Schema config'
    exec(@sql)
    print ' - Config schema created'
end
Else
```

```
begin
    print ' - Config schema already exists.'
end
print ''
go

print 'Config.Applications table'
If Not Exists(Select [schemas].[name] + '.' + [tables].[name]
As [Schema.Table]
    From [sys].[tables]
    Join [sys].[schemas]
        On [schemas].[schema_id] = [tables].[schema_id]
    Where [schemas].[name] = N'config'
        And [tables].[name] = N'Applications')
begin
    print ' - Create config.Applications table'
    Create Table [config].[Applications]
    (
        ApplicationId int identity(1, 1)
            Constraint PK_config_Applications Primary Key Clustered
        , ApplicationName nvarchar(255) Not NULL
            Constraint UQ_config_Applications_ApplicationName
            Unique
    )
    print ' - Config.Applications table created'
end
Else
begin
    print ' - Config.Applications table already exists.'
end
print ''
go
```

CHAPTER 9 DEPLOY A SIMPLE, CUSTOM, FILE-BASED AZURE-SSIS FRAMEWORK

```
print 'Config.Packages table'
If Not Exists(Select [schemas].[name] + '.' + [tables].[name]
As [Schema.Table]
    From [sys].[tables]
    Join [sys].[schemas]
        On [schemas].[schema_id] = [tables].[schema_id]
    Where [schemas].[name] = N'config'
        And [tables].[name] = N'Packages')

begin
    print ' - Create config.Packages table'
    Create Table [config].[Packages]
    (
        PackageId int identity(1, 1)
        Constraint PK_config_Packages Primary Key Clustered
        , PackageLocation nvarchar(255) Not NULL
        , PackageName nvarchar(255) Not NULL
        , Constraint UQ_config_Packages_PackageName
            Unique(PackageLocation, PackageName)
    )
    print ' - Config.Packages table created'
end
Else
begin
    print ' - Config.Packages table already exists.'
end
print ''
go

print 'Config.ApplicationPackages table'
If Not Exists(Select [schemas].[name] + '.' + [tables].[name]
As [Schema.Table]
    From [sys].[tables]
    Join [sys].[schemas]
        On [schemas].[schema_id] = [tables].[schema_id]
    Where [schemas].[name] = N'config'
        And [tables].[name] = N'ApplicationPackages')
```

```
begin
    print ' - Create config.ApplicationPackages table'
    Create Table [config].[ApplicationPackages]
    (
        ApplicationPackageId int identity(1, 1)
        Constraint PK_config_ApplicationPackages Primary Key Clustered
        , ApplicationId int Not NULL
            Constraint FK_config_ApplicationPackages_config_Applications
                Foreign Key References [config].[Applications](ApplicationId)
        , PackageId int Not NULL
            Constraint FK_config_ApplicationPackages_config_Packages
                Foreign Key References [config].[Packages](PackageId)
        , ExecutionOrder int Not NULL
            Constraint DF_config_ApplicationPackages_ExecutionOrder
                Default(10)
        , ApplicationPackageEnabled bit Not NULL
            Constraint DF_config_ApplicationPackages_ApplicationPackageEnabled
                Default(1)
        , FailApplicationOnPackageFailure bit Not NULL
            Constraint DF_config_ApplicationPackages_
                FailApplicationOnPackageFailure
                Default(1)
        , Constraint UQ_config_ApplicationPackages_ApplicationId_PackageId_
            ExecutionOrder
            Unique(ApplicationId, PackageId, ExecutionOrder)
    )
    print ' - Config.ApplicationPackages table created'
end
Else
begin
    print ' - Config.ApplicationPackages table already exists.'
end
print ''
go
```

CHAPTER 9 DEPLOY A SIMPLE, CUSTOM, FILE-BASED AZURE-SSIS FRAMEWORK

```
Set NoCount ON

declare @ApplicationName nvarchar(255) = N'Framework Test'

print @ApplicationName
declare @ApplicationId int = (Select [Applications].[ApplicationId]
                                From [config].[Applications]
                                Where [Applications].[ApplicationName] =
                                      @ApplicationName)

If (@ApplicationId Is NULL)
begin
    print ' - Adding ' + @ApplicationName + ' application to config.
          Applications table'

    declare @AppTbl table(ApplicationId int)

    Insert Into [config].[Applications]
    (ApplicationName)
    Output inserted.ApplicationId into @AppTbl
    Values (@ApplicationName)

    Set @ApplicationId = (Select ApplicationId
                           From @AppTbl)

    print ' - ' + @ApplicationName + ' application added to config.
          Applications table'
end
Else
begin
    print ' - ' + @ApplicationName + ' application already exists in the
          config.Applications table.'
end

Select @ApplicationId As ApplicationId
print ''
go

Set NoCount ON
```

```
declare @PackageLocation nvarchar(255) = N'E:\Projects\TestSSISolution\TestSSISProject\'  
declare @PackageName nvarchar(255) = N'ReportAndSucceed.dtsx'  
  
print @PackageLocation + @PackageName  
declare @PackageId int = (Select [Packages].[PackageId]  
                           From [config].[Packages]  
                           Where [Packages].[PackageLocation] = @  
                                 PackageLocation  
                                 And [Packages].[PackageName] = @PackageName)  
If (@PackageId Is NULL)  
begin  
    print ' - Adding ' + @PackageName + ' package to config.Packages table'  
  
    declare @PkgTbl table(PackageId int)  
  
    Insert Into [config].[Packages]  
    (PackageLocation, PackageName)  
    Output inserted.PackageId into @PkgTbl  
    Values (@PackageLocation, @PackageName)  
  
    Set @PackageId = (Select PackageId  
                      From @PkgTbl)  
  
    print ' - ' + @PackageName + ' package added to config.Packages table'  
end  
Else  
begin  
    print ' - ' + @PackageName + ' package already exists in the config.  
          Packages table.'  
end  
  
Select @PackageId As PackageId  
print ''
```

CHAPTER 9 DEPLOY A SIMPLE, CUSTOM, FILE-BASED AZURE-SSIS FRAMEWORK

```
set @PackageName = N'ReportAndFail.dtsx'

print @PackageLocation + @PackageName
set @PackageId = (Select [Packages].[PackageId]
                  From [config].[Packages]
                  Where [Packages].[PackageLocation] = @PackageLocation
                        And [Packages].[PackageName] = @PackageName)
If (@PackageId Is NULL)
begin
    print ' - Adding ' + @PackageName + ' application to config.Packages
          table'

Delete @PkgTbl

Insert Into [config].[Packages]
(PackageLocation, PackageName)
Output inserted.PackageId into @PkgTbl
Values (@PackageLocation, @PackageName)

Set @PackageId = (Select PackageId
                  From @PkgTbl)

print ' - ' + @PackageName + ' package added to config.Packages table'
end
Else
begin
    print ' - ' + @PackageName + ' package already exists in the config.
          Packages table.'
end

Select @PackageId As PackageId
print ''
go

Set NoCount ON

declare @ApplicationName nvarchar(255) = N'Framework Test'
declare @PackageLocation nvarchar(255) = N'E:\Projects\TestSSISSolution\
TestSSISProject'
declare @PackageName nvarchar(255) = N'ReportAndSucceed.dtsx'
```

```
declare @ExecutionOrder int = 10

print @ApplicationName + ' - ' + @PackageLocation + @PackageName

declare @ApplicationId int = (Select [Applications].[ApplicationId]
                                From [config].[Applications]
                                Where [Applications].[ApplicationName] =
                                      @ApplicationName)

declare @PackageId int = (Select [Packages].[PackageId]
                                From [config].[Packages]
                                Where [Packages].[PackageLocation] =
                                      @PackageLocation
                                And [Packages].[PackageName] = @PackageName)

declare @ApplicationPackageId int = (Select ApplicationPackageId
                                         From config.ApplicationPackages
                                         Where ApplicationId = @ApplicationId
                                         And PackageId = @PackageId
                                         And ExecutionOrder = @

ExecutionOrder)

If (@ApplicationPackageId Is NULL)
begin
    print ' - Assigning ' + @PackageName + ' package to '
          + @ApplicationName + ' application'
          + ' in config.ApplicationPackages table'
          + ' at ExecutionOrder ' + Convert(varchar(9), @ExecutionOrder)

    Insert Into [config].[ApplicationPackages]
    (ApplicationId
    , PackageId
    , ExecutionOrder)
    Values (@ApplicationId
            , @PackageId
            , @ExecutionOrder)
```

CHAPTER 9 DEPLOY A SIMPLE, CUSTOM, FILE-BASED AZURE-SSIS FRAMEWORK

```
print ' - ' + @PackageName + ' package assigned to '
      + @ApplicationName + ' application'
      + ' in config.ApplicationPackages table'
      + ' at ExecutionOrder ' + Convert(varchar(9), @ExecutionOrder)
end
Else
begin
print ' - ' + @PackageName + ' package already'
      + ' assigned to ' + @ApplicationName
      + ' application in config.ApplicationPackages table'
      + ' at ExecutionOrder ' + Convert(varchar(9), @ExecutionOrder)
      + '.'
end
print ''

set @PackageName = N'ReportAndFail.dtsx'
set @ExecutionOrder = 20

print @ApplicationName + ' - ' + @PackageLocation + @PackageName

set @ApplicationId = (Select [Applications].[ApplicationId]
                      From [config].[Applications]
                      Where [Applications].[ApplicationName] =
                            @ApplicationName)

set @PackageId = (Select [Packages].[PackageId]
                      From [config].[Packages]
                      Where [Packages].[PackageLocation] = @PackageLocation
                            And [Packages].[PackageName] = @PackageName)

set @ApplicationPackageId = (Select ApplicationPackageId
                                From config.ApplicationPackages
                                Where ApplicationId = @ApplicationId
                                      And PackageId = @PackageId
                                      And ExecutionOrder = @ExecutionOrder)
```

```
If (@ApplicationPackageId Is NULL)
begin
    print ' - Assigning ' + @PackageName + ' package to '
    + @ApplicationName + ' application'
    + ' in config.ApplicationPackages table'
    + ' at ExecutionOrder ' + Convert(varchar(9), @ExecutionOrder)

    Insert Into [config].[ApplicationPackages]
    (ApplicationId
    , PackageId
    , ExecutionOrder)
    Values (@ApplicationId
            , @PackageId
            , @ExecutionOrder)

    print ' - ' + @PackageName + ' package assigned to '
    + @ApplicationName + ' application'
    + ' in config.ApplicationPackages table'
    + ' at ExecutionOrder ' + Convert(varchar(9), @ExecutionOrder)
end
Else
begin
    print ' - ' + @PackageName + ' package already'
    + ' assigned to ' + @ApplicationName
    + ' application in config.ApplicationPackages table'
    + ' at ExecutionOrder ' + Convert(varchar(9), @ExecutionOrder)
    + '.'.

end
print ''
go

print 'Log schema'
If Not Exists(Select [schemas].[name]
    From [sys].[schemas]
    Where [schemas].[name] = N'log')
```

```

begin
    print ' - Create log schema'
    declare @sql nvarchar(100) = N'Create Schema log'
    exec(@sql)
    print ' - Log schema created'
end
Else
begin
    print ' - Log schema already exists.'
end
print ''
go

print 'Log.ApplicationInstance table'
If Not Exists(Select [schemas].[name] + '.' + [tables].[name]
As [Schema.Table]
    From [sys].[tables]
    Join [sys].[schemas]
        On [schemas].[schema_id] = [tables].[schema_id]
    Where [schemas].[name] = N'log'
        And [tables].[name] = N'ApplicationInstance')

begin
    print ' - Create log.ApplicationInstance table'
    Create Table [log].[ApplicationInstance]
    (
        ApplicationInstanceId int identity(1, 1)
            Constraint PK_log_ApplicationInstance Primary Key Clustered
        , ApplicationId int Not NULL
            Constraint FK_log_ApplicationInstance_config_Applications
                Foreign Key References [config].[Applications](ApplicationId)
        , ApplicationStartTime datetimeoffset(7) Not NULL
            Constraint DF_log_ApplicationInstance_ApplicationStartTime
                Default(sysdatetimeoffset())
        , ApplicationEndTime datetimeoffset(7) NULL
        , ApplicationStatus nvarchar(25) Not NULL
            Constraint DF_log_ApplicationInstance_ApplicationStatus

```

```

    Default(N'Running')
)
print ' - Log.ApplicationInstance table created'
end
Else
begin
    print ' - Log.ApplicationInstance table already exists.'
end
print ''
go

print 'Log.ApplicationPackageInstance table'
If Not Exists(Select [schemas].[name] + '.' + [tables].[name]
As [Schema.Table]
    From [sys].[tables]
    Join [sys].[schemas]
        On [schemas].[schema_id] = [tables].[schema_id]
    Where [schemas].[name] = N'log'
        And [tables].[name] = N'ApplicationPackageInstance')
begin
    print ' - Create log.ApplicationPackageInstance table'
    Create Table [log].[ApplicationPackageInstance]
    (
        ApplicationPackageInstanceId int identity(1, 1)
            Constraint PK_log_ApplicationPackageInstance Primary Key Clustered
        , ApplicationInstanceId int Not NULL
            Constraint FK_log_ApplicationPackageInstance_log_ApplicationInstance
                Foreign Key References [log].[ApplicationInstance]
                    (ApplicationInstanceId)
        , ApplicationPackageId int Not NULL
            Constraint FK_log_ApplicationPackageInstance_config_
                ApplicationPackages
                Foreign Key References [config].[ApplicationPackages]
                    (ApplicationPackageId)
        , ApplicationPackageStartTime datetimeoffset(7) Not NULL
    )
)

```

```

    Constraint DF_log_ApplicationPackageInstance_
    ApplicationPackageStartTime
        Default(sysdatetimeoffset())
    , ApplicationPackageEndTime datetimeoffset(7) NULL
    , ApplicationPackageStatus nvarchar(25) Not NULL
        Constraint DF_log_ApplicationPackageInstance_
        ApplicationPackageStatus
        Default(N'Running')
)
print ' - Log.ApplicationPackageInstance table created'
end
Else
begin
    print ' - Log.ApplicationPackageInstance table already exists.'
end
print ''
go

```

If all goes as planned, the script returns results shown in Figure 9-14.

Results		Messages	
ApplicationId			
1			
PackageId			
1			
PackageId			
1			

Figure 9-14. Results of query execution

The Messages returned from executing the T-SQL query in Listing 9-2 are shown in Listing 9-3.

Listing 9-3. Messages returned from SSISConfig artifact and metadata creation

Started executing query at Line 1

Config schema

- Create config schema
- Config schema created

Started executing query at Line 17

Config.Applications table

- Create config.Applications table
- Config.Applications table created

Started executing query at Line 43

Config.Packages table

- Create config.Packages table
- Config.Packages table created

Started executing query at Line 70

Config.ApplicationPackages table

- Create config.ApplicationPackages table
- Config.ApplicationPackages table created

Started executing query at Line 110

Framework Test

- Adding Framework Test application to config.Applications table
- Framework Test application added to config.Applications table

Started executing query at Line 143

E:\Projects\TestSSISolution\TestSSISProject\ReportAndSucceed.dtsx

- Adding ReportAndSucceed.dtsx package to config.Packages table
- ReportAndSucceed.dtsx package added to config.Packages table

E:\Projects\TestSSISolution\TestSSISProject\ReportAndFail.dtsx

- Adding ReportAndFail.dtsx application to config.Packages table
- ReportAndFail.dtsx package added to config.Packages table

CHAPTER 9 DEPLOY A SIMPLE, CUSTOM, FILE-BASED AZURE-SSIS FRAMEWORK

Started executing query at Line 209

Framework Test - E:\Projects\TestSSISolution\TestSSISProject\
ReportAndSucceed.dtsx

- Assigning ReportAndSucceed.dtsx package to Framework Test application ➤
in config.ApplicationPackages table at ExecutionOrder 10
- ReportAndSucceed.dtsx package assigned to Framework Test application ➤
in config.ApplicationPackages table at ExecutionOrder 10

Framework Test - E:\Projects\TestSSISolution\TestSSISProject\
ReportAndFail.dtsx

- Assigning ReportAndFail.dtsx package to Framework Test application ➤
in config.ApplicationPackages table at ExecutionOrder 20
- ReportAndFail.dtsx package assigned to Framework Test application ➤
in config.ApplicationPackages table at ExecutionOrder 20

Log.ApplicationInstance table

- Create log.ApplicationInstance table
- Log.ApplicationInstance table created

Log.ApplicationPackageInstance table

- Create log.ApplicationPackageInstance table
- Log.ApplicationPackageInstance table created

Total execution time: 00:00:00.557

Test the deployment and test metadata insert using the T-SQL query shown in Listing 9-4.

Listing 9-4. Retrieve SSIS application metadata

```
Select a.ApplicationName  
    , p.PackageLocation + p.PackageName As PackagePath  
    , ap.ExecutionOrder  
    , ap.FailApplicationOnPackageFailure  
From [config].[ApplicationPackages] ap  
Join [config].[Applications] a  
    On a.ApplicationId = ap.ApplicationId  
Join [config].Packages p  
    On p.PackageId = ap.PackageId
```

```
Where a.ApplicationName = N'Framework Test'
And ap.ApplicationPackageEnabled = 1
Order By ap.ExecutionOrder
```

If all has gone according to plan, your results should appear similar to Figure 9-15.

ApplicationName	PackagePath	ExecutionOrder	FailApplicationOnPackageFailure
Framework Test	E:\Projects\TestSSISolution\TestSSISProject\ReportAndSucceed.dtsx	10	1
Framework Test	E:\Projects\TestSSISolution\TestSSISProject\ReportAndFail.dtsx	20	1

Figure 9-15. A Test SSIS application in SSISConfig in Azure SQL database

If your SSISConfig database query results appear as shown in Figure 9-10, congratulations! You followed the instructions correctly. But there's an issue with the instructions: how will the Azure Data Factory version of the SSIS execution engine find my – or your – local drive (my E drive, in this case)? The short answer is, “That would be difficult.”

Before we update the PackagePath metadata, let's first provision an Azure File Share.

Provision an Azure File Share

Begin by downloading Microsoft Azure Storage Explorer at storageexplorer.com, as shown in Figure 9-16.



Figure 9-16. Browse to StorageExplorer.com to download Azure Storage Explorer

Install Azure Storage Explorer, and connect to the Azure Storage account provisioned in the previous chapter. When connected, storage explorer surfaces an Explorer window that appears similar to Figure 9-17.

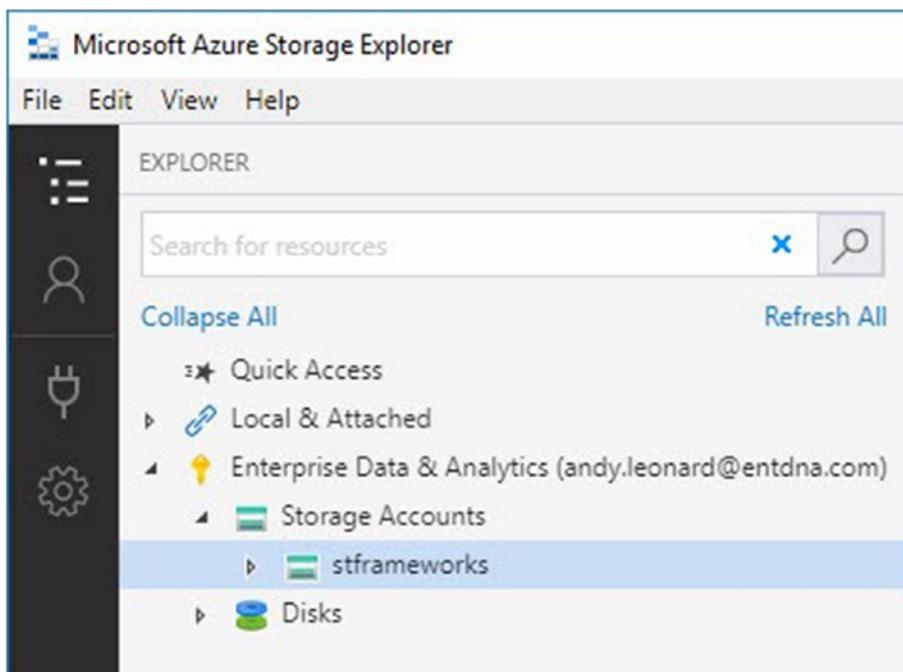


Figure 9-17. Viewing Explorer in Azure Storage Explorer

In Chapter 8, we provisioned a storage account named “stframeworks” which we see in the screenshot of Microsoft Azure Storage Explorer in Figure 9-17. Expand the stframeworks storage account to surface virtual folders for Blob Containers, File Shares, Queues, and Tables, as shown in Figure 9-18.

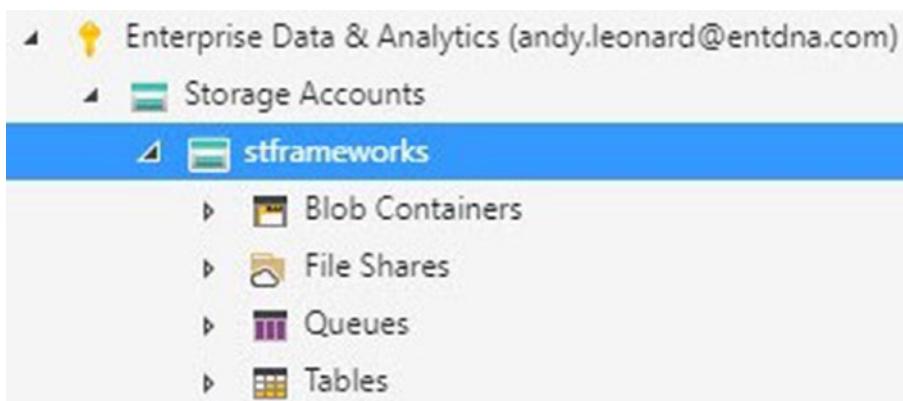


Figure 9-18. stframeworks virtual folders

Right-click the File Shares virtual folder, and click “Create File Share” as shown in Figure 9-19.

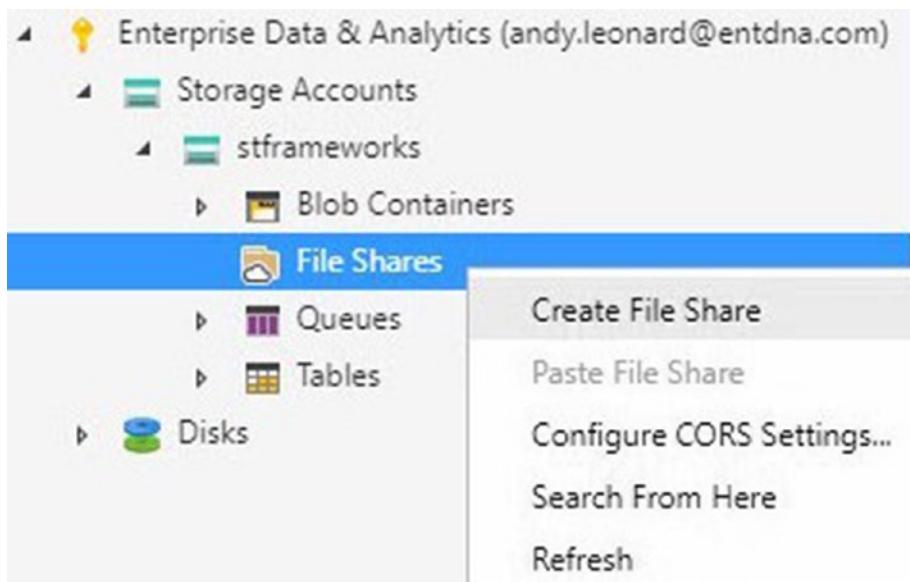


Figure 9-19. Creating a new File Share

When the new file share displays, enter a name for the file share – such as “fs-ssis,” as shown in Figure 9-20.

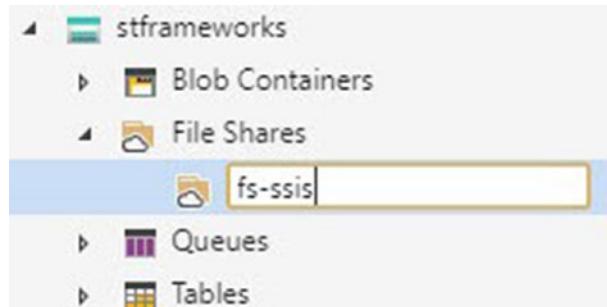


Figure 9-20. Naming the new file share

Press the Enter key to finish creating the file share.

Our new file share opens and is now ready to store SSIS packages in Azure, as shown in Figure 9-21.

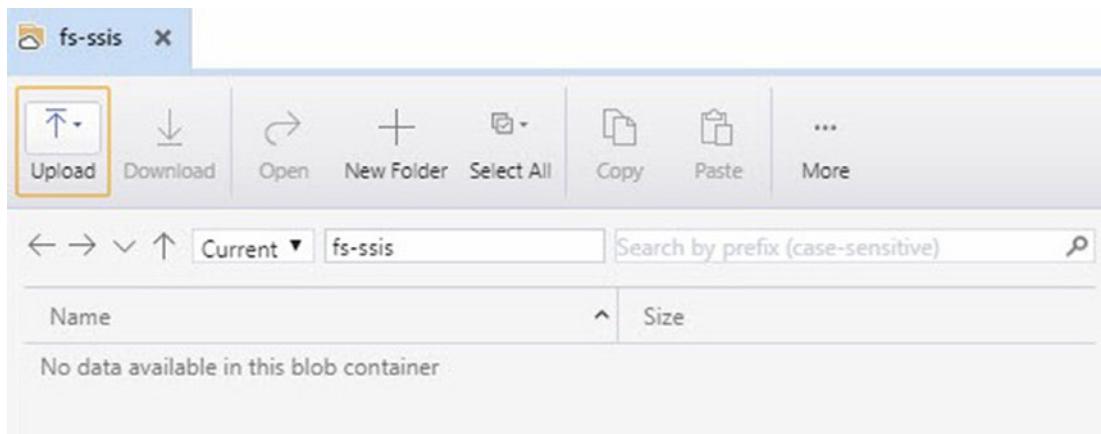


Figure 9-21. The new fs-ssis file share

Upload SSIS Packages

Uploading SSIS packages is straightforward using Storage Explorer. Click the Upload button to begin, as shown in Figure 9-22.

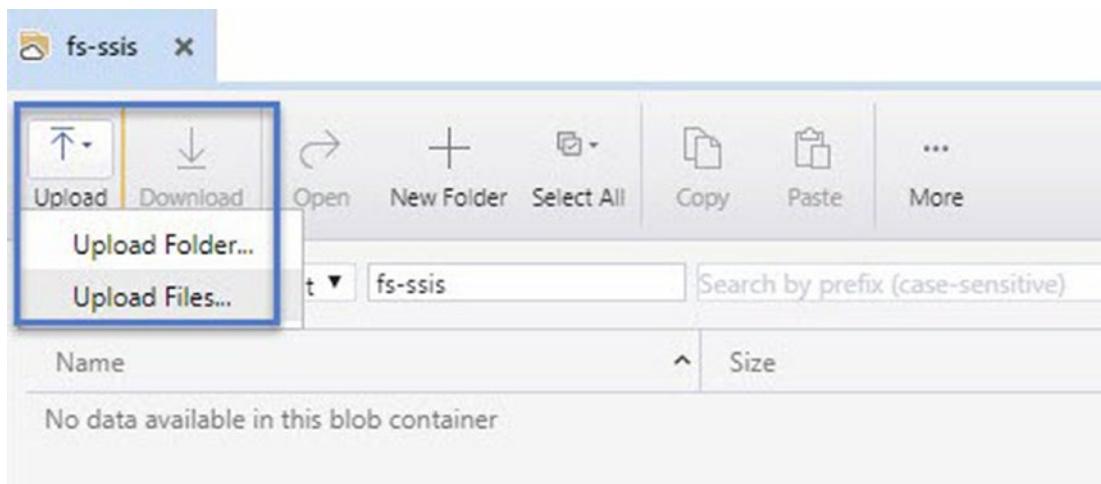


Figure 9-22. Preparing to upload SSIS packages to the file share

You may upload a folder or files, as shown in Figure 9-22.

Click “Upload Files...” to open the Upload Files dialog. First, select a file or files to be uploaded – such as the ReportAndSucceed.dtsx and ReportAndFail.dtsx SSIS packages developed as described in Chapter 5. Second, accept the default Destination directory (“/”). Third, click the Upload button, as shown in Figure 9-23.

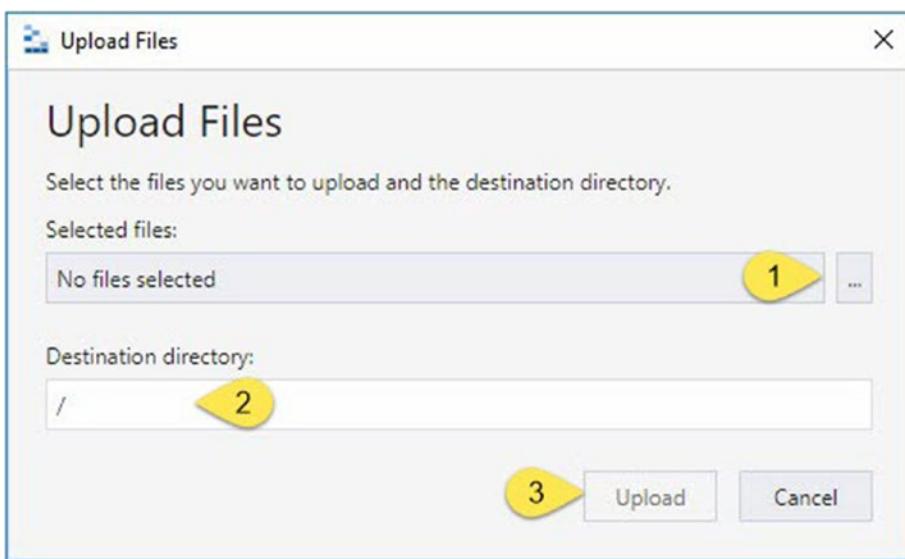


Figure 9-23. Configuring the upload to the file share

When the upload completes, you may view the results of the operation in the Azure Storage Explorer Activities window, as shown in Figure 9-24.

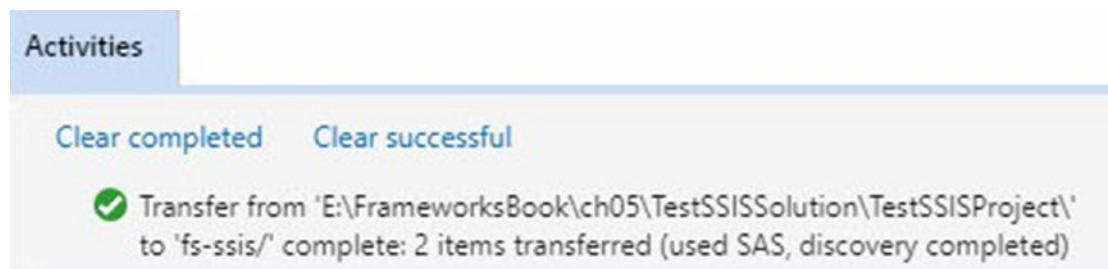


Figure 9-24. Viewing activities associated with the upload

ReportAndSucceed.dtsx and ReportAndFail.dtsx SSIS packages have been uploaded and are now visible in the fs-ssis file share, as shown in Figure 9-25.

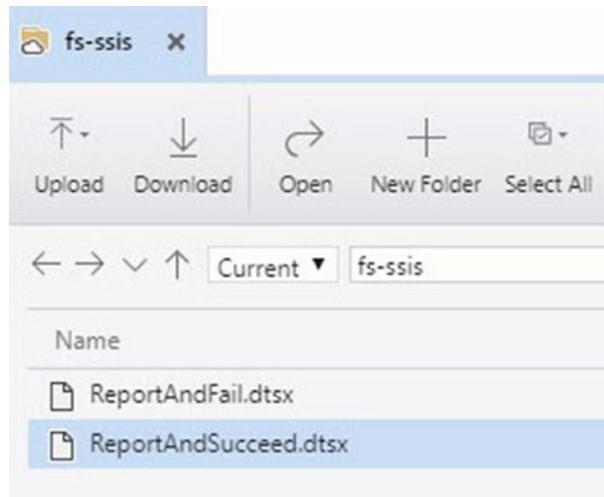


Figure 9-25. ReportAndSucceed.dtsx and ReportAndFail.dtsx SSIS packages uploaded

The next step is to update the metadata in the SSISConfig config.Packages table to reflect the file share location of the packages.

Update PackageLocation Values

Question: What is the difference between a simple, custom, file-based SSIS framework that executes SSIS packages stored in an on-premises file store and a simple, custom, file-based SSIS framework that executes SSIS packages stored in an Azure file share?

Answer: One metadata value.

Execute the T-SQL Update statement shown in Listing 9-5 to update the PackageLocation values stored in the config.Packages table.

Listing 9-5. Updating the config.Packages' PackageLocation values

```
Select p.PackageLocation + p.PackageName As PackagePath
From [config].[Packages] p

Update [config].[Packages]
Set PackageLocation = N'\\stframeworks.file.core.windows.net\fs-ssis\' 
Where PackageLocation = N'E:\Projects\TestSSISSolution\TestSSISProject\' 

Select p.PackageLocation + p.PackageName As PackagePath
From [config].[Packages] p
```

Results from executing the T-SQL shown in Listing 9-5 should appear similar to Figure 9-26.

	Results	Messages
	PackagePath	
1	E:\Projects\TestSSISolution\TestSSISProject\ReportAndFail.dtsx	
2	E:\Projects\TestSSISolution\TestSSISProject\ReportAndSucceed.dtsx	
	PackagePath	
1	\stframeworks.file.core.windows.net\fs-ssis\ReportAndFail.dtsx	
2	\stframeworks.file.core.windows.net\fs-ssis\ReportAndSucceed.dtsx	

Figure 9-26. Updating the config.Packages metadata

We are now ready to build the SSIS framework execution engine, Azure Data Factory edition.

Build the SSIS Framework ADF Execution Engine

Until this point, deploying a simple, custom, file-based SSIS framework that executes SSIS packages stored in an Azure file share required no changes to the underlying architecture of the SSIS framework metadata database, only metadata was changed, and then only the values stored in a single field.

The SSIS framework execution engine is an Azure Data Factory pipeline, and pipelines are different from SSIS packages. These differences will drive changes in the database and elsewhere.

Retrieve a List of SSIS Packages

To begin building the SSIS framework execution engine in ADF, connect to the Author and Monitor site for the Azure Data Factory instance, click the “+” beside the “Filter resources by name” textbox, and then click “Pipeline” to create a new pipeline, as shown in Figure 9-27.

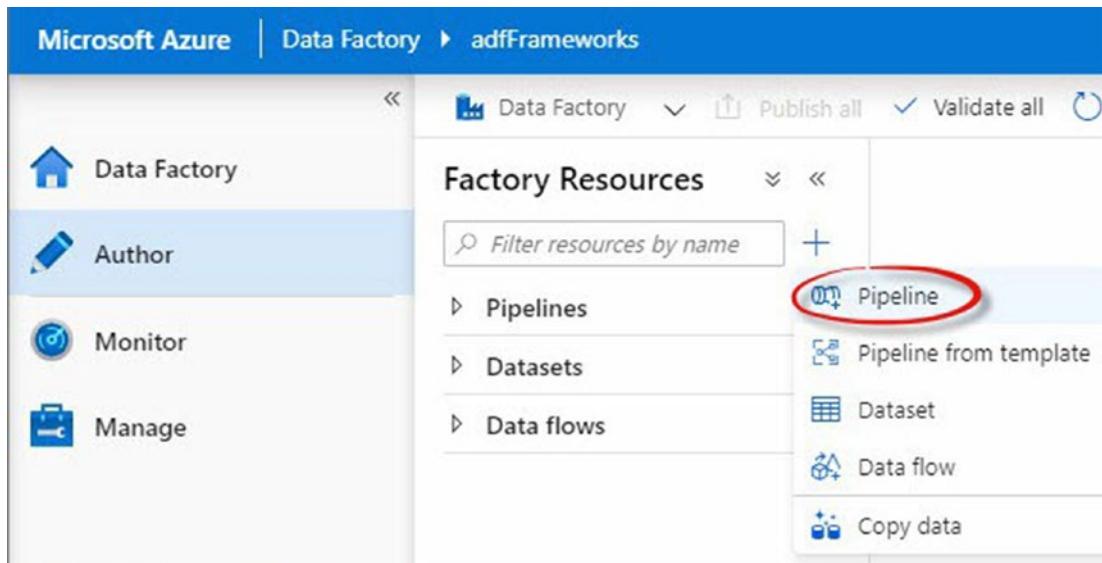


Figure 9-27. Creating a new ADF pipeline

When the new pipeline opens, change the name to “parent,” as shown in Figure 9-28.



Figure 9-28. Renaming the new pipeline “parent”

In the Activities blade, expand the General category, and then drag a Lookup activity onto the pipeline canvas. Rename the Lookup activity to “Get Application Packages,” as shown in Figure 9-29.

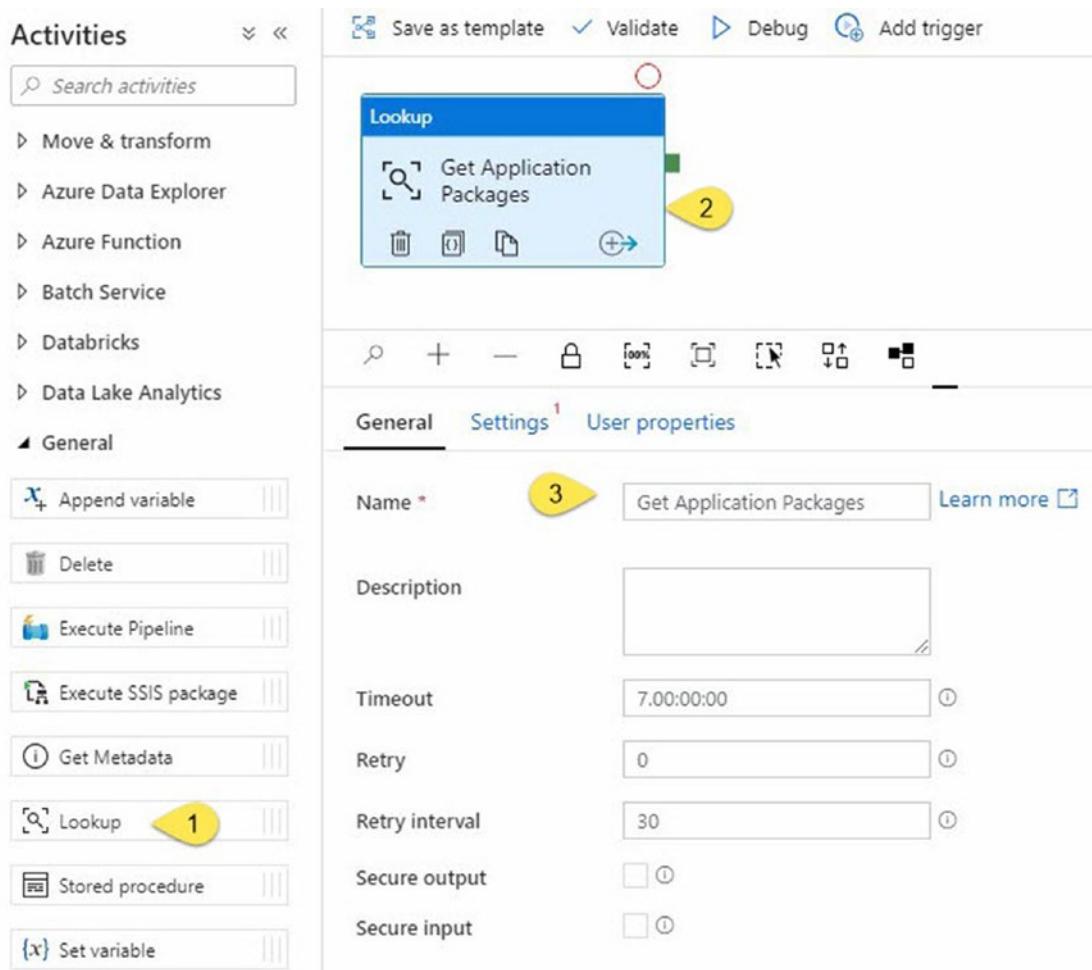


Figure 9-29. Adding and renaming a *Lookup* activity to the parent pipeline

One nice feature of Azure Data Factory is flexible order of authoring, or “FOA.” FOA means one may approach ADF development from either direction, bottom up or top down.

Bottom-up ADF development involves building Linked Services *first* – to allow pipeline artifacts to communicate with data stores and services that reside outside the pipeline. Once Linked Services are built in bottom-up fashion, Datasets are developed. Datasets consume Linked Services and surface collections of data to pipeline artifacts. If you are familiar with SSIS, ADF Linked Services are analogous to SSIS Connection Managers, and ADF Datasets are (less) analogous to SSIS Data Flow source and destination adapters.

For this example, I will leverage FOA for top-down development. What does top-down ADF development look like? I will click and select “New” buttons, links, and options to construct ADF artifacts like Datasets and Linked Services in the next few sections.

On the “Get Application Packages” lookup activity’s Settings tab, click the “+ New” link to configure the “Source dataset” property, as shown in Figure 9-30.

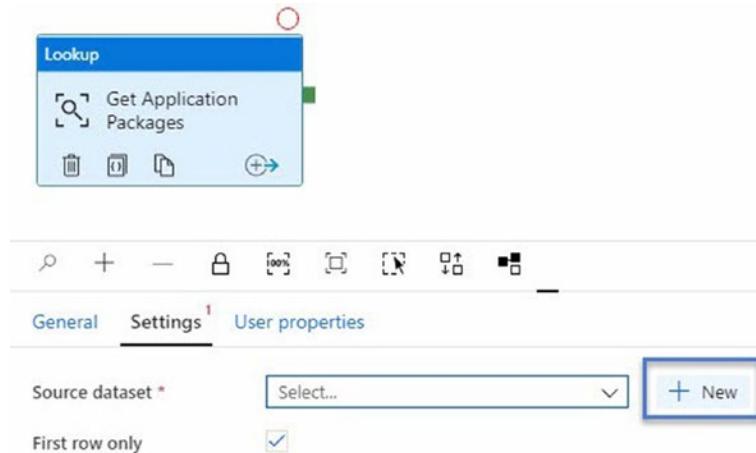


Figure 9-30. Add a new Source dataset

When the New dataset blade displays, select “Azure SQL Database,” and then click the Continue button, as shown in Figure 9-31.

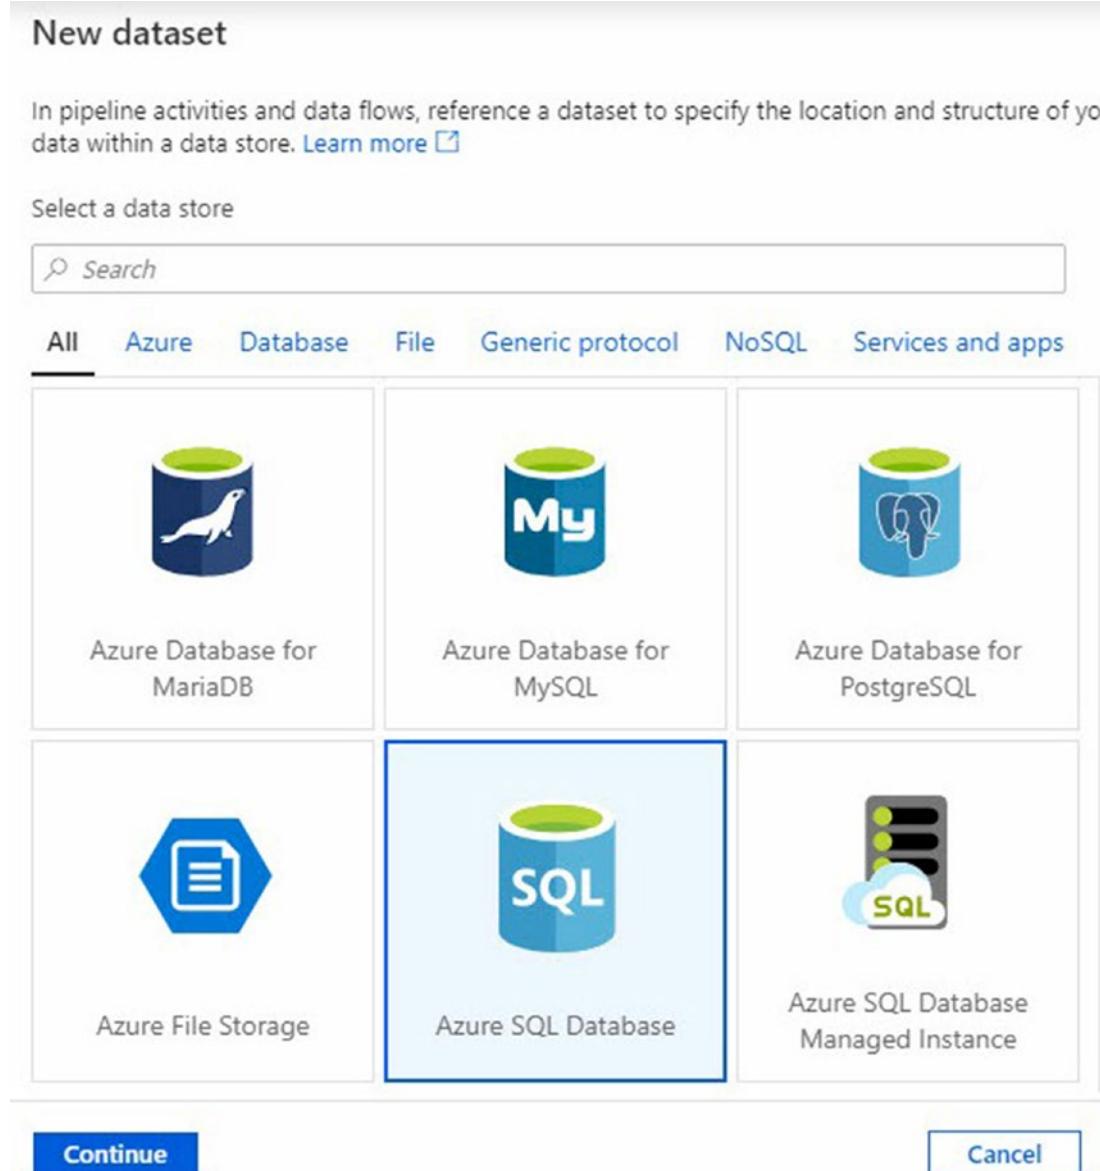


Figure 9-31. Configuring the New dataset's data store type

Set the dataset's Name property; I named my dataset "ssisFrameworkDataSet". Click the Linked Service drop-down, and then click the "+ New" option as shown in Figure 9-32.

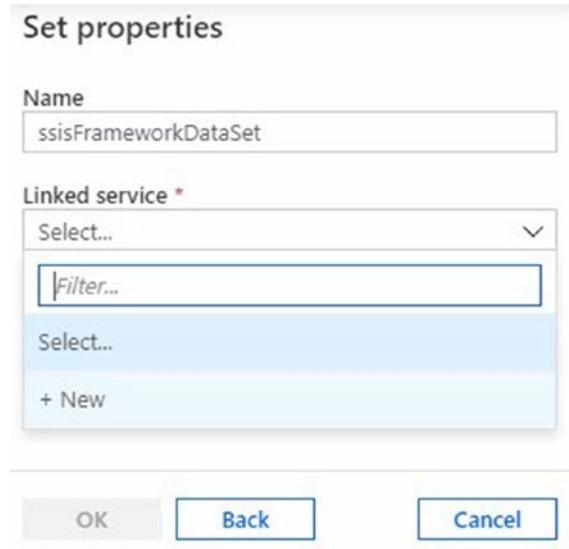


Figure 9-32. Configuring the dataset Name and Linked service properties

When the “New linked service (Azure SQL Database)” page displays, enter a Name and optional Description. Leave the “Connect via integration runtime” option set to its default value, which is “AutoResolveIntegrationRuntime”.

For the purposes of this example, select the “Connection string” connection source, and then select the “From Azure subscription” option for the “Account selection method” property. Select your “Azure subscription” in the subscription drop-down, select the name of the server you created when provisioning the Azure SQL database from the “Server name” drop-down, and then select the name of the Azure SQL database you provisioned from the “Database name” drop-down.

Select “SQL authentication” in the “Authentication type” drop-down, and enter your Azure SQL database user name and password in the “User name” and “Password” textboxes, respectively.

You are now ready to test the connection. Click the “Test connection” link to test connectivity. The connection test should fail, as shown in Figure 9-33.

New linked service (Azure SQL Database)

Name *

Description

Connect via integration runtime * ⓘ

Connection string **Azure Key Vault**

Account selection method From Azure subscription Enter manually ⓘ

Azure subscription ⏺

Server name * ⏺

Database name *

Authentication type *

User name *

Password **Azure Key Vault**

Password *

Additional connection properties

Annotations

Connection failed

Cannot connect to SQL Database: 'svssis.database.windows.net', Database: 'SSISConfig', User: 'adminandy'. Check the linked service configuration is correct, and make sure the SQL Database firewall allows the integration runtime to access. Cannot open server 'svssis' requested by the login. Client with IP address '20.44.17.80' is not allowed to access the server. To enable access, use the Windows Azure Management Portal or run sp_set_firewall_rule on the master database to create a firewall rule for this IP address or address range. It may take up to five minutes for this change to take effect.
SqlErrorNumber=40615,Class=14,State=1, Activity ID:

Is this helpful? Yes No

✖ Connection failed. [More](#)

🔗 Test connection Cancel

Create

Figure 9-33. Failed connection test

The reason this error is included in this example is because it is such a common error. To resolve the error, open the Azure Portal in a new browser tab and hover over “SQL databases” in the Azure left menu to prompt the display of the SQL databases “hover card,” as shown in Figure 9-34.

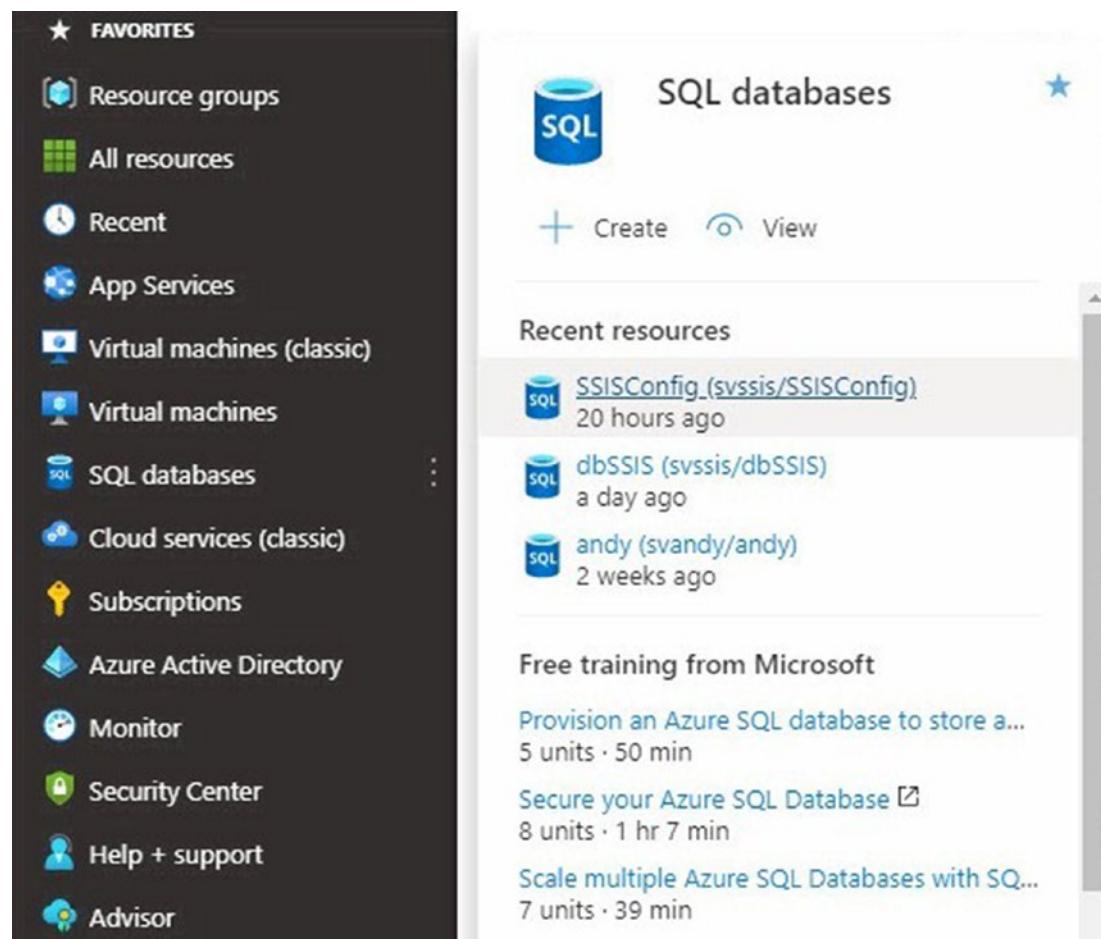


Figure 9-34. Viewing SQL databases

Click the recently provisioned SSISConfig database to open the SSISConfig page, as shown in Figure 9-35.

The screenshot shows the Azure portal interface for the SSISConfig database. The left sidebar lists 'Overview', 'Activity log', 'Tags', 'Diagnose and solve problems', 'Quick start', and 'Query editor (preview)'. The main content area displays the following details:

Resource group (change)	: rgFrameworks	Server name	: svssis.database.windows.net
Status	: Online	Elastic pool	: No elastic pool
Location	: East US 2	Connection strings	: Show database connection strings
Subscription (change)	: Enterprise Data & Analytics	Pricing tier	: General Purpose: Gen5, 2 vCores
Subscription ID	:	Earliest restore point	: 2020-06-02 12:17 UTC
Tags (change)	: Click here to add tags		

Figure 9-35. The SSISConfig Azure SQL database blade

On the Overview page, click the Server name link (“svssis.database.windows.net” in this case) – shown in Figure 9-35 – to open the server blade, as shown in Figure 9-36.

The screenshot shows the Azure portal interface for the SVSSIS server. The left sidebar lists 'Overview', 'Activity log', 'Access control (IAM)', 'Tags', and 'Diagnose and solve problems'. The main content area displays the following details:

Resource group (change)	: rgFrameworks	Server admin	: adminandy
Status	: Available	Firewalls and virtual netw...	: Show firewall settings
Location	: East US 2	Active Directory admin	: Not configured
Subscription (change)	: Enterprise Data & Analytics	Server name	: svssis.database.windows.net
Subscription ID	:		
Tags (change)	: Click here to add tags		

Figure 9-36. The server that hosts the SSISConfig Azure SQL database

When the server overview page displays, click the “Show firewall settings” link to open the “Firewalls and virtual networks” page. Change the “Allow Azure services and resources to access this server” to “Yes,” as shown in Figure 9-37.

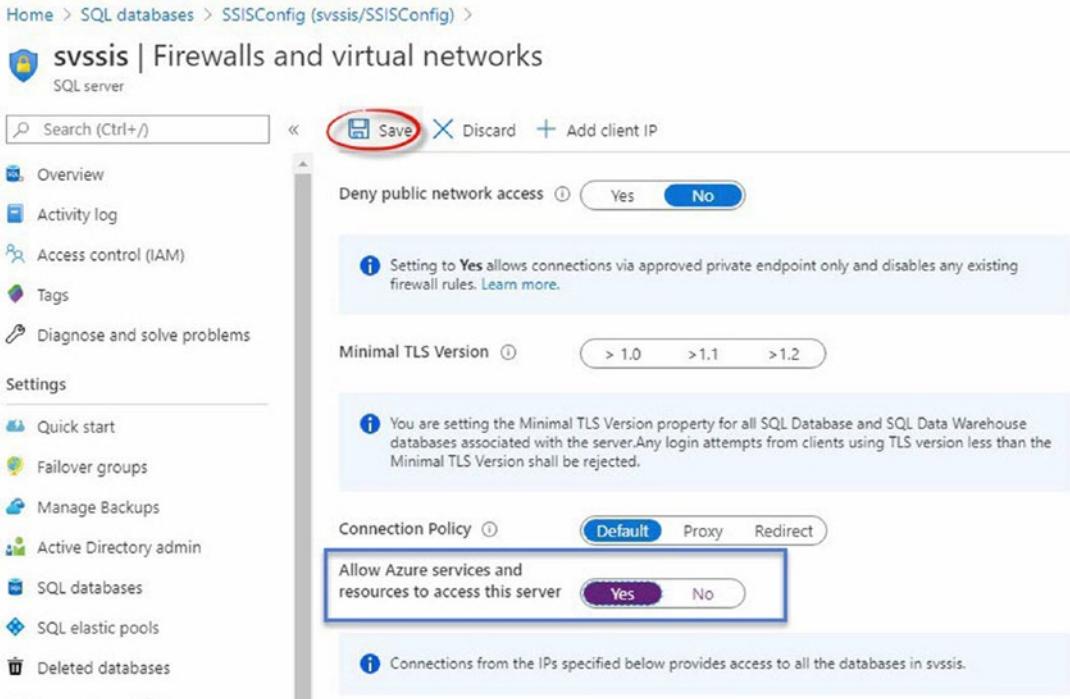


Figure 9-37. Updating “Allow Azure services and resources to access this server”

Click the Save button shown in Figure 9-37 to store the update to “Allow Azure services and resources to access this server,” and then return to the “New linked service” configuration in Azure Data Factory. Make sure you complete the Linked Service configuration (in the following image, the “User name” property value is redacted). Click “Test connection” to retest the connection to the SSISConfig database, similar to Figure 9-38.

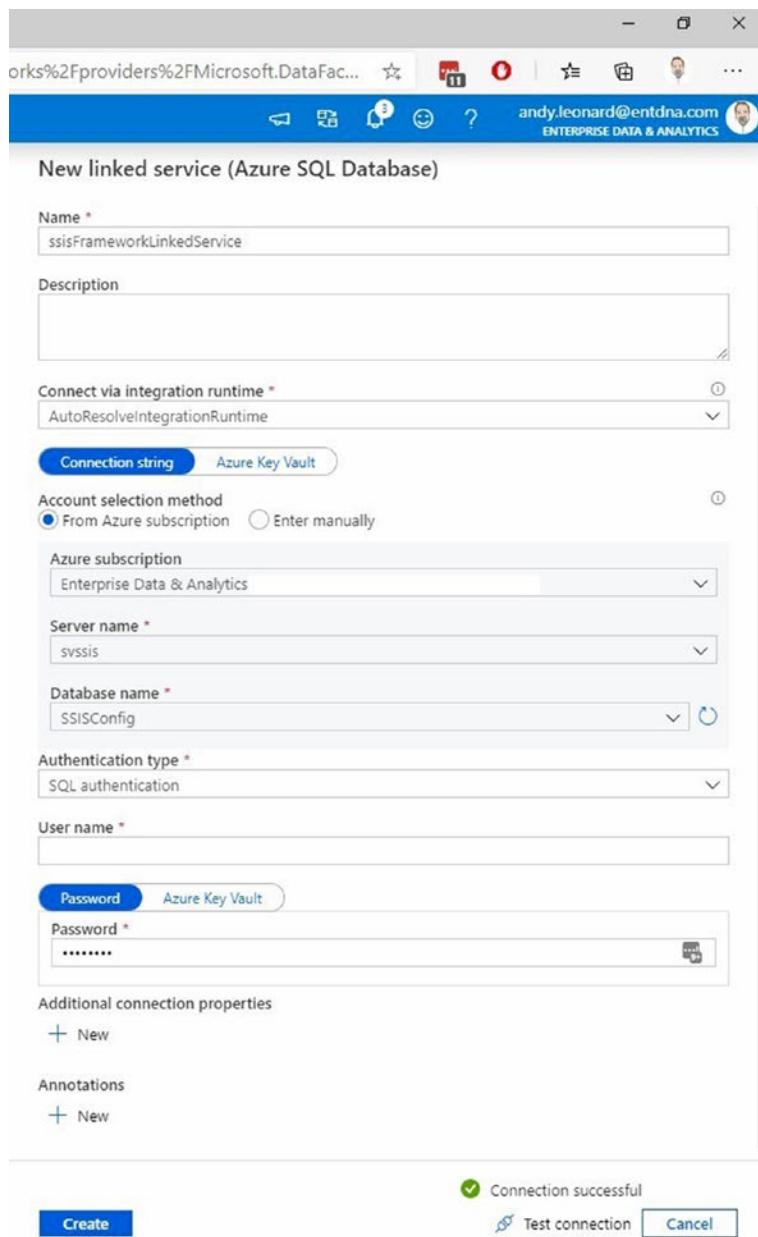


Figure 9-38. Successful connection test

The preceding section has not only served as an example of FOA (flexible order of authoring), the section also provided one solution to a very common Azure security obstacle.

Continue configuring the “Get Application Packages” lookup activity’s Settings tab by changing the “Use query” property to “Query.” Initialize the “Query” property with the T-SQL statement shown in Listing 9-6.

Listing 9-6. T-SQL to select “Framework Test” SSIS application packages

```
Select a.ApplicationName  
    , p.PackageLocation + p.PackageName As PackagePath  
    , ap.ExecutionOrder  
    , ap.FailApplicationOnPackageFailure  
From [config].[ApplicationPackages] ap  
Join [config].[Applications] a  
    On a.ApplicationId = ap.ApplicationId  
Join [config].Packages p  
    On p.PackageId = ap.PackageId  
Where a.ApplicationName = N'Framework Test'  
    And ap.ApplicationPackageEnabled = 1  
Order By ap.ExecutionOrder
```

Leave the “Query timeout (minutes)” and “Isolation level” properties set to their defaults for the purposes of this example. Make sure the “First row only” checkbox is unchecked, as shown in Figure 9-39.

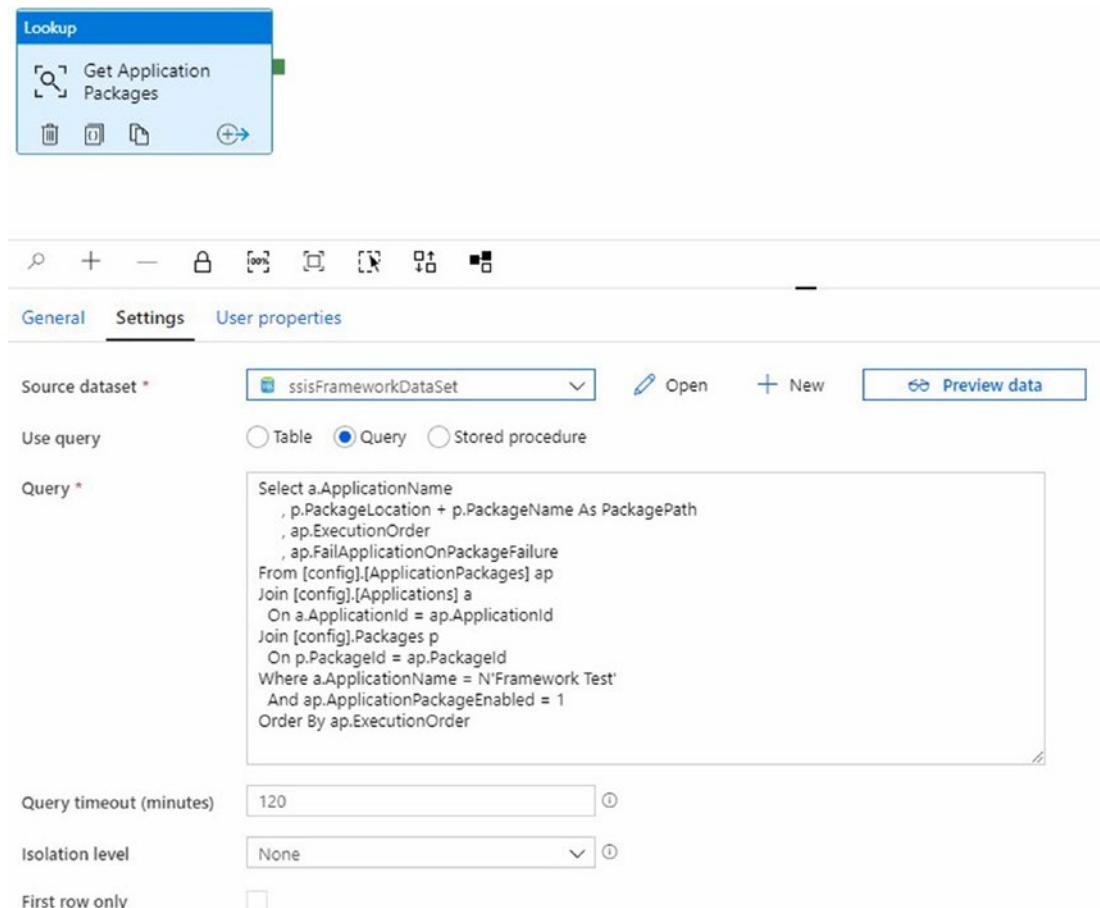


Figure 9-39. Completing “Get Application Packages” lookup activity settings

The “Get Application Packages” lookup activity may now be tested by clicking the Debug item in the toolbar, as shown in Figure 9-40.

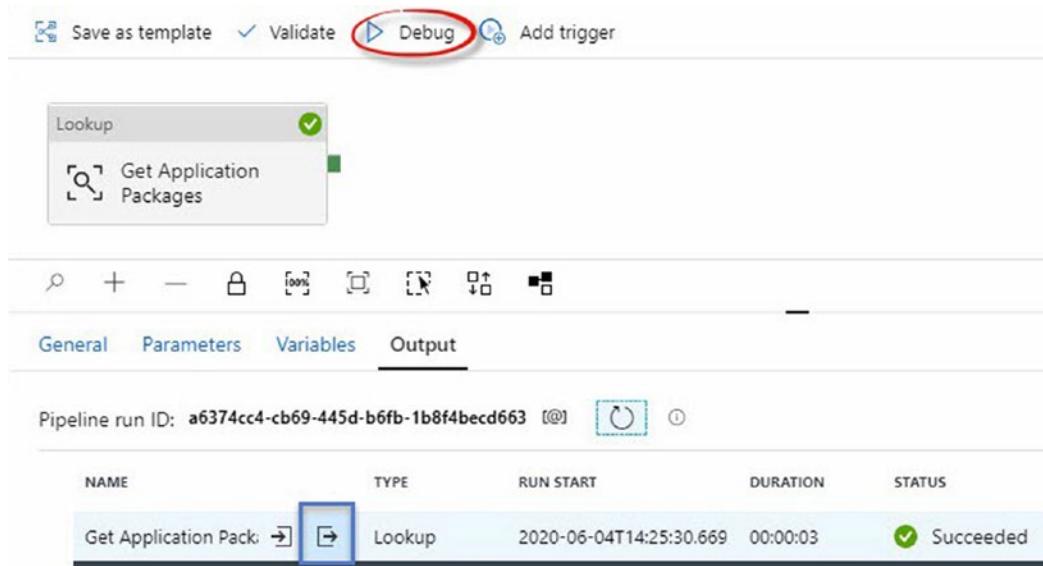


Figure 9-40. View the output (blue box) after a Debug execution (red circle)

After the Debug execution completes, click the Output viewer shown on the “Get Application Packages” lookup activity’s Output tab (“blue box” in Figure 9-40) to view the JSON returned by the Lookup activity, shown in Figure 9-41.

```
{
  "count": 2,
  "value": [
    {
      "ApplicationName": "Framework Test",
      "PackagePath": "\\\stframeworks.file.core.windows.net\\fs-ssis\\ReportAndSucceed.dtsx",
      "ExecutionOrder": 10,
      "FailApplicationOnPackageFailure": true
    },
    {
      "ApplicationName": "Framework Test",
      "PackagePath": "\\\stframeworks.file.core.windows.net\\fs-ssis\\ReportAndFail.dtsx",
      "ExecutionOrder": 20,
      "FailApplicationOnPackageFailure": true
    }
  ],
  "effectiveIntegrationRuntime": "DefaultIntegrationRuntime (East US 2)",
  "billingReference": {
    "activityType": "PipelineActivity",
    "billableDuration": [
      {
        "meterType": "AzureIR",
        "duration": 0.016666666666666666,
        "unit": "DIUHours"
      }
    ]
  },
  "durationInQueue": {
    "integrationRuntimeQueue": 0
  }
}
```

Figure 9-41. JSON returned by “Get Application Packages” lookup activity

Now that we have a list of SSIS packages, the next step is to execute the SSIS packages contained in the list.

Execute the Retrieved SSIS Packages

In the SSIS version of the simple, custom SSIS framework, a Foreach Loop Container enumerates the collection of SSIS packages. In Azure Data Factory, a ForEach activity performs the same function.

Expand the Activities category named “Iteration & conditionals,” and drag a ForEach activity onto the pipeline canvas. Connect a Success pipe from the “Get Application Packages” lookup activity to the ForEach activity, and then rename the ForEach activity to “ForEach Application Package,” as shown in Figure 9-42.

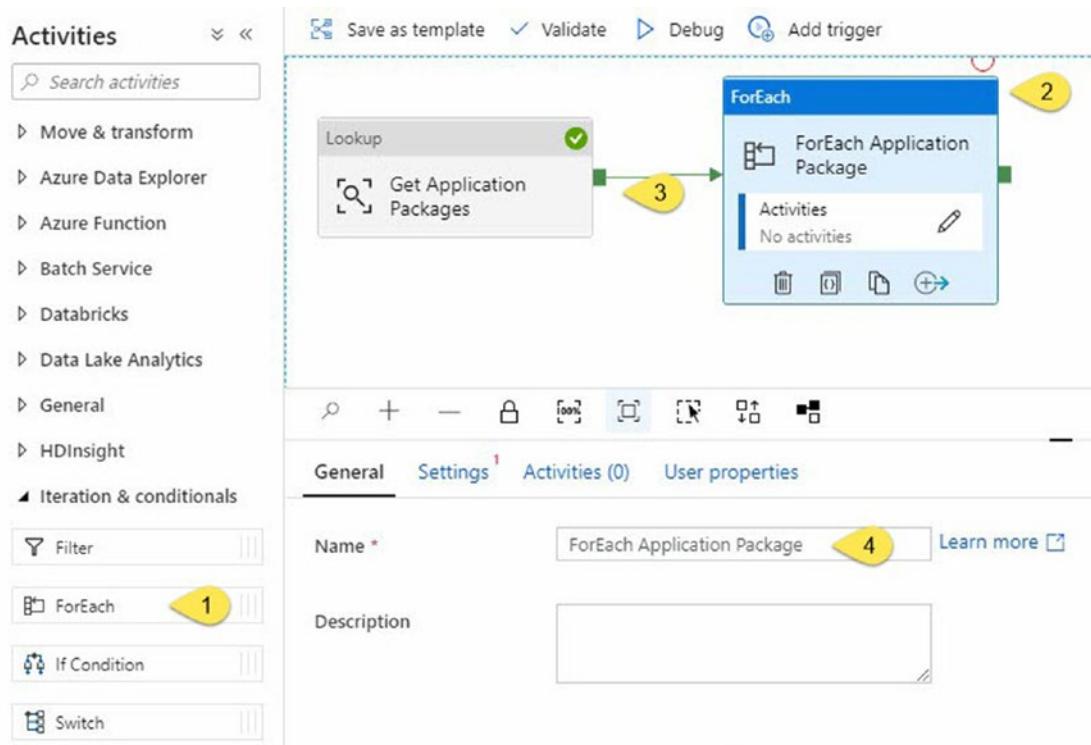


Figure 9-42. Adding a ForEach activity

On the Settings tab, click inside the Items property textbox. The “Add dynamic content [Alt + P]” link is displayed just below the Items textbox, as shown in Figure 9-43.

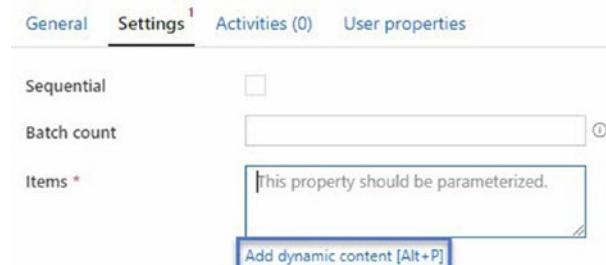


Figure 9-43. The “Add dynamic content [Alt + P]” link on the “ForEach Application Package” ForEach activity’s Settings tab

Click the “Add dynamic content [Alt + P]” link to open the “Add dynamic content” blade. In the “Activity outputs” category, click Get Application Packages. @activity('Get Application Packages').output appears in the Add dynamic content textbox. Append “.value” to the expression so that the Azure Data Factory expression reads @activity('Get Application Packages').output.value, as shown in Figure 9-44.

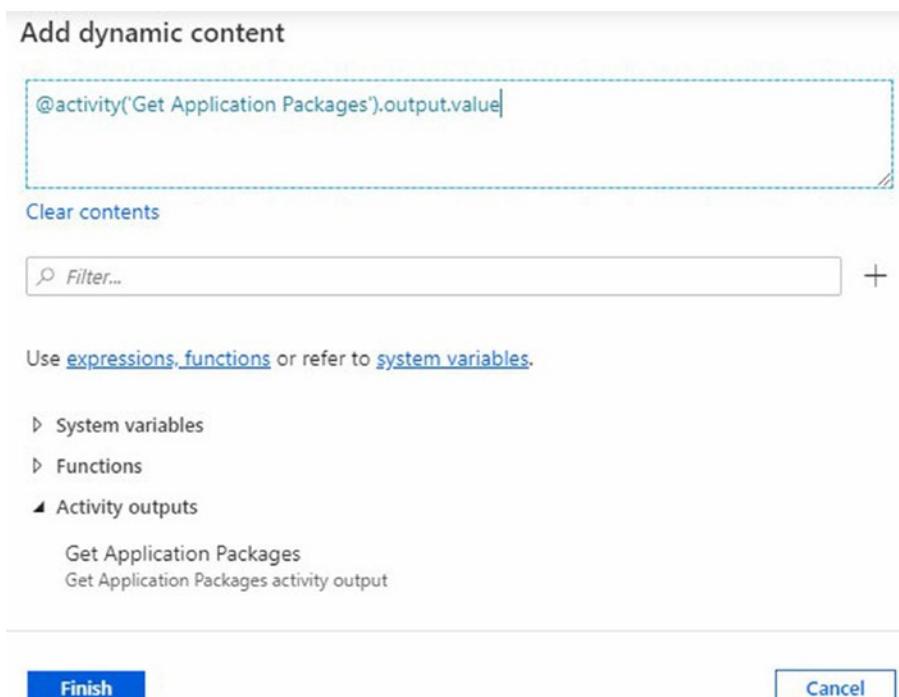


Figure 9-44. Setting the items property dynamic value for the ForEach Application Packages activity

Click the Finish button to complete the property configuration. The ForEach Application Packages Items property should now appear similar to Figure 9-45.

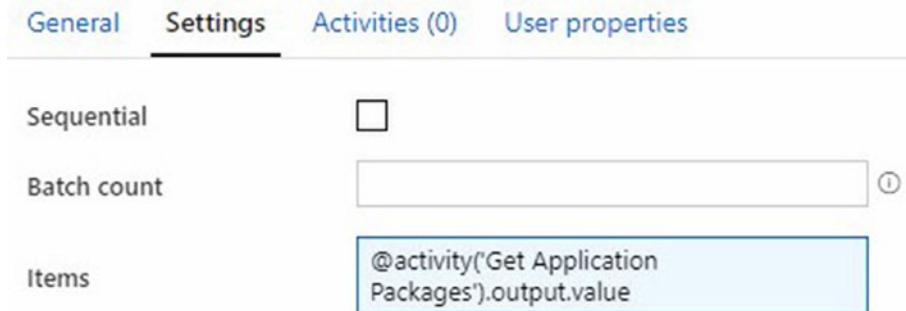


Figure 9-45. The ForEach Application Packages Items property, configured

The enumerator is now configured. The next step is to configure the activities we wish the parent pipeline to perform for each application package.

Click ForEach Application Package's activities Edit icon (the pencil), shown in Figure 9-46.



Figure 9-46. Edit the ForEach Application Package's activities

When the parent ► ForEach Application Package ► Activities page displays, drag an “Execute SSIS package” activity from the Activities ► General pipeline activities, and then rename the new Execute SSIS package activity to “Execute Application Package,” as shown in Figure 9-47.

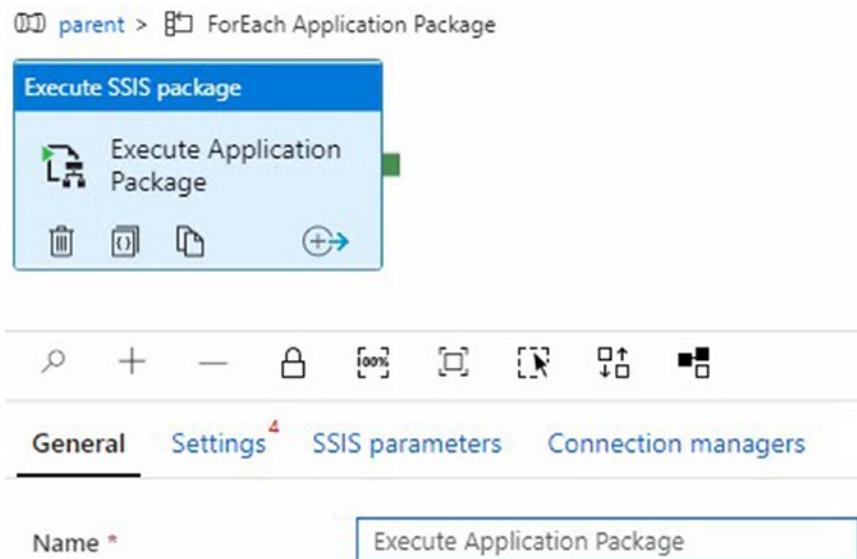


Figure 9-47. Adding and renaming the “Execute Application Package” execute SSIS package activity

Configure the “Execute Application Package” execute SSIS activity’s Settings tab Azure-SSIS IR property, selecting the “Azure-SSIS-Files” integration runtime from the drop-down, as shown in Figure 9-48.

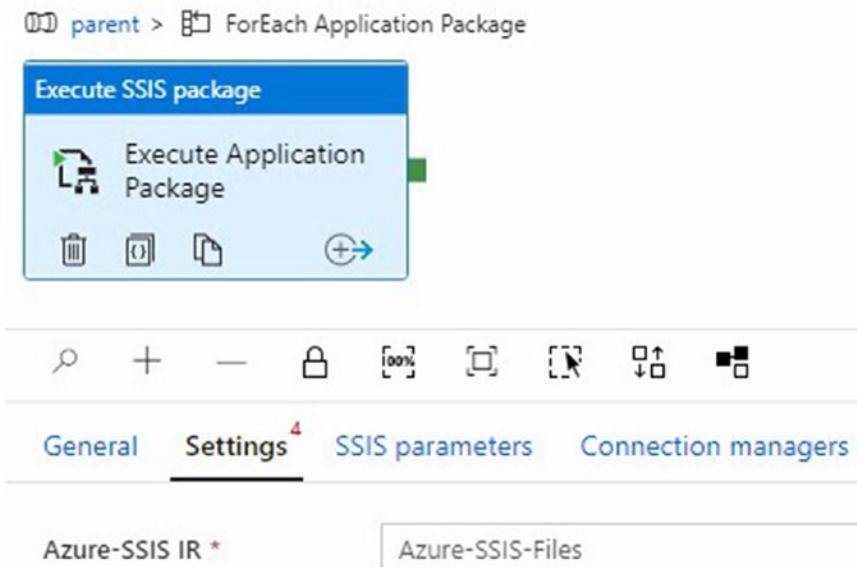


Figure 9-48. Selecting the Azure-SSIS-Files integration runtime

For the purposes of this example, ensure the Windows authentication and 32-bit runtime property checkboxes are unchecked. Click inside the “Package path” property’s textbox, and then click the “Add dynamic content [Alt + P]” link to surface the Add dynamic content blade, expand the ForEach iterator category, click the ForEach Application Package Current Item option, and then append “.PackagePath” to “@item()” in the dynamic content textbox, as shown in Figure 9-49.

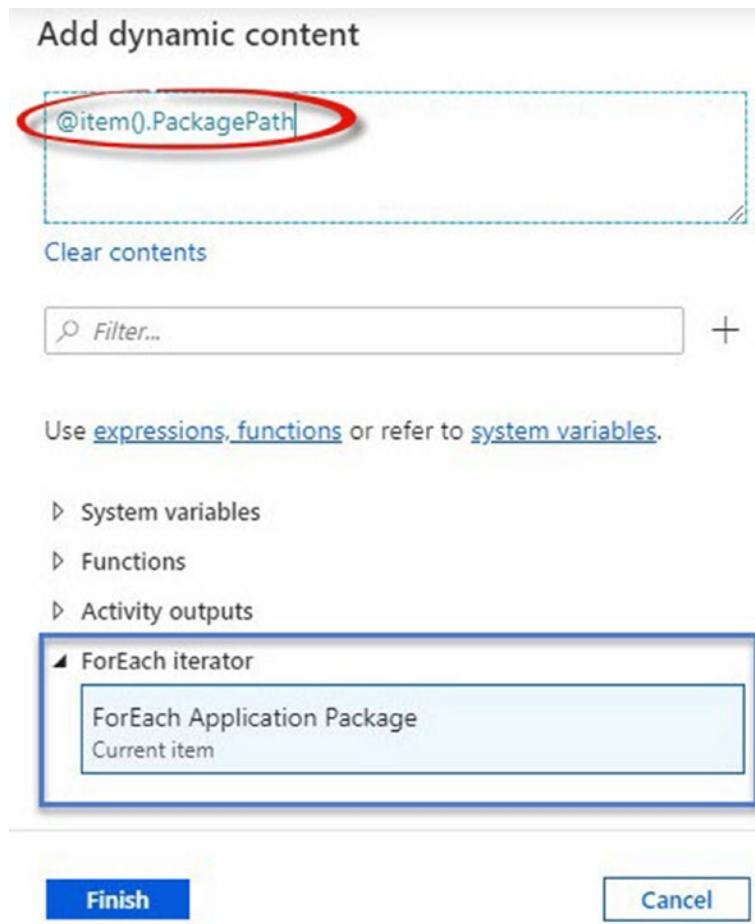


Figure 9-49. Configuring the PackagePath value in ForEach Application Package activities

Click the Finish button to return to the Settings tab, as shown in Figure 9-50.

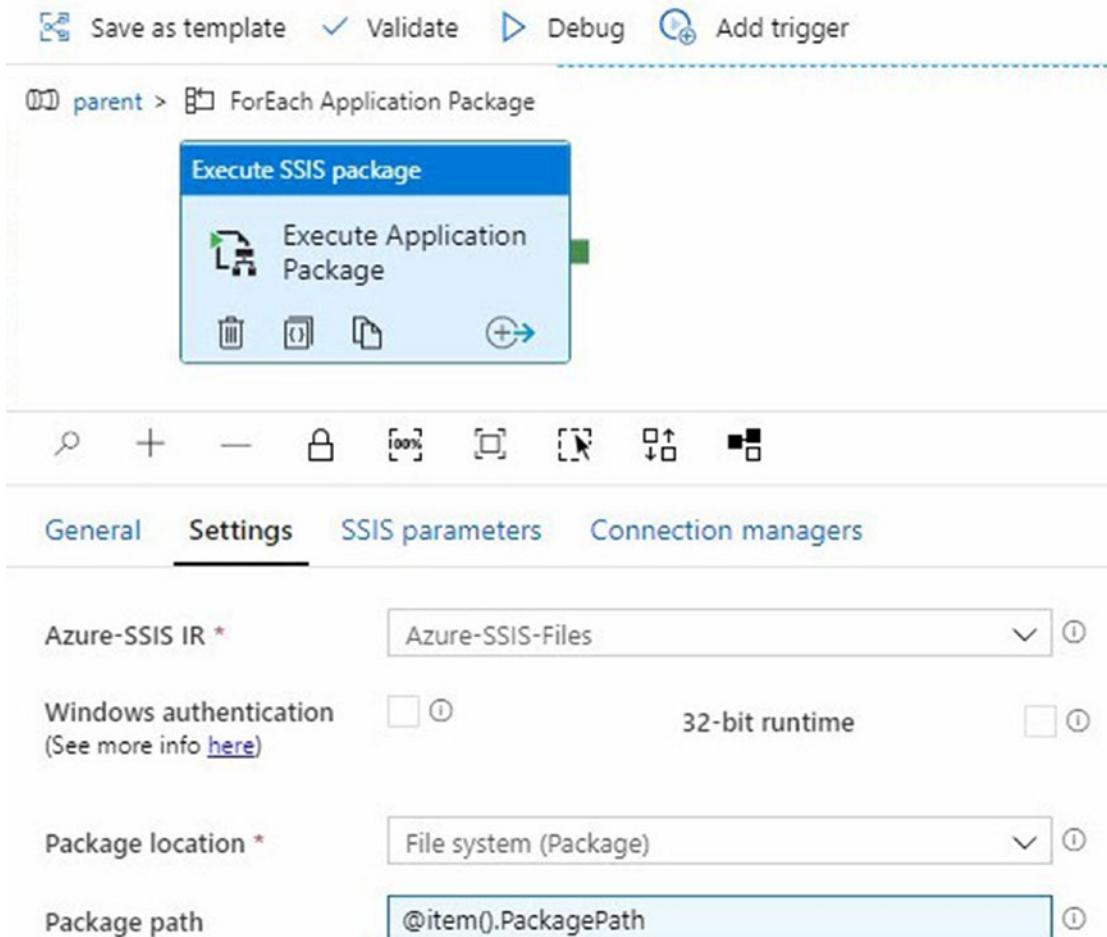


Figure 9-50. Execute Application Package's Package path property, configured

For the purposes of this example, leave the “Configuration path” property empty, although it contains “sample” configuration metadata, set the Domain property to “Azure,” and set the Username property to the name of the Azure File Share – “stframeworks” in this case, as shown in Figure 9-51.

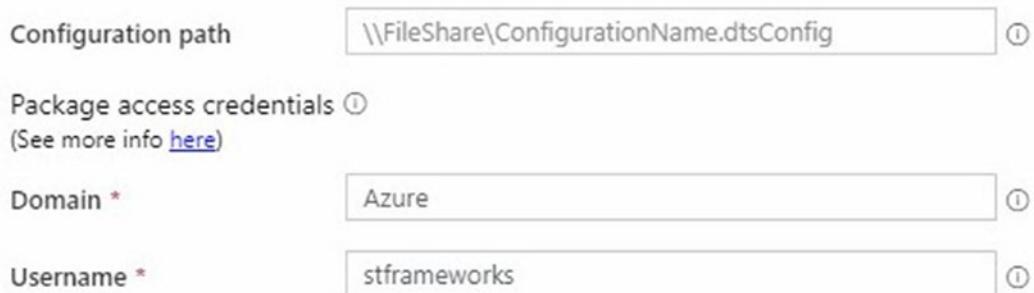


Figure 9-51. Configuring the Execute SSIS package Configuration path, Domain, and Username properties

Configuring the Password property is the next step, and it is perhaps the most unintuitive property of all. The value required for the Execute SSIS package activity is not a *password*, it is a *key* for accessing the Azure file share. To obtain this value, open Microsoft Azure Storage Explorer and connect to the Azure file share. In the Explorer pane, right-click the Azure file share, and then click “Copy Primary Key,” as shown in Figure 9-52.

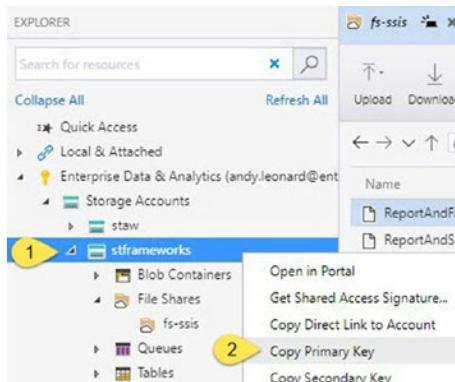


Figure 9-52. Copying the Azure file share's primary key for access

The Azure file share primary key is the value required for the password property. Return to the Azure Data Factory portal, and paste the Azure file share primary key into the Password property textbox, as shown in Figure 9-53.

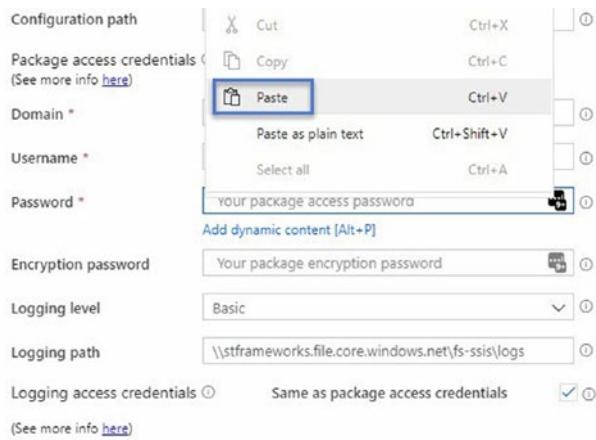


Figure 9-53. Pasting the Azure file share primary key into the Execute SSIS package activity's Password property

If the SSIS package's Package Protection Level property is set to EncryptAllWithPassword or EncryptSensitiveWithPassword, the password in the SSIS package's Package Password property must be supplied to the Execute SSIS package activity's "Encryption password" property (#1 in Figure 9-54). Since our test SSIS project packages use the default Protection Level, EncryptSensitiveWithUserKey, we may leave the Encryption password property empty. For the purposes of this example, accept the default "Logging level" property setting (Basic - #2 in Figure 9-54) and set the "Logging path" property to the Azure file share path plus "\logs" (#3 in Figure 9-54). Set the "Logging access credentials" property by checking the "Same as package access credentials" checkbox (#4 in Figure 9-54).

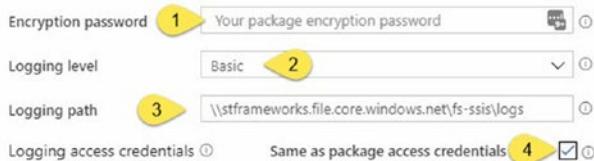


Figure 9-54. Configuring the Encryption password, Logging level, Logging path, and Logging access credentials properties for the Execute SSIS package activity

The parent pipeline is now minimally configured, enough so we can click the Debug menu item and perform a test execution (make sure the Azure-SSIS-Files integration runtime is running first!). By design (at this point), the pipeline debug execution should fail, as shown in Figure 9-55.

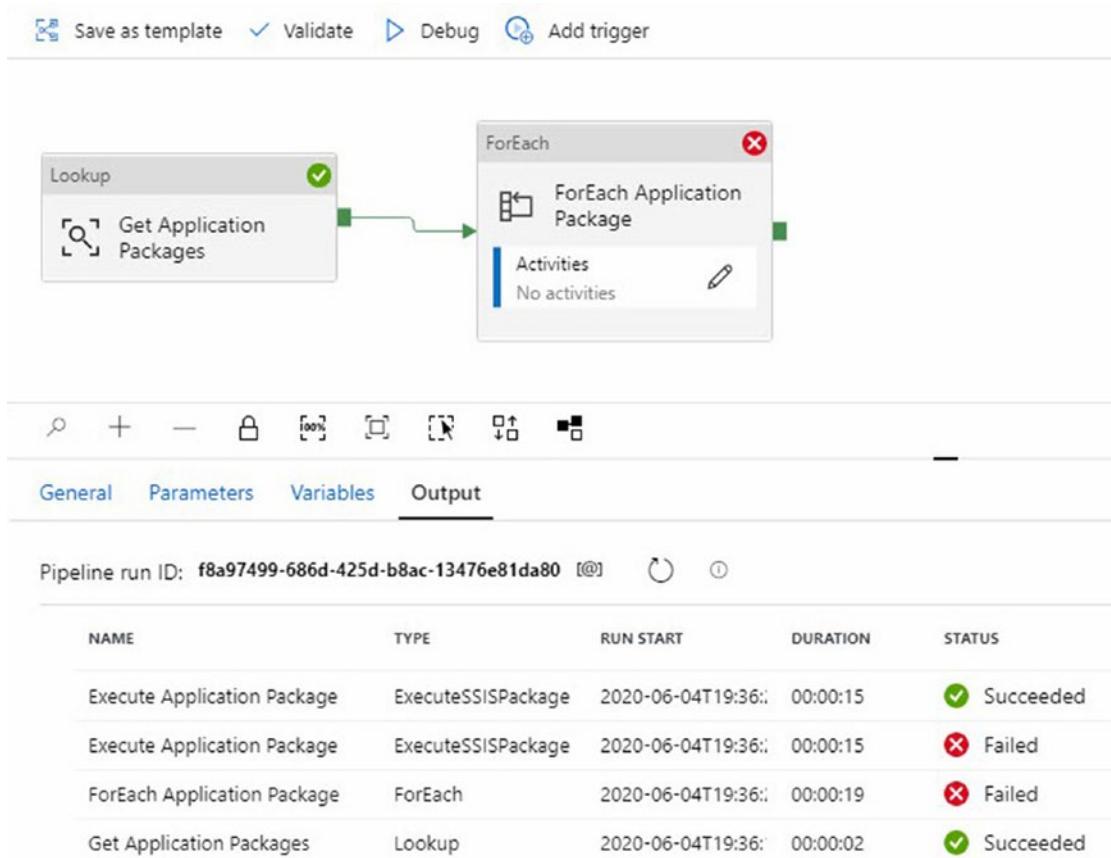


Figure 9-55. Failed test execution or parent pipeline

The test execution failed, but why?

In Figure 9-55, the first execution of the “Execute Application Package” execute SSIS activity failed, but the second execution of the “Execute Application Package” execute SSIS activity succeeded. You may see *both* executions fail. The most likely cause is the Azure-SSIS-Files integration runtime is not running.

The next step is to view the log file generated in the Azure file share’s logs folder.

View Test Execution Logs

Returning to storage explorer, the logs folder is likely *not* displayed. Refresh the Azure file share from the “More” drop-down menu item. Click “Refresh,” as shown in Figure 9-56.

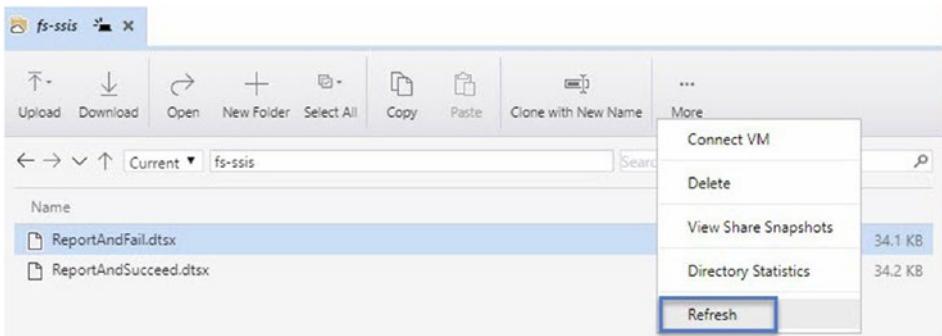


Figure 9-56. Clicking More ► Refresh to view the logs folder

The logs folder, shown in Figure 9-57, should have been created (if the logs folder did not exist beforehand) when the test execution occurred.

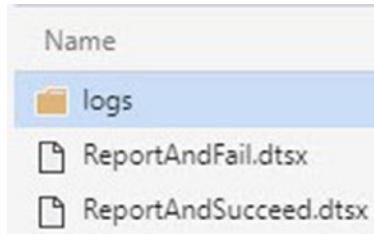


Figure 9-57. The logs folder, surfaced

Double-click the logs folder to view the contents. Each SSIS package execution generates a run ID visible in the pipeline’s Output tab in ADF, as shown in Figure 9-58, along with the corresponding logs\folder name in storage explorer.

The screenshot shows the SSIS Pipeline execution's Output tab and the Azure Storage Explorer. The Pipeline execution tab displays a table of tasks with their names, types, statuses, and run IDs. The 'Execute Application Package' task appears twice: once with status 'Succeeded' and once with status 'Failed'. The Storage Explorer shows a folder structure under 'fs-ssis\logs'. The top-level folder contains subfolders named after the pipeline run ID ('64e7e9521fb4432bb82007a2405a0df4'). Inside these subfolders are several log files, including 'execution_data_statistics_2020-06-04_19-36-23_700674_000000.log', 'execution_component_phases_2020-06-04_19-36-23_700674_000000.log', 'executable_statistics_2020-06-04_19-36-23_700674_000000.log', 'event_messages_2020-06-04_19-36-23_700674_000000.log', and 'event_message_context_2020-06-04_19-36-23_700674_000000.log'.

Figure 9-58. Pipeline execution's Output tab aligned with Storage Explorer's logs folder contents

“Execute Application Package” executed twice, once to run the SSIS package named “ReportAndSucceed.dtsx” and again for the SSIS package named “ReportAndFail.dtsx”. Figure 9-55 showed the topmost instance of “Execute Application Package” succeeded, and the next listing of “Execute Application Package” failed.

Figure 9-58 is a screenshot of the parent pipeline’s Output tab *aligned* with the contents of the logs folder in the Azure file share. Please note the pipeline’s Output Run ID value is similar to the name of the logs subfolder. Entering the topmost logs subfolder in storage explorer reveals a collection of log files, as shown in Figure 9-59.

The screenshot shows the Azure Storage Explorer interface. The current view is 'fs-ssis\logs'. Inside the 'logs' folder, there is a subfolder named after the pipeline run ID ('64e7e9521fb4432bb82007a2405a0df4'). This subfolder contains five log files: 'execution_data_statistics_2020-06-04_19-36-23_700674_000000.log', 'execution_component_phases_2020-06-04_19-36-23_700674_000000.log', 'executable_statistics_2020-06-04_19-36-23_700674_000000.log', 'event_messages_2020-06-04_19-36-23_700674_000000.log', and 'event_message_context_2020-06-04_19-36-23_700674_000000.log'. The 'event_messages...' file is currently selected.

Figure 9-59. Log files stored in the Azure file share fs-ssis\logs folder

Right-click the file that begins with “event_messages_” and then click “Open,” as shown in Figure 9-60.

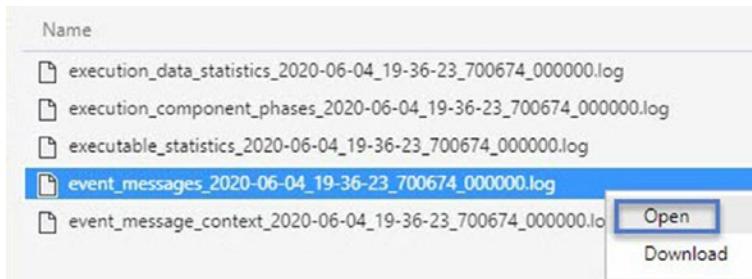


Figure 9-60. Opening the event_messages log

Since my server is configured to open files with the “.log” extension using Notepad, the event_messages file opens in Notepad, as shown in Figure 9-61.

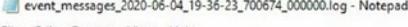
```

event_message_id,operation_id,message_time,message_type,message_source_type,message_text,extended_info_id,package_name,event_name,message
b9ca7a78-2103-4811-bfb5-f9f124414049,1,2020-06-04 19:36:28 +00:00,70,20,No web proxy set. Use system default proxy.,,,OnInfo
271720a1-ec5f-4a65-ad01-bb307d9d6c5c,1,2020-06-04 19:36:29 +00:00,10,30,ReportAndSucceed:Validation has started.,,,OnPreValidate
eac23b2d-c672-4e96-bfda-19a06a985b34,1,2020-06-04 19:36:29 +00:00,10,40,SCR Log Values:Validation has started.,,,OnPostValidate
16f2ab63-fb12-42e9-8fec-11f07b04cc92,1,2020-06-04 19:36:29 +00:00,20,40,SCR Log Values:Validation is complete.,,,OnPostValidate
18ece38e-bee2-4652-bb19-bdb137784176,1,2020-06-04 19:36:29 +00:00,20,30,ReportAndSucceed:Validation is complete.,,,OnPostValidate
0800647a-d26e-484a-980b-3c90543cc1c5,1,2020-06-04 19:36:29 +00:00,30,30,"ReportAndSucceed:Start, 7:36:29 PM.",,,OnPreExecute
afe81fb2-a8bd-4953-bf26-bca43c8e40bf,1,2020-06-04 19:36:29 +00:00,40,30,"ReportAndSucceed:Finished, 7:36:29 PM, Elapsed time
edf0e8c4-3398-434d-a482-1a73318c994c,1,2020-06-04 19:36:29 +00:00,30,40,"SCR Log Values:Start, 7:36:29 PM.",,,OnPreExecute,
ce859125-89be-4d38-aaef0-b769387242b9,1,2020-06-04 19:36:29 +00:00,10,40,SCR Log Values:Validation has started.,,,OnPostValidate
e7a5ffa6-5991-427c-ae84-7dd674e71329,1,2020-06-04 19:36:29 +00:00,20,40,SCR Log Values:Validation is complete.,,,OnPostValidate
d721292b-3974-4012-933a-eb55d2d062bb,1,2020-06-04 19:36:29 +00:00,70,40,SCR Log Values:Information: I am ReportAndSucceed.,
ab46d828-5029-478d-bb60-42a9b8905bdd,1,2020-06-04 19:36:29 +00:00,40,40,"SCR Log Values:Finished, 7:36:29 PM, Elapsed time:

```

Figure 9-61. Viewing the event_messages log file for ReportAndSucceed.dtsx test execution

This is the event_messages log file for the execution of the ReportAndSucceed.dtsx SSIS package. Opening the event_messages log file from the other folder displays the event_messages from the execution of the ReportAndFail.dtsx SSIS package, as shown in Figure 9-62.



```

event_messages_2020-06-04_19-36-23_700674_000000.log - Notepad
File Edit Format View Help
EventMessageId,OperationId,MessageTime,MessageType,MessageSourceType,MessageText,ExtendedInfoId,PackageName,EventName,MessageSo
132eefc8-de7d-4a6a-9e2f-8e862e37159b,1,2020-06-04 19:36:28 +00:00,70,20,No web proxy set. Use system default proxy.,,,OnInfor
214cf0ba-6fffb-4755-922e-adef15154fd16,1,2020-06-04 19:36:29 +00:00,10,30,ReportAndFail:Validation has started.,,,OnPreValidate,R
b1430748-a515-470e-9093-1c40521a8add,1,2020-06-04 19:36:29 +00:00,10,40,SCR Log Values:Validation has started.,,,OnPreValidate,
335bb3be-2782-45fd-ae7d-1313c4bf382e,1,2020-06-04 19:36:29 +00:00,20,40,SCR Log Values:Validation is complete.,,,OnPostValidate
6dc5508e-d3e0-4325-9e4f-8f058cb3d0e6,1,2020-06-04 19:36:29 +00:00,20,30,ReportAndFail:Validation is complete.,,,OnPostValidate,
78fe5a00-ad5a-4425-bf1c-bbd72e46195b,1,2020-06-04 19:36:29 +00:00,30,30,"ReportAndFail:Start, 7:36:29 PM.",,,OnPostExecute,Repor
8d61c66d-356e-44d4-b052-933cff096c11,1,2020-06-04 19:36:29 +00:00,110,30,"ReportAndFail:Warning: SSIS Warning Code DTS_W_MAXIMU
695f060b-9416-422c-a477-b01edadee1c2,1,2020-06-04 19:36:29 +00:00,40,30,"ReportAndFail:Finished, 7:36:29 PM, Elapsed time: 00:0
d8c121a3-fd33-4f55-8c99-aa337137f32c,1,2020-06-04 19:36:29 +00:00,30,40,"SCR Log Values:Start, 7:36:29 PM.",,,OnPostExecute,SCR
d78979ef-37ab-4839-81af-8beb00d5e9e9,1,2020-06-04 19:36:29 +00:00,10,40,SCR Log Values:Validation has started.,,,OnPreValidate,
6b45c231-3428-43b3-af99-3593497e0e55,1,2020-06-04 19:36:29 +00:00,20,40,SCR Log Values:Validation is complete.,,,OnPostValidate
e4fc00eb-fd2d-4019-9275-51375590cf07,1,2020-06-04 19:36:29 +00:00,120,40,SCR Log Values:Error: ReportAndFail execution failed.,
d1296801-f67c-4833-9aa3-0677bc61e80d,1,2020-06-04 19:36:29 +00:00,110,40,"SCR Log Values:Warning: SSIS Warning Code DTS_W_MAXIM
93e42c01-fb54-4e2f-87c5-7d17b7a054bc,1,2020-06-04 19:36:29 +00:00,130,40,"SCR Log Values:The task, 'SCR Log Values', failed.",,
a4f07233-f9c1-4221-80c8-d703dacfc943,1,2020-06-04 19:36:29 +00:00,40,40,"SCR Log Values:Finished, 7:36:29 PM, Elapsed time: 00:

```

Figure 9-62. Viewing the event_messages log file for ReportAndFail.dtsx test execution

Conclusion

This chapter covered the following:

- Provisioning an Azure SQL Database instance
 - Deploying the SSISConfig database to the Azure SQL DB instance
- Provisioning an Azure File Share
 - Deploying (copying) test SSIS packages to the Azure file share
- Building the SSIS framework execution engine
 - Test-executing the parent pipeline and viewing execution logs

The next step is persisting application and application package operational data in the SSISConfig database.

CHAPTER 10

Framework Logging in ADF

Chapters 5–7 covered designing and building an SSIS framework using a SQL Server database named SSISConfig and an SSIS package as the execution engine (the Parent.dtsx SSIS package) to execute SSIS packages stored in the file system of an on-premises server. Chapter 8 covered provisioning an instance of Azure Data Factory (ADF), an Azure storage account, and an Azure-SSIS integration runtime configured to execute SSIS packages stored in an Azure file share. Chapter 9 covered provisioning an Azure SQL database, deploying the SSISConfig database to the Azure SQL DB, provisioning an Azure file share (fs-ssis), and then building the execution engine (the parent pipeline) in ADF.

In this chapter, we add logging functionality to the parent ADF pipeline, much like the functionality we added to the on-premises version of the SSIS framework in Chapter 7. The log schema and associated artifacts have already been deployed, but we will make a few adjustments and modifications to the logging-related tables.

Although they perform similar operations, Azure Data Factory pipelines are *different* from SSIS packages. The differences will drive changes in the SSIS framework design.

Add the ApplicationName Parameter

The parent pipeline will execute SSIS framework applications. To make the ADF pipeline generic and able to execute *any* application stored in the SSISConfig database's metadata, the pipeline requires an ApplicationName parameter.

Open the parent pipeline, click the whitespace of the pipeline, and then click the Parameters tab. When the Parameters tab displays, click the “+ New” link highlighted in Figure 10-1.

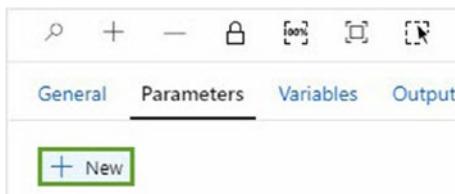


Figure 10-1. Beginning to add a new parameter

When the Parameter configuration displays, set the following parameter properties, which are shown in Figure 10-2.

- Name: ApplicationName
- Type: String
- Default Value: Framework Test

Parameter Configuration			
	NAME	TYPE	DEFAULT VALUE
<input type="button" value="New"/>	ApplicationName	String	Framework Test

Figure 10-2. Configuring the ApplicationName parent pipeline parameter

The ApplicationName parameter is now accessible by activities in the parent ADF pipeline.

A Quick Review of SSIS Framework Applications and Packages

An SSIS framework application is a collection of SSIS packages that execute in a prescribed order (*one* SSIS application may be configured to execute *one or many* SSIS packages). Metadata for the SSIS application is stored in the config.Applications table and metadata for SSIS package location is stored in the config.Packages table. Because the cardinality between SSIS applications and SSIS packages is many-to-many (*one* SSIS package may be configured to run in *one or many* SSIS applications), the config.ApplicationPackages table exists to resolve relationships between applications and packages.

Add Application Instance Logging

In the introduction to this chapter, I mentioned there are differences between the way ADF and SSIS accomplish similar functions, and these differences will drive changes in the SSIS framework design in Azure Data Factory. We encounter the first difference here as we configure Application Instance logging.

The SSIS framework collects operational data about instances of application execution in the log.ApplicationInstance table. Operational data regarding instances of application package execution are stored in the log.ApplicationPackageInstance table.

Modifying the Log.ApplicationInstance Table

An Azure Data Factory (ADF) pipeline is made up of many constituent objects and artifacts, including ADF activities. When ADF activities execute, a *RunId* value is generated for use in operational logging. When an Azure Data Factory pipeline executes, a RunId value is generated and used to connect internal ADF pipeline logs to pipeline activity logs.

Modify the log.ApplicationInstance table by executing the T-SQL in Listing 10-1.

Listing 10-1. Adding the ApplicationRunId column to the ApplicationInstance table

```
print 'Log.ApplicationInstance.ApplicationRunId column'
If Not Exists(Select s.[name]
    + '.' + t.[name]
    + '.' + c.[name]
    From [sys].[schemas] s
    Join [sys].[tables] t
    On t.[schema_id] = s.[schema_id]
    Join [sys].[columns] c
    On c.[object_id] = t.[object_id]
    Where s.[name] = N'log'
        And t.[name] = N'ApplicationInstance'
        And c.[name] = N'ApplicationRunId')
begin
    print ' - Adding log.ApplicationInstance.ApplicationRunId column'
    Alter Table log.ApplicationInstance
```

```

    Add ApplicationRunId nvarchar(55) NULL
    print ' - Log.ApplicationInstance.ApplicationRunId column added'
    end
Else
begin
    print ' - Log.ApplicationInstance.ApplicationRunId column already
exists.'
end

```

Let's next encapsulate the T-SQL for adding Application Instance rows in a stored procedure.

Adding the Log.InsertApplicationInstance Stored Procedure

In the SSIS execution engine (the Parent.dtsx SSIS package), an Execute SQL Task was used to insert the initial ApplicationInstance record into the log.ApplicationInstance table and return an ApplicationInstanceId value. In the Azure Data Factory execution engine (the parent pipeline), a Lookup activity will call a stored procedure that inserts the record and returns an ApplicationInstanceId value. The ApplicationInstanceId value will be used later in the pipeline to update the ApplicationInstance record.

In Azure Data Studio (or SSMS), connect to the Azure SQL database named SSISConfig. Starting with the T-SQL code from Listing 7-4 in Chapter 7, execute the T-SQL to create a new stored procedure named “log.InsertApplicationInstance” in an idempotent (re-executable) manner, as shown in Listing 10-2.

Listing 10-2. Idempotent T-SQL to create the log.InsertApplicationInstance stored procedure

```

print 'log.InsertApplicationInstance stored procedure'
If Exists(Select s.[name] + '.' + p.[name]
        From [sys].[procedures] p
        Join [sys].[schemas] s
        On s.[schema_id] = p.[schema_id]
        Where s.[name] = N'log'
        And p.[name] = N'InsertApplicationInstance')

```

```

begin
    print ' - Dropping log.InsertApplicationInstance stored procedure'
    Drop Procedure log.InsertApplicationInstance
    print ' - Log.InsertApplicationInstance stored procedure dropped'
end

print ' - Creating log.InsertApplicationInstance stored procedure'
go

Create Procedure log.InsertApplicationInstance
    @ApplicationName nvarchar(255)
    , @ApplicationRunId nvarchar(55) = NULL
As

declare @ApplicationId int = (Select ApplicationId
                                From config.Applications
                                Where ApplicationName = @ApplicationName)

Insert Into [log].ApplicationInstance (ApplicationId, ApplicationRunId)
Output inserted.ApplicationInstanceId
Values (@ApplicationId, @ApplicationRunId)

go

print ' - Log.InsertApplicationInstance stored procedure created'
go

```

The log.InsertApplicationInstance stored procedure requires the name of an SSIS framework application be sent to the @ApplicationName string [nvarchar(255)] parameter. An internal T-SQL int parameter named @ApplicationId looks up the value of the ApplicationId in the config.Applications table for the given @ApplicationName value.

The log.InsertApplicationInstance stored procedure also accepts a parameter named @ApplicationRunId which defaults to NULL when no value is supplied. @ApplicationRunId is used to store the Azure Data Factory pipeline RunId value. Storing the RunId in the SSIS framework allows enterprise DevOps teams to relate operational log data in the SSIS framework to operational data captured in ADF logs, which is *extremely* useful when something unfortunate happens during ADF pipeline execution.

The T-SQL Insert statement initializes the ApplicationInstance record by inserting new @ApplicationId and @ApplicationRunId values into the log.ApplicationInstance table. Defaults configured on the log.ApplicationInstance table insert the following values in the row:

- ApplicationStartTime is set to sysdatetimeoffset() by DF_log_ApplicationInstance_ApplicationStartTime.
- ApplicationStatus is set to N'Running' by DF_log_ApplicationInstance_ApplicationStatus.

If all goes as planned, Azure Data Studio should return Messages that appear as shown in Figure 10-3.

Messages

```
Started executing query at Line 1
log.InsertApplicationInstance stored procedure
  - Creating log.InsertApplicationInstance stored procedure
Started executing query at Line 16
Commands completed successfully.
Started executing query at Line 31
  - Log.InsertApplicationInstance stored procedure created
Total execution time: 00:00:00.110
```

Figure 10-3. Messages from creating the log.InsertApplicationInstance stored procedure

After the log.InsertApplicationInstance stored procedure has been created, the next step is to add an activity to the pipeline to execute the stored procedure.

Logging Application Instance

The ADF execution engine needs to execute a stored procedure, log.InsertApplicationInstance, that returns a value named ApplicationInstanceId. A Lookup activity is built for this!

Return to the parent ADF pipeline. Drag a new Lookup activity onto the surface. Add a Success output from the new Lookup activity to the Get Application Packages lookup activity. Rename the new Lookup activity to "Log Application Instance Start," as shown in Figure 10-4.

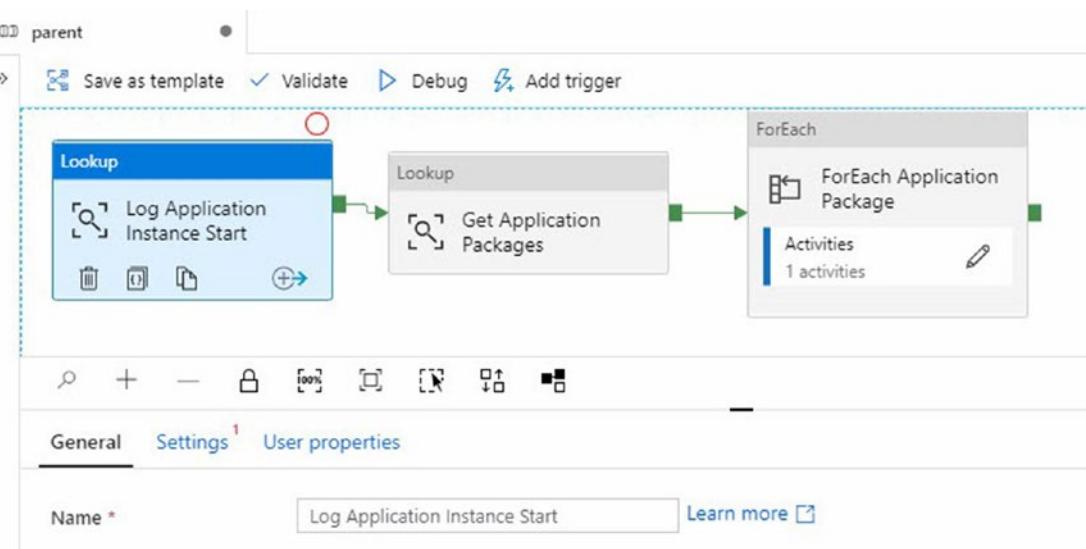


Figure 10-4. Adding the Log Application Instance Start lookup activity

Click the Log Application Instance Start lookup activity’s Settings tab to continue configuration. Click the “Source dataset” property drop-down, and select “ssisFrameworkDataSet,” as shown in Figure 10-5.

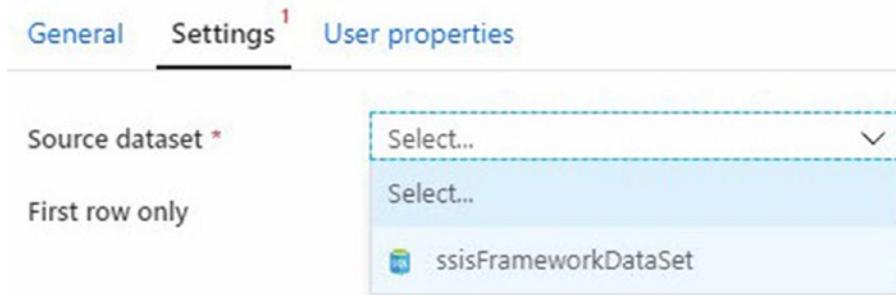


Figure 10-5. Selecting “ssisFrameworkDataSet” for the Log Application Instance Start lookup activity’s “Source dataset” property

Next, set the Log Application Instance Start lookup activity’s “Use query” property option to “Stored procedure.” Select “log.InsertApplicationInstance” from the “Name” property drop-down, as shown in Figure 10-6.

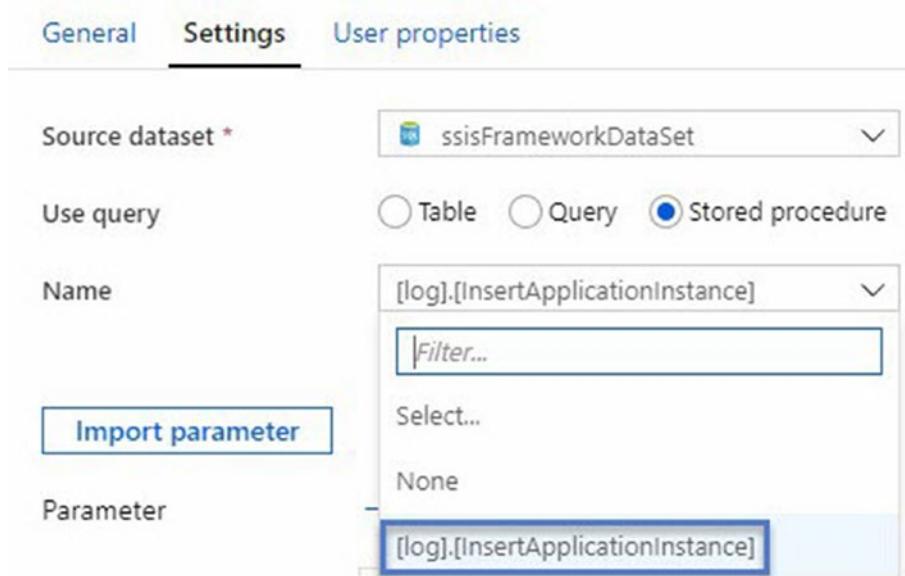


Figure 10-6. Selecting the `log.InsertApplicationInstance` stored procedure

Next, click the “Import parameter” button to import the `@ApplicationName` and `@ApplicationRunId` String parameters from the `log.InsertApplicationInstance` stored procedure. Click inside the `@ApplicationName` parameter’s Value property textbox, and then click the “Add dynamic content [Alt + P]” link, as shown in Figure 10-7.

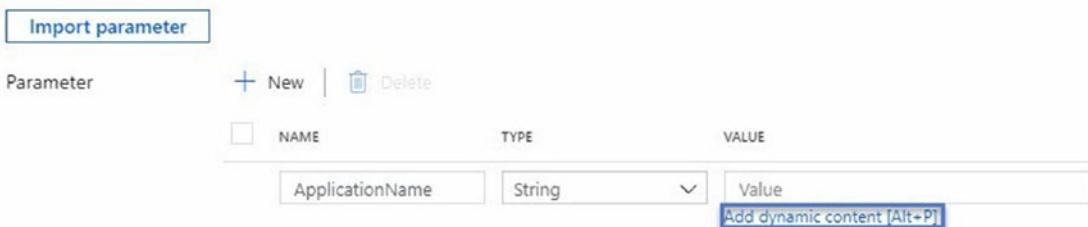


Figure 10-7. Preparing to add dynamic content to the parameter value property

When the “Add dynamic content” blade displays, expand the Parameters expression category, and click “ApplicationName” to set the parameter value ADF expression to `@pipeline.parameters.ApplicationName`, as shown in Figure 10-8.

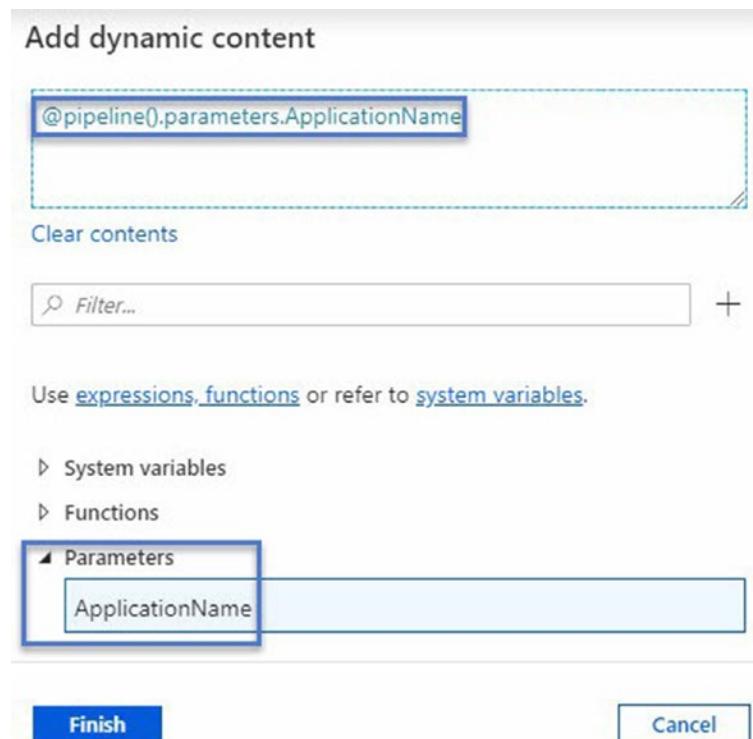


Figure 10-8. Setting the ApplicationName parameter value expression

Click the Finish button to proceed.

Click inside the @ApplicationRunId parameter's Value property textbox, and then click the “Add dynamic content [Alt + P]” link. When the “Add dynamic content” blade displays, expand the “System variables” expression category, and click “Pipeline run ID” to set the parameter value ADF expression to @pipeline().RunId, as shown in Figure 10-9.

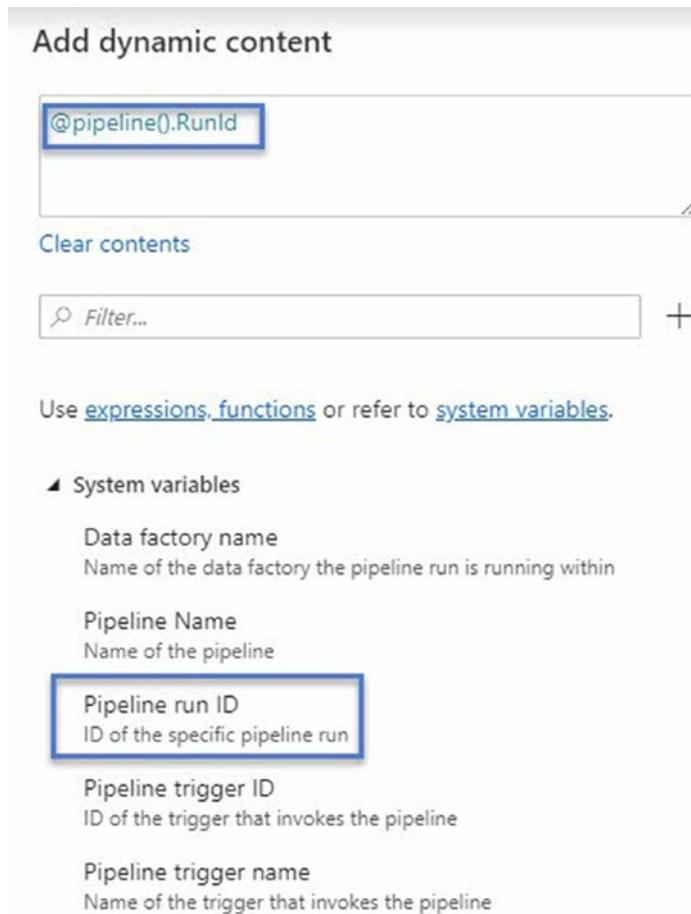


Figure 10-9. Setting the ApplicationRunId parameter value expression

Leave the “Query timeout (minutes)” property set to its default (120 minutes), the “Isolation level” property set to “None” (default), and the “First row only” property checked (default), as shown in Figure 10-10.

Query timeout (minutes)	<input type="text" value="120"/>	(i)
Isolation level	<input type="text" value="None"/>	(i)
First row only	<input checked="" type="checkbox"/>	

Figure 10-10. Leaving default values for Query timeout, Isolation level, and First row only properties

The Log Application Instance Start lookup activity is now configured to call the log.InsertApplicationInstance stored procedure, passing the stored procedure the pipeline's ApplicationName and ApplicationRunId parameter values and returning the ApplicationInstanceId value.

Adding the Log.UpdateApplicationInstanceState Stored Procedure

As configured, the parent pipeline fails execution because the ReportAndFail.dtsx SSIS package fails (by design). The log.ApplicationInstance row should be updated to reflect the SSIS framework application instance failed.

Begin by adding a new stored procedure named “log.UpdateApplicationInstanceState” to the SSISConfig database, as shown in Listing 10-3.

Listing 10-3. Idempotent T-SQL to create the log.UpdateApplicationInstanceState stored procedure

```
print 'log.UpdateApplicationInstanceState stored procedure'
If Exists(Select s.[name] + '.' + p.[name]
          From [sys].[procedures] p
          Join [sys].[schemas] s
            On s.[schema_id] = p.[schema_id]
          Where s.[name] = N'log'
            And p.[name] = N'UpdateApplicationInstanceState')
begin
    print ' - Dropping log.UpdateApplicationInstanceState stored procedure'
    Drop Procedure log.UpdateApplicationInstanceState
    print ' - Log.UpdateApplicationInstanceState stored procedure dropped'
end

print ' - Creating log.UpdateApplicationInstanceState stored procedure'
go

Create Procedure log.UpdateApplicationInstanceState
    @ApplicationInstanceId int
    , @ApplicationStatus nvarchar(55) = 'Succeeded'
As
```

```

Update [log].ApplicationInstance
Set ApplicationEndTime = sysdatetimeoffset()
    , ApplicationStatus = @ApplicationStatus
Where ApplicationInstanceId = @ApplicationInstanceId

go

print ' - Log.UpdateApplicationInstanceStatus stored procedure created'
go

```

If all goes as planned, Azure Data Studio messages should appear similar to Figure 10-11.

Messages

```

Started executing query at Line 1
log.UpdateApplicationInstanceStatus stored procedure
- Creating log.UpdateApplicationInstanceStatus stored procedure
Started executing query at Line 17
Commands completed successfully.
Started executing query at Line 29
- Log.UpdateApplicationInstanceStatus stored procedure created
Total execution time: 00:00:00.657

```

Figure 10-11. Azure Data Studio messages reflecting successful creation of the log.UpdateApplicationInstanceStatus stored procedure

The next step is to add a Stored Procedure Activity that calls the log.UpdateApplicationInstanceStatus stored procedure on failure.

Updating Application Instance

The ADF execution engine needs to execute a stored procedure, log.UpdateApplicationInstanceStatus that accepts ApplicationInstanceId and ApplicationStatus values and then updates the status of the current Application Instance. A Stored Procedure Activity is perfect for this!

Return to the parent ADF pipeline. Drag a Stored Procedure Activity onto the surface. Add a Failure output from the ForEach Application Package activity to the new Stored Procedure Activity. Rename the new Stored Procedure Activity to “Log Application Instance Failure,” as shown in Figure 10-12.

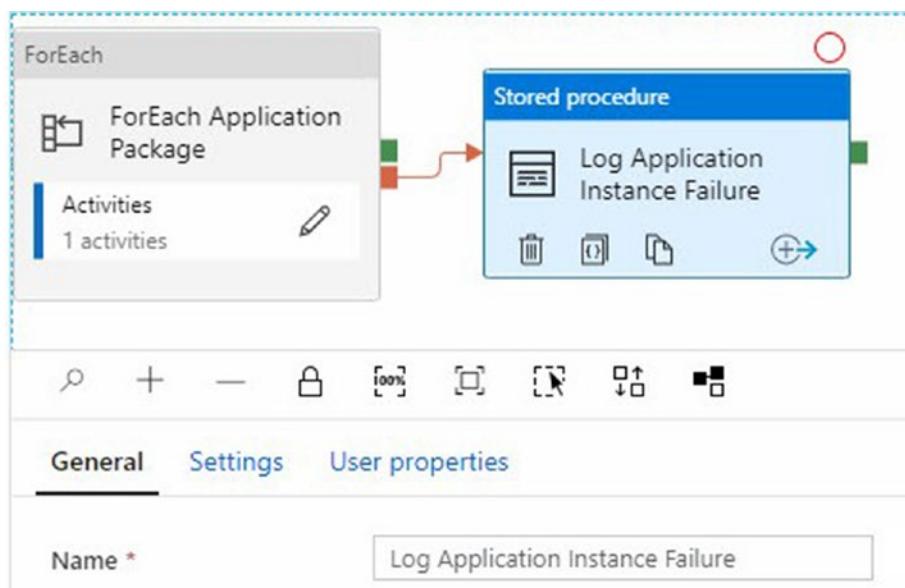


Figure 10-12. Adding the “Log Application Instance Failure” stored procedure activity

On the Settings tab, select “ssisFrameworkLinkedService” from the “Linked service” property drop-down, and then select “[log].[UpdateApplicationInstanceStatus]” from the “Stored procedure name” drop-down, as shown in Figure 10-13.

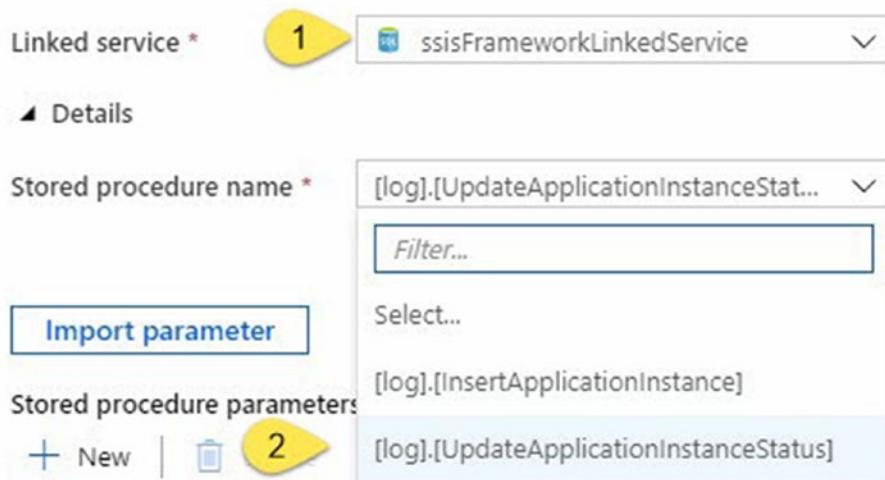


Figure 10-13. Configuring the “Linked service” and “Stored procedure name” properties

Click the “Import parameter” button to import parameters from the log. UpdateApplicationInstanceStatus stored procedure. Click inside the value textbox for the ApplicationInstanceId parameter, and then click the “Add dynamic content [Alt + P]” link, as shown in Figure 10-14.

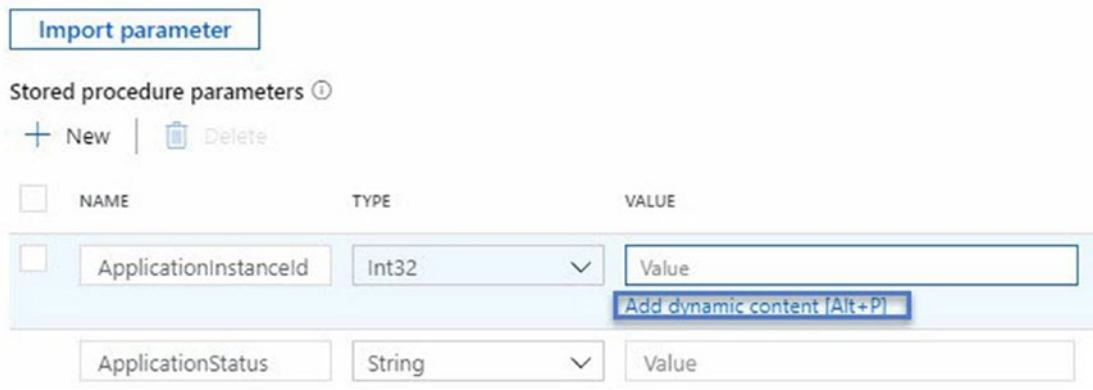


Figure 10-14. Preparing to enter dynamic content for the ApplicationInstanceId parameter

When the “Add dynamic content” blade displays, expand the “Activity outputs” expressions category, and click “Log Application Instance Start.” The dynamic content textbox will display @activity('Log Application Instance Start').output.Append “.firstrow.ApplicationInstanceId” to the expression to map the value of the ApplicationInstanceId returned from the execution of the “Log Application Instance Start” lookup activity into the ApplicationInstanceId parameter’s value property, as shown in Figure 10-15.

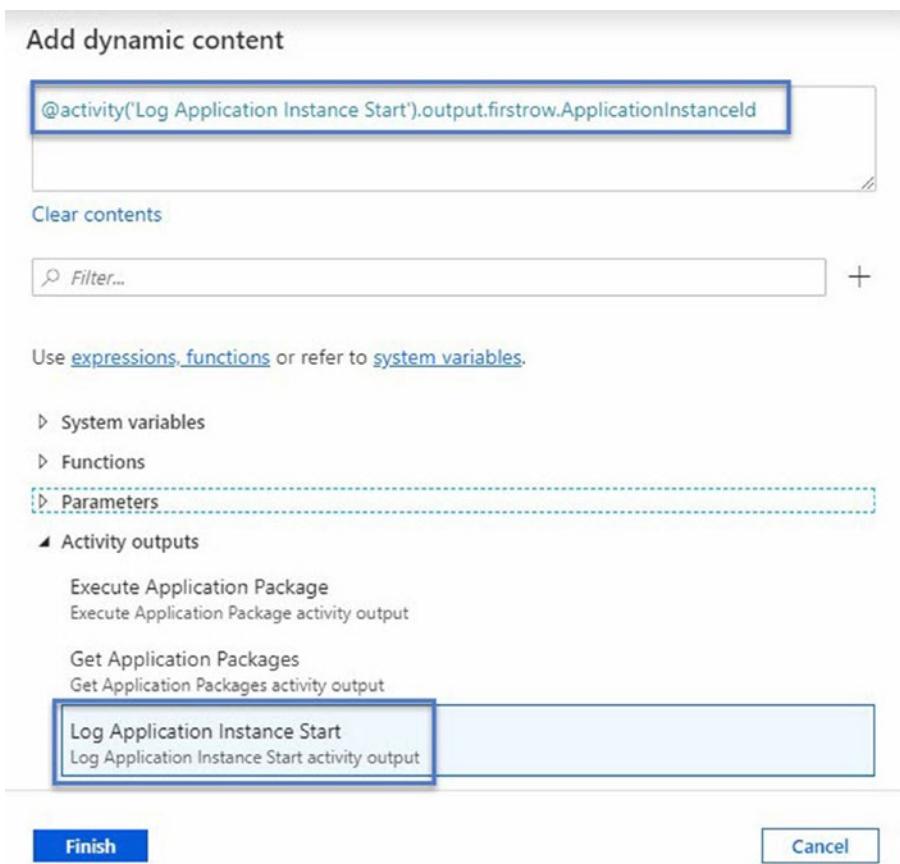


Figure 10-15. Mapping the “Log Application Instance Start” lookup activity into the ApplicationInstanceId parameter’s value property

Click the Finish button to proceed.

Type “Failed” into the ApplicationStatus parameter’s value property, as shown in Figure 10-16.

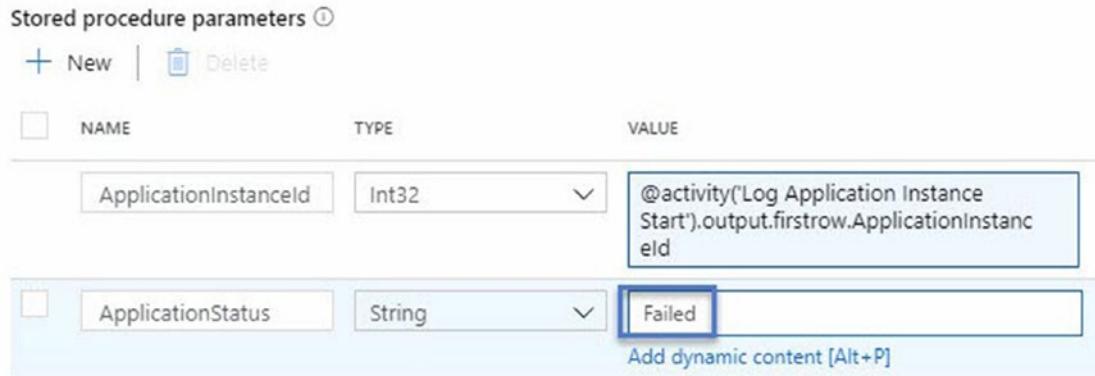


Figure 10-16. Setting the *ApplicationStatus* property to Failed

The ADF version of the SSIS framework execution engine (the parent pipeline) is now configured to execute an SSIS application and record a failed application instance.

The next step is to test-execute the pipeline.

Let's Test It!

Click the Debug item on the parent pipeline's toolbar. The Output tab displays and, after some time, reveals the parent pipeline activities executed, as shown in Figure 10-17.

NAME	TYPE	DURATION	STATUS
Log Application Instance Failure	SqlServerStoredProcedure	00:00:13	✓ Succeeded
Execute Application Package	ExecuteSSISPackage	00:00:22	✓ Succeeded
Execute Application Package	ExecuteSSISPackage	00:00:26	✗ Failed
ForEach Application Package	ForEach	00:00:32	✗ Failed
Get Application Packages	Lookup	00:00:13	✓ Succeeded
Log Application Instance Start	Lookup	00:00:03	✓ Succeeded

Figure 10-17. Results of a test execution of the parent pipeline

Execute the T-SQL query shown in Listing 10-4 to examine the Application Instance results.

Listing 10-4. Application Instance results

```
Select a.ApplicationName
    , ai.ApplicationStartTime
    , ai.ApplicationStatus
From log.ApplicationInstance ai
Join config.Applications a
    On a.ApplicationId = ai.ApplicationId
Order By ApplicationInstanceId Desc
```

The results should appear similar to Figure 10-18.

Results			Messages
	ApplicationName	ApplicationStartTime	ApplicationStatus
1	Framework Test	2020-06-08 16:22:53...	Failed

Figure 10-18. Application Instance results

While the results match what was expected, something is missing: there is no Stored Procedure Activity that updates the Application Instance on Success.

Updating Application Instance on Success

The parent pipeline needs a way to log Application Instance success. Begin by clicking the Clone icon on the “Log Application Instance Failure” stored procedure activity as shown in Figure 10-19.



Figure 10-19. Preparing to clone the “Log Application Instance Success” stored procedure activity from the “Log Application Instance Failure” stored procedure activity

Cloning the “Log Application Instance Failure” stored procedure activity makes sense because the new Stored Procedure Activity will be calling the *same* stored procedure – log.UpdateApplicationInstanceStatus – and *only one* parameter value configuration will change.

Connect a Success output from the “ForEach Application Package” ForEach activity to the cloned Stored Procedure Activity, and then rename the cloned Stored Procedure Activity to “Log Application Instance Success,” as shown in Figure 10-20.

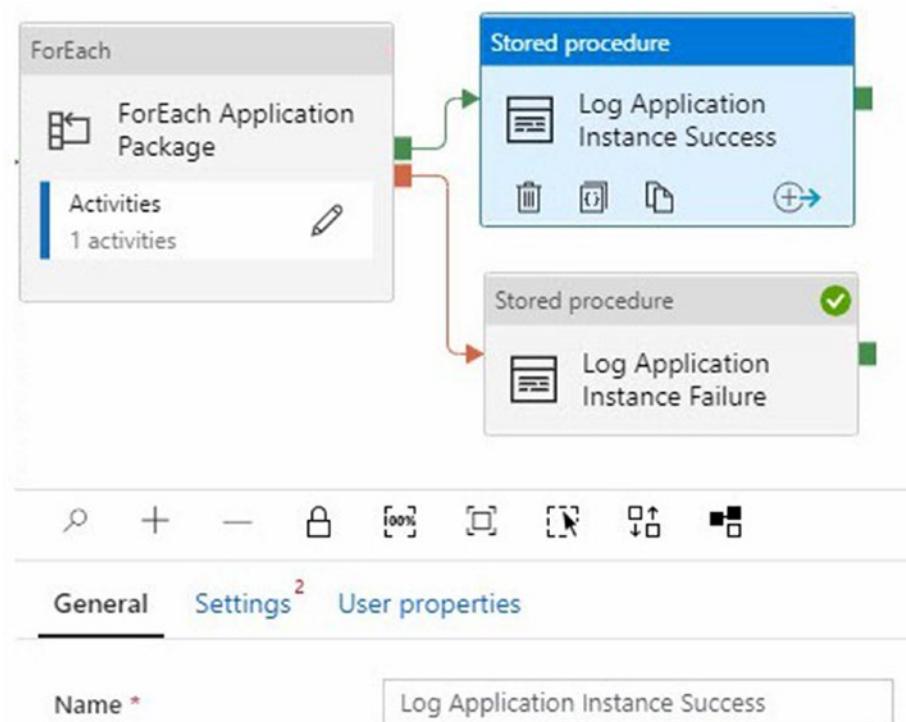


Figure 10-20. Adding the “Log Application Instance Success” stored procedure activity

Next, click the Settings tab for the “Log Application Instance Success” stored procedure activity. The only change required on the Settings tab is the value of the ApplicationStatus parameter’s value property, which needs to be updated from “Failed” to “Succeeded,” as shown in Figure 10-21.

NAME	TYPE	VALUE
ApplicationInstanceId	Int32	@activity('Log Application Instance Start').output.firstrow.ApplicationInstanceId
ApplicationStatus	String	Succeeded

Figure 10-21. Updating the ApplicationStatus parameter value

A fresh test execution confirms the parent pipeline logs SSIS framework application failures as designed. The Azure Data Factory version of the SSIS framework needs similar instance logging at the application package scope.

The next step is implementing Application Package Instance logging.

Before proceeding, we need to make two changes to the parent pipeline:

- Add the ApplicationPackageId field to the query in the “Get Application Packages” lookup activity.
- Add the ApplicationInstanceId variable (and store the ApplicationInstanceId value in the ApplicationInstanceId variable).

Adding the ApplicationPackageId Field

To add the ApplicationPackageId field to the metadata dataset returned from the SSISConfig database, select the “Get Application Packages” lookup activity, and then click the Settings tab. Modify the Query property using the T-SQL in Listing 10-5.

Listing 10-5. Modified query to retrieve application package metadata

```
Select a.ApplicationName
    , p.PackageLocation + p.PackageName As PackagePath
    , ap.ExecutionOrder
    , ap.FailApplicationOnPackageFailure
    , ap.ApplicationPackageId
From [config].[ApplicationPackages] ap
Join [config].[Applications] a
    On a.ApplicationId = ap.ApplicationId
Join [config].Packages p
    On p.PackageId = ap.PackageId
Where a.ApplicationName = N'Framework Test'
    And ap.ApplicationPackageEnabled = 1
Order By ap.ExecutionOrder
```

After the T-SQL has been modified, the next step is to add a variable to hold the ApplicationInstanceId pipeline variable.

Adding the ApplicationInstanceId Pipeline Variable

To add the ApplicationInstanceId pipeline variable, click the pipeline, and then select the Variables tab. Click the “+ New” button and name the new variable “ApplicationInstanceId” as shown in Figure 10-22.

NAME	TYPE	DEFAULT VALUE
GetApplicationPackages	String	Value
CancelCommand	String	Value
ApplicationInstanceId	String	0

Figure 10-22. Adding the ApplicationInstanceId pipeline variable

“Why is the ApplicationInstanceId variable – an integer – declared as a String variable?” That’s a fair question. At the time of this writing, the other variable type options are Boolean and Array. The ApplicationInstanceId is neither a Boolean nor an Array.

The next step is to add a Set variable activity to set the value of the ApplicationInstanceId pipeline parameter to the ApplicationInstanceId value returned from the “Log Application Instance Start” lookup activity. To begin, select the Success output of the “Log Application Instance Start” lookup activity, right-click the Success output, and then click Delete, as shown in Figure 10-23.

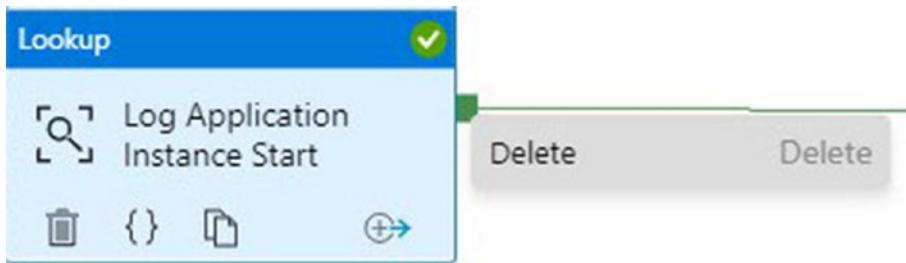


Figure 10-23. Deleting the Success output from the “Log Application Instance Start” lookup activity

To add a new Success output to the “Log Application Instance Start” lookup activity, click the “Add Output” icon, and then click “Success,” as shown in Figure 10-24.

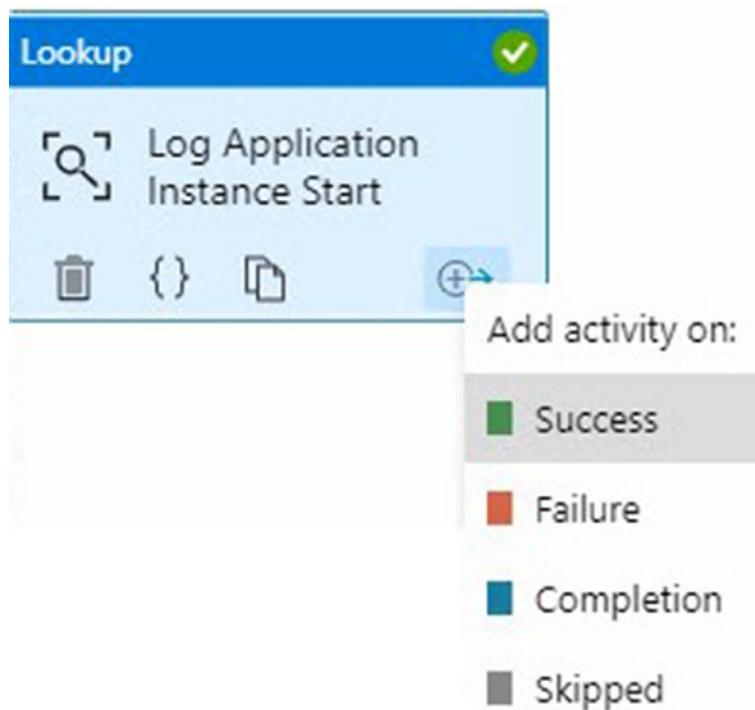


Figure 10-24. Adding a new Success output to the “Log Application Instance Start” lookup activity

Drag a Set variable activity onto the pipeline canvas, and rename it to “Set ApplicationInstanceId.” Click the Variables tab, and select ApplicationInstanceId from the Name property drop-down. Click inside the Value property textbox, and then click the “Add dynamic content [Alt + P]” link to open the Set dynamic content blade. Scroll to the Activity outputs category, and click “Log Application Instance Start” to add @activity('Log Application Instance Start').output to the expression textbox. Edit @activity(, to read @string(activity(.Append .firstrow.ApplicationInstanceId) to the expression to assign the ApplicationInstanceId value returned in the first row of the output from the “Log Application Instance Start” lookup activity, as shown in Figure 10-25.

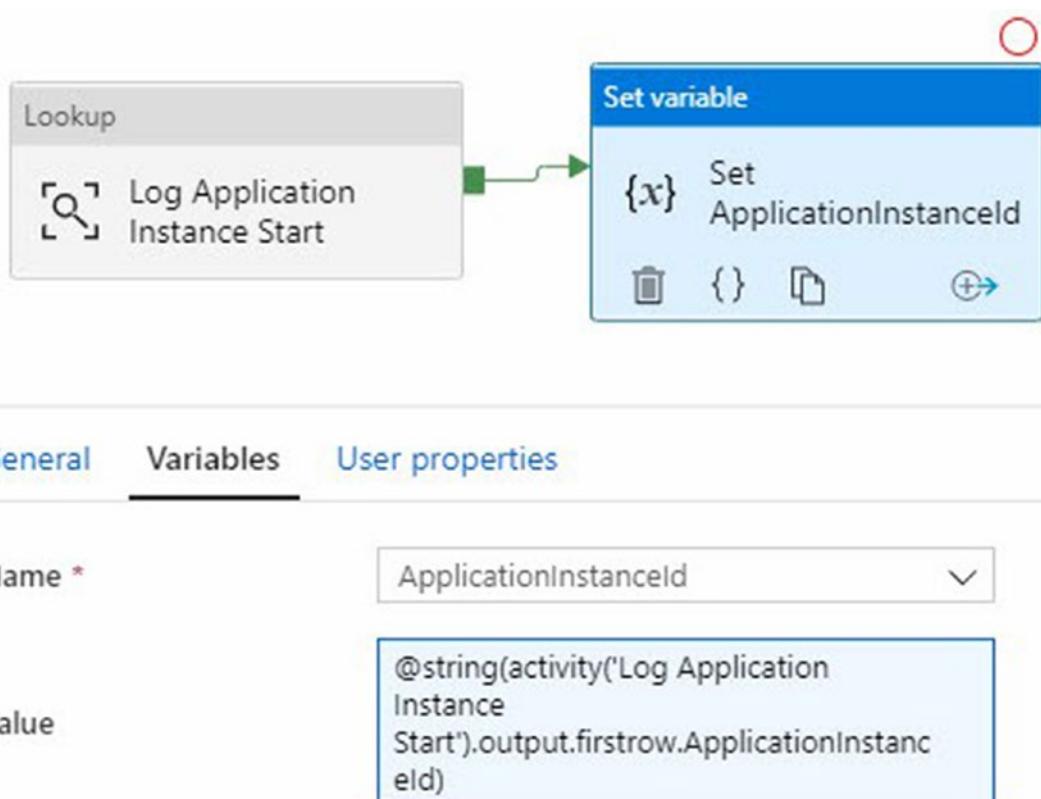


Figure 10-25. Configuring the “Set ApplicationInstanceId” set variable activity

Connect a Success output from the “Set ApplicationInstanceId” set variable activity to the “Get Application Packages” lookup activity.

The “Set ApplicationInstanceId” set variable activity is now configured to read the ApplicationInstanceId value returned from the “Log Application Instance Start” lookup activity and store the value in the ApplicationInstanceId pipeline variable.

Finally, update the “Log Application Instance Success” and “Log Application Instance Failure” stored procedure activities to use @int(variables('ApplicationInstanceId')) for the ApplicationInstanceId stored procedure parameters instead of the output from the “Log Application Instance Start” lookup activity, as shown in Figure 10-26.

The screenshot shows the 'Settings' tab in the Azure Data Factory interface. It is configuring a stored procedure named '[log].[UpdateApplicationInstanceStatus]'. Under 'Stored procedure parameters', there are two entries:

- ApplicationInstanceId: Type Int32, Value: @int(variables('ApplicationInstanceId'))
- ApplicationStatus: Type String, Value: Succeeded, with a checked 'Treat as null' option.

Figure 10-26. Using `@variables('ApplicationInstanceId')`

As stated earlier, the ApplicationInstanceId pipeline variable is a string. At the time of this writing, the value would be implicitly converted to an integer here and elsewhere. That said, please be *intentional* in your coding. Here, “intentional” means explicit conversion to an integer value.

The next step is to add a child pipeline to execute the application packages.

Add the Child Pipeline

As mentioned previously, the differences between the on-premises and the Azure Data Factory versions of an SSIS framework drive design differences. The first step in the redesign is migrating application package execution from the parent pipeline to a “child” pipeline.

Begin by creating a new ADF pipeline named “child,” as shown in Figure 10-27.

The screenshot shows the Azure Data Factory pipeline editor. A new pipeline named 'child' is being created under the 'parent' pipeline. The 'Properties' pane on the right shows the 'Name' field set to 'child'. The pipeline structure is visible on the left, with a dashed line indicating the scope of the new child pipeline.

Figure 10-27. Adding the child pipeline

In the parent pipeline, the “Execute Application Package” execute SSIS activity has direct access to the properties – populated from SSISConfig metadata – required to dynamically execute application packages related to an SSIS application. The child pipeline will use parameters to access this same metadata. On the Parameters tab of the child pipeline, create the following pipeline parameter / type / default value combinations:

- PackagePath / String / [Empty String]
- ApplicationPackageId / Int / 0
- ApplicationInstanceId / Int / 0
- FailApplicationOnPackageFailure / Bool / true
- ParentRunId / String / [Empty String]

When configured, child pipeline parameters should appear similar to Figure 10-28.

NAME	TYPE	DEFAULT VALUE
PackagePath	String	Value
ApplicationPackageId	Int	0
ApplicationInstanceId	Int	0
FailApplicationOnPacka	Bool	true
ParentRunId	String	Value

Figure 10-28. Child pipeline parameters, configured

Drag an Execute SSIS package activity onto the child canvas, and rename the new activity to “Execute Application Package,” as shown in Figure 10-29.

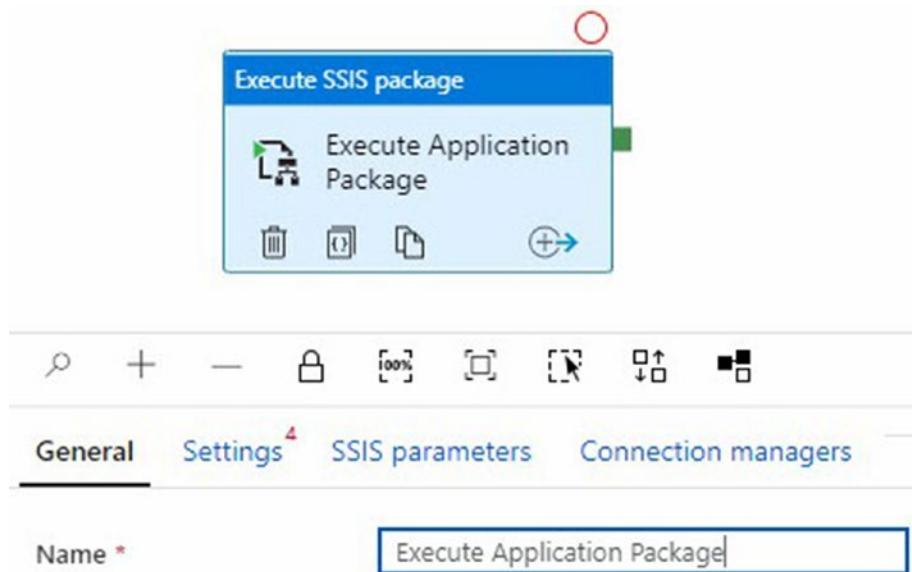


Figure 10-29. Adding the “Execute Application Package” execute SSIS activity

Configure the “Execute Application Package” execute SSIS activity’s properties on the Settings tab as follows:

- Azure-SSIS IR: Azure-SSIS-Files
- Windows authentication: [Unchecked]
- 32-bit runtime: [Unchecked]
- Package location: File system (Package)
- Package path: @pipeline().parameters.PackagePath
 - Add dynamic content [Alt + P], then parameters ► PackagePath.

- Configuration path: [Empty]
- Domain: Azure
- Username: stframeworks
 - The name of the Azure file share that contains the SSIS packages
- Password: [Your file share primary key]
 - The Primary Key to access the Azure file share
- Encryption password: [Empty]
- Logging level: Basic
- Logging path: [Your file share logging path]
- Logging access credentials
 - Same as package access credentials: [Checked if logging folder resides in the same Azure file share as SSIS packages]

When configured, the “Execute Application Package” execute SSIS activity’s properties on the Settings tab should appear similar to Figure 10-30.

General	Settings	SSIS parameters	Connection managers
Azure-SSIS IR *	Azure-SSIS-Files ▼ ⓘ		
Windows authentication (See more info here)	<input type="checkbox"/> ⓘ	32-bit runtime	<input type="checkbox"/> ⓘ
Package location *	File system (Package) ▼ ⓘ		
Package path	@pipeline().parameters.PackagePath ⓘ		
Configuration path	\FileShare\ConfigurationName.dtsConfig ⓘ		
Package access credentials ⓘ (See more info here)			
Domain *	Azure ⓘ		
Username *	stframeworks ⓘ		
Password *	***** ⓘ		
Encryption password	Your package encryption password ⓘ		
Logging level	Basic ▼ ⓘ		
Logging path	\stframeworks.file.core.windows.net\fs-ssis\logs ⓘ		
Logging access credentials ⓘ (See more info here)	Same as package access credentials		<input checked="" type="checkbox"/> ⓘ

Figure 10-30. “Execute Application Package” execute SSIS activity’s properties on the Settings tab, configured

In the parent pipeline, navigate to the “ForEach Application Package” ForEach activity’s Activity surface, and click the Delete icon (trash can) to delete the “Execute Application Package” execute SSIS activity, as shown in Figure 10-31.



Figure 10-31. Deleting the “Execute Application Package” execute SSIS activity

Drag an Execute Pipeline activity onto the “ForEach Application Package” ForEach activity’s Activity surface, and rename it to “Execute Child Pipeline,” as shown in Figure 10-32.

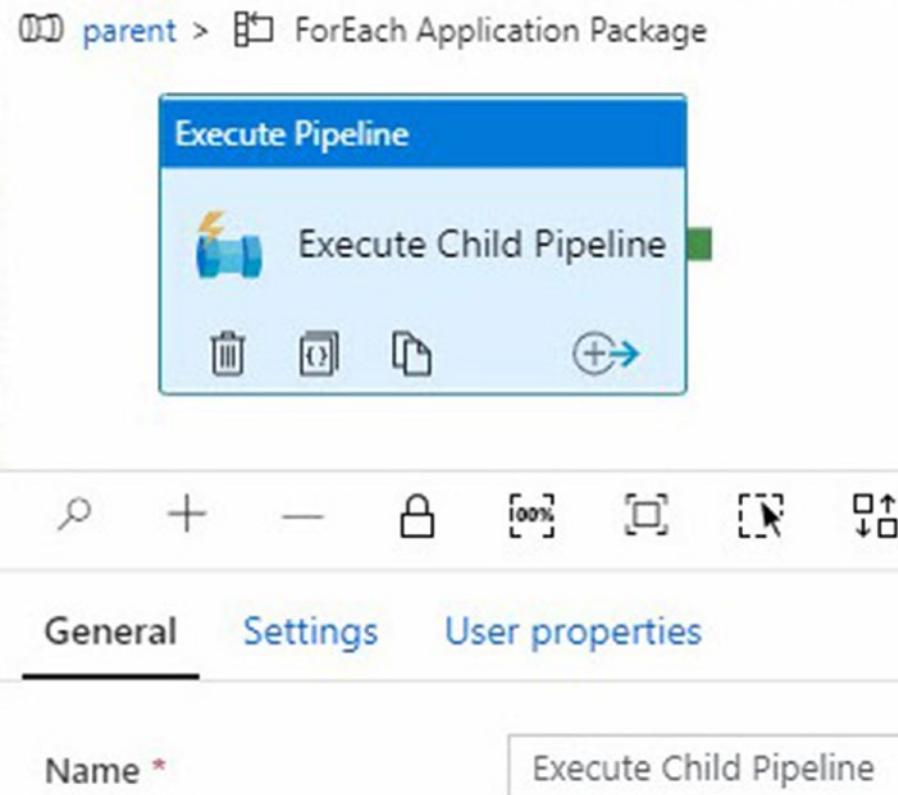


Figure 10-32. Adding the “Execute Child Pipeline” execute pipeline activity

Click the “Execute Child Pipeline” execute pipeline activity’s Setting tab, and configure the following properties:

- Invoked pipeline: child
- Wait on completion: [Checked]
- Parameters:
 - PackagePath / string / @item().PackagePath
 - Add dynamic content [Alt + P], ForEach iterator ➤ ForEach Application Package, and then append “PackagePath,” as shown in Figure 10-33.

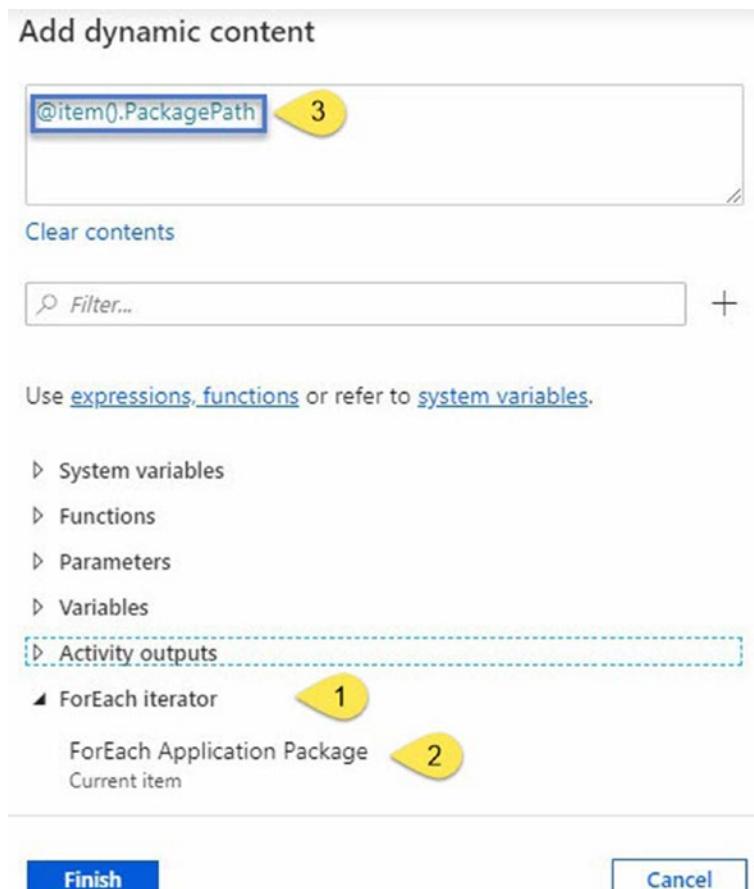


Figure 10-33. Configuring the PackagePath parameter for the child pipeline

- ApplicationPackageId / int / @item().ApplicationPackageId
- ApplicationInstanceId / int / @variables('ApplicationInstanceId')
- FailApplicationOnPackageFailure / bool / @item().FailApplicationOnPackageFailure
- ParentRunId / string / @pipeline().RunId
 - Add dynamic content [Alt + P], and then System variables ► Pipeline run ID, as shown in Figure 10-34.

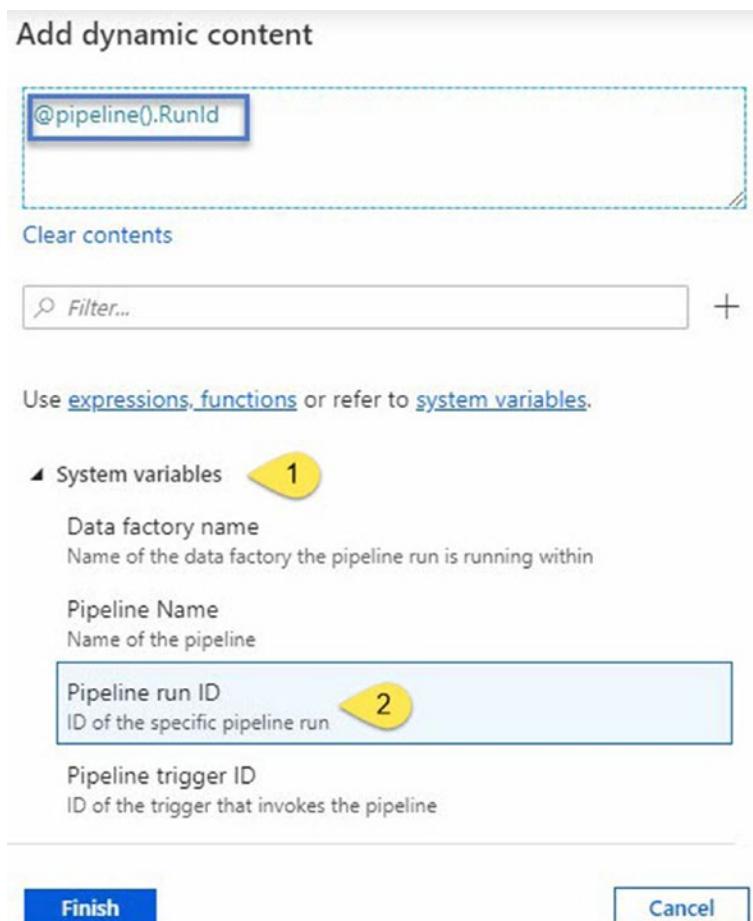


Figure 10-34. Configuring the ParentRunId parameter for the child pipeline

When configured, the “Execute Child Pipeline” execute pipeline activity’s properties on the Settings tab should appear similar to Figure 10-35.

NAME	TYPE	VALUE	DEFAULT VALUE
PackagePath	string	@item().PackagePath	
ApplicationPackageId	int	@item().ApplicationPackageId	0
ApplicationInstanceId	int	@variables('ApplicationInstanceId')	0
FailApplicationOnPackageFailure	bool	@item().FailApplicationOnPackageFailure	true
parentRunId	string	@pipeline().RunId	

Figure 10-35. “Execute Child Pipeline” execute pipeline activity’s properties on the Settings, configured

The “Execute Child Pipeline” execute pipeline activity is now configured to call the child pipeline and pass the parameter values required to execute an application package for which metadata is stored in the SSIS framework.

Let’s Test It!

Click Debug to test-execute the parent pipeline. The “Execute Child Pipeline” execute pipeline activity should execute twice, failing once and succeeding once, and that’s what we see when we view the Output tab of the parent pipeline post-execution, as shown in Figure 10-36.

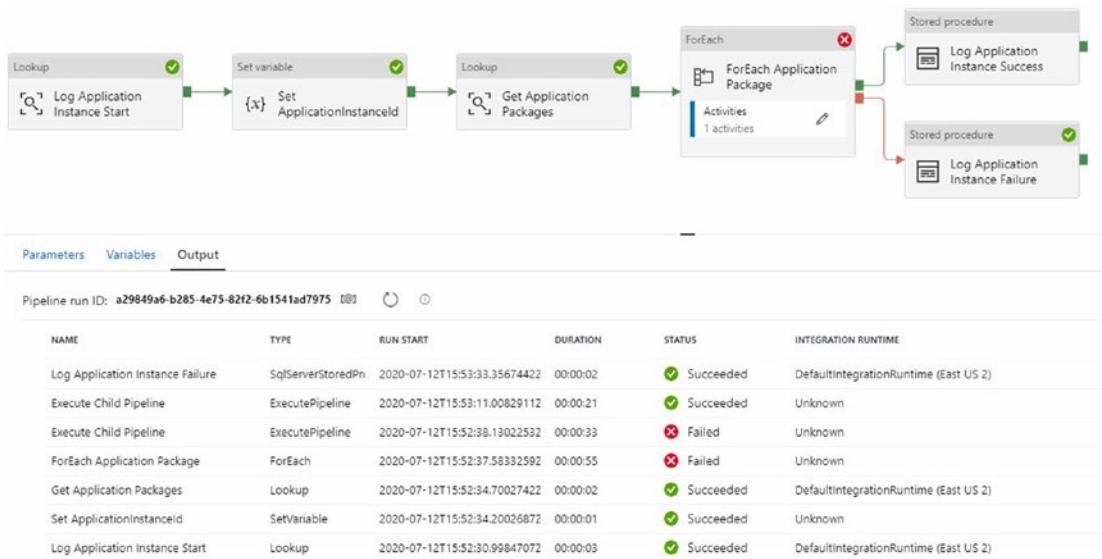


Figure 10-36. Viewing the parent pipeline's debug execution output

The updated SSIS framework execution engine works as designed, but there is more to do to implement the fault tolerance for FailApplicationOnPackageFailure, which we will manage after the next step, adding application instance logging.

Add Application Package Instance Logging

Application instance logging happens in the parent pipeline. The pattern for application instance logging is

1. Log application instance start
2. Do some stuff
3. Log application instance status, succeeded or failed

Application instance logging is a solid pattern. Application package instance logging should follow the same pattern.

Modifying the Log.ApplicationPackageInstance Table

As mentioned earlier in the log.ApplicationInstance section, when ADF activities execute, a *RunId* value is generated for use in operational logging. As with logging the parent pipeline's Run Id when an Azure Data Factory parent pipeline executes, logging the child pipeline's RunId is a great idea. The child Run Id value may be used to connect internal ADF pipeline logs to pipeline activity logs.

Modify the log.ApplicationPackageInstance table by executing the T-SQL in Listing 10-6.

Listing 10-6. Adding the ApplicationPackageRunId column to the ApplicationPackageInstance table

```
print 'Log.ApplicationPackageInstance.ApplicationPackageRunId column'
If Not Exists(Select s.[name]
    + '.' + t.[name]
    + '.' + c.[name]
    From [sys].[schemas] s
    Join [sys].[tables] t
    On t.[schema_id] = s.[schema_id]
    Join [sys].[columns] c
    On c.[object_id] = t.[object_id]
    Where s.[name] = N'log'
        And t.[name] = N'ApplicationPackageInstance'
        And c.[name] = N'ApplicationPackageRunId')
begin
    print ' - Adding log.ApplicationPackageInstance.ApplicationPackageRunId
column'
    Alter Table log.ApplicationPackageInstance
        Add ApplicationPackageRunId nvarchar(55) NULL
    print ' - Log.ApplicationPackageInstance.ApplicationPackageRunId column
added'
end
```

```

Else
begin
    print ' - Log.ApplicationPackageInstance.ApplicationPackageRunId column
already exists.'
end

```

Let's next encapsulate the T-SQL for adding Application Package Instance rows in a stored procedure.

Adding the Log.InsertApplicationPackageInstance Stored Procedure

In the SSIS execution engine (the Parent.dtsx SSIS package), an Execute SQL Task was used to insert the initial ApplicationPackageInstance record into the log. ApplicationPackageInstance table and return an ApplicationPackageInstanceId value. In the Azure Data Factory execution engine (*now the parent and child pipelines*), a Lookup activity will call a stored procedure that inserts the record and returns an ApplicationPackageInstanceId value. The ApplicationPackageInstanceId value will be used later in the pipeline to update the ApplicationPackageInstance record.

In Azure Data Studio (or SSMS), connect to the Azure SQL database named SSISConfig. Starting with the T-SQL code from Listing 7-7 in Chapter 7, edit the T-SQL to create a new stored procedure named “log.InsertApplicationPackageInstance” in an idempotent (re-executable) manner, as shown in Listing 10-7.

Listing 10-7. Idempotent T-SQL to create the log.InsertApplicationPackageInstance stored procedure

```

print 'log.InsertApplicationPackageInstance stored procedure'
If Exists(Select s.[name] + '.' + p.[name]
        From [sys].[procedures] p
        Join [sys].[schemas] s
        On s.[schema_id] = p.[schema_id]
        Where s.[name] = N'log'
        And p.[name] = N'InsertApplicationPackageInstance')
begin
    print ' - Dropping log.InsertApplicationPackageInstance stored procedure'
    Drop Procedure log.InsertApplicationPackageInstance

```

```

print ' - Log.InsertApplicationPackageInstance stored procedure dropped'
end

print ' - Creating log.InsertApplicationPackageInstance stored procedure'
go

Create Procedure log.InsertApplicationPackageInstance
    @ApplicationInstanceId int
    , @ApplicationPackageId int
    , @ApplicationPackageRunId nvarchar(55) = NULL
As

Insert Into [log].ApplicationPackageInstance
( ApplicationInstanceId
, ApplicationPackageId
, ApplicationPackageRunId)
Output inserted.ApplicationPackageInstanceId
Values
( @ApplicationInstanceId
, @ApplicationPackageId
, @ApplicationPackageRunId)

go

print ' - Log.InsertApplicationPackageInstance stored procedure created'
go

```

The `log.InsertApplicationPackageInstance` stored procedure requires the application instance id of the application instance be sent to the `@ApplicationInstanceId int` parameter in order to *correlate* the application instance with the application package instance. Remember, `ApplicationInstanceId` is a child pipeline parameter sent to the child pipeline from the parent pipeline. Another child pipeline parameter passed from the parent pipeline is `ApplicationPackageId`. We pass the value of the `ApplicationPackageId` pipeline parameter to the `log.InsertApplicationPackageInstance` stored procedure.

The `log.InsertApplicationPackageInstance` stored procedure also accepts a parameter named `@ApplicationPackageRunId` which defaults to `NULL` when no value is supplied. `@ApplicationPackageRunId` is used to store the Azure Data Factory child

pipeline's RunId value. As stated earlier, storing the RunId in the SSIS framework allows enterprise DevOps teams to relate operational log data in the SSIS framework to operational data captured in ADF logs, which is *extremely* useful when something unfortunate happens during ADF pipeline execution.

The T-SQL Insert statement initializes the ApplicationPackageInstance record by inserting new @ApplicationInstanceId, @ApplicationPackageId, and @ApplicationPackageRunId values into the log.ApplicationPackageInstance table. Defaults configured on the log.ApplicationPackageInstance table insert the following values in the row:

- ApplicationPackageStartTime is set to sysdatetimeoffset() by DF_log_ApplicationPackageInstance_ApplicationPackageStartTime.
- ApplicationPackageStatus is set to N'Running' by DF_log_ApplicationPackageInstance_ApplicationPackageStatus.

If all goes as planned, Azure Data Studio should return Messages that appear as shown in Figure 10-37.

Messages

```
Started executing query at Line 1
log.InsertApplicationPackageInstance stored procedure
  - Creating log.InsertApplicationPackageInstance stored procedure
Started executing query at Line 16
Commands completed successfully.
Started executing query at Line 34
  - Log.InsertApplicationPackageInstance stored procedure created
Total execution time: 00:00:00.136
```

Figure 10-37. Messages from creating the log.InsertApplicationPackageInstance stored procedure

After the log.InsertApplicationPackageInstance stored procedure has been created, the next step is to add an activity to the pipeline to execute the stored procedure.

Logging Application Package Instance

The ADF execution engine needs to execute a stored procedure, log.InsertApplicationPackageInstance, that returns a value named ApplicationPackageInstanceId. A Lookup activity is built for this!

Return to the child ADF pipeline. Drag a new Lookup activity onto the surface and rename the new Lookup activity to “Log Application Package Instance Start,” as shown in Figure 10-38.

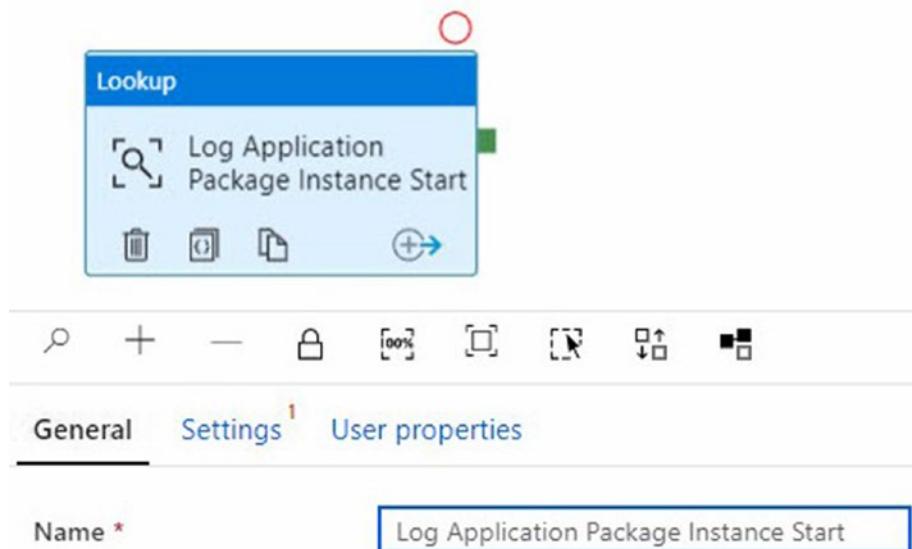


Figure 10-38. Adding the Log Application Package Instance Start lookup activity

Click the Log Application Instance Start lookup activity’s Settings tab to continue configuration. Click the “Source dataset” property drop-down, and select “ssisFrameworkDataSet,” and then set the Log Application Package Instance Start lookup activity’s “Use query” property option to “Stored procedure.” Select “log. InsertApplicationPackageInstance” from the “Name” property drop-down, as shown in Figure 10-39.

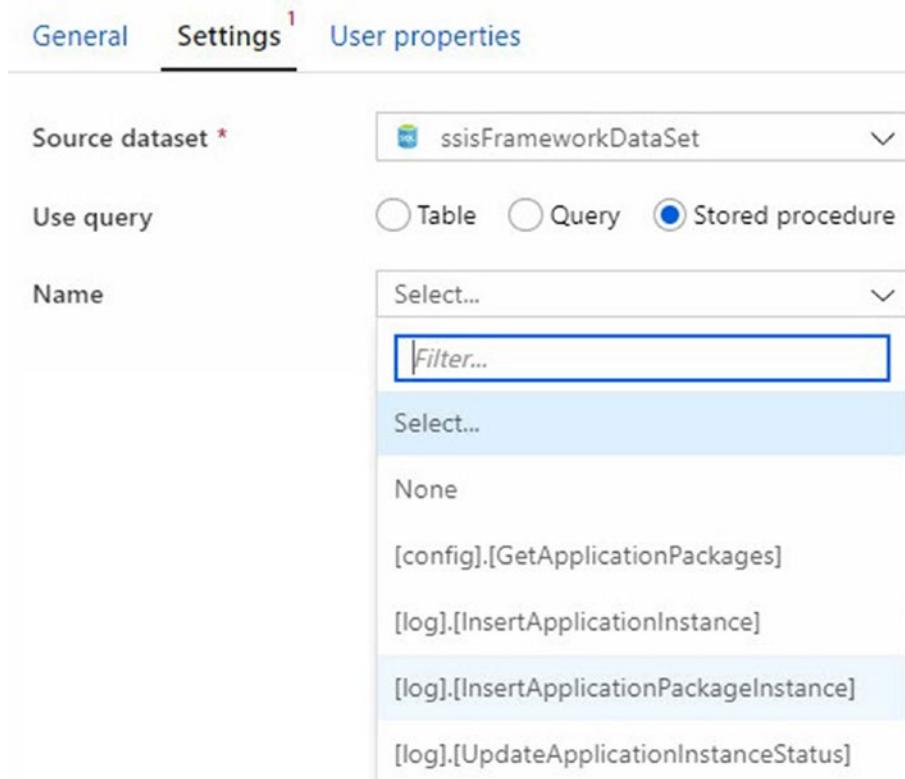


Figure 10-39. Selecting the `log.InsertApplicationPackageInstance` stored procedure

Next, click the “Import parameter” button to import the `@ApplicationInstanceId`, `@ApplicationPackageId`, and `@ApplicationPackageRunId` Int32 and String parameters from the `log.InsertApplicationPackageInstance` stored procedure. Click inside the `@ApplicationInstanceId` parameter’s Value property textbox, and then click the “Add dynamic content [Alt + P]” link, as shown in Figure 10-40.

Import parameter			
Parameter	New	Delete	
	NAME	TYPE	VALUE
	ApplicationInstanceId	Int32	Value
			Add dynamic content [Alt+P]
	ApplicationPackageId	Int32	Value
	ApplicationPackageRu	String	Value

Figure 10-40. Preparing to add dynamic content to the parameter value property

When the “Add dynamic content” blade displays, expand the Parameters expression category and click “ApplicationInstanceId” to set the parameter value ADF expression to @pipeline.parameters.ApplicationInstanceId, as shown in Figure 10-41.

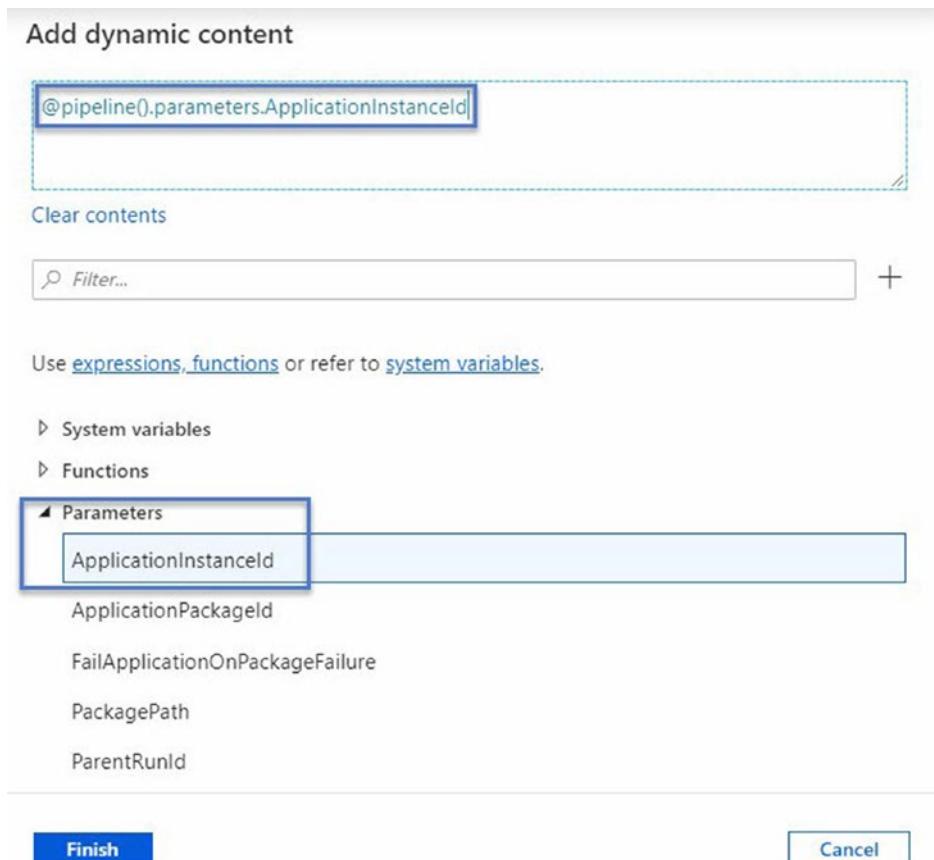


Figure 10-41. Setting the ApplicationInstanceId parameter value expression

Click the Finish button to proceed.

Repeat the preceding procedure to assign the @ApplicationPackageId stored procedure parameter to the child pipeline's ApplicationPackageId parameter as shown in Figure 10-42.

NAME	TYPE	VALUE
ApplicationInstanceId	Int32	@pipeline().parameters.ApplicationInstanceId
ApplicationPackageId	Int32	@pipeline().parameters.ApplicationPackageId
ApplicationPackageRunId	String	Value

Figure 10-42. Setting the ApplicationPackageId parameter value expression

Click inside the @ApplicationPackageRunId parameter's Value property textbox, and then click the “Add dynamic content [Alt + P]” link. When the “Add dynamic content” blade displays, expand the “System variables” expression category, and click “Pipeline run ID” to set the parameter value ADF expression to @pipeline().RunId, as shown in Figure 10-43.

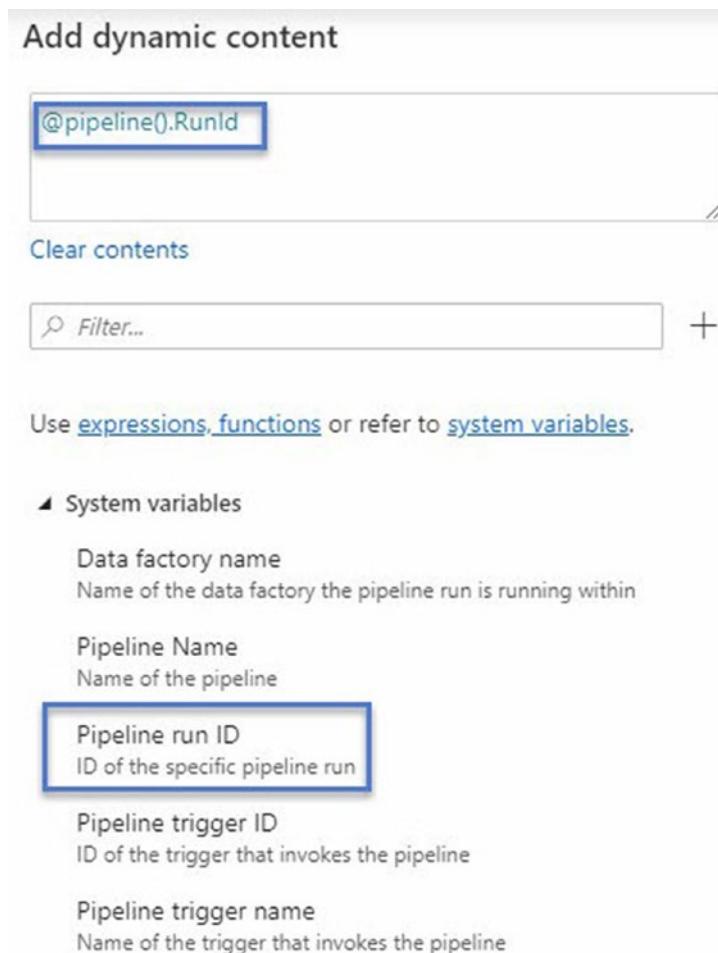


Figure 10-43. Setting the `ApplicationPackageRunId` parameter value expression

Leave the “Query timeout (minutes)” property set to its default (120 minutes), the “Isolation level” property set to “None” (default), and the “First row only” property checked (default).

The Log Application Package Instance Start lookup activity is now configured to call the `log.InsertApplicationPackageInstance` stored procedure, passing the stored procedure the child pipeline’s `ApplicationInstanceId`, `ApplicationPackageId`, and `ApplicationPackageRunId` parameter values and returning the `ApplicationPackageInstanceId` value.

Connect a Success output from the Log Application Package Instance Start lookup activity to the Execute Application Package execute SSIS activity.

Adding the Log.UpdateApplicationPackageInstanceStatus Stored Procedure

As configured, the child pipeline fails execution when called to execute the ReportAndFail.dtsx SSIS package because the ReportAndFail.dtsx SSIS package is designed to fail. The log.ApplicationPackageInstance row needs to be updated to reflect the SSIS framework application package instance failed.

Begin by adding a new stored procedure named “log.UpdateApplicationPackageInstanceStatus” to the SSISConfig database, as shown in Listing 10-8.

Listing 10-8. Idempotent T-SQL to create the log.

UpdateApplicationPackageInstanceStatus stored procedure

```
print 'log.UpdateApplicationPackageInstanceStatus stored procedure'
If Exists(Select s.[name] + '.' + p.[name]
          From [sys].[procedures] p
          Join [sys].[schemas] s
            On s.[schema_id] = p.[schema_id]
          Where s.[name] = N'log'
            And p.[name] = N'UpdateApplicationPackageInstanceStatus')
begin
    print ' - Dropping log.UpdateApplicationPackageInstanceStatus
          stored procedure'
    Drop Procedure log.UpdateApplicationPackageInstanceStatus
    print ' - Log.UpdateApplicationPackageInstanceStatus stored procedure
          dropped'
end

print ' - Creating log.UpdateApplicationPackageInstanceStatus stored
      procedure'
go

Create Procedure log.UpdateApplicationPackageInstanceStatus
    @ApplicationPackageInstanceId int
    , @ApplicationPackageStatus nvarchar(55) = 'Succeeded'
```

As

```

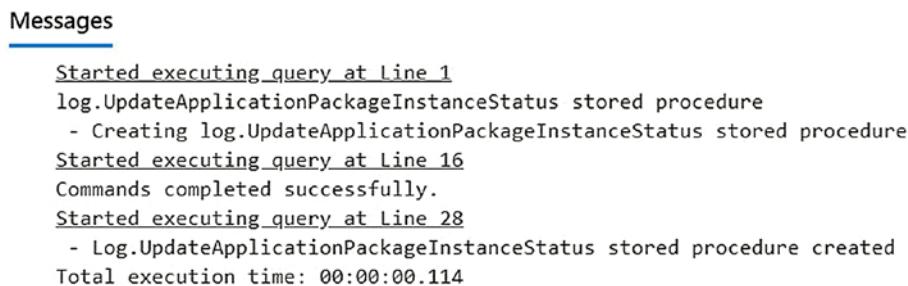
Update [log].ApplicationPackageInstance
Set ApplicationPackageEndTime = sysdatetimeoffset()
    , ApplicationPackageStatus = @ApplicationPackageStatus
Where ApplicationPackageInstanceId = @ApplicationPackageInstanceId

go

print ' - Log.UpdateApplicationPackageInstanceStatus stored procedure
created'
go

```

If all goes as planned, Azure Data Studio messages should appear similar to Figure 10-44.



The screenshot shows the 'Messages' tab in Azure Data Studio. It displays the following log entries:

- Started executing query at Line 1
- log.UpdateApplicationPackageInstanceStatus stored procedure
 - Creating log.UpdateApplicationPackageInstanceStatus stored procedure
- Started executing query at Line 16
- Commands completed successfully.
- Started executing query at Line 28
- Log.UpdateApplicationPackageInstanceStatus stored procedure created
- Total execution time: 00:00:00.114

Figure 10-44. Azure Data Studio messages reflecting successful creation of the log.UpdateApplicationPackageInstanceStatus stored procedure

The next step is to add a Stored Procedure Activity that calls the log.UpdateApplicationPackageInstanceStatus stored procedure on failure.

Updating Application Package Instance

The ADF execution engine needs to execute a stored procedure, log.UpdateApplicationPackageInstanceStatus, that accepts ApplicationPackageInstanceId and ApplicationPackageStatus values and then updates the status of the current Application Package Instance. A Stored Procedure Activity is perfect for this!

Return to the child ADF pipeline. Drag a Stored Procedure Activity onto the surface. Add a Failure output from the Execute Application Package activity to the new Stored Procedure Activity. Rename the new Stored Procedure Activity to “Log Application Package Instance Failure,” as shown in Figure 10-45.

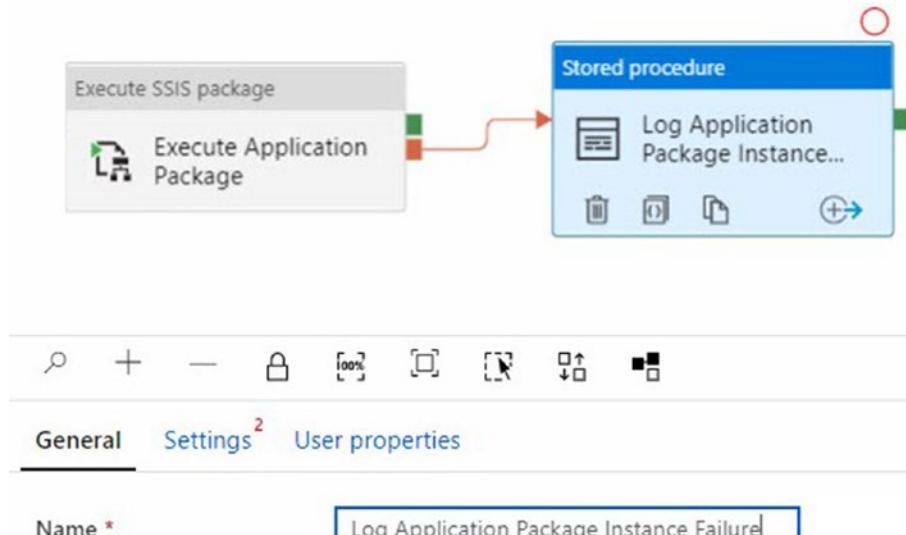


Figure 10-45. Adding the “Log Application Package Instance Failure” stored procedure activity

On the Settings tab, select “ssisFrameworkLinkedService” from the “Linked service” property drop-down, and then select “[log].[UpdateApplicationPackageInstanceState]” from the “Stored procedure name” drop-down. Click the “Import parameter” button to import parameters from the log.UpdateApplicationPackageInstanceState stored procedure. Click inside the value textbox for the ApplicationPackageInstanceId parameter, and then click the “Add dynamic content [Alt + P]” link. When the “Add dynamic content” blade displays, expand the “Activity outputs” expressions category, and click “Log Application Package Instance Start.” The dynamic content textbox will display @activity('Log Application Package Instance Start').output.Append .firstrow.ApplicationPackageInstanceId to the expression to map the value of the ApplicationPackageInstanceId returned from the execution of the “Log Application Package Instance Start” lookup activity into the ApplicationPackageInstanceId parameter’s value property, as shown in Figure 10-46.

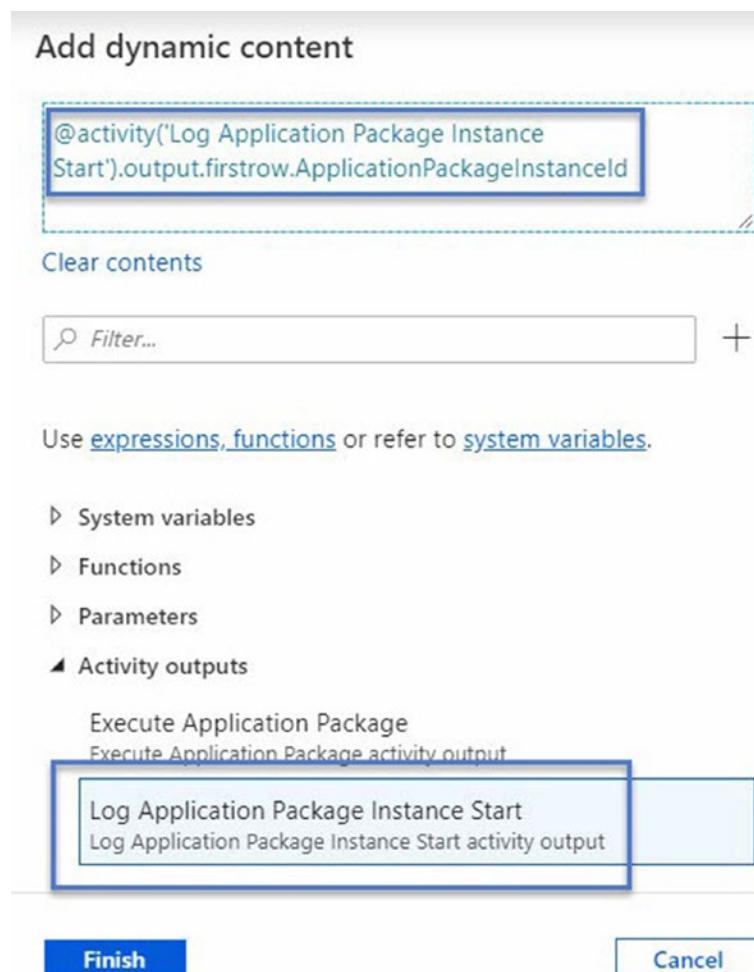


Figure 10-46. Mapping the “Log Application Package Instance Start” lookup activity into the ApplicationPackageInstanceId parameter’s value property

Click the Finish button to proceed.

Type “Failed” into the ApplicationStatus parameter’s value property, as shown in Figure 10-47.

NAME	TYPE	VALUE
ApplicationPackageIdns	Int32	@activity('Log Application Package Instance Start').output.firstrow.ApplicationPackageInstanceId
ApplicationPackageStatus	String	Failed

Add dynamic content [Alt+P]

Figure 10-47. Setting the *ApplicationPackageStatus* property to Failed

The child pipeline for the ADF version of the SSIS framework execution engine is now configured to execute an SSIS application package and record a failed application package instance.

The next step is to test-execute the pipeline.

Let's Test It!

Return to the parent pipeline to test the execution engine. Click the Debug item on the parent pipeline's toolbar. The Output tab displays and, after some time, reveals the parent pipeline activities executed, as shown in Figure 10-48.

Parameters	Variables	Output	
NAME	TYPE	DURATION	STATUS
Log Application Instance Failure	SqlServerStoredProcedure	00:00:02	✓ Succeeded
Execute Child Pipeline	ExecutePipeline	00:00:33	✓ Succeeded
Execute Child Pipeline	ExecutePipeline	00:00:29	✗ Failed
ForEach Application Package	ForEach	00:01:06	✗ Failed
Get Application Packages	Lookup	00:00:04	✓ Succeeded
Log Application Instance Start	Lookup	00:00:02	✓ Succeeded

Figure 10-48. Results of a text execution of the parent pipeline

Execute the T-SQL query shown in Listing 10-9 to examine the Application Instance results.

Listing 10-9. Application execution results

```
Select top 1
    a.ApplicationName
    , ai.ApplicationInstanceId
    , ai.ApplicationStatus
From log.ApplicationInstance ai
Join config.Applications a
    On a.ApplicationId = ai.ApplicationId
Order By ApplicationInstanceId Desc

declare @ApplicationInstanceId int = (Select top 1 ApplicationInstanceId
                                         From log.ApplicationInstance
                                         Order By ApplicationInstanceId DESC)

Select p.PackageName
    , ai.ApplicationInstanceId
    , api.ApplicationPackageInstanceId
    , api.ApplicationPackageStatus
    , ap.FailApplicationOnPackageFailure
From log.ApplicationPackageInstance api
Join log.ApplicationInstance ai
    On ai.ApplicationInstanceId = api.ApplicationInstanceId
Join config.ApplicationPackages ap
    On ap.ApplicationPackageId = api.ApplicationPackageId
Join config.Packages p
    On p.PackageId = ap.PackageId
Where ai.ApplicationInstanceId = @ApplicationInstanceId
Order By ApplicationPackageInstanceId
```

The results should appear similar to Figure 10-49:

Results		Messages		
ApplicationName	ApplicationInstanceId	ApplicationStatus		
1 Framework Test	61	Succeeded		
PackageName	ApplicationInstanceId	ApplicationPackageInstanceId	ApplicationPackageStatus	FailApplicationOnPackageFailure
1 ReportAndFail.dtsx	61	73	Failed	1
2 ReportAndSucceed.dtsx	61	74	Running	1

Figure 10-49. Application execution results

While the results match what was expected, something is missing: there is currently no Stored Procedure Activity that updates the Application Package Instance on Success, which is why the results of the query in Listing 10-9, shown in Figure 10-49, display the ReportAndSucceed.dtsx SSIS ApplicationPackageStatus as “Running.”

Updating Application Package Instance on Success

The child pipeline needs a way to log Application Package Instance success. Begin by clicking the Clone icon on the “Log Application Package Instance Failure” stored procedure activity as shown in Figure 10-50.

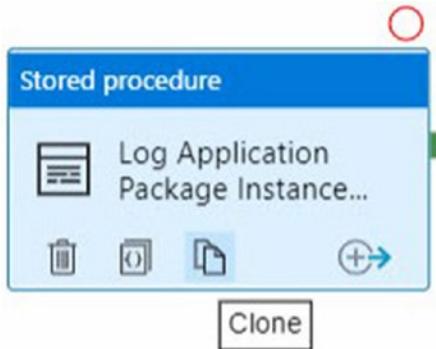


Figure 10-50. Preparing to clone the “Log Application Package Instance Success” stored procedure activity from the “Log Application Package Instance Failure” stored procedure activity

As with cloning the “Log Application Instance Failure” stored procedure activity earlier, cloning the “Log Application Package Instance Failure” stored procedure activity makes sense because the new Stored Procedure Activity will be calling the *same* stored procedure – log.UpdateApplicationPackageInstanceStatus – and *only one* parameter value configuration will change.

Connect a Success output from the “Execute Application Package” execute SSIS package activity to the cloned Stored Procedure Activity, and then rename the cloned Stored Procedure Activity to “Log Application Package Instance Success,” as shown in Figure 10-51.

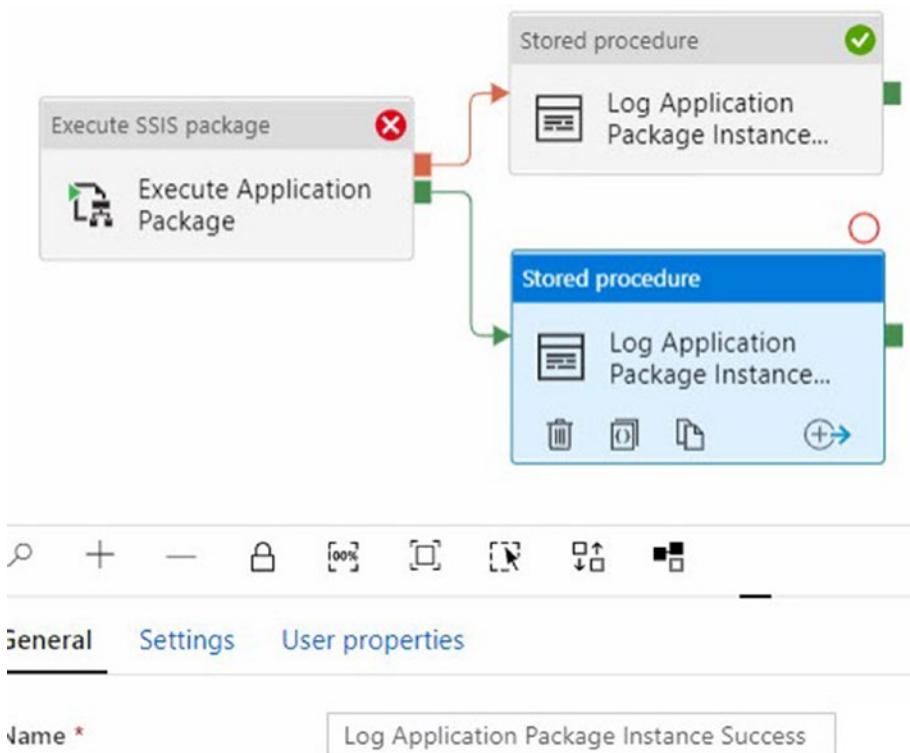


Figure 10-51. Adding the “Log Application Package Instance Success” stored procedure activity

Next, click the Settings tab for the “Log Application Package Instance Success” stored procedure activity. The only change required on the Settings tab is the value of the ApplicationPackageStatus parameter’s value property, which needs to be updated from “Failed” to “Succeeded,” as shown in Figure 10-52.

The screenshot shows the 'Settings' tab of the ADF interface. At the top, there are tabs for 'General', 'Settings' (which is selected), and 'User properties'. Below the tabs, there's a section for 'Linked service *' with a dropdown set to 'ssisFrameworkLinkedService' and a 'Test connection' button. Under 'Details', the 'Stored procedure name *' is set to '[log].[UpdateApplicationPackageInstanceSt...]' with an 'Edit' link. A 'Import parameter' button is highlighted with a blue border. Below this, 'Stored procedure parameters' are listed with a 'New' and 'Delete' button. There are two parameters: 'ApplicationPackageInstns' (Int32) with a value of '@activity('Log Application Package Instance Start').output.firstrow.ApplicationPackageInstanceId' and 'ApplicationPackageStat...' (String) with a value of 'Succeeded'. A 'Add dynamic content [Alt+P]' link is also present.

Figure 10-52. Updating the ApplicationPackageStatus parameter value

A fresh test execution confirms the parent pipeline logs SSIS framework application failures as designed, and re-executing the query in Listing 10-9 confirms – as shown in Figure 10-53.

Results		Messages		
	ApplicationName	ApplicationInstanceId	ApplicationStatus	
1	Framework Test	62	Running	
	PackageName	ApplicationInstanceId	ApplicationPackageInstanceId	ApplicationPackageStatus
1	ReportAndFail.dtsx	62	75	Failed
2	ReportAndSucceed.dtsx	62	76	Succeeded
				FailApplicationOnPackageFailure

Figure 10-53. ReportAndSucceed.dtsx application package succeeds

In our examples thus far, each application instance fails because a child package (named “ReportAndFail.dtsx”) fails each time it is executed. The next step is to add fault tolerance to the SSIS framework, ADF version, and that is the topic of the next chapter.

Conclusion

In this chapter, we added logging functionality to the parent pipeline, modified the SSISConfig log tables, and encapsulated logging logic in stored procedures. We decoupled the execution of SSIS framework application packages by adding a child pipeline to perform application package execution.

CHAPTER 11

Fault Tolerance in the ADF Framework

The previous chapter covered logging functionality for the Azure Data Factory version of the SSIS framework's execution engine. Differences in functionality available to SSIS packages and functionality available to Azure Data Factory pipelines drove design changes.

In this chapter, we complete ADF execution engine functionality by implementing fault tolerance to programmatically stop pipeline execution based on SSISConfig metadata configurations.

A Brief Introduction to Fault Tolerance

In the author's humble opinion, *fault tolerance* is a different way of describing graceful failure. Thinking "what happens when this fails?" is one difference between a technician and an engineer. Technicians complete projects that accomplish some task, while engineers build solutions that manage issue domains. It comes down to what an individual considers *done*. Technicians work through failures and stop when the project works. Engineers do not stop until the solution does not fail; and – when the solution *does fail* – engineers make certain the failure is graceful.

You may recall fault tolerance was the most complex portion of the design of the on-premises SSIS framework execution engine (the Parent.dtsx package) designed in Chapters 5-7. The design of the Azure Data Factory version of the SSIS framework execution engine (the parent and child pipelines) is no different. To achieve fault tolerance, the ADF execution strategy requires radical redesign.

The example has already implemented some of the required redesign by building the child pipeline. Why did application package execution need to move to the child pipeline? We answer this question in this chapter.

Add ADF Managed Identity to Contributor Role

In this section, we will use the Azure Data Factory REST API. If that concerns you, don't worry; using the ADF REST API is interesting and will level up your Azure Data Factory automation skills!

The first step is to add the Azure Data Factory Managed Identity to the Contributor role in the Azure Data Factory Access control blade. Begin by logging into the Azure portal and navigating to Azure Data Factory blade. Click "Access control (IAM)," as shown in Figure 11-1.

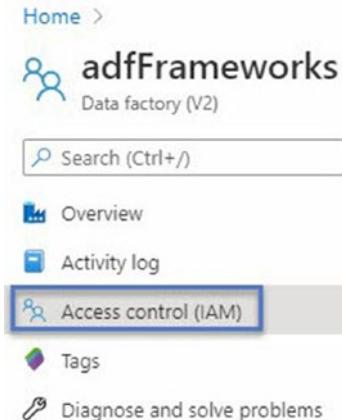


Figure 11-1. Navigating to the Access control (IAM) blade

When the Access control (IAM) blade displays, click the "Add a role assignment" Add button, as shown in Figure 11-2.

The screenshot shows two blades from the Azure portal. On the left, the 'Check access' blade is open, featuring a 'Find' search bar and a dropdown menu set to 'Azure AD user, group, or service principal'. On the right, the 'Add a role assignment' blade is open, displaying a brief description of what role assignments do, an 'Add' button, and a 'Learn more' link.

Figure 11-2. Preparing to assign a role

When the “Add role assignment” blade displays, select “Contributor” from the Role drop-down. Leave the “Assign access to” drop-down set to the default value (“Azure AD user, group, or service principal”), and enter the name of your Azure Data Factory in the Select textbox, as shown in Figure 11-3.

The screenshot shows the 'Add role assignment' blade. It has three input fields: 'Role' (set to 'Contributor'), 'Assign access to' (set to 'Azure AD user, group, or service principal'), and 'Select' (containing the text 'adfFrameworks'). Below the 'Select' field is a list box showing a single item, 'adfFrameworks', which is highlighted with a blue border.

Figure 11-3. Finding the Azure Data Factory managed identity

In the “Add role assignment” blade, click the Azure Data Factory managed identity. The “Selected members” list reflects the current selection, and the Save button is enabled indicating the role assignment is ready to store, as shown in Figure 11-4.

Add role assignment

Role ⓘ
Contributor ⓘ

Assign access to ⓘ
Azure AD user, group, or service principal

Select ⓘ
adfFrameworks

No users, groups, or service principals found.

Selected members:

adfFrameworks Remove

Save **Discard**

Figure 11-4. Ready to complete the role assignment

Click the “Role assignments” tab to view role assignments for ADF, as shown in Figure 11-5.

Name	Type	Role	Scope
adfFrameworks	Data Factory	Contributor ⓘ	This resource

Figure 11-5. ADF managed identity assigned to Contributor role

Now that the Azure Data Factory managed identity is assigned to the Contributor role, ADF Web activities may interact with the many methods hosted within the Azure Data Factory REST API.

Add Application Package Fault Tolerance

How is fault tolerance implemented within the Azure Data Factory version of the SSIS framework?

About Fail Application on Package Failure

`FailApplicationOnPackageFailure` is stored along with the application package metadata in the `SSISConfig` database. The value of `FailApplicationOnPackageFailure` is returned to the “Get Application Packages” lookup activity in the parent pipeline, and then `FailApplicationOnPackageFailure` is passed to the child pipeline parameter (named `FailApplicationOnPackageFailure`) in the “Execute Child Pipeline” execute pipeline activity inside the “ForEach Application Package” foreach activity’s inner Activities.

Implementing Application Package Fault Tolerance

Begin implementing application package fault tolerance in the child pipeline by right-clicking the “Execute Application Package” execute SSIS package activity’s Failure output and then clicking Delete, as shown in Figure 11-6.

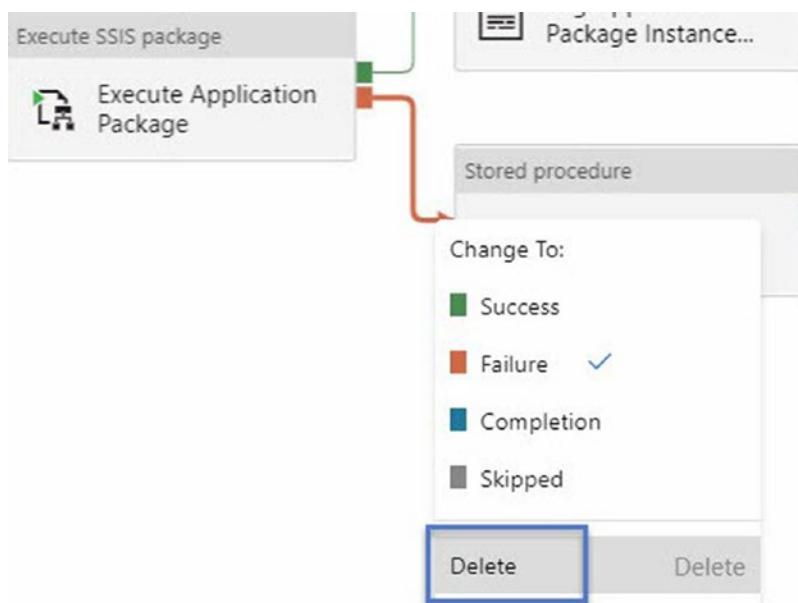


Figure 11-6. Deleting the “Execute Application Package” execute SSIS package activity’s Failure output

Drag an If Condition activity onto the child pipeline canvas and rename the If Condition “If Fail Application On Package Failure,” as shown in Figure 11-7.

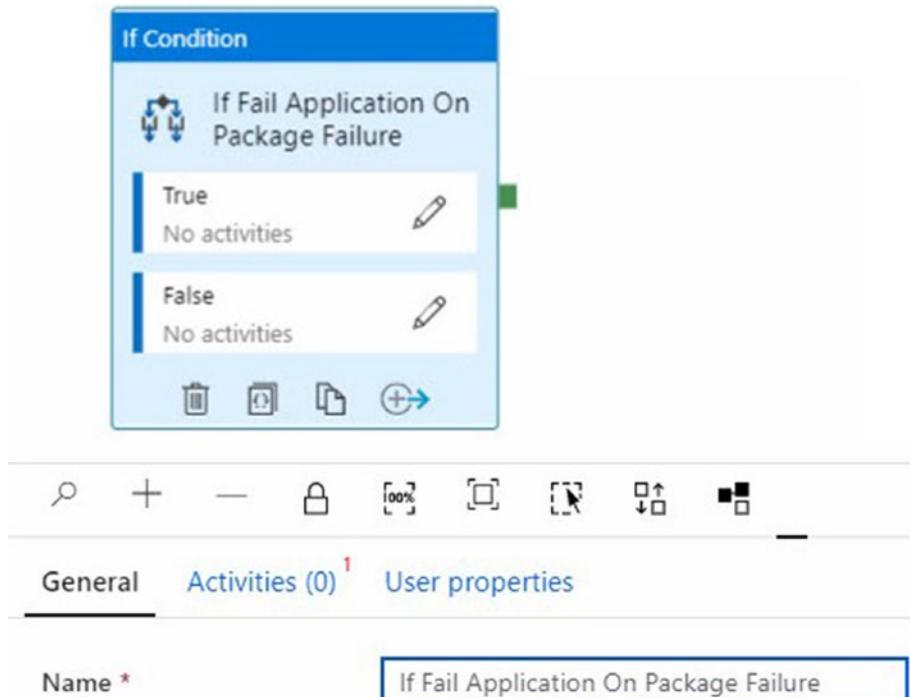


Figure 11-7. Adding the “If Fail Application On Package Failure” if condition activity

Connect a Failure output from the “Execute Application Package” execute SSIS activity to the “If Fail Application On Package Failure” if condition activity, as shown in Figure 11-8.

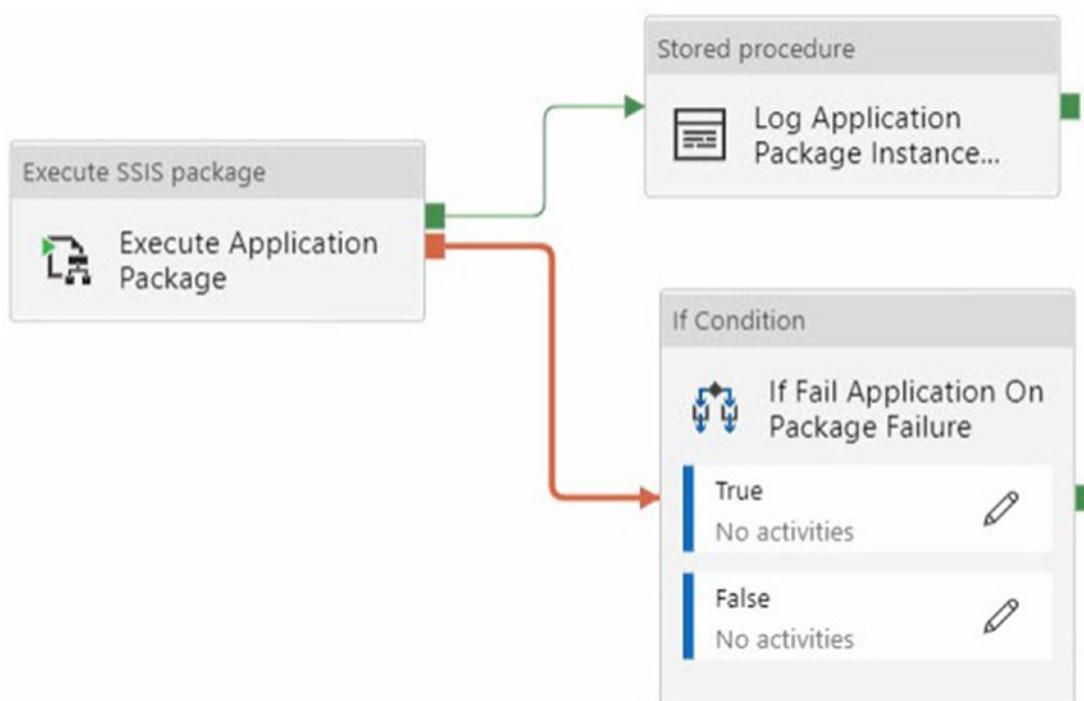


Figure 11-8. Connecting a Failure output to the “If Fail Application On Package Failure” if condition activity

Click the “If Fail Application On Package Failure” if condition activity’s Activities tab, and then click inside the expression textbox. Click the “Add dynamic content [Alt + P]” link beneath the expression textbox. When the “Add dynamic content” blade displays, expand the Parameters category and select the FailApplicationOnPackageFailure child pipeline parameter. The expression textbox will display “pipeline().parameters.FailApplicationOnPackageFailure”. Prefix the expression with “@bool(“ and add a closing parenthesis – “)” – to complete the expression, as shown in Figure 11-9.

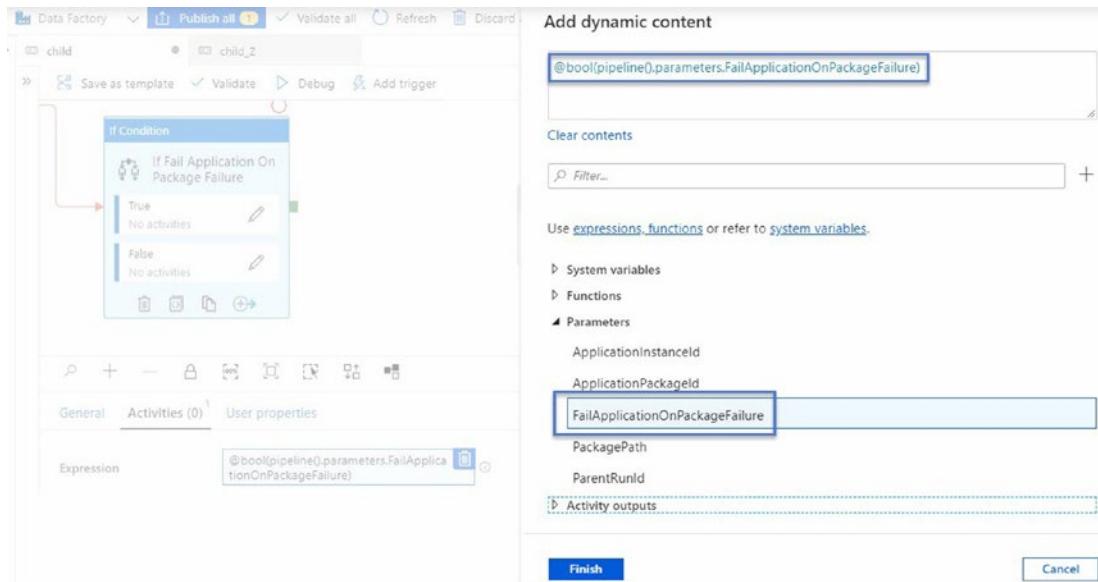


Figure 11-9. Configuring the expression property for the “If Fail Application On Package Failure” if condition activity

Click the Finish to complete this portion of the configuration.

Checking the Logic

Let’s pause here and think through the logic so far. There are two application package execution possible outcomes: the application package execution will either succeed or fail.

If the application package execution succeeds, the “Log Application Package Instance Success” stored procedure activity executes and updates the log. ApplicationPackageInstance table’s ApplicationPackageStatus value to “Succeeded” for the current ApplicationPackageInstanceId.

The on-success – first use case – is already implemented. Fault tolerance logic is required *if and only if* the application package execution fails.

If the application package execution fails, the “If Fail Application On Package Failure” if condition activity evaluates the value of the child package pipeline’s FailApplicationOnPackageFailure parameter (converted by the expression to a Boolean value).

The second use case occurs when the application package execution fails *and* FailApplicationOnPackageFailure is false. In this case, we want to

- Log the status (failure with information) of the application package execution
- *Continue* the execution of the SSIS framework application's application packages

The third use case occurs when the application package execution fails *and* FailApplicationOnPackageFailure is true. In this case, we want to

- Log the status (failure with information) of the application package execution
- Log the status (failure) of the application execution
- *Stop* the execution of application packages

Implementing the Fault Tolerance Logic

To begin implementing the fault tolerance logic, click the “If Fail Application On Package Failure” if condition activity’s False activities editor, as shown in Figure 11-10.

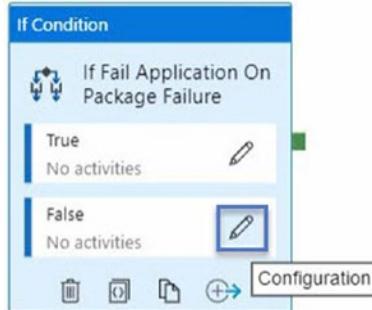


Figure 11-10. Opening the “If Fail Application On Package Failure” if condition activity’s False activities editor

Drag a Stored Procedure Activity onto the canvas and rename the new Stored Procedure Activity “Log Application Package Instance Failed 0,” as shown in Figure 11-11.

child > If Fail Application On Package Failure > False activities

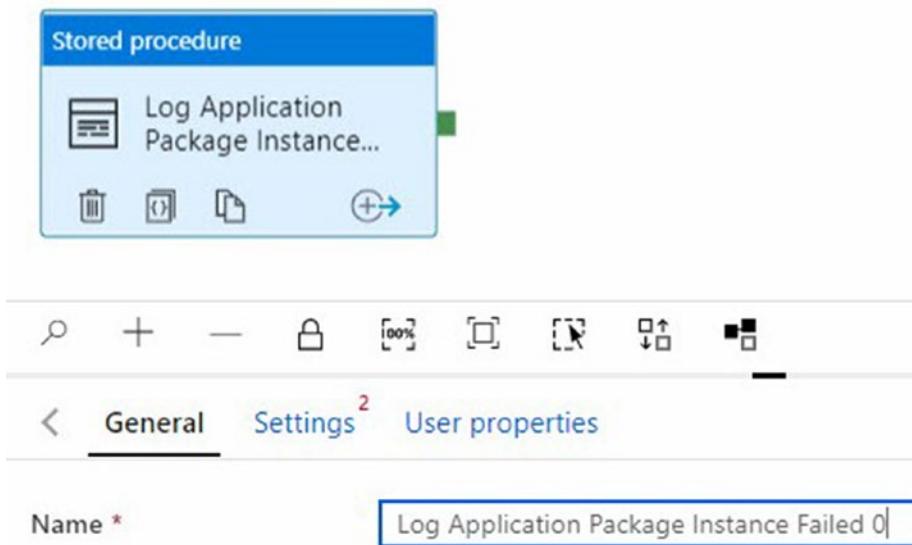


Figure 11-11. Adding the “Log Application Package Instance Failed 0” stored procedure activity

Click the Settings tab and set the Linked service property to ssisFrameworkLinkedService. Select log.UpdateApplicationPackageInstanceId from the Stored procedure name drop-down, as shown in Figure 11-12.

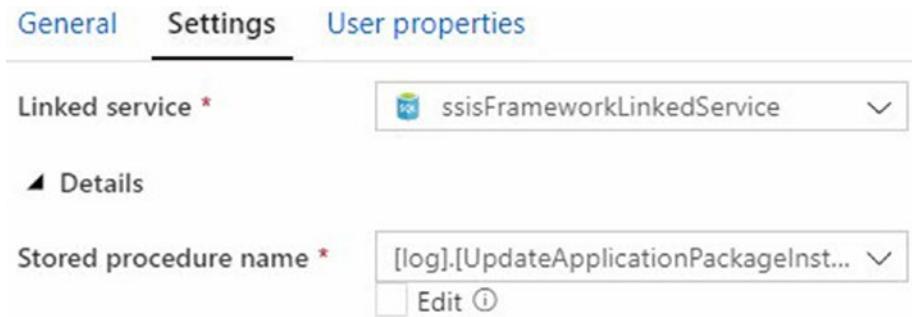


Figure 11-12. Configuring Linked service and Stored procedure name

Click the Import parameter button to begin configuring the log. UpdateApplicationPackageInstanceId stored procedure parameter values. Configure the ApplicationPackageInstanceId stored procedure parameter by clicking the “Add dynamic content [Alt + P]” link and selecting the “Log Application Package

“Instance Start” activity, which will enter the expression @activity('Log Application Package Instance Start').output into the expression textbox. Append “.firstrow.ApplicationPackageInstanceId” to the expression. In the ApplicationPackageStatus value textbox, enter “Failed (FailAppOnPkgFail: 0),” as shown in Figure 11-13.

NAME	TYPE	VALUE
ApplicationPackageInstanceId	Int32	@activity('Log Application Package Instance Start').output.firstrow.ApplicationPackageInstanceId
ApplicationPackageStatus	String	Failed (FailAppOnPkgFail: 0)

Figure 11-13. Configuring the ApplicationPackageInstanceId and ApplicationPackageStatus parameter values

Once stored procedure parameter values are set, configuration of the “Log Application Package Instance Failed 0” stored procedure activity is complete. Please note the ApplicationPackageStatus includes additional information: “(FailAppOnPkgFail: 0)”. The “Log Application Package Instance Failed 0” stored procedure addresses use case 2.

The ApplicationPackageStatus column in the log.ApplicationPackageInstance table is currently configured as nvarchar(25), which is not large enough to contain the text “Failed (FailAppOnPkgFail: 0)”. Execute the T-SQL in Listing 11-1 to expand the column size.

Listing 11-1. Expanding ApplicationPackageStatus

```
Alter Table [log].[ApplicationPackageInstance]
    Alter Column ApplicationPackageStatus nvarchar(55) Not NULL
```

Before discussing and adding code to support use case 3, a quick review of the ForEach activity’s default behavior is a good idea.

ForEach Activity Default Behavior

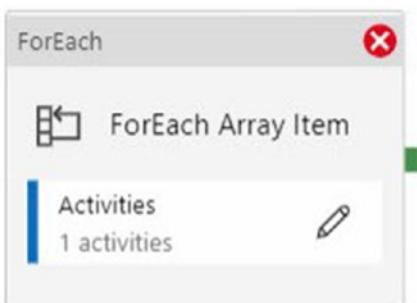
The ForEach activity's default iterating behavior is to continue iterating over the configured collection, executing "inner activities," until the items in the collection have all been iterated. One side effect is that errors occurring in the "inner activities" do *not* stop iteration. "Inner activities" are executed as the collection is iterated regardless of the outcome of said execution.

After iteration completes, any failures that occurred during iteration cause the ForEach activity to fail.

Assume a ForEach activity named "ForEach Array Item" contains a single "inner activity," a Wait activity named "Wait" (creative, I know). Assume an array variable containing four items is supplied to the ForEach activity's Items property. Consider the following scenario for ForEach "inner activities" execution in which the second iteration of the Wait activity fails, numbered by iteration:

1. Wait executes successfully.
2. Wait execution fails.
3. Wait executes successfully.
4. Wait executes successfully.

Each iteration will execute, resulting in four executions of the wait activity. The ForEach activity will fail. A Debug output will appear as shown in Figure 11-14.



The screenshot shows the configuration of a ForEach activity named 'ForEach'. It contains one child activity, 'ForEach Array Item'. Below this, a table displays the execution history of the 'Wait' activity and the overall ForEach execution.

NAME	TYPE	DURATION	STATUS
Wait	Wait	00:00:04	✓ Succeeded
Wait	Wait	00:00:03	✓ Succeeded
Wait	Wait	00:00:00	✗ Failed
Wait	Wait	00:00:02	✓ Succeeded
ForEach Array Item	ForEach	00:00:13	✗ Failed

Figure 11-14. Output of a ForEach activity execution

Figure 11-14 shows four executions of the Wait activity. Three succeeded, one failed. Please note each Wait activity executed, and two Wait activity executions fired *after* the Wait activity in the second iteration *failed*. Also note the ForEach activity failed.

ForEach Activity Default Behavior, Applied

Applied to the SSIS framework's fault tolerance, the “ForEach Application Package” ForEach activity – whose “inner activities” are configured in the parent pipeline to execute the child pipeline – is already configured to continue executing the child package regardless of the outcome of the child package execution. The next step, therefore, is to add logic to *stop* application package execution when our use case calls for it – when the application package execution fails *and* FailApplicationOnPackageFailure is true (use case 3).

To configure the “If Fail Application On Package Failure” if condition activity’s “True activities,” click the Edit icon (pencil) beside “True activities,” as shown in Figure 11-15.

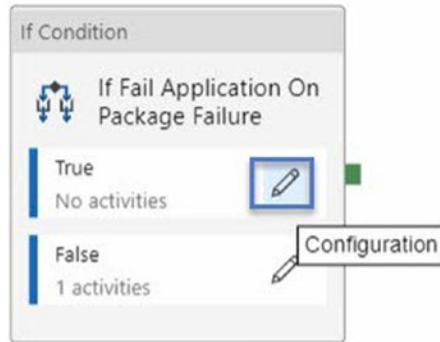


Figure 11-15. Preparing to configure the “True activities”

The order of operations for responding to a failed child package execution when the FailApplicationOnPackageFailure bit is configured to true (1, which is the default) is as follows:

1. Log application package instance failure
2. Log application instance cancelled
3. Set the command to cancel the parent pipeline
4. Execute the command to cancel the parent pipeline

The first step for responding to a failed child package execution when the FailApplicationOnPackageFailure bit is configured to true (1) is as follows: log application package instance failure.

To begin adding the functionality to log application package instance failure, add a Stored Procedure Activity to the “If Fail Application On Package Failure” if condition activity’s “True activities” canvas and rename the new Stored Procedure Activity “Log Application Package Instance Failed 1,” as shown in Figure 11-16.

00 child > ⚡ If Fail Application On Package Failure

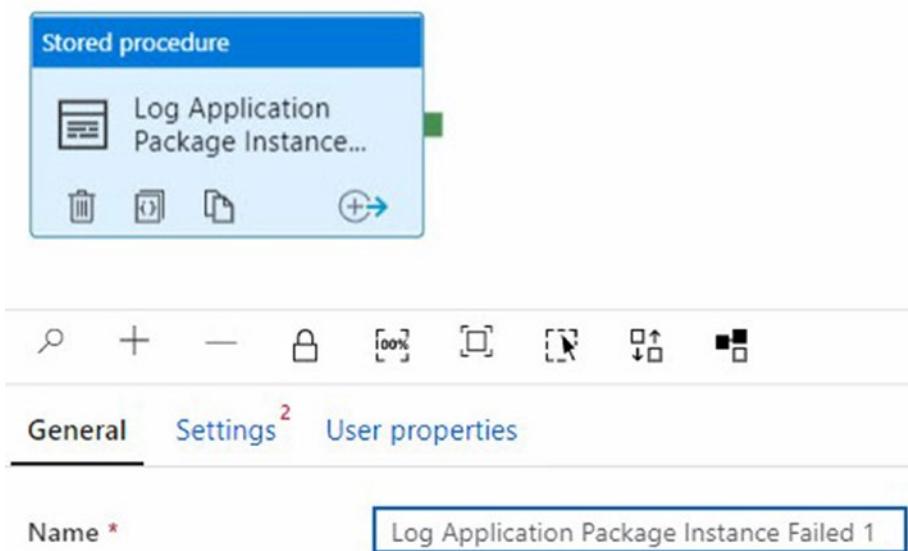


Figure 11-16. Adding the “Log Application Package Instance Failed 1” stored procedure activity

On the “Log Application Package Instance Failed 1” stored procedure activity’s Settings tab, click the “Linked service” property’s drop-down, and select “ssisFrameworkLinkedService”. In the “Stored procedure name” drop-down, select “[log].[UpdateApplicationPackageInstanceState]”, as shown in Figure 11-17.

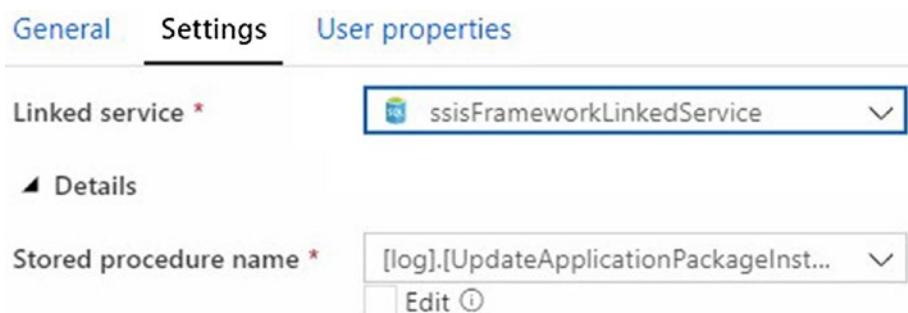


Figure 11-17. Configuring “Log Application Package Instance Failed 1” stored procedure activity Settings

Click the Import parameter button. When the ApplicationPackageInstanceId and ApplicationPackageStatus parameters are displayed, configure the ApplicationPackageInstanceId by clicking inside the value textbox, and then clicking the “Add dynamic content [Alt + P]” link, and then entering the expression “@activity('Log Application Package Instance Start').output.firstrow.ApplicationPackageInstanceId”. In the ApplicationPackageStatus parameter value textbox, enter the text “Failed (FailAppOnPkgFail: 1),” as shown in Figure 11-18.

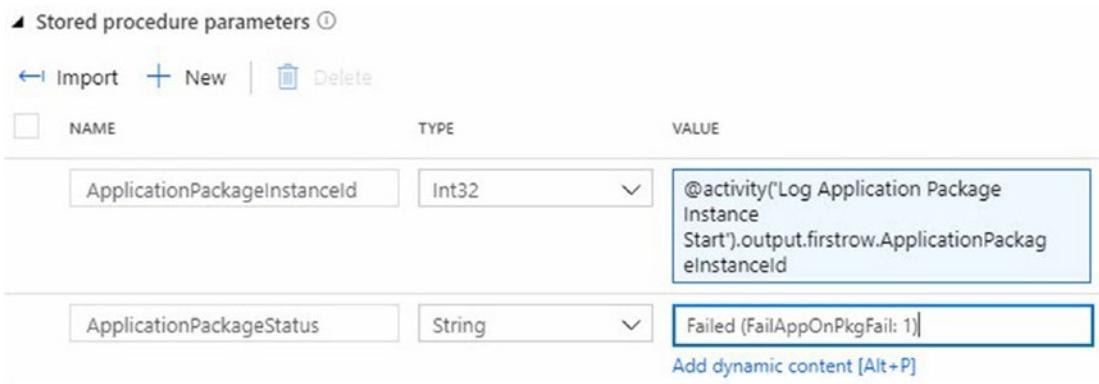


Figure 11-18. Configuring “Log Application Package Instance Failed 1” stored procedure activity Settings parameter values

The preceding completes the first step for responding to a failed child package execution when the FailApplicationOnPackageFailure bit is configured to true (1): log application package instance failure.

The second step for responding to a failed child package execution when the FailApplicationOnPackageFailure bit is configured to true (1) is as follows: log application instance cancelled.

To begin adding the functionality to log application instance cancelled, add a Stored Procedure Activity to the “If Fail Application On Package Failure” if condition activity’s “True activities” canvas, connect a Success output from the “Log Application Package Instance Failed 1” stored procedure activity and the new Stored Procedure Activity, and then rename the new Stored Procedure Activity “Log Application Instance Cancelled,” as shown in Figure 11-19.

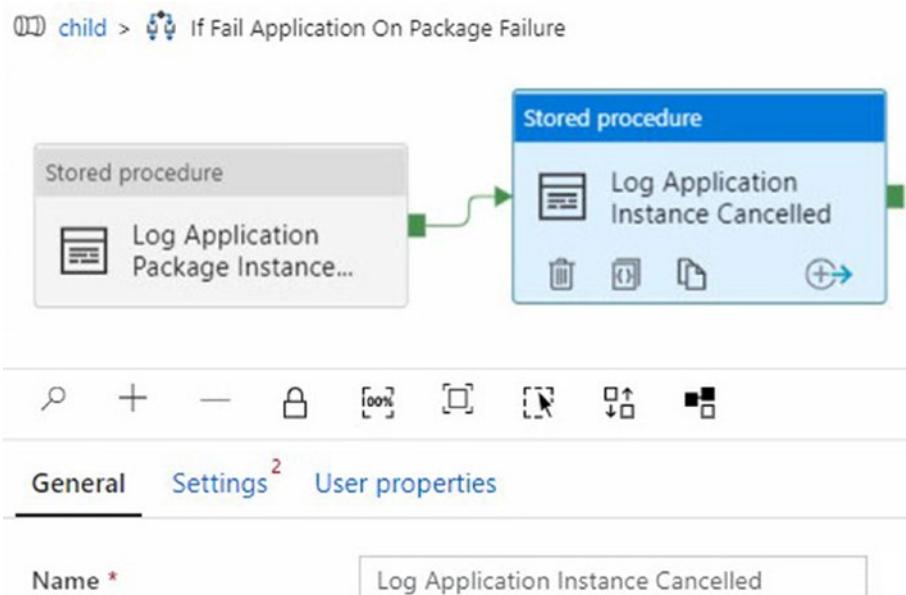


Figure 11-19. Adding the “Log Application Instance Cancelled” stored procedure activity

Connect a Success output from the “Log Application Package Failed 1” stored procedure activity to the “Log Application Instance Cancelled” stored procedure activity. On the “Log Application Instance Cancelled” stored procedure activity’s Settings tab, click the “Linked service” property’s drop-down, and select “ssisFrameworkLinkedService”. In the “Stored procedure name” drop-down, select “[log].[UpdateApplicationIntanceStatus]”, as shown in Figure 11-20.

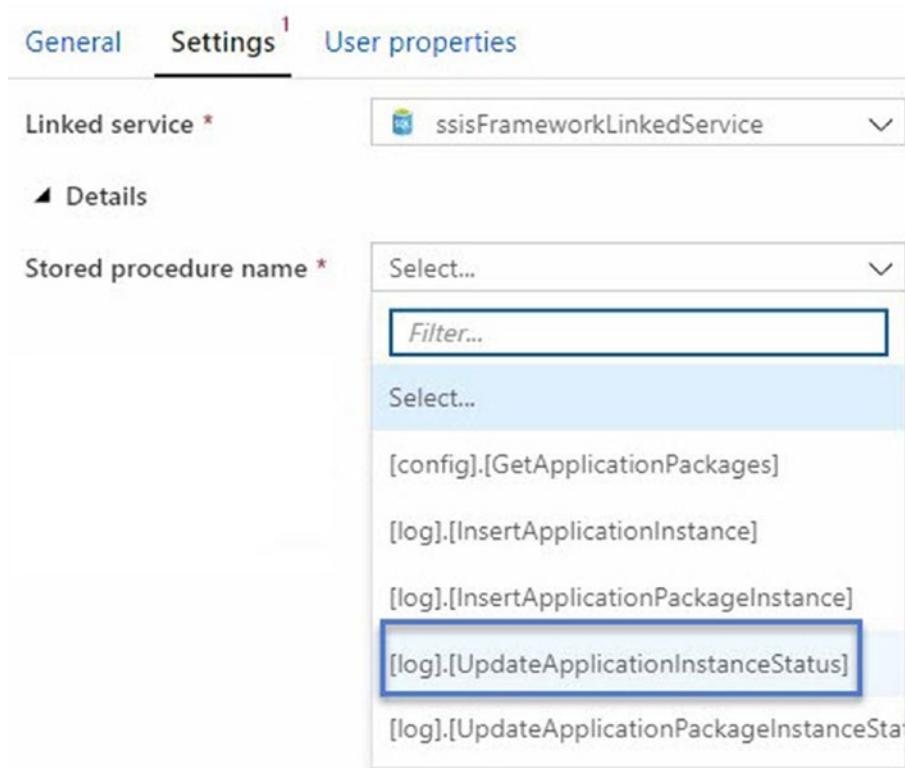


Figure 11-20. Configuring “Log Application Instance Cancelled” stored procedure activity Settings

Click the Import parameter button. When the ApplicationInstanceId and ApplicationStatus parameters are displayed, configure the ApplicationInstanceId by clicking inside the value textbox, and then clicking the “Add dynamic content [Alt + P]” link, and then entering the expression “@pipeline().parameters.ApplicationInstanceId”. In the ApplicationStatus parameter value textbox, enter the text “Cancelled,” as shown in Figure 11-21.

<input type="checkbox"/>	NAME	TYPE	VALUE
	ApplicationInstanceId	Int32	@pipeline().parameters.ApplicationInstanceId
	ApplicationStatus	String	Cancelled

Figure 11-21. Configuring “Log Application Instance Cancelled” stored procedure activity Settings parameter values

The preceding completes the second step for responding to a failed child package execution when the FailApplicationOnPackageFailure bit is configured to true (1): log application instance cancelled.

Before we begin this portion of the procedure, navigate to the child pipeline settings, and click the Variables tab. Click the “+ New” button, and add a String variable named “StopParentPipelineRunIdString,” as shown in Figure 11-22.

<input type="checkbox"/>	NAME	TYPE	DEFAULT VALUE
	StopParentPipelineRunIdString	String	Value

Figure 11-22. Adding the StopParentPipelineRunIdString variable to the child pipeline

To begin adding the functionality to build the cancel-pipeline run command, add a set variable activity to the “If Fail Application On Package Failure” if condition activity’s “True activities” canvas, connect a Success output from the “Log Application Instance Cancelled” stored procedure activity and the new set variable activity, and then rename the new set variable activity “Set Cancel Parent Pipeline Command,” as shown in Figure 11-23.

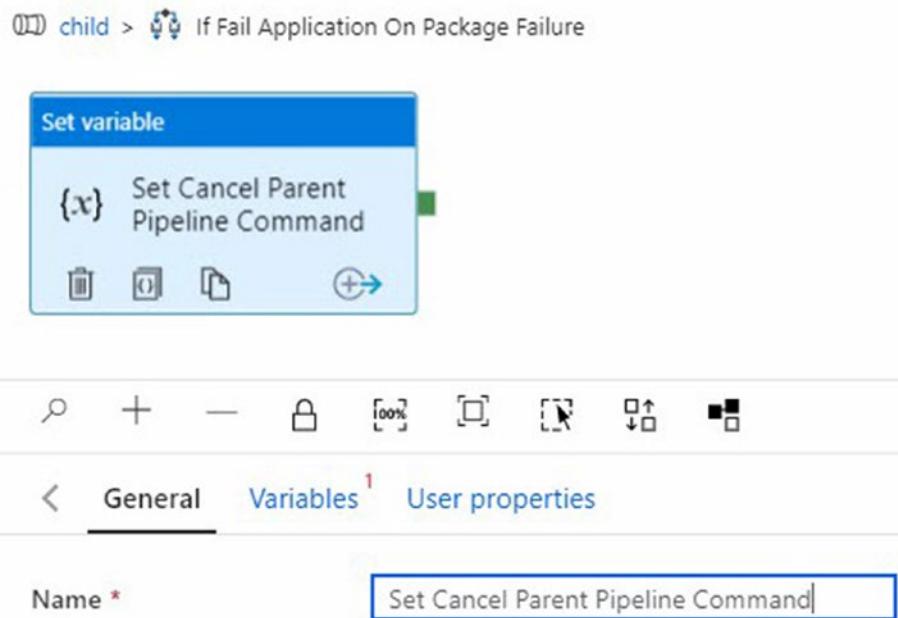


Figure 11-23. Adding the “Set Cancel Parent Pipeline Command” set variable activity

On the “Set Cancel Parent Pipeline Command” set variable activity’s Variables tab, click the “Name” property’s drop-down, and select “StopParentPipelineRunIdString”. Click inside the “Value” textbox, and then click the “Add dynamic content [Alt + P]” link. When the “Add dynamic content” blade displays, enter an expression similar to “@concat('https://management.azure.com/subscriptions/<subscription id>/resourcegroups/<resource group name>/providers/Microsoft.DataFactory/factories',pipeline().DataFactory,'/pipelineruns/',pipeline().parameters.ParentRunId,'/cancel?api-version=2018-06-01')”, as shown in Figure 11-24.

In the expression shown in Figure 11-24, replace <subscription id> with your Subscription Id value and <resource group name> with your Azure Data Factory’s Resource Group Name.

child_Z > If Fail Application On Package Failure > True activities

The screenshot shows the 'Variables' tab of the Azure Data Factory Variables blade. A modal window titled 'Set variable' is open, showing a single activity: '{x} Set Stop Parent Pipeline Command'. Below the modal, the 'Variables' section has two rows:

Name *	Value
StopParentPipelineRunIdString	<pre>@concat('https://management.azure.co m/subscriptions/<subscription id>/resourcegroups/<resource group name>/providers/Microsoft.DataFactory /factories',pipeline().DataFactory,'/pipeli neruns/',pipeline().parameters.ParentRu nId,'/cancel?api-version=2018-06-01')</pre>

Figure 11-24. Building the value of the “StopParentPipelineRunIdString” variable

The preceding completes the third step for responding to a failed child package execution when the FailApplicationOnPackageFailure bit is configured to true (1): set the command to cancel the parent pipeline.

The fourth step for responding to a failed child package execution when the FailApplicationOnPackageFailure bit is configured to true (1) is as follows: execute the command to cancel the parent pipeline.

To begin adding the functionality to execute the cancel-pipeline run command, add a web activity to the “If Fail Application On Package Failure” if condition activity’s “True activities” canvas, connect a Success output from the “Set Cancel Parent Pipeline Command” set variable activity and the new web activity, and then rename the new web activity “Stop Parent Execution,” as shown in Figure 11-25.

child > If Fail Application On Package Failure > True activities

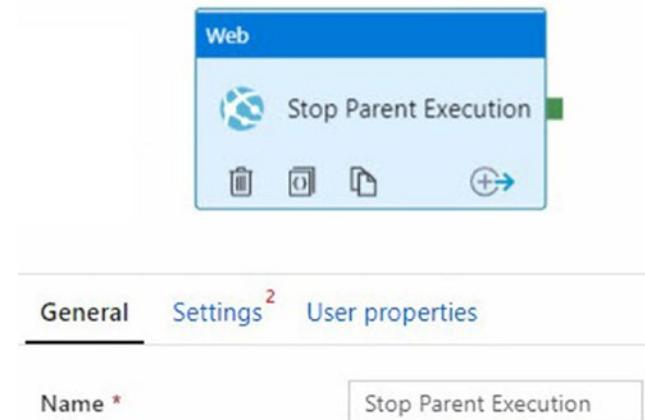


Figure 11-25. Adding the “Stop Parent Execution” web activity

On the “Stop Parent Execution” web activity’s Settings tab, click inside the “URL” property’s textbox and click the “Add dynamic content [Alt + P]” link. When the “Add dynamic content” blade displays, enter the expression “@variables('StopParentPipelineRunIdString')”, as shown in Figure 11-26.

child > If Fail Application On Package Failure > True activities

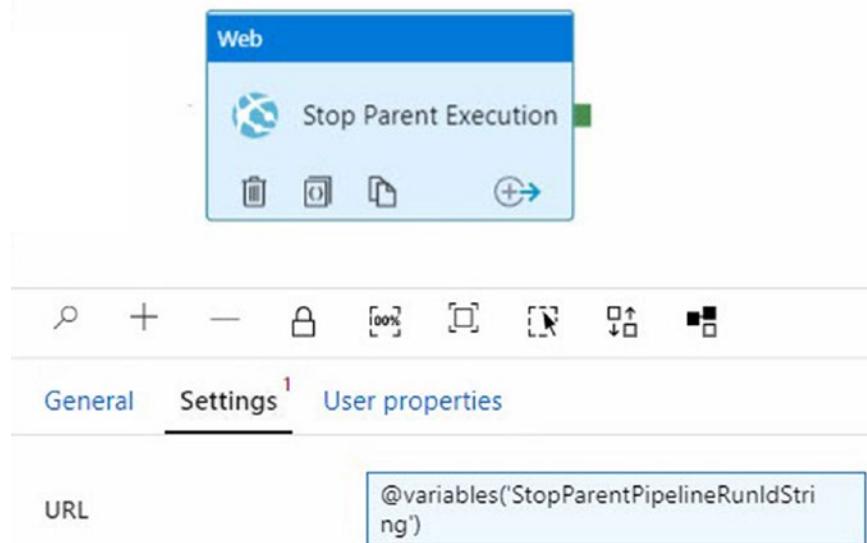


Figure 11-26. Configuring the URL property of the “Stop Parent Execution” web activity

Complete configuration of the properties on the “Stop Parent Execution” web activity’s Settings tab as follows (and as shown in Figure 11-27):

- Method: POST
- Body: {"message": "Stopping the parent pipeline"}
- Advanced ▶ Authentication: MSI
- Advanced ▶ Resource: <https://management.azure.com>

The screenshot shows the 'Settings' tab for a 'Stop Parent Execution' web activity. The configuration includes:

- URL:** @variables('StopParentPipelineRunIdString')
- Method:** POST
- Headers:** New
- Body:** {"message": "Stopping the parent pipeline"}
- Datasets:** Select...
- Linked services:** Select...
- Integration runtime:** AutoResolveIntegrationRuntime
- Authentication:** MSI (selected)
- Resource:** <https://management.azure.com>

Figure 11-27. “Stop Parent Execution” web activity Settings configurations

The preceding completes the fourth and final step for responding to a failed child package execution when the FailApplicationOnPackageFailure bit is configured to true (1): set the command to cancel the parent pipeline.

The activities inside the “If Fail Application On Package Failure” if condition activity’s “True activities,” as configured, work together to update the statuses of the application package instance and application instance and then invoke the Azure Data Factory’s REST API Pipeline Runs Cancel method for the ongoing execution of the parent pipeline.

Before testing the changes, delete the (now orphaned) “Log Application Package Instance Failure” stored procedure activity, as shown in Figure 11-28.

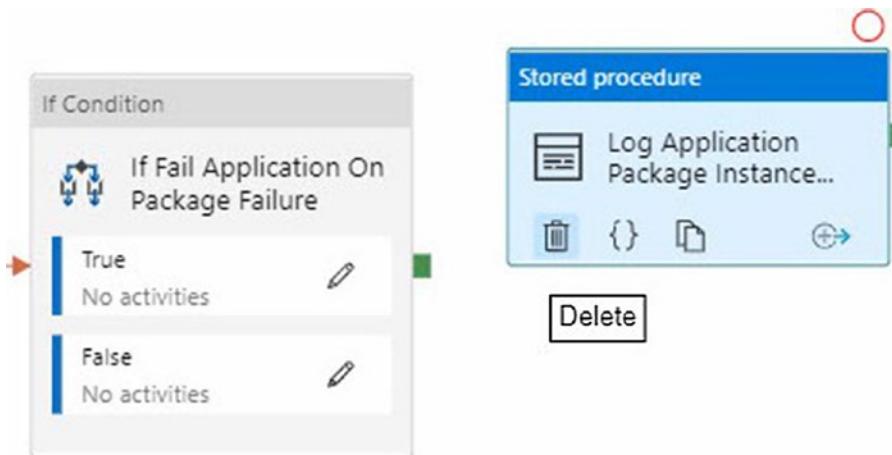


Figure 11-28. Deleting the old “Log Application Package Instance Failure” stored procedure activity

Let’s Test It!

Deploy the parent and child pipelines by clicking the Publish all button in the Data Factory toolbar, as shown in Figure 11-29.



Figure 11-29. Publishing changes to the pipelines and ADF artifacts

In the parent pipeline, click the “Add trigger” menu item, and then click “Trigger now,” as shown in Figure 11-30.



Figure 11-30. Triggering the parent pipeline

When the parent pipeline execution completes, the Monitor page should display results similar to Figure 11-31.

 A screenshot of the ADF Monitor page. On the left, there is a sidebar with icons for Dashboards, Pipeline runs (which is selected and highlighted in blue), Trigger runs, Integration runtimes, Alerts & metrics, and Application packages. The main area is titled 'Pipeline runs' and shows a table of executed pipelines. The table has columns: PIPELINE NAME, DURATION, TRIGGERED BY, STATUS, PARAMETERS, and ERROR. It displays two rows: one for 'child' which failed (status 'Failed') and one for 'parent' which was cancelled (status 'Cancelled').

PIPELINE NAME	DURATION	TRIGGERED BY	STATUS	PARAMETERS	ERROR
child	00:00:30	a0d61d47-8652-4cb6-82	Failed		
parent	00:00:49	Manual trigger	Cancelled		

Figure 11-31. Parent and child pipeline execution results

Execute the T-SQL query shown in Listing 11-2 to confirm the application instance and application package instance results.

Listing 11-2. Application and application package execution results

```
Select top 1
    a.ApplicationName
    , ai.ApplicationInstanceId
    , ai.ApplicationStatus
From log.ApplicationInstance ai
Join config.Applications a
    On a.ApplicationId = ai.ApplicationId
Order By ApplicationInstanceId Desc

declare @ApplicationInstanceId int = (Select top 1 ApplicationInstanceId
                                         From log.ApplicationInstance
                                         Order By ApplicationInstanceId DESC)
```

```

Select p.PackageName
    , ai.ApplicationInstanceId
    , api.ApplicationPackageInstanceId
    , api.ApplicationPackageStatus
    , ap.FailApplicationOnPackageFailure
From log.ApplicationPackageInstance api
Join log.ApplicationInstance ai
    On ai.ApplicationInstanceId = api.ApplicationInstanceId
Join config.ApplicationPackages ap
    On ap.ApplicationPackageId = api.ApplicationPackageId
Join config.Packages p
    On p.PackageId = ap.PackageId
Where ai.ApplicationInstanceId = @ApplicationInstanceId
Order By ApplicationPackageInstanceId

```

The results should appear similar to Figure 11-32.

Results		Messages	
	ApplicationName	ApplicationInstanceId	ApplicationStatus
1	Framework Test	67	Cancelled
PackageName	ApplicationInstanceId	ApplicationPackageInstanceId	ApplicationPackageStatus
1	ReportAndFail.dtsx	67	82
			Failed (FailAppOnPkgFail: 1)
			1

Figure 11-32. Application and application package execution results

The “Framework Test” SSIS framework application contains two application packages, ReportAndFail.dtsx and ReportAndSucceed.dtsx. As advertised, ReportAndFail.dtsx execution fails every single time it is executed. In the SSIS framework metadata, the FailApplicationOnPackageFailure bit is set to 1 (true) for both application packages, as seen when executing the T-SQL query in Listing 11-3 (with results shown in Figure 11-33).

Results		Messages			
	ApplicationPackageId	ApplicationName	PackageName	ExecutionOrder	FailApplicationOnPackageFailure
1	2	Framework Test	ReportAndFail.dtsx	10	1
2	1	Framework Test	ReportAndSucceed.dtsx	20	1

Figure 11-33. Viewing SSIS framework metadata T-SQL query results

Listing 11-3. Viewing SSIS framework metadata

```
Select ap.ApplicationPackageId
, a.ApplicationName
, p.PackageName
, ap.ExecutionOrder
, ap.FailApplicationOnPackageFailure
From config.Applications a
Join config.ApplicationPackages ap
    On ap.ApplicationId = a.ApplicationId
Join config.Packages p
    On p.PackageId = ap.PackageId
Where a.ApplicationName = N'Framework Test'
Order By ap.ExecutionOrder
```

When the parent pipeline is executed, both application packages – ReportAndFail.dtsx and ReportAndSucceed.dtsx – are retrieved from the SSISConfig database and dumped into the parent pipeline’s “ForEach Application Package” ForEach activity’s iterator. Because the ExecutionOrder of ReportAndFail.dtsx is 10 and the ExecutionOrder of ReportAndSucceed.dtsx is 20, ReportAndFail.dtsx executes first. By design, when the execution of ReportAndFail.dtsx fails in the child pipeline and the FailApplicationOnPackageFailure bit is true, the parent pipeline is cancelled.

What happens if ReportAndFail.dtsx fails in the child pipeline and the FailApplicationOnPackageFailure bit is *false*? To test, update the FailApplicationOnPackageFailure bit for the ReportAndFail.dtsx application package using the T-SQL in Listing 11-4.

Listing 11-4. Updating the FailApplicationOnPackageFailure bit for the ReportAndFail.dtsx application package

```
Select ap.ApplicationPackageId
, a.ApplicationName
, p.PackageName
, ap.ExecutionOrder
, ap.FailApplicationOnPackageFailure
From config.Applications a
Join config.ApplicationPackages ap
```

```

On ap.ApplicationId = a.ApplicationId
Join config.Packages p
  On p.PackageId = ap.PackageId
Where a.ApplicationName = N'Framework Test'
Order By ap.ExecutionOrder

Update config.ApplicationPackages
Set FailApplicationOnPackageFailure = 0
Where ApplicationPackageId = 2

Select ap.ApplicationPackageId
, a.ApplicationName
, p.PackageName
, ap.ExecutionOrder
, ap.FailApplicationOnPackageFailure
From config.Applications a
Join config.ApplicationPackages ap
  On ap.ApplicationId = a.ApplicationId
Join config.Packages p
  On p.PackageId = ap.PackageId
Where a.ApplicationName = N'Framework Test'
Order By ap.ExecutionOrder

```

The results of executing the T-SQL query in Listing 11-4 should appear similar to Figure 11-34.

	ApplicationPackageId	ApplicationName	PackageName	ExecutionOrder	FailApplicationOnPackageFailure
1	2	Framework Test	ReportAndFail.dtsx	10	1
2	1	Framework Test	ReportAndSucceed.dtsx	20	1
	ApplicationPackageId	ApplicationName	PackageName	ExecutionOrder	FailApplicationOnPackageFailure
1	2	Framework Test	ReportAndFail.dtsx	10	0
2	1	Framework Test	ReportAndSucceed.dtsx	20	1

Figure 11-34. FailApplicationOnPackageFailure bit for the ReportAndFail.dtsx application package, updated

To test the implications of changing the FailApplicationOnPackageFailure bit for the ReportAndFail.dtsx application package, re-trigger the parent pipeline. The Monitor page surfaces the fault-tolerant results: ReportAndFail.dtsx failed but did not stop execution of the SSIS framework application. ReportAndSucceed.dtsx executed and succeeded, as shown in Figure 11-35.

Pipeline runs				
	All status	Rerun	Cancel	Refresh
Showing 1 - 13 items				
<input type="checkbox"/> PIPELINE NAME	RUN START ↑↓	DURATION	TRIGGERED BY	STATUS
<input type="checkbox"/> child	6/17/20, 3:32:55 PM	00:00:22	91493b3f-b66b-47ee-9c	✓ Succeeded
<input type="checkbox"/> child	6/17/20, 3:32:07 PM	00:00:45	4519d0d5-0a8a-4eb5-a6	✗ Failed
<input type="checkbox"/> parent	6/17/20, 3:31:52 PM	00:01:39	Manual trigger	✗ Failed

Figure 11-35. Fault tolerance in action; an application package failure does not stop application execution

My only complaint is that the parent pipeline reports failure.

Conclusion

In this chapter, we added fault tolerance to the Azure Data Factory version of the SSIS framework. In the process, we learned about configuring security for the Azure Data Factory Managed Identity so web activities in the data factory's pipelines could invoke methods in the Azure Data Factory REST API. We also learned more about the default behavior of the ForEach activity.

Index

A, B

Azure Data Factory (ADF) framework, 209
IAM
 add role assignment, 357
 complete task, 358
 Contributor role, 358
 Select textbox, 357
Azure-SSIS framework
 Azure Storage Explorer
 create File Share, 269
 explorer window, 268
 fs-ssis, 269, 270
 install, 267
 stframeworks storage account, 268
 update PackageLocation Values, 272
 upload SSIS packages, 270
 SSISConfig database
 complete portal, 249
 “Compute + storage” property, 244
 configure database, 245
 connection details, 250
 create Database, 242
 create SSISConfig database, 251
 data details, 244
 enable advanced data security
 property, 247
 reconfiguring database, 246
 results of, query execution, 264
 Review + create page, 248
 SSISConfig DDL, 252–264

subscriptions and resource group
 property, 243
T-SQL query, 267
 updated settings, 247
Azure-SSIS integration
 ADF
 Azure dashboard, 215
 Azure Data Factory page, 216
 create button to provision, 215
 create resource, 209
 Data Factory Overview page, 216
 data factory page, 210
 enable GIT checkbox, 214
 new data factory page, 211
 resource group drop-down, 212
 searching data factory, 209
 select location, 213
 Azure Storage
 Access tier property, 220
 configure account kind
 property, 219
 configure replication
 property, 220
 configuring storage account, 218
 create storage account, 217,
 221, 223
 instance details, 219
 performance property, 219
 storage account name, 218
 storage account setting, 222

INDEX

Azure-SSIS IR
 advanced settings, 233
 configuration choices, 234
 configure name, description and location, 226
 create SSIS catalog, 232
 dashboard, 238
 deployment settings, 231
 edition/license property, 228, 229
 integration runtime setup, 224
 integration runtime
 setup blade, 225
 monitor Integration runtimes, 238
 node number property, 228
 node size option, 226, 227, 230
 opening ADF connections, 224
 reduce BYOL, 229
 running status, 235
 starting status, 235
 status dialog, 239
 stop button, 236, 239
 survey dialog trigger, 237

ConnectionString Override, 148
ConnectionString property, 145
Expression Builder, 146
Expression Editor, 147
Expressions property, 144
file connection manager editor, 139
file system, 137
new connection, 138
Parent.dtsx package tab, 150
ReportAndSucceed.dtsx file, 140
ReportAndSucceed.dtsx File, 142
Task Editor, 137
Test Execution Results, 149
TestSSISSolutionTestSSISProject
 ReportAndSucceed.dtsx file, 140

C, D

Controller pattern
 FWDemo.DailyProcessController, 19
 Monthly Process Controller, 23–25

E

Execute Package Task (EPT)
 adding Execute Child Package, 136
 ChildPackage, 142
 ChildPackage file connection
 manager, 143
 configure EPT Execute Child
 Package, 141

Fault Tolerance Logic
 check logic, 362
 condition activity, 360, 362
 connecting Failure output, 361
 execute application package, 359
 implementation
 adding Log Application Package
 Instance, 369
 add trigger, 379
 application and application
 package, 380
 configuring Log Application
 Instance Cancelled, 372
 deleting Log Application Package
 Instance Failure, 378
 ForEach Array Item, 366
 if condition activity, 363
 linked service property, 364
 Log Application Instance
 Cancelled, 373

- parent and child pipeline
 - execution, [379](#)
- [ReportAndFail.dtsx](#), [382](#)
- Set Cancel Parent Pipeline
 - Command, [374](#)
- stop parent execution, [377](#)
- StopParentPipeline
 - RunIdString, [373, 375](#)
- True activity, [368](#)
- T-SQL query results, [380](#)
- URL property, [376](#)
- web activity, [376](#)
- Framework
 - analysis of, [4](#)
 - create Schema and log table, [5](#)
 - daily Process, [7–12](#)
 - monthly Process, [13–18](#)
- I, J, K**
 - Integrated development environment (IDE), [77](#)
- L**
 - Log.ApplicationInstance Table
 - add Execute Child Pipeline, [330](#)
 - adding child pipeline, [325](#)
 - adding log.UpdateApplication
 - InstanceState, [311](#)
 - adding lookup activity, [307](#)
 - ApplicationInstanceId parameter
 - value, [315](#)
 - ApplicationInstanceId variable, [321](#)
 - Application Instance results, [317](#)
 - ApplicationName parameter, [309](#)
 - ApplicationPackageId metadata, [320](#)
 - ApplicationStatus property Failed, [316](#)
 - Azure Data Studio message, [312](#)
- Child pipeline parameters, [326](#)
- configuring ParentRunId parameter, [332](#)
- deleting Execute Application Package, [330](#)
- Execute Application Package, [327, 329](#)
- Execute Child Pipeline, [333](#)
- import parameter, [308, 314](#)
- linked service, [313](#)
- Log Application Instance Failure, [318](#)
- Log Application Instance Start, [322, 323](#)
- Log Application Instance Success, [319](#)
- log.InsertApplicationInstance, [304, 308](#)
- log.UpdateApplicationInstance
 - Status, [313](#)
- modify table, [303](#)
- PackagePath parameter, [331](#)
- query timeout property, [310](#)
- records, [306](#)
- set ApplicationInstanceId, [324](#)
- setting ApplicationRunId parameter, [310](#)
- ssisFrameworkDataSet property, [307](#)
- text execution, [316](#)
- viewing parent pipeline, [334](#)
- Log Application Package
 - adding, [131](#)
 - add.Net code, [133](#)
 - instrumentation messages, [135](#)
 - logging, [132](#)
 - Parent.dtsx control flow, [134](#)
 - Script Task Editor, [132](#)
- Log execution values
 - adding and renaming, [99](#)
 - ApplicationName parameter, [101](#)
 - debug execution, [103](#)
 - information event, [103](#)
 - information event code, [104](#)
 - ReadOnlyVariables script, [100](#)
 - variables and parameter, [101](#)

M, N

Metadata

- controllers
 - @DailyProcName, 41
 - monthly controllers, 40
 - @MonthlyProcName, 44
 - procedure completion, 42
- create table, 33
- Daily Process Application, 37
- DISABLE procedure, 51
- FDemo.Application, 37
- FDemo.
 - ApplicationProcessProcedure, 35
 - FDemo.Application table, 34
 - FDemo.ProcessProcedure, 33
 - FDemo.ProcessProcedure table, 36
 - monthly process application, 37
 - orchestrator, 47–50
 - ProcessLogID, 51
 - ProcessLogMessage, 54
 - Process log output, 50, 53

O, P, Q, R

- Orchestrator process
 - execute statement, 30
 - FDemo.ProcessLog, 30, 31
 - FDemo.ProcessOrchestrator, 27
 - @RunDate, 30

S, T, U, V, W, X, Y, Z

- SQL Get Application Packages
 - adding, 105
 - adding SQLStatement property, 112
 - ADO.NET, 106
 - configure ADO.NET Connection Manager, 108–110
- execute SQL Task Editor, 113
- framework Test, 118
- full result set, 115
- new Connection, 107
- new variable, 116
- parameter Mapping, 114
- parent container, 117
- SQL task editor, 111, 119
- SSISConfig database, 111
- task editor, 106
- SQL Server Management Studio (SSMS), 63, 160
- SSIS framework
 - ApplicationName parameter
 - add parameter, 302
 - parameter configuration, 302
 - add ReportAndFail, 82
 - Application Packages
 - attributes of, execution, 73
 - config. ApplicationPackages table, 69, 71, 72
 - config.Packages table, 67, 69
 - default constraint, 73
 - keys management, 74
 - components, 61
 - config.Applications table, 66, 67
 - Config Schema, 63, 65
 - create TestSSISSolution, 78
 - debug execution, 81
 - empathy, 60
 - error event, 84
 - functionality, 60
 - information event, 82
 - install Integration Services extension, 77
 - install Visual Studio 2019, 76
 - log.ApplicationPackageInstance table
 - add dynamic content, 340

adding the Application
 PackageRunId, 335
 ApplicationInstanceId parameter,
 341, 342
 ApplicationPackageRunId
 parameter, 343
 ApplicationStatus parameter, 348
 failure, 346
 Log Application Package Instance
 Failure, 350
 Log Application Package Instance
 Success, 351
 log.InsertApplicationPackage
 Instance, 336, 338, 340
 log.UpdateApplicationPackage
 InstanceStatus, 344
 mapping parameter, 347
 parent pipeline activity, 348
 ReportAndSucceed.dtsx, 352
 results, 350
 start lookup activity, 339
 updating
 ApplicationPackageStatus, 352
 Main() method, 83
 metadata-driven execution
 management, 62
 principles, 60
 ReportAndSucceed.dtsx, 78
 Script Task Editor, 80
 SCR Log Values, 79
 simplicity, 60
 SQL Server Integration Services, 77
 SSISConfig database, 62, 63

SSIS Framework ADF Execution Engine
 activity outputs, 288
 add dynamic content, 288
 add ForEach activity, 287
 add Source dataset, 276

Azure-SSIS-Files integration, 290
 change name, 274
 configuration path, 293
 connection test, 283
 copy Primary Key, 293
 create new pipeline, 274
 event_messages log, 297
 Execute Application Package, 290,
 292, 297
 failed connection test, 279
 ForEach Application Packages Item, 289
 Get Application Packages, 275
 JSON return, 286
 logging level, 294
 name property, 278
 new dataset, 277
 PackagePath property, 291
 parent pipeline, 295
 Password property textbox, 294
 ReportAndSucceed.dtsx, 298, 299
 server name link, 281
 show firewall settings, 282
 SSISConfig database, 281
 T-SQL statement, 284
 view logs folder, 296
 view output, 286
 view SQL database, 280

SSIS framework application
 FailApplicationOnPackageFailure
 functionality, 197
 FailApplicationOnPackageFailure
 logic, 195
 metadata
 ApplicationPackageStatus result, 206
 FailPackageOnFailure property, 200
 fault tolerance, 201
 foreach loop container, 198
 MaximumErrorCount property, 199

INDEX

- SSIS framework application (*cont.*)
- running
 - ApplicationPaackageStatus, 204
 - SQL Log Application Package
 - Instance Failure, 205
 - update ReportAndFail.dtsx, 202
 - package update results, 194
 - Parent.dtsx execution fails, 195
 - view instance log data, 191
 - view instance result, 191, 193
- SSIS framework logging
- create log schema, 159
 - log.ApplicationInstance
 - table, 161, 162
 - log. ApplicationPackageInstance
 - table, 162, 164
 - SSMS, 160
- SQL Log Application Instance
- adding, 165
 - adding T-SQL, 172
 - add Variable, 169
 - ApplicationInstanceId, 176
 - ConnectionType property, 174
 - create new variable, 168
 - Event Handlers tab, 173
 - OnError Event handler, 174
 - parameter Mapping, 168, 172
 - parent Executable, 173
 - ResultSet property, 166
 - SQL Log Application Instance
 - Success, 171
 - SQL Query dialog, 167, 175
 - SQL Task Editor, 170
- SQL Log Application Package Instance
- adding, 177
 - adding T-SQL, 185
 - add Variable dialog, 182
 - ConnectionType property, 178
- failure script task, 187
- FOREACH Application Package, 189, 190
- foreach loop container, 183
- new variable, 181
- parameter Mapping, 180, 185
- SQL Query dialog, 179, 188
- success script task, 184
- SSIS Framework Metadata Management
- add Framework Test, 87
 - add SSIS application, 85
 - add SSIS packages, 88
 - assign SSIS Application Packages, 91, 94
 - config schema, 95
 - framework application, 95
 - re-executing SSIS Application Packages, 94, 95
- ReportAndSucceed.dtsx, 90
- Test ApplicationId, 87
- T-SQL query, 96
- SSIS package
- EPT
 - adding Execute Child Package, 136
 - ChildPackage, 142
 - ChildPackage file connection manager, 143
 - Configure EPT Execute Child Package, 141
 - ConnectionString Override, 148
 - ConnectionString property, 145
 - Expression Builder, 146
 - Expression Editor, 147
 - Expressions property, 144
 - file Connection Manager editor, 139
 - file system, 137
 - new connection, 138
- Parent.dtsx package tab, 150

- ReportAndSucceed.dtsx file, [140](#)
- ReportAndSucceed.dtsx File, [142](#)
- Task Editor, [137](#)
 - test execution results, [149](#)
- iterate collections
 - ADO.Net dataset, [126](#)
 - ADO object source variable, [123](#)
 - assign User, [125](#)
 - changing ordinal, [126](#)
 - collection page, [124](#)
 - ExecutionOrder column, [128](#)
 - Foreach ADO Enumerator, [122](#)
 - FOREACH Application Package, [120](#)
 - PackagePath variable, [125](#)
 - Parent.dtsx control flow, [130](#)
 - SSIS variable, [127](#)
 - user, [129](#)
 - variable Mappings page, [124](#)
- Log Application Package
 - adding, [131](#)
 - add.Net code, [133](#)
 - instrumentation messages, [135](#)
 - Logging, [132](#)
 - Parent.dtsx control flow, [134](#)
 - Script Task Editor, [132](#)
- log execution values
 - adding and renaming, [99](#)
 - ApplicationName parameter, [101](#)
 - debug execution, [103](#)
 - Information event, [103](#)
- Information event code, [104](#)
- ReadOnlyVariables script, [100](#)
- variables and parameter, [101](#)
- Log Package Execution
 - adding script, [151](#)
 - application package, [152](#)
 - failure script task, [154](#)
 - information event, [153](#)
 - .Net code, [155, 156](#)
 - Parent.dtsx execution fails, [157](#)
- ReadOnlyVariables
 - property, [151, 155](#)
 - success script task, [153](#)
 - viewing application package, [158](#)
- SQL Get Application Packages
 - adding, [105](#)
 - adding SQLStatement property, [112](#)
- ADO.NET, [106](#)
- Configure ADO.NET Connection Manager, [108–110](#)
- Execute SQL Task Editor, [113](#)
- framework Test, [118](#)
- full result set, [115](#)
- new Connection, [107](#)
- new variable, [116](#)
- parameter mapping, [114](#)
- parent container, [117](#)
- SQL task editor, [111, 119](#)
- SSISConfig database, [111](#)
- task editor, [106](#)